

Your grade: 100%

Your latest: 100% • Your highest: 100% • To pass you need at least 80%. We keep your highest score.

Next item →

1.

A developer is building an AI agent that interacts with a user over multiple steps. The agent begins making inconsistent decisions because it does not remember actions it previously took (e.g., API calls and results). What is the most effective way to fix this issue?

Store past interactions externally and include only the most relevant messages in each new prompt

Include previous user and assistant messages, including past actions, in the prompt for each new interaction

Increase the randomness of the model's responses to improve adaptability across interactions

Nice work

LLMs are stateless, so including prior messages (especially assistant actions and responses) allows the model to maintain continuity and make informed decisions.

1 / 1 point
2.

A developer uses a single prompt to ask an LLM to:

generate application code

write documentation

and create test cases

The output is inconsistent and incomplete across all three tasks.

Which approach would most effectively improve the quality of the results?

Context Management

Progressive Enhancement

Prompt Chaining

Output Processing

Nice work

Breaking the task into sequential steps allows the model to focus on one aspect at a time, improving quality and coherence.

1 / 1 point
3.

Why are the results of API or tool calls typically represented as “user” messages in an AI agent system?

To provide the model with the result as new input so it can decide the next action

To allow the model to store results internally for future use

To keep the conversation consistent by treating tool outputs as part of the ongoing interaction

Nice work

Tool outputs are passed back as input so the model can interpret them and determine what to do next.

1 / 1 point
4.

A developer is building an AI agent that needs to trigger actions like API calls and database updates. However, the model's responses vary in format and include extra explanations, making them difficult to execute programmatically. What is the most effective way to address this issue?

Ask the model to generate more detailed explanations so the system has more context to interpret

Constrain the model to produce responses in a consistent, structured format that can be parsed into actions

Allow the model to generate free-form responses and rely on heuristics to extract the intended action

Increase the number of examples in the prompt so the model learns to behave more consistently over time

Nice work

Structured outputs allow software to reliably interpret the model's response and map it to executable actions.

1 / 1 point
5.

In a multi-message AI interaction, which message type is primarily used to define the model's behavior and constraints before it processes user input?

User message, because it contains the task or question

System message, because it establishes instructions that guide how the model responds throughout the interaction

All message types equally contribute to defining behavior

Nice work

System messages set the rules and behavior for the model, influencing how it interprets and responds to all subsequent messages.

1 / 1 point