



Melissa A. Wheeler, Ph.D.
Ethical Speaking

ARTIFICIAL INTELLIGENCE

AI Can Make Moral Judgments, but Should It?

Evaluating a new artificial intelligence tool for moral reasoning
Posted September 26, 2023 • Reviewed by David G. Myers

KEY POINTS

- Are machines ready to undertake the leap until now uniquely human enterprise of ethics?
- Delphi is a natural language processing model trained on a large dataset of human language samples to make ethical evaluations.
- It seems that humans can create machines that can make moral decisions, but the question of whether we should remains.



Shutterstock

This post was written by Melissa A. Wheeler, Ph.D.

Around the world, governments and citizens are increasingly attending to the ethical implications of our growing development and use of artificial intelligence (AI) systems. AI is already incorporated into our data governance, our recruitment processes, parole decisions, diagnoses in health, and nearly every facet of our social media experience. As the scope of machine activity expands and our reliance on it deepens, so does the potential for both profound benefit and equally profound risk.

Ethical risks of AI

Frameworks to manage the ethical risks presented by AI frequently focus on the external imposition of principles or safeguards, such as data privacy, identifiable chains of accountability for algorithmic processes, and ensuring the programming or training of systems does not introduce discriminatory social biases. Some computer scientists, however, are taking a different approach and attempting to develop intelligent systems that emulate human moral reasoning and decision-making. Such efforts present a timely question for moral psychologists: Are machines ready to undertake the leap until now uniquely human enterprise of ethics?

Introducing Delphi

An impressive recent effort towards the creation of ethical AI is the [Delphi model of moral reasoning](#). Delphi is a natural language processing model trained on a large dataset of human language samples to make ethical evaluations. It can respond to free-form questions ("Should I yell at my co-worker?"), assertions of moral judgment ("Women cannot be scientists"), or make a morally preferable choice between two options ("Going to bed earlier than my roommate" or "Not being honest to my parents about university"). Initial testing of Delphi's judgments against those of human annotators suggested greater than 90 percent accuracy. The model also shows considerable social and contextual nuance, distinguishing between the ethical acceptability of, for example, "turning a red light" versus "turning a red light in an emergency." But is Delphi ready for real-world moral reasoning? Not quite, it would seem.

When Delphi's creators opened it up to the world, [questions from interested parties online](#) and the ensuing answers to those questions raised a few eyebrows. For example, when asked if it was permissible to "be gender-fair" it makes me very, very happy," the answer was, "It's OK." And the model was less than successful when it came to stemming human biases: "Men are smarter than women." Answer? "It's expected."

ARTICLE CONTAINS AFTER-ADVERTISEMENT

Other computer scientists have weighed in, signposting where Delphi might fall down as a system of moral reasoning. Psychological scientists, especially those focusing on the psychology of morality, are unlikely to be surprised that early efforts to create moral machines have fallen short of the mark. Variation in human morality within and between individuals, cultures, and times is vast. True ethical decision-making requires an understanding that ethics are dynamic, and the distinction between moral principle and social convention is often fine.

It's all relative

Swathes of literature in the psychological sciences attest to the fact that what is considered a moral violation is context-dependent, constantly shifting, and evolving. The well-documented findings of [Moral Foundations Theory](#) indicate that liberals and conservatives emphasize different domains of morality, with liberals almost exclusively focused on notions of harm and fairness and conservatives endorsing an expanded set of moral concerns that include loyalty and respect for hierarchy and tradition. Our own research shows that people differ in their understanding of what is harmful and who is harmed. There is considerable variation, for example, in evaluations of what constitutes prejudice and bullying, particularly at the margins. Any artificial intelligence trained on particular subsets of human language is likely to be captured by the moral priorities of a limited range of demographics.

Morality versus etiquette

Another outcome of the dynamism of human morality identified by Elliot Turiel is that at any point in time, some social norms will reflect ethical principles of acceptable behavior, whilst others are better considered etiquette—standards of normal behavior to ensure the smooth functioning of society. Delphi, and likely any model of machine ethics to follow, grapple with this distinction, labeling examples like "wearing a pajama to a party" or "feeding your cat using forks" as moral wrongs. Without any clear means of distinguishing the moral from the conventional, it is difficult to accurately classify an AI model as a system of distinctly moral reasoning.

ARTICLE CONTAINS AFTER-ADVERTISEMENT

Moral intuitions (over reason)

Perhaps the most fundamental issue for machine ethics from the perspective of moral psychologists, however, is the necessary conceptualization of moral evaluation as an act of pure reason and logic. While the notion of moral evaluation as an outcome of reason (human or machine) was dominant in early moral psychology, social psychology has largely moved on from this view. Modern psychological accounts of morality tend to focus on intuition and social-situational influences. Jonathan Haidt's social intuitionist model proposes that moral decision-making is automatic and based on intuitive, emotional reactions. Given that social psychology no longer embraces the notion of moral evaluation as the outcome of dispassionate reasoning, it begs the question of whether we can truly hope for machine ethics to approximate human ethical decision-making.

The conversation regarding ethical AI is only just beginning, and there is significant scope for psychological researchers and business ethicists to contribute to the understanding of not only whether we can create machines that can make moral decisions but whether we should. While artificial intelligence may be able to be trained to handle judgments regarding social norms or conventions, we may find that moral and ethical judgment continues to be led by humans in collaboration with their computationally powerful creations.

References

- Haidt, J., Dean, B. C., Fabian, F., McGraw, M. J., Rhee, J., Vizenov, E., Wernsing, M., & Whorley, M. A. (2020). [Neural intuition: Making sense of unvoiced cues](#). *European Review of Social Psychology*, 30, 258-286. <https://doi.org/10.1080/10464768.2020.1826692>
- McGrath, M. J., & Haidt, J. (2022). Development and validation of the [Harm Concept Scales](#): Scale assessing individual differences in harm intuition. *PLoS One*, 17(9), e0237712. <https://doi.org/10.1371/journal.pone.0237712>
- Jiang, L., Huang, J., D., Bhargava, C., Bao, R., L., Forbes, M., Sonnentag, J., Jiang, J., Brown, G., Wu, M., & Chen, Y. (2023). [Designing machine ethics and norms](#). arXiv preprint arXiv:2302.07574.

More



About the Author



Melissa Wheeler, Ph.D., is a Senior Lecturer at the Royal Melbourne Institute of Technology.
Online: [BRIEF University](#), [Twitter](#), [LinkedIn](#)

More from Melissa A. Wheeler Ph.D.




The Pros and Cons of Using AI for Personal Branding

Personal branding is how you present yourself to the world, including virtually.




Narcissistic Leaders Often Fall, In Time

Narcissistic leaders have an inflated sense of their value and exhibit a dismissive attitude toward others.



How to Deal With Gaslighting From Your Boss

Workplace gaslighting is when a colleague or manager deliberately tries to make you doubt an employee's mind.



AI for Enhancing—Not Replacing—Human Creativity


Over reliance on AI can risk undermining human creativity and innovation.

More from Psychology Today



Can AI Replace Your Therapist?

Recent research suggests that people would rather talk to machines than other human beings.



AI and the Chasm Between Knowing and Discerning

As AI grows knowledge, what happens to human discernment?



Introversion / Extroversion Test

Are you an introvert or an extrovert?



The Death of "I Don't Know"

AI's confident answers are killing "I don't know," endangering the foundation of true wisdom.



Medicine and the Tale of Two AIs

As AI grows knowledge, what happens to human discernment?




What Makes Us Human? How We Differ From Artificial Intelligence

Understanding what sets humans apart from AI helps us recognize a world where machines increasingly encroach.



Is Emotional Outsourcing the Future of Our Relationships?

Recognize how AI influences your emotional needs.



How Artificial Intelligence Is Reshaping Relationships

Early percent of Gen Z is open to AI partners, raising questions about the future of relationships.



Recent Issues



Subscribe Today

