Hide menu
About this Course
Key Components of Llamafile
Developing Portable CLI with Cosmopolitan
Running Llamafile
Conclusion and Next Steps
Reading: Key Terms
1 min
Reading: Llamafile server
5 min
Reading: Course Conclusion
5 min
Graded Assignment: Final Quiz-Llamafile
Submitted
Reading: Next Steps
1 min
Lab: Interactive Llamafile Sandbox
1h

Your grade: 100%

Your latest: **100%** • Your highest: **100%** • To pass you need at least 80%. We keep your highest score.

Next item

Final Quiz-Llamafile

Review Learning Objectives

Instructions

coach

Ready to review what you've learned before starting the assignment? I'm here to help.

Help me practice Let's chat

Assignment details

Due Aug 16, 11:59 PM IST
Attempts Unlimited
Try again

Submitted
Aug 12, 3:17 AM IST

Your grade
To pass you need at least 80%. We keep your highest score.
View submission See feedback
100%

Like Dislike Report an issue

1. What does the llama.cpp example server demonstrate? **1 / 1 point**

- How to train a language model from scratch
- How to deploy a language model on mobile devices
- How to expose a language model through an HTTP API
- How to optimize a language model for better performance

Correct
Correct. The primary purpose of the example server is to demonstrate serving the model via a web API.

2. Which of the following is NOT an API endpoint provided by the example server? **1 / 1 point**

- /train
- /completion
- /tokenize
- /embedding

Correct
Correct. The example server does not include an endpoint for training the model.

3. How can the model and generation settings be customized when running the server? **1 / 1 point**

- By modifying the server source code
- By providing a config file in YAML format
- By passing command line options
- By setting environment variables

Correct
Correct. The server accepts various command line flags to specify things like model path, context size, etc.

4. What does the /completion API endpoint accept to generate text? **1 / 1 point**

- An image file to be captioned
- A text prompt to continue
- An audio clip to be transcribed
- A numerical vector representing text

Correct
Correct. The endpoint takes a prompt parameter containing the text to generate a completion for.

5. What task does the /tokenize API endpoint perform? **1 / 1 point**

- Generating text completions
- Splitting text into individual tokens
- Converting tokens back into readable text
- Extracting embeddings from text

Correct
Correct. The /tokenize endpoint takes a text input and returns the corresponding list of tokens.