## Your grade: 100%

Your latest: 100% • Your highest: 100% • To pass you need at least 80%. We keep your highest score.

Next item →

1. In an RNN encoder-decoder model without attention, what does the encoder typically produce? 1 / 1 point

- (•) A single vector representing the input sequence
- ( ) A sequence of images
- ( ) A probability distribution over classes
- ( ) A sequence of words

Nice work

2. True or False: Sequence-to-sequence models can be based on LSTMs. 1 / 1 point

- (•) True
- ( ) False

Nice work

3. Which of the following is a popular variant of the attention mechanism used in transformers? 1 / 1 point

- ( ) Local attention
- ( ) Gated attention
- ( ) Global attention
- (•) Scaled dot-product attention

Nice work

4. True or False: Transformer-based language models are typically encoder-decoder models. 1 / 1 point

- (•) False
- ( ) True

Nice work

5. What can be a reason for overfitting? 1 / 1 point

- (•) Too much model capacity for the training set
- ( ) Not training for enough epochs
- ( ) A hidden size that is small compared to the output dimension
- ( ) All of the above

Nice work