

Advances in Neural Information Processing Systems 35 (NeurIPS 2022)

Edited by: S. Koyejo and S. Mohamed and A. Agarwal and D. Belgrave and K. Cho and A. Oh

ISBN: 9781713871088

[Purchase Printed Proceeding](#)

● Main Conference Track ● Datasets and Benchmarks Track

- [Federated Submodel Optimization for Hot and Cold Data Features](#) Yucheng Ding, Chaoyue Niu, Fan Wu, Shaojie Tang, Chengfei Lyu, yanghe feng, Guihai Chen
- [On Kernelized Multi-Armed Bandits with Constraints](#) Xingyu Zhou, Bo Ji
- [Geometric Order Learning for Rank Estimation](#) Seon-Ho Lee, Nyeong Ho Shin, Chang-Su Kim
- [Structured Recognition for Generative Models with Explaining Away](#) Changmin Yu, Hugo Soulat, Neil Burgess, Maneesh Sahani
- [NAS-Bench-Graph: Benchmarking Graph Neural Architecture Search](#) Yijian Qin, Ziwei Zhang, Xin Wang, Zeyang Zhang, Wenwu Zhu
- [Fast Bayesian Coresets via Subsampling and Quasi-Newton Refinement](#) Cian Naik, Judith Rousseau, Trevor Campbell
- [What You See is What You Classify: Black Box Attributions](#) Steven Stalder, Nathanael Perraudin, Radhakrishna Achanta, Fernando Perez-Cruz, Michele Volpi
- [Adaptive Interest for Emphatic Reinforcement Learning](#) Martin Klissarov, Rasool Fakoor, Jonas W. Mueller, Kavosh Asadi, Taesup Kim, Alexander J. Smola
- [Scaling & Shifting Your Features: A New Baseline for Efficient Model Tuning](#) Dongze Lian, Daquan Zhou, Jiashi Feng, Xinchao Wang
- [Zero-Shot Video Question Answering via Frozen Bidirectional Language Models](#) Antoine Yang, Antoine Miech, Josef Sivic, Ivan Laptev, Cordelia Schmid
- [Active Learning with Neural Networks: Insights from Nonparametric Statistics](#) Yinglun Zhu, Robert Nowak
- [IM-Loss: Information Maximization Loss for Spiking Neural Networks](#) Yufei Guo, Yuanpei Chen, Liwen Zhang, Xiaode Liu, Yinglei Wang, Xuhui Huang, Zhe Ma
- [Using natural language and program abstractions to instill human inductive biases in machines](#) Sreejan Kumar, Carlos G. Correa, Ishita Dasgupta, Raja Marjeh, Michael Y Hu, Robert Hawkins, Jonathan D Cohen, nathaniel daw, Karthik Narasimhan, Tom Griffiths
- [Second Thoughts are Best: Learning to Re-Align With Human Values from Text Edits](#) Ruibo Liu, Chenyan Jia, Ge Zhang, Ziyu Zhuang, Tony Liu, Soroush Vosoughi
- [SatMAE: Pre-training Transformers for Temporal and Multi-Spectral Satellite Imagery](#) Yezhen Cong, Samar Khanna, Chenlin Meng, Patrick Liu, Erik Rozi, Yutong He, Marshall Burke, David Lobell, Stefano Ermon
- [On Sample Optimality in Personalized Collaborative and Federated Learning](#) Mathieu Even, Laurent Massoulié, Kevin Scaman
- [Offline Multi-Agent Reinforcement Learning with Knowledge Distillation](#) Wei-Cheng Tseng, Tsun-Hsuan Johnson Wang, Yen-Chen Lin, Phillip Isola
- [Decentralized Gossip-Based Stochastic Bilevel Optimization over Communication Networks](#) Shuoguang Yang, Xuezhou Zhang, Mengdi Wang
- [Conditional Meta-Learning of Linear Representations](#) Giulia Denevi, Massimiliano Pontil, Carlo Ciliberto
- [Theory and Approximate Solvers for Branched Optimal Transport with Multiple Sources](#) Peter Lippmann, Enrique Fita Sanmartín, Fred A. Hamprecht
- [CHIMLE: Conditional Hierarchical IMLE for Multimodal Conditional Image Synthesis](#) Shichong Peng, Seyed Alireza Moazenipourasil, Ke Li
- [Active Ranking without Strong Stochastic Transitivity](#) Hao Lou, Tao Jin, Yue Wu, Pan Xu, Quanquan Gu, Farzad Farnoud
- [Offline Goal-Conditioned Reinforcement Learning via \$f\$ -Advantage Regression](#) Jason Yecheng Ma, Jason Yan, Dinesh Jayaraman, Osbert Bastani
- [Rethinking and Improving Robustness of Convolutional Neural Networks: a Shapley Value-based Approach in Frequency Domain](#) Yiting Chen, Qibing Ren, Junchi Yan
- [Adversarial Style Augmentation for Domain Generalized Urban-Scene Segmentation](#) Zhun Zhong, Yuyang Zhao, Gim Hee Lee, Nicu Sebe
- [Fully Sparse 3D Object Detection](#) Lue Fan, Feng Wang, Naiyan Wang, ZHAO-XIANG ZHANG
- [Diffusion Visual Counterfactual Explanations](#) Maximilian Augustin, Valentyn Boreiko, Francesco Croce, Matthias Hein
- [Recurrent Video Restoration Transformer with Guided Deformable Attention](#) Jingyun Liang, Yuchen Fan, Xiaoyu Xiang, Rakesh Ranjan, Eddy Ilg, Simon Green, Jiezhong Cao, Kai Zhang, Radu Timofte, Luc V Gool
- [A Consolidated Cross-Validation Algorithm for Support Vector Machines via Data Reduction](#) Boxiang Wang, Archer Yang
- [On-Demand Sampling: Learning Optimally from Multiple Distributions](#) Nika Haghtalab, Michael Jordan, Eric Zhao
- [Asynchronous SGD Beats Minibatch SGD Under Arbitrary Delays](#) Konstantin Mishchenko, Francis Bach, Mathieu Even, Blake E. Woodworth
- [Coresets for Relational Data and The Applications](#) Jiaxiang Chen, Qingyuan Yang, Ruomin Huang, Hu Ding
- [Model-Based Offline Reinforcement Learning with Pessimism-Modulated Dynamics Belief](#) Kaiyang Guo, Shao Yunfeng, Yanhui Geng
- [Generating Training Data with Language Models: Towards Zero-Shot Language Understanding](#) Yu Meng, Jiaxin Huang, Yu Zhang, Jiawei Han
- [Wavelet Score-Based Generative Modeling](#) Florentin Guth, Simon Coste, Valentin De Bortoli, Stephane Mallat

[Report an Issue](#) | [Name Change Policy](#)

- [Why do tree-based models still outperform deep learning on typical tabular data?](#) Leo Grinsztajn, Edouard Oyallon, Gael Varoquaux
- [Generalizing Consistent Multi-Class Classification with Rejection to be Compatible with Arbitrary Losses](#) Yuzhou Cao, Tianchi Cai, Lei Feng, Lihong Gu, Jinjie GU, Bo An, Gang Niu, Masashi Sugiyama
- [Markovian Interference in Experiments](#) Vivek Farias, Andrew Li, Tianyi Peng, Andrew Zheng
- [Identifiability and generalizability from multiple experts in Inverse Reinforcement Learning](#) Paul Rolland, Luca Viano, Norman Schürhoff, Boris Nikolov, Volkan Cevher
- [Parallel Tempering With a Variational Reference](#) Nikola Surjanovic, Saifuddin Syed, Alexandre Bouchard-Côté, Trevor Campbell
- [Provably Efficient Reinforcement Learning in Partially Observable Dynamical Systems](#) Masatoshi Uehara, Ayush Sekhari, Jason D. Lee, Nathan Kallus, Wen Sun
- [Off-Policy Evaluation for Episodic Partially Observable Markov Decision Processes under Non-Parametric Models](#) Rui Miao, Zhengling Qi, Xiaoke Zhang
- [Efficient Knowledge Distillation from Model Checkpoints](#) Chaofei Wang, Qisen Yang, Rui Huang, Shiji Song, Gao Huang
- [Decoupled Self-supervised Learning for Graphs](#) Teng Xiao, Zhengyu Chen, Zhimeng Guo, Zeyang Zhuang, Suhang Wang
- [Shadow Knowledge Distillation: Bridging Offline and Online Knowledge Transfer](#) Lujun Li, ZHE JIN
- [ComENet: Towards Complete and Efficient Message Passing for 3D Molecular Graphs](#) Limei Wang, Yi Liu, Yuchao Lin, Haoran Liu, Shuiwang Ji
- [VLMbench: A Compositional Benchmark for Vision-and-Language Manipulation](#) Kaizhi Zheng, Xiaotong Chen, Odest Chadwicke Jenkins, Xin Wang
- [Tiered Reinforcement Learning: Pessimism in the Face of Uncertainty and Constant Regret](#) Jiawei Huang, Li Zhao, Tao Qin, Wei Chen, Nan Jiang, Tie-Yan Liu
- [Between Stochastic and Adversarial Online Convex Optimization: Improved Regret Bounds via Smoothness](#) Sarah Sachs, Hedi Hadiji, Tim van Erven, Cristóbal Guzmán
- [Differentially Private Learning Needs Hidden State \(Or Much Faster Convergence\)](#) Jiayuan Ye, Reza Shokri
- [BR-SNIS: Bias Reduced Self-Normalized Importance Sampling](#) Gabriel Cardoso, Sergey Samsonov, Achille Thin, Eric Moulines, Jimmy Olsson
- [Learning to Configure Computer Networks with Neural Algorithmic Reasoning](#) Luca Beurer-Kellner, Martin Vechev, Laurent Vanbever, Petar Veličković
- [Early Stage Convergence and Global Convergence of Training Mildly Parameterized Neural Networks](#) Mingze Wang, Chao Ma
- [On Divergence Measures for Bayesian Pseudocoresets](#) Balhae Kim, Jungwon Choi, Seanie Lee, Yoonho Lee, Jung-Woo Ha, Juho Lee
- [Unsupervised Learning of Equivariant Structure from Sequences](#) Takeru Miyato, Masanori Koyama, Kenji Fukumizu
- [Multi-Class \$H\$ -Consistency Bounds](#) Pranjal Awasthi, Anqi Mao, Mehryar Mohri, Yutao Zhong
- [On the Frequency-bias of Coordinate-MLPs](#) Sameera Ramasinghe, Lachlan E. MacDonald, Simon Lucey
- [DC-BENCH: Dataset Condensation Benchmark](#) Justin CUI, Ruochen Wang, Si Si, Cho-Jui Hsieh
- [Mask Matching Transformer for Few-Shot Segmentation](#) siyu jiao, Gengwei Zhang, Shant Navasardyan, Ling Chen, Yao Zhao, Yunchao Wei, Humphrey Shi
- [Queue Up Your Regrets: Achieving the Dynamic Capacity Region of Multiplayer Bandits](#) Ilai Bistritz, Nicholas Bambos
- [Differentially Private Covariance Revisited](#) Wei Dong, Yuting Liang, Ke Yi
- [Trimmed Maximum Likelihood Estimation for Robust Generalized Linear Model](#) Pranjal Awasthi, Abhimanyu Das, Weihao Kong, Rajat Sen
- [Causal Discovery in Linear Latent Variable Models Subject to Measurement Error](#) Yuqin Yang, AmirEmad Ghassami, Mohamed Nafea, Negar Kiyavash, Kun Zhang, Ilya Shpitser
- [Density-driven Regularization for Out-of-distribution Detection](#) Wenjian Huang, Hao Wang, Jiahao Xia, Chengyan Wang, Jianguo Zhang
- [Sparsity in Continuous-Depth Neural Networks](#) Hananeh Aliee, Till Richter, Mikhail Solonin, Ignacio Ibarra, Fabian Theis, Niki Kilbertus
- [Environment Diversification with Multi-head Neural Network for Invariant Learning](#) Bo-Wei Huang, Keng-Te Liao, Chang-Sheng Kao, Shou-De Lin
- [Learning Probabilistic Models from Generator Latent Spaces with Hat EBM](#) Mitch Hill, Erik Nijkamp, Jonathan Mitchell, Bo Pang, Song-Chun Zhu
- [Learning Best Combination for Efficient N:M Sparsity](#) Yuxin Zhang, Mingbao Lin, ZhiHang Lin, Yiting Luo, Ke Li, Fei Chao, Yongjian Wu, Rongrong Ji
- [Why Do Artificially Generated Data Help Adversarial Robustness](#) Yue Xing, Qifan Song, Guang Cheng
- [Neural Surface Reconstruction of Dynamic Scenes with Monocular RGB-D Camera](#) Hongrui Cai, Wanquan Feng, Xuetao Feng, Yan Wang, Juyong Zhang
- [Global Optimal K-Medoids Clustering of One Million Samples](#) Jiayang Ren, Kaixun Hua, Yankai Cao
- [Batch Multi-Fidelity Active Learning with Budget Constraints](#) Shibo Li, Jeff M Phillips, Xin Yu, Robert Kirby, Shandian Zhe
- [UniCLIP: Unified Framework for Contrastive Language-Image Pre-training](#) Janghyeon Lee, Jongsuk Kim, Hyounguk Shon, Bumsoo Kim, Seung Hwan Kim, Honglak Lee, Junmo Kim
- [Efficient Multi-agent Communication via Self-supervised Information Aggregation](#) Cong Guan, Feng Chen, Lei Yuan, Chenghe Wang, Hao Yin, Zongzhang Zhang, Yang Yu
- [Accelerated Training of Physics-Informed Neural Networks \(PINNs\) using Meshless Discretizations](#) Ramansh Sharma, Varun Shankar

- [DOPE: Doubly Optimistic and Pessimistic Exploration for Safe Reinforcement Learning](#) Archana Bura, Aria HasanzadeZonuzy, Dileep Kalathil, Srinivas Shakkottai, Jean-Francois Chamberland
- [Improved Regret Analysis for Variance-Adaptive Linear Bandits and Horizon-Free Linear Mixture MDPs](#) Yeoneung Kim, Insoon Yang, Kwang-Sung Jun
- [Communication-Efficient Topologies for Decentralized Learning with \$O\(1\)\$ Consensus Rate](#) Zhuoqing Song, Weijian Li, Kexin Jin, Lei Shi, Ming Yan, Wotao Yin, Kun Yuan
- [Moderate-fitting as a Natural Backdoor Defender for Pre-trained Language Models](#) Biru Zhu, Yujia Qin, Ganqu Cui, Yangyi Chen, Weilin Zhao, Chong Fu, Yangdong Deng, Zhiyuan Liu, Jingang Wang, Wei Wu, Maosong Sun, Ming Gu
- [Dataset Distillation via Factorization](#) Songhua Liu, Kai Wang, Xingyi Yang, Jingwen Ye, Xinchao Wang
- [Adaptive Sampling for Discovery](#) Ziping Xu, Eunjae Shim, Ambuj Tewari, Paul Zimmerman
- [A Large Scale Search Dataset for Unbiased Learning to Rank](#) Lixin Zou, Haitao Mao, Xiaokai Chu, Jiliang Tang, Wenwen Ye, Shuaiqiang Wang, Dawei Yin
- [SegNeXt: Rethinking Convolutional Attention Design for Semantic Segmentation](#) Meng-Hao Guo, Cheng-Ze Lu, Qibin Hou, Zhengning Liu, Ming-Ming Cheng, Shi-min Hu
- [Understanding Hyperdimensional Computing for Parallel Single-Pass Learning](#) Tao Yu, Yichi Zhang, Zhiru Zhang, Christopher M. De Sa
- [Syndicated Bandits: A Framework for Auto Tuning Hyper-parameters in Contextual Bandit Algorithms](#) QIN DING, Yue Kang, Yi-Wei Liu, Thomas Chun Man Lee, Cho-Jui Hsieh, James Sharpnack
- [Benign, Tempered, or Catastrophic: Toward a Refined Taxonomy of Overfitting](#) Neil Mallinar, James Simon, Amirhesam Abedsoltan, Parthe Pandit, Misha Belkin, Preetum Nakkiran
- [Pre-trained Adversarial Perturbations](#) Yuanhao Ban, Yinpeng Dong
- [An Empirical Study on Disentanglement of Negative-free Contrastive Learning](#) Jinkun Cao, Ruiqian Nai, Qing Yang, Jialei Huang, Yang Gao
- [MABSplit: Faster Forest Training Using Multi-Armed Bandits](#) Mo Tiwari, Ryan Kang, Jaeyong Lee, Chris Piech, Ilan Shomorony, Sebastian Thrun, Martin J. Zhang
- [Counterfactual Fairness with Partially Known Causal Graph](#) Aoqi Zuo, Susan Wei, Tongliang Liu, Bo Han, Kun Zhang, Mingming Gong
- [Controlled Sparsity via Constrained Optimization or: How I Learned to Stop Tuning Penalties and Love Constraints](#) Jose Gallego-Posada, Juan Ramirez, Akram Erraqabi, Yoshua Bengio, Simon Lacoste-Julien
- [Algorithms and Hardness for Learning Linear Thresholds from Label Proportions](#) Rishi Saket
- [Predictive Coding beyond Gaussian Distributions](#) Luca Pinchetti, Tommaso Salvatori, Yordan Yordanov, Beren Millidge, Yuhang Song, Thomas Lukasiewicz
- [Semi-supervised Active Linear Regression](#) Nived Rajaraman, Fnu Devvrit, Pranjal Awasthi
- [Boosting Barely Robust Learners: A New Perspective on Adversarial Robustness](#) Avrim Blum, Omar Montasser, Greg Shakhnarovich, Hongyang Zhang
- [Decision-Focused Learning without Decision-Making: Learning Locally Optimized Decision Losses](#) Sanket Shah, Kai Wang, Bryan Wilder, Andrew Perrault, Milind Tambe
- [Neural Stochastic PDEs: Resolution-Invariant Learning of Continuous Spatiotemporal Dynamics](#) Cristopher Salvi, Maud Lemerrier, Andris Gerasimovics
- [Okapi: Generalising Better by Making Statistical Matches Match](#) Myles Bartlett, Sara Romiti, Viktoriia Sharmanska, Novi Quadrianto
- [Revisiting Heterophily For Graph Neural Networks](#) Sitao Luan, Chenqing Hua, Qincheng Lu, Jiaqi Zhu, Mingde Zhao, Shuyuan Zhang, Xiao-Wen Chang, Doina Precup
- [Museformer: Transformer with Fine- and Coarse-Grained Attention for Music Generation](#) Botao Yu, Peiling Lu, Rui Wang, Wei Hu, Xu Tan, Wei Ye, Shikun Zhang, Tao Qin, Tie-Yan Liu
- [Emergent Communication: Generalization and Overfitting in Lewis Games](#) Mathieu Rita, Corentin Tallec, Paul Michel, Jean-Bastien Grill, Olivier Pietquin, Emmanuel Dupoux, Florian Strub
- [Towards Efficient Post-training Quantization of Pre-trained Language Models](#) Haoli Bai, Lu Hou, Lifeng Shang, Xin Jiang, Irwin King, Michael R Lyu
- [TweetNERD - End to End Entity Linking Benchmark for Tweets](#) Shubhanshu Mishra, Aman Saini, Raheleh Makki, Sneha Mehta, Aria Haghighi, Ali Mollahosseini
- [Riemannian Neural SDE: Learning Stochastic Representations on Manifolds](#) Sung Woo Park, Hyomin Kim, Kyungjae Lee, Junseok Kwon
- [OccGen: Selection of Real-world Multilingual Parallel Data Balanced in Gender within Occupations](#) Marta Costa-jussà, Christine Basta, Oriol Domingo, André Rubungo
- [Learning in Observable POMDPs, without Computationally Intractable Oracles](#) Noah Golowich, Ankur Moitra, Dhruv Rohatgi
- [Cross-Image Context for Single Image Inpainting](#) Tingliang Feng, Wei Feng, Weiqi Li, Di Lin
- [Efficient and Effective Augmentation Strategy for Adversarial Training](#) Sravanti Addepalli, Samyak Jain, Venkatesh Babu R
- [Multi-Sample Training for Neural Image Compression](#) Tongda Xu, Yan Wang, Dailan He, Chenjian Gao, Han Gao, Kunzan Liu, Hongwei Qin
- [Adaptive Data Debiasing through Bounded Exploration](#) Yifan Yang, Yang Liu, Parinaz Naghizadeh
- [Learning to Navigate Wikipedia by Taking Random Walks](#) Manzil Zaheer, Kenneth Marino, Will Grathwohl, John Schultz, Wendy Shang, Shailesh Babayan, Arun Abhishek, Ishita Dasgupta, Christine Kaeffer, Chen, Bob Fergus

[Report an Issue](#) | [Name Change Policy](#)

- [When does return-conditioned supervised learning work for offline reinforcement learning?](#) David Brandfonbrener, Alberto Bietti, Jacob Buckman, Romain Laroche, Joan Bruna
- [Provable Subspace Identification Under Post-Nonlinear Mixtures](#) Qi Lyu, Xiao Fu
- [S³-NeRF: Neural Reflectance Field from Shading and Shadow under a Single Viewpoint](#) Wenqi Yang, Guanying Chen, Chaofeng Chen, Zhenfang Chen, Kwan-Yee K. Wong
- [AdaFocal: Calibration-aware Adaptive Focal Loss](#) Arindam Ghosh, Thomas Schaaf, Matthew Gormley
- [PDEBench: An Extensive Benchmark for Scientific Machine Learning](#) Makoto Takamoto, Timothy Praditia, Raphael Leiteritz, Daniel MacKinlay, Francesco Alesiani, Dirk Pflüger, Mathias Niepert
- [Learning Robust Dynamics through Variational Sparse Gating](#) Arnav Kumar Jain, Shivakanth Sujit, Shruti Joshi, Vincent Michalski, Danijar Hafner, Samira Ebrahimi Kahou
- [Where to Pay Attention in Sparse Training for Feature Selection?](#) Ghada Sokar, Zahra Atashgahi, Mykola Pechenizkiy, Decebal Constantin Mocanu
- [Maximizing Revenue under Market Shrinkage and Market Uncertainty](#) Maria-Florina F. Balcan, Siddharth Prasad, Tuomas Sandholm
- [General Cutting Planes for Bound-Propagation-Based Neural Network Verification](#) Huan Zhang, Shiqi Wang, Kaidi Xu, Linyi Li, Bo Li, Suman Jana, Cho-Jui Hsieh, J. Zico Kolter
- [StrokeRehab: A Benchmark Dataset for Sub-second Action Identification](#) Aakash Kaku, Kangning Liu, Avinash Parnandi, Haresh Rengaraj Rajamohan, Kannan Venkataramanan, Anita Venkatesan, Audre Wirtanen, Natasha Pandit, Heidi Schambra, Carlos Fernandez-Granda
- [An \$\alpha\$ -regret analysis of Adversarial Bilateral Trade](#) Yossi Azar, Amos Fiat, Federico Fusco
- [LDSA: Learning Dynamic Subtask Assignment in Cooperative Multi-Agent Reinforcement Learning](#) Mingyu Yang, Jian Zhao, Xunhan Hu, Wengang Zhou, Jiangcheng Zhu, Houqiang Li
- [Mildly Conservative Q-Learning for Offline Reinforcement Learning](#) Jiafei Lyu, Xiaoteng Ma, Xiu Li, Zongqing Lu
- [Iterative Feature Matching: Toward Provable Domain Generalization with Logarithmic Environments](#) Yining Chen, Elan Rosenfeld, Mark Sellke, Tengyu Ma, Andrej Risteski
- [Certifying Robust Graph Classification under Orthogonal Gromov-Wasserstein Threats](#) Hongwei Jin, Zishun Yu, Xinhua Zhang
- [Functional Ensemble Distillation](#) Coby Penso, Idan Achituve, Ethan Fetaya
- [Use-Case-Grounded Simulations for Explanation Evaluation](#) Valerie Chen, Nari Johnson, Nicholay Topin, Gregory Plumb, Ameet Talwalkar
- [Lethal Dose Conjecture on Data Poisoning](#) Wenxiao Wang, Alexander Levine, Soheil Feizi
- [Online Decision Mediation](#) Daniel Jarrett, Alihan Hüyük, Mihaela van der Schaar
- [Neural-Symbolic Entangled Framework for Complex Query Answering](#) Zezhong Xu, Wen Zhang, Peng Ye, Hui Chen, Huajun Chen
- [Reinforcement Learning with Automated Auxiliary Loss Search](#) Tairan He, Yuge Zhang, Kan Ren, Minghuan Liu, Che Wang, Weinan Zhang, Yuqing Yang, Dongsheng Li
- [FinRL-Meta: Market Environments and Benchmarks for Data-Driven Financial Reinforcement Learning](#) Xiao-Yang Liu, Ziyi Xia, Jingyang Rui, Jiechao Gao, Hongyang Yang, Ming Zhu, Christina Wang, Zhaoran Wang, Jian Guo
- [TempEL: Linking Dynamically Evolving and Newly Emerging Entities](#) Klim Zaporozets, Lucie-Aimée Kaffee, Johannes Deleu, Thomas Demeester, Chris Develder, Isabelle Augenstein
- [M⁴I: Multi-modal Models Membership Inference](#) Pingyi Hu, Zihan Wang, Ruoxi Sun, Hu Wang, Minhui Xue
- [Best of Both Worlds Model Selection](#) Aldo Pacchiano, Christoph Dann, Claudio Gentile
- [The Unreasonable Effectiveness of Fully-Connected Layers for Low-Data Regimes](#) Peter Kocsis, Peter Sűkeník, Guillem Braso, Matthias Niessner, Laura Leal-Taixé, Ismail Elezi
- [Augmentations in Hypergraph Contrastive Learning: Fabricated and Generative](#) Tianxin Wei, Yuning You, Tianlong Chen, Yang Shen, Jingrui He, Zhangyang Wang
- [On the Global Convergence Rates of Decentralized Softmax Gradient Play in Markov Potential Games](#) Runyu Zhang, Jincheng Mei, Bo Dai, Dale Schuurmans, Na Li
- [Sym-NCO: Leveraging Symmetry for Neural Combinatorial Optimization](#) Minsu Kim, Junyoung Park, Jinkyoo Park
- [Few-Shot Parameter-Efficient Fine-Tuning is Better and Cheaper than In-Context Learning](#) Haokun Liu, Derek Tam, Mohammed Muqeeth, Jay Mohta, Tenghao Huang, Mohit Bansal, Colin A. Raffel
- [HF-NeuS: Improved Surface Reconstruction Using High-Frequency Details](#) Yiqun Wang, Ivan Skorokhodov, Peter Wonka
- [On the Epistemic Limits of Personalized Prediction](#) Lucas Monteiro Paes, Carol Long, Berk Ustun, Flavio Calmon
- [DeepInteraction: 3D Object Detection via Modality Interaction](#) Zeyu Yang, Jiaqi Chen, Zhenwei Miao, Wei Li, Xiatian Zhu, Li Zhang
- [Deep Differentiable Logic Gate Networks](#) Felix Petersen, Christian Borgelt, Hilde Kuehne, Oliver Deussen
- [Maximizing and Satisficing in Multi-armed Bandits with Graph Information](#) Parth Thaker, Mohit Malu, Nikhil Rao, Gautam Dasarathy
- [MoGDE: Boosting Mobile Monocular 3D Object Detection with Ground Depth Estimation](#) Yunsong Zhou, Quan Liu, Hongzi Zhu, Yunzhe Li, Shan Chang, Minyi Guo
- [Causality Preserving Chaotic Transformation and Classification using Neurochaos Learning](#) Harikrishnan N B, Aditi Kathpalia, Nithin Nagaraj
- [GOOD: A Graph Out-of-Distribution Benchmark](#) Shurui Gui, Xiner Li, Limei Wang, Shuiwang Ji
- [CascadeXML: Rethinking Transformers for End-to-end Multi-resolution Training in Extreme Multi-label Classification](#) Siddhant Kharbanda, Atmadeep Banerjee, Erik Schultheis, Rohit Babbar

- [VF-PS: How to Select Important Participants in Vertical Federated Learning, Efficiently and Securely?](#) Jiawei Jiang, Lukas Burkhalter, Fangcheng Fu, Bolin Ding, Bo Du, Anwar Hithnawi, Bo Li, Ce Zhang
- [PopArt: Efficient Sparse Regression and Experimental Design for Optimal Sparse Linear Bandits](#) Kyoungseok Jang, Chicheng Zhang, Kwang-Sung Jun
- [Augmenting Online Algorithms with \$\epsilon\$ -Accurate Predictions](#) Anupam Gupta, Debmalya Panigrahi, Bernardo Subercaseaux, Kevin Sun
- [Unsupervised Multi-Object Segmentation by Predicting Probable Motion Patterns](#) Laurynas Karazija, Subhabrata Choudhury, Iro Laina, Christian Rupprecht, Andrea Vedaldi
- [Incrementality Bidding via Reinforcement Learning under Mixed and Delayed Rewards](#) Ashwinkumar Badanidiyuru Varadaraja, Zhe Feng, Tianxi Li, Haifeng Xu
- [Masked Generative Adversarial Networks are Data-Efficient Generation Learners](#) Jiaxing Huang, Kaiwen Cui, Dayan Guan, Aoran Xiao, Fangneng Zhan, Shijian Lu, Shengcai Liao, Eric Xing
- [What You See is What You Get: Principled Deep Learning via Distributional Generalization](#) Bogdan Kulynych, Yao-Yuan Yang, Yaodong Yu, Jarosław Błasiok, Preetum Nakkiran
- [Towards Understanding the Condensation of Neural Networks at Initial Training](#) Hanxu Zhou, Zhou Qixuan, Tao Luo, Yaoyu Zhang, Zhi-Qin Xu
- [CoNT: Contrastive Neural Text Generation](#) Chenxin An, Jiangtao Feng, Kai Lv, Lingpeng Kong, Xipeng Qiu, Xuanjing Huang
- [GAPX: Generalized Autoregressive Paraphrase-Identification X](#) Yifei Zhou, Renyu Li, Hayden Housen, Ser Nam Lim
- [Scalable Infomin Learning](#) Yanzhi Chen, weihao sun, Yingzhen Li, Adrian Weller
- [Learning to Accelerate Partial Differential Equations via Latent Global Evolution](#) Tailin Wu, Takashi Maruyama, Jure Leskovec
- [Monte Carlo Augmented Actor-Critic for Sparse Reward Deep Reinforcement Learning from Suboptimal Demonstrations](#) Albert Wilcox, Ashwin Balakrishna, Jules Dedieu, Wyame Benslimane, Daniel Brown, Ken Goldberg
- [Not too little, not too much: a theoretical analysis of graph \(over\)smoothing](#) Nicolas Keriven
- [Kernel similarity matching with Hebbian networks](#) Kyle Luther, Sebastian Seung
- [HumanLiker: A Human-like Object Detector to Model the Manual Labeling Process](#) Haoran Wei, Ping Guo, Yangguang Zhu, Chenglong Liu, Peng Wang
- [Scalable Representation Learning in Linear Contextual Bandits with Constant Regret Guarantees](#) Andrea Tirinzoni, Matteo Papini, Ahmed Touati, Alessandro Lazaric, Matteo Pirota
- [DevFly: Bio-Inspired Development of Binary Connections for Locality Preserving Sparse Codes](#) Tianqi Wei, Rana Alkhoury Maroun, Qinghai Guo, Barbara Webb
- [Why neural networks find simple solutions: The many regularizers of geometric complexity](#) Benoit Dherin, Michael Munn, Mihaela Rosca, David Barrett
- [Efficiently Computing Local Lipschitz Constants of Neural Networks via Bound Propagation](#) Zhouxing Shi, Yihan Wang, Huan Zhang, J. Zico Kolter, Cho-Jui Hsieh
- [A Causal Analysis of Harm](#) Sander Beckers, Hana Chockler, Joseph Halpern
- [Seeing the forest and the tree: Building representations of both individual and collective dynamics with transformers](#) Ran Liu, Mehdi Azabou, Max Dabagia, Jingyun Xiao, Eva Dyer
- [Semi-Supervised Learning with Decision Trees: Graph Laplacian Tree Alternating Optimization](#) Arman Zharmagambetov, Miguel A. Carreira-Perpinan
- [Riemannian Score-Based Generative Modelling](#) Valentin De Bortoli, Emile Mathieu, Michael Hutchinson, James Thornton, Yee Whye Teh, Arnaud Doucet
- [Intra-agent speech permits zero-shot task acquisition](#) Chen Yan, Federico Carnevale, Petko I Georgiev, Adam Santoro, Aurelia Guy, Alistair Muldal, Chia-Chun Hung, Joshua Abramson, Timothy Lillicrap, Gregory Wayne
- [Free Probability for predicting the performance of feed-forward fully connected neural networks](#) Reda CHHAIBI, Tariq Daouda, Ezechiell Kahn
- [The Minority Matters: A Diversity-Promoting Collaborative Metric Learning Algorithm](#) Shilong Bao, Qianqian Xu, Zhiyong Yang, Yuan He, Xiaochun Cao, Qingming Huang
- [Open-Ended Reinforcement Learning with Neural Reward Functions](#) Robert Meier, Asier Mujika
- [A Reduction to Binary Approach for Debiasing Multiclass Datasets](#) Ibrahim M. Alabdulmohsin, Jessica Schrouff, Sanmi Koyejo
- [Obj2Seq: Formatting Objects as Sequences with Class Prompt for Visual Tasks](#) Zhiyang Chen, Yousong Zhu, Zhaowen Li, Fan Yang, Wei Li, Haixin Wang, Chaoyang Zhao, Liwei Wu, Rui Zhao, Jinqiao Wang, Ming Tang
- [Learn to Explain: Multimodal Reasoning via Thought Chains for Science Question Answering](#) Pan Lu, Swaroop Mishra, Tanglin Xia, Liang Qiu, Kai-Wei Chang, Song-Chun Zhu, Oyvind Tafjord, Peter Clark, Ashwin Kalyan
- [Few-Shot Audio-Visual Learning of Environment Acoustics](#) Sagnik Majumder, Changan Chen, Ziad Al-Halah, Kristen Grauman
- [The Phenomenon of Policy Churn](#) Tom Schaul, Andre Barreto, John Quan, Georg Ostrovski
- [Molecule Generation by Principal Subgraph Mining and Assembling](#) Xiangzhe Kong, Wenbing Huang, Zhixing Tan, Yang Liu
- [Implicit Neural Representations with Levels-of-Experts](#) Zekun Hao, Arun Mallya, Serge Belongie, Ming-Yu Liu
- [Planning for Sample Efficient Imitation Learning](#) Zhao-Heng Yin, Weirui Ye, Qifeng Chen, Yang Gao
- [Concept Activation Regions: A Generalized Framework For Concept-Based Explanations](#) Jonathan Crabbé, Mihaela van der Schaar

- [Understanding Cross-Domain Few-Shot Learning Based on Domain Similarity and Few-Shot Difficulty](#) Jaehoon Oh, Sungnyun Kim, Namgyu Ho, Jin-Hwa Kim, Hwanjun Song, Se-Young Yun
- [Ambiguous Images With Human Judgments for Robust Visual Event Classification](#) Kate Sanders, Reno Kriz, Anqi Liu, Benjamin Van Durme
- [Sustainable Online Reinforcement Learning for Auto-bidding](#) Zhiyu Mou, Yusen Huo, Rongquan Bai, Mingzhou Xie, Chuan Yu, Jian Xu, Bo Zheng
- [Uni-Perceiver-MoE: Learning Sparse Generalist Models with Conditional MoEs](#) Jinguo Zhu, Xizhou Zhu, Wenhai Wang, Xiaohua Wang, Hongsheng Li, Xiaogang Wang, Jifeng Dai
- [Improved Coresets for Euclidean \$k\$ -Means](#) Vincent Cohen-Addad, Kasper Green Larsen, David Saulpic, Chris Schwiegelshohn, Omar Ali Sheikh-Omar
- [Accelerated Linearized Laplace Approximation for Bayesian Deep Learning](#) Zhijie Deng, Feng Zhou, Jun Zhu
- [Learning to Reason with Neural Networks: Generalization, Unseen Data and Boolean Measures](#) Emmanuel Abbe, Samy Bengio, Elisabetta Cornacchia, Jon Kleinberg, Aryo Lotfi, Maithra Raghu, Chiyuan Zhang
- [Using Partial Monotonicity in Submodular Maximization](#) Loay Mualem, Moran Feldman
- [Enhanced Meta Reinforcement Learning via Demonstrations in Sparse Reward Environments](#) Desik Rengarajan, Sapana Chaudhary, Jaewon Kim, Dileep Kalathil, Srinivas Shakkottai
- [Riemannian Diffusion Models](#) Chin-Wei Huang, Milad Aghajohari, Joey Bose, Prakash Panangaden, Aaron C. Courville
- [Training and Inference on Any-Order Autoregressive Models the Right Way](#) Andy Shih, Dorsa Sadigh, Stefano Ermon
- [Lazy and Fast Greedy MAP Inference for Determinantal Point Process](#) Shinichi Hemmi, Taihei Oki, Shinsaku Sakaue, Kaito Fujii, Satoru Iwata
- [Generalization Bounds for Stochastic Gradient Descent via Localized \$\epsilon\$ -Covers](#) Sejun Park, Umut Simsekli, Murat A. Erdogdu
- [Semi-Supervised Semantic Segmentation via Gentle Teaching Assistant](#) Ying Jin, Jiaqi Wang, Dahua Lin
- [Anytime-Valid Inference For Multinomial Count Data](#) Michael Lindon, Alan Malek
- [What I Cannot Predict, I Do Not Understand: A Human-Centered Evaluation Framework for Explainability Methods](#) Julien Colin, Thomas FEL, Remi Cadene, Thomas Serre
- [S4ND: Modeling Images and Videos as Multidimensional Signals with State Spaces](#) Eric Nguyen, Karan Goel, Albert Gu, Gordon Downs, Preety Shah, Tri Dao, Stephen Baccus, Christopher Ré
- [Understanding Programmatic Weak Supervision via Source-aware Influence Function](#) Jieyu Zhang, Haonan Wang, Cheng-Yu Hsieh, Alexander J. Ratner
- [Approximation with CNNs in Sobolev Space: with Applications to Classification](#) Guohao Shen, Yuling Jiao, Yuanyuan Lin, Jian Huang
- [Sample Complexity of Learning Heuristic Functions for Greedy-Best-First and A* Search](#) Shinsaku Sakaue, Taihei Oki
- [TransTab: Learning Transferable Tabular Transformers Across Tables](#) Zifeng Wang, Jimeng Sun
- [Perceptual Attacks of No-Reference Image Quality Models with Human-in-the-Loop](#) Weixia Zhang, Dingquan Li, Xionghuo Min, Guangtao Zhai, Guodong Guo, Xiaokang Yang, Kede Ma
- [Sketch-GNN: Scalable Graph Neural Networks with Sublinear Training Complexity](#) Mucong Ding, Tahseen Rabbani, Bang An, Evan Wang, Furong Huang
- [Tree Mover's Distance: Bridging Graph Metrics and Stability of Graph Neural Networks](#) Ching-Yao Chuang, Stefanie Jegelka
- [DivBO: Diversity-aware CASH for Ensemble Learning](#) Yu Shen, Yupeng Lu, Yang Li, Yaofeng Tu, Wentao Zhang, Bin CUI
- [Revisiting Graph Contrastive Learning from the Perspective of Graph Spectrum](#) Nian Liu, Xiao Wang, Deyu Bo, Chuan Shi, Jian Pei
- [Functional Indirection Neural Estimator for Better Out-of-distribution Generalization](#) Kha Pham, Thai Hung Le, Man Ngo, Truyen Tran
- [Combinatorial Bandits with Linear Constraints: Beyond Knapsacks and Fairness](#) Qingsong Liu, Weihang Xu, Siwei Wang, Zhixuan Fang
- [Will Bilevel Optimizers Benefit from Loops](#) Kaiyi Ji, Mingrui Liu, Yingbin Liang, Lei Ying
- [Combining Explicit and Implicit Regularization for Efficient Learning in Deep Networks](#) Dan Zhao
- [MBW: Multi-view Bootstrapping in the Wild](#) Mosam Dabhi, Chaoyang Wang, Tim Clifford, László Jeni, Ian Fasel, Simon Lucey
- [On A Mallows-type Model For \(Ranked\) Choices](#) Yifan Feng, Yuxuan Tang
- [\(De-\)Randomized Smoothing for Decision Stump Ensembles](#) Miklós Horváth, Mark Müller, Marc Fischer, Martin Vechev
- [Learning to Break the Loop: Analyzing and Mitigating Repetitions for Neural Text Generation](#) Jin Xu, Xiaojiang Liu, Jianhao Yan, Deng Cai, Huayang Li, Jian Li
- [Debiased Machine Learning without Sample-Splitting for Stable Estimators](#) Qizhao Chen, Vasilis Syrgkanis, Morgane Austern
- [Near-Optimal Sample Complexity Bounds for Constrained MDPs](#) Sharan Vaswani, Lin Yang, Csaba Szepesvari
- [Integral Probability Metrics PAC-Bayes Bounds](#) Ron Amit, Baruch Epstein, Shay Moran, Ron Meir
- [Bellman Residual Orthogonalization for Offline Reinforcement Learning](#) Andrea Zanette, Martin J Wainwright
- [Quantum Speedups of Optimizing Approximately Convex Functions with Applications to Logarithmic Regret Stochastic Convex Bandits](#) Tongyang Li, Ruizhe Zhang
- [Learning Neural Acoustic Fields](#) Andrew Luo, Yilun Du, Michael Tarr, Josh Tenenbaum, Antonio Torralba, Chuang Gan
- [A Universal Error Measure for Input Predictions Applied to Online Graph Problems](#) Giulia Bernardini, Alexander Lindermayr, Alberto Marchetti-Spaccamela, Nicole Megow, Leen Stougie, Michelle Sweering
- [Online Reinforcement Learning for Mixed Policy Scopes](#) Junzhe Zhang, Elias Bareinboim

- [XTC: Extreme Compression for Pre-trained Transformers Made Simple and Efficient](#) Xiaoxia Wu, Zhewei Yao, Minjia Zhang, Conglong Li, Yuxiong He
- [Meta-Album: Multi-domain Meta-Dataset for Few-Shot Image Classification](#) Ihsan Ullah, Dustin Carrión-Ojeda, Sergio Escalera, Isabelle Guyon, Mike Huisman, Felix Mohr, Jan N. van Rijn, Haozhe Sun, Joaquin Vanschoren, Phan Anh Vu
- [S3GC: Scalable Self-Supervised Graph Clustering](#) Fnu Devvrit, Aditya Sinha, Inderjit Dhillon, Prateek Jain
- [Contrastive Neural Ratio Estimation](#) Benjamin K Miller, Christoph Weniger, Patrick Forré
- [An Information-Theoretic Framework for Deep Learning](#) Hong Jun Jeon, Benjamin Van Roy
- [Uncoupled Learning Dynamics with \$O\(\log T\)\$ Swap Regret in Multiplayer Games](#) Ioannis Anagnostides, Gabriele Farina, Christian Kroer, Chung-Wei Lee, Haipeng Luo, Tuomas Sandholm
- [Robust Semi-Supervised Learning when Not All Classes have Labels](#) Lan-Zhe Guo, Yi-Ge Zhang, Zhi-Fan Wu, Jie-Jing Shao, Yu-Feng Li
- [Private Multiparty Perception for Navigation](#) Hui Lu, Mia Chiquier, Carl Vondrick
- [Improving Task-Specific Generalization in Few-Shot Learning via Adaptive Vicinal Risk Minimization](#) Long-Kai Huang, Ying Wei
- [EgoTaskQA: Understanding Human Tasks in Egocentric Videos](#) Baoxiong Jia, Ting Lei, Song-Chun Zhu, Siyuan Huang
- [C-Mixup: Improving Generalization in Regression](#) Huaxiu Yao, Yiping Wang, Linjun Zhang, James Y. Zou, Chelsea Finn
- [Generalised Mutual Information for Discriminative Clustering](#) Louis Ohl, Pierre-Alexandre Mattei, Charles Bouveyron, Warith HARCHAOUI, Mickaël Leclercq, Arnaud Droit, Frederic Precioso
- [Consistent Interpolating Ensembles via the Manifold-Hilbert Kernel](#) Yutong Wang, Clay Scott
- [Geo-Neus: Geometry-Consistent Neural Implicit Surfaces Learning for Multi-view Reconstruction](#) Qiancheng Fu, Qingshan Xu, Yew Soon Ong, Wenbing Tao
- [Sublinear Algorithms for Hierarchical Clustering](#) Arpit Agarwal, Sanjeev Khanna, Huan Li, Prathamesh Patil
- [Is Sortition Both Representative and Fair?](#) Soroush Ebadian, Gregory Kehne, Evi Micha, Ariel D. Procaccia, Nisarg Shah
- [Beyond Rewards: a Hierarchical Perspective on Offline Multiagent Behavioral Analysis](#) Shayegan Omidshafiei, Andrei Kapishnikov, Yannick Assogba, Lucas Dixon, Been Kim
- [Dynamic pricing and assortment under a contextual MNL demand](#) Noemie Perivier, Vineet Goyal
- [DGD²: A Linearly Convergent Distributed Algorithm For High-dimensional Statistical Recovery](#) Marie Maros, Gesualdo Scutari
- [Pseudo-Riemannian Graph Convolutional Networks](#) Bo Xiong, Shichao Zhu, Nico Potyka, Shirui Pan, Chuan Zhou, Steffen Staab
- [CroCo: Self-Supervised Pre-training for 3D Vision Tasks by Cross-View Completion](#) Philippe Weinzaepfel, Vincent Leroy, Thomas Lucas, Romain BRÉGIER, Yohann Cabon, Vaibhav ARORA, Leonid Antsfeld, Boris Chidlovskii, Gabriela Csurka, Jerome Revaud
- [Sound and Complete Verification of Polynomial Networks](#) Elias Abad Rocamora, Mehmet Fatih Sahin, Fanghui Liu, Grigorios Chrysos, Volkan Cevher
- [Better SGD using Second-order Momentum](#) Hoang Tran, Ashok Cutkosky
- [Learning Predictions for Algorithms with Predictions](#) Misha Khodak, Maria-Florina F. Balcan, Ameet Talwalkar, Sergei Vassilvitskii
- [Unsupervised Point Cloud Completion and Segmentation by Generative Adversarial Autoencoding Network](#) Changfeng Ma, Yang Yang, Jie Guo, Fei Pan, Chongjun Wang, Yanwen Guo
- [CalFAT: Calibrated Federated Adversarial Training with Label Skewness](#) Chen Chen, Yuchen Liu, Xingjun Ma, Lingjuan Lyu
- [Rethinking Generalization in Few-Shot Classification](#) Markus Hiller, Rongkai Ma, Mehrtash Harandi, Tom Drummond
- [Stimulative Training of Residual Networks: A Social Psychology Perspective of Loafing](#) Peng Ye, Shengji Tang, Baopu Li, Tao Chen, Wanli Ouyang
- [EGSDE: Unpaired Image-to-Image Translation via Energy-Guided Stochastic Differential Equations](#) Min Zhao, Fan Bao, Chongxuan Li, Jun Zhu
- [Cryptographic Hardness of Learning Halfspaces with Massart Noise](#) Ilias Diakonikolas, Daniel Kane, Pasin Manurangsi, Lisheng Ren
- [Frank-Wolfe-based Algorithms for Approximating Tyler's M-estimator](#) Lior Danon, Dan Garber
- [Reinforcement Learning with Non-Exponential Discounting](#) Matthias Schultheis, Constantin A. Rothkopf, Heinz Koeppel
- [Picking on the Same Person: Does Algorithmic Monoculture lead to Outcome Homogenization?](#) Rishi Bommasani, Kathleen A. Creel, Ananya Kumar, Dan Jurafsky, Percy S. Liang
- [Causal Identification under Markov equivalence: Calculus, Algorithm, and Completeness](#) Amin Jaber, Adele Ribeiro, Jiji Zhang, Elias Bareinboim
- [Unravelling the Performance of Physics-informed Graph Neural Networks for Dynamical Systems](#) Abishek Thangamuthu, Gunjan Kumar, Suresh Bishnoi, Ravinder Bhattoo, N M Anoop Krishnan, Sayan Ranu
- [Dynamic Fair Division with Partial Information](#) Gerdus Benade, Daniel Halpern, Alexandros Psomas
- [Generalized Variational Inference in Function Spaces: Gaussian Measures meet Bayesian Deep Learning](#) Veit David Wild, Robert Hu, Dino Sejdinovic
- [Communicating Natural Programs to Humans and Machines](#) Sam Acquaviva, Yewen Pu, Marta Kryven, Theodoros Sechopoulos, Catherine Wong, Gabrielle Ecanow, Maxwell Nye, Michael Tessler, Josh Tenenbaum
- [SCAMPS: Synthetics for Camera Measurement of Physiological Signals](#) Daniel McDuff, Miah Wander, Xin Liu, Brian Hill, Javier Hernandez, Jonathan Lester, Tadas Baltrusaitis
- [A Closer Look at Learned Optimization: Stability, Robustness, and Inductive Biases](#) James Harrison, Luke Metz, Jascha Sohl-Dickstein
- ["Lossless" Compression of Deep Neural Networks: A High-dimensional Neural Tangent Kernel Approach](#) lingyu gu, Yongqi Du, yuan

- [Privacy of Noisy Stochastic Gradient Descent: More Iterations without More Privacy Loss](#) Jason Altschuler, Kunal Talwar
- [Theseus: A Library for Differentiable Nonlinear Optimization](#) Luis Pineda, Taosha Fan, Maurizio Monge, Shobha Venkataraman, Paloma Sodhi, Ricky T. Q. Chen, Joseph Ortiz, Daniel DeTone, Austin Wang, Stuart Anderson, Jing Dong, Brandon Amos, Mustafa Mukadam
- [K-Radar: 4D Radar Object Detection for Autonomous Driving in Various Weather Conditions](#) Dong-Hee Paek, SEUNG-HYUN KONG, Kevin Tirta Wijaya
- [Asymmetric Temperature Scaling Makes Larger Networks Teach Well Again](#) Xin-Chun Li, Wen-shu Fan, Shaoming Song, Yinchuan Li, bingshuai Li, Shao Yunfeng, De-Chuan Zhan
- [Solving Quantitative Reasoning Problems with Language Models](#) Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, Yuhuai Wu, Behnam Neyshabur, Guy Gur-Ari, Vedant Misra
- [Structural Knowledge Distillation for Object Detection](#) Philip de Rijk, Lukas Schneider, Marius Cordts, Darius Gavrila
- [Thompson Sampling Efficiently Learns to Control Diffusion Processes](#) Mohamad Kazem Shirani Faradonbeh, Mohamad Sadegh Shirani Faradonbeh, Mohsen Bayati
- [Discrete Compositional Representations as an Abstraction for Goal Conditioned Reinforcement Learning](#) Riashat Islam, Hongyu Zang, Anirudh Goyal, Alex M. Lamb, Kenji Kawaguchi, Xin Li, Romain Laroche, Yoshua Bengio, Remi Tachet des Combes
- [Graph Convolution Network based Recommender Systems: Learning Guarantee and Item Mixture Powered Strategy](#) Leyan Deng, Defu Lian, Chenwang Wu, Enhong Chen
- [Local Metric Learning for Off-Policy Evaluation in Contextual Bandits with Continuous Actions](#) Haanvid Lee, Jongmin Lee, Yunseon Choi, Wonseok Jeon, Byung-Jun Lee, Yung-Kyun Noh, Kee-Eung Kim
- [Flare7K: A Phenomenological Nighttime Flare Removal Dataset](#) Yuekun Dai, Chongyi Li, Shangchen Zhou, Ruicheng Feng, Chen Change Loy
- [USB: A Unified Semi-supervised Learning Benchmark for Classification](#) Yidong Wang, Hao Chen, Yue Fan, Wang SUN, Ran Tao, Wenxin Hou, Renjie Wang, Linyi Yang, Zhi Zhou, Lan-Zhe Guo, Heli Qi, Zhen Wu, Yu-Feng Li, Satoshi Nakamura, Wei Ye, Marios Savvides, Bhiksha Raj, Takahiro Shinozaki, Bernt Schiele, Jindong Wang, Xing Xie, Yue Zhang
- [Nocturne: a scalable driving benchmark for bringing multi-agent learning one step closer to the real world](#) Eugene Vinitzky, Nathan Lichtlé, Xiaomeng Yang, Brandon Amos, Jakob Foerster
- [Quasi-Newton Methods for Saddle Point Problems](#) Chengchang Liu, Luo Luo
- [Self-Supervised Contrastive Pre-Training For Time Series via Time-Frequency Consistency](#) Xiang Zhang, Ziyuan Zhao, Theodoros Tsiligkaridis, Marinka Zitnik
- [Uncalibrated Models Can Improve Human-AI Collaboration](#) Kailas Vodrahalli, Tobias Gerstenberg, James Y. Zou
- [Learning Dynamical Systems via Koopman Operator Regression in Reproducing Kernel Hilbert Spaces](#) Vladimir Kostic, Pietro Novelli, Andreas Maurer, Carlo Ciliberto, Lorenzo Rosasco, Massimiliano Pontil
- [Self-supervised surround-view depth estimation with volumetric feature fusion](#) Jung-Hee Kim, Junhwa Hur, Tien Phuoc Nguyen, Seong-Gyun Jeong
- [CARLANE: A Lane Detection Benchmark for Unsupervised Domain Adaptation from Simulation to multiple Real-World Domains](#) Bonifaz Stuhr, Johann Haselberger, Julian Gebele
- [On Enforcing Better Conditioned Meta-Learning for Rapid Few-Shot Adaptation](#) Markus Hiller, Mehrtash Harandi, Tom Drummond
- [Oracle-Efficient Online Learning for Smoothed Adversaries](#) Nika Haghtalab, Yanjun Han, Abhishek Shetty, Kunhe Yang
- [A Policy-Guided Imitation Approach for Offline Reinforcement Learning](#) Haoran Xu, Li Jiang, Li Jianxiong, Xianyuan Zhan
- [Sample-Efficient Learning of Correlated Equilibria in Extensive-Form Games](#) Ziang Song, Song Mei, Yu Bai
- [VectorAdam for Rotation Equivariant Geometry Optimization](#) Selena Zihan Ling, Nicholas Sharp, Alec Jacobson
- [Matrix Multiplicative Weights Updates in Quantum Zero-Sum Games: Conservation Laws & Recurrence](#) Rahul Jain, Georgios Piliouras, Ryann Sim
- [On the Convergence Theory for Hessian-Free Bilevel Algorithms](#) Daouda Sow, Kaiyi Ji, Yingbin Liang
- [Equivariant Networks for Crystal Structures](#) Oumar Kaba, Siamak Ravanbakhsh
- [A General Framework for Auditing Differentially Private Machine Learning](#) Fred Lu, Joseph Munoz, Maya Fuchs, Tyler LeBlond, Elliott Zaresky-Williams, Edward Raff, Francis Ferraro, Brian Testa
- [Generalization Analysis on Learning with a Concurrent Verifier](#) Masaaki Nishino, Kengo Nakamura, Norihito Yasuda
- [Spartan: Differentiable Sparsity via Regularized Transportation](#) Kai Sheng Tai, Taipeng Tian, Ser Nam Lim
- [Focal Modulation Networks](#) Jianwei Yang, Chunyuan Li, Xiyang Dai, Jianfeng Gao
- [HSurf-Net: Normal Estimation for 3D Point Clouds by Learning Hyper Surfaces](#) Qing Li, Yu-Shen Liu, Jin-San Cheng, Cheng Wang, Yi Fang, Zhizhong Han
- [Robust Streaming PCA](#) Daniel Bienstock, Minchan Jeong, Apurv Shukla, Se-Young Yun
- [NeMF: Neural Motion Fields for Kinematic Animation](#) Chengan He, Jun Saito, James Zachary, Holly Rushmeier, Yi Zhou
- [Global Normalization for Streaming Speech Recognition in a Modular Framework](#) Ehsan Variani, Ke Wu, Michael D Riley, David Rybach, Matt Shannon, Cyril Allauzen
- [Resource-Adaptive Federated Learning with All-In-One Neural Composition](#) Yiqun Mei, Pengfei Guo, Mo Zhou, Vishal Patel
- [SoteriaFL: A Unified Framework for Private Federated Learning with Communication Compression](#) Zhize Li, Haoyu Zhao, Boyue Li, Yuejie Ch:

[Report an Issue](#) | [Name Change Policy](#)

- [Your Transformer May Not be as Powerful as You Expect](#) Shengjie Luo, Shanda Li, Shuxin Zheng, Tie-Yan Liu, Liwei Wang, Di He
- [Redundancy-Free Message Passing for Graph Neural Networks](#) Rongqin Chen, Shenghui Zhang, Leong Hou U, Ye Li
- [Diffusion-LM Improves Controllable Text Generation](#) Xiang Li, John Thickstun, Ishaan Gulrajani, Percy S. Liang, Tatsunori B. Hashimoto
- [Making Sense of Dependence: Efficient Black-box Explanations Using Dependence Measure](#) Paul Novello, Thomas FEL, David Vigouroux
- [Energy-Based Contrastive Learning of Visual Representations](#) Beomsu Kim, Jong Chul Ye
- [Why Robust Generalization in Deep Learning is Difficult: Perspective of Expressive Power](#) Binghui Li, Jikai Jin, Han Zhong, John Hopcroft, Liwei Wang
- [Asynchronous Actor-Critic for Multi-Agent Reinforcement Learning](#) Yuchen Xiao, Weihao Tan, Christopher Amato
- [Polynomial Neural Fields for Subband Decomposition and Manipulation](#) Guandao Yang, Sagie Benaïm, Varun Jampani, Kyle Genova, Jonathan Barron, Thomas Funkhouser, Bharath Hariharan, Serge Belongie
- [On the Generalizability and Predictability of Recommender Systems](#) Duncan McElfresh, Sujay Khandagale, Jonathan Valverde, John Dickerson, Colin White
- [Optimal Rates for Regularized Conditional Mean Embedding Learning](#) Zhu Li, Dimitri Meunier, Mattes Mollenhauer, Arthur Gretton
- [Divert More Attention to Vision-Language Tracking](#) Mingzhe Guo, Zhipeng Zhang, Heng Fan, Liping Jing
- [Rethinking Image Restoration for Object Detection](#) Shangquan Sun, Wenqi Ren, Tao Wang, Xiaochun Cao
- [Optimal Brain Compression: A Framework for Accurate Post-Training Quantization and Pruning](#) Elias Frantar, Dan Alistarh
- [Challenging Common Assumptions in Convex Reinforcement Learning](#) Mirco Mutti, Riccardo De Santi, Piersilvio De Bartolomeis, Marcello Restelli
- [Association Graph Learning for Multi-Task Classification with Category Shifts](#) Jiayi Shen, Zehao Xiao, Xiantong Zhen, Cees Snoek, Marcel Worring
- [Delving into Sequential Patches for Deepfake Detection](#) Jiazhi Guan, Hang Zhou, Zhibin Hong, Errui Ding, Jingdong Wang, Chengbin Quan, Youjian Zhao
- [On the detrimental effect of invariances in the likelihood for variational inference](#) Richard Kurle, Ralf Herbrich, Tim Januschowski, Yuyang (Bernie) Wang, Jan Gasthaus
- [Self-Supervised Learning Through Efference Copies](#) Franz Scherr, Qinghai Guo, Timoleon Moraitis
- [Precise Learning Curves and Higher-Order Scalings for Dot-product Kernel Regression](#) Lechao Xiao, Hong Hu, Theodor Misiakiewicz, Yue Lu, Jeffrey Pennington
- [A Unified Model for Multi-class Anomaly Detection](#) Zhiyuan You, Lei Cui, Yujun Shen, Kai Yang, Xin Lu, Yu Zheng, Xinyi Le
- [Constrained GPI for Zero-Shot Transfer in Reinforcement Learning](#) Jaekyeom Kim, Seohong Park, Gunhee Kim
- [UnfoldML: Cost-Aware and Uncertainty-Based Dynamic 2D Prediction for Multi-Stage Classification](#) Yanbo Xu, Alind Khare, Glenn Matlin, Monish Ramadoss, Rishikesan Kamaleswaran, Chao Zhang, Alexey Tumanov
- [Generative Status Estimation and Information Decoupling for Image Rain Removal](#) Di Lin, Xin WANG, Jia Shen, Renjie Zhang, Ruonan Liu, Miaohui Wang, Wuyuan Xie, Qing Guo, Ping Li
- [Learning with little mixing](#) Ingvar Ziemann, Stephen Tu
- [ComGAN: Unsupervised Disentanglement and Segmentation via Image Composition](#) Rui Ding, Kehua Guo, Xiangyuan Zhu, Zheng Wu, Liwei Wang
- [Sparse Winning Tickets are Data-Efficient Image Recognizers](#) Mukund Varma T, Xuxi Chen, Zhenyu Zhang, Tianlong Chen, Subhashini Venugopalan, Zhangyang Wang
- [VAEL: Bridging Variational Autoencoders and Probabilistic Logic Programming](#) Eleonora Misino, Giuseppe Marra, Emanuele Sansone
- [The alignment property of SGD noise and how it helps select flat minima: A stability analysis](#) Lei Wu, Mingze Wang, Weijie Su
- [EvenNet: Ignoring Odd-Hop Neighbors Improves Robustness of Graph Neural Networks](#) Runlin Lei, Zhen Wang, Yaliang Li, Bolin Ding, Zhewei Wei
- [Differentiable Analog Quantum Computing for Optimization and Control](#) Jiaqi Leng, Yuxiang Peng, Yi-Ling Qiao, Ming Lin, Xiaodi Wu
- [Promising or Elusive? Unsupervised Object Segmentation from Real-world Single Images](#) Yafei YANG, Bo Yang
- [Efficiently Factorizing Boolean Matrices using Proximal Gradient Descent](#) Sebastian Dalleiger, Jilles Vreeken
- [Learning from Future: A Novel Self-Training Framework for Semantic Segmentation](#) Ye Du, Yujun Shen, Haochen Wang, Jingjing Fei, Wei Li, Liwei Wu, Rui Zhao, Zehua Fu, Qingjie LIU
- [A Simple and Provably Efficient Algorithm for Asynchronous Federated Contextual Linear Bandits](#) Jiafan He, Tianhao Wang, Yifei Min, Quanquan Gu
- [How Powerful are K-hop Message Passing Graph Neural Networks](#) Jiarui Feng, Yixin Chen, Fuhai Li, Anindya Sarkar, Muhan Zhang
- [Dynamic Tensor Product Regression](#) Aravind Reddy, Zhao Song, Lichen Zhang
- [Generalization Analysis of Message Passing Neural Networks on Large Random Graphs](#) Sohir Maskey, Ron Levie, Yunseok Lee, Gitta Kutyniok
- [On Infinite Separations Between Simple and Optimal Mechanisms](#) Alexandros Psomas, Ariel Schwartzman Cohenca, S. Weinberg
- [Online Learning and Pricing for Network Revenue Management with Reusable Resources](#) Huiwen Jia, Cong Shi, Siqian Shen
- [Learning to Find Proofs and Theorems by Learning to Refine Search Strategies: The Case of Loop Invariant Synthesis](#) Jonathan Laurent, André Platzer
- [Exploitability Minimization in Games and Beyond](#) Denizalp Goktas, Amy Greenwald

- [Where2comm: Communication-Efficient Collaborative Perception via Spatial Confidence Maps](#) Yue Hu, Shaoheng Fang, Zixing Lei, Yiqi Zhong, Siheng Chen
- [Your Out-of-Distribution Detection Method is Not Robust!](#) Mohammad Azizmalayeri, Arshia Soltani Moakhar, Arman Zarei, Reihaneh Zohrabi, Mohammad Manzuri, Mohammad Hossein Rohban
- [Discovering and Overcoming Limitations of Noise-engineered Data-free Knowledge Distillation](#) Piyush Raikwar, Deepak Mishra
- [NaturalProver: Grounded Mathematical Proof Generation with Language Models](#) Sean Welleck, Jiacheng Liu, Ximing Lu, Hannaneh Hajishirzi, Yejin Choi
- [Smoothed Online Convex Optimization Based on Discounted-Normal-Predictor](#) Lijun Zhang, Wei Jiang, Jinfeng Yi, Tianbao Yang
- [One for All: Simultaneous Metric and Preference Learning over Multiple Users](#) Gregory Canal, Blake Mason, Ramya Korlakai Vinayak, Robert Nowak
- [Few-Shot Fast-Adaptive Anomaly Detection](#) Ze Wang, Yipin Zhou, Rui Wang, Tsung-Yu Lin, Ashish Shah, Ser Nam Lim
- [SegViT: Semantic Segmentation with Plain Vision Transformers](#) Bowen Zhang, Zhi Tian, Quan Tang, Xiangxiang Chu, Xiaolin Wei, Chunhua Shen, Yifan liu
- [Unsupervised Learning From Incomplete Measurements for Inverse Problems](#) Julián Tachella, Dongdong Chen, Mike Davies
- [Redeeming intrinsic rewards via constrained optimization](#) Eric Chen, Zhang-Wei Hong, Joni Pajarinen, Pulkit Agrawal
- [A Unified Evaluation of Textual Backdoor Learning: Frameworks and Benchmarks](#) Ganqu Cui, Lifan Yuan, Bingxiang He, Yangyi Chen, Zhiyuan Liu, Maosong Sun
- [MissDAG: Causal Discovery in the Presence of Missing Data with Continuous Additive Noise Models](#) Erdun Gao, Ignavier Ng, Mingming Gong, Li Shen, Wei Huang, Tongliang Liu, Kun Zhang, Howard Bondell
- [Escaping Saddle Points with Bias-Variance Reduced Local Perturbed SGD for Communication Efficient Nonconvex Distributed Learning](#) Tomoya Murata, Taiji Suzuki
- [AttCAT: Explaining Transformers via Attentive Class Activation Tokens](#) Yao Qiang, Deng Pan, Chengyin Li, Xin Li, Rhongho Jang, Dongxiao Zhu
- [A Theoretical Study on Solving Continual Learning](#) Gyuhak Kim, Changnan Xiao, Tatsuya Konishi, Zixuan Ke, Bing Liu
- [BiMLP: Compact Binary Architectures for Vision Multi-Layer Perceptrons](#) Yixing Xu, Xinghao Chen, Yunhe Wang
- [Understanding Deep Neural Function Approximation in Reinforcement Learning via \$\epsilon\$ -Greedy Exploration](#) Fanghui Liu, Luca Viano, Volkan Cevher
- [Misspecified Phase Retrieval with Generative Priors](#) Zhaoqiang Liu, Xinshao Wang, Jiulong Liu
- [Data-Efficient Augmentation for Training Neural Networks](#) Tian Yu Liu, Baharan Mirzasoleiman
- [Divide and Contrast: Source-free Domain Adaptation via Adaptive Contrastive Learning](#) Ziyi Zhang, Weikai Chen, Hui Cheng, Zhen Li, Siyuan Li, Liang Lin, Guanbin Li
- [Towards Human-Level Bimanual Dexterous Manipulation with Reinforcement Learning](#) Yuanpei Chen, Tianhao Wu, Shengjie Wang, Xidong Feng, Jiechuan Jiang, Zongqing Lu, Stephen McAleer, Hao Dong, Song-Chun Zhu, Yaodong Yang
- [Outlier-Robust Sparse Mean Estimation for Heavy-Tailed Distributions](#) Ilias Diakonikolas, Daniel Kane, Jasper Lee, Ankit Pensia
- [Local-Global MCMC kernels: the best of both worlds](#) Sergey Samsonov, Evgeny Lagutin, Marylou Gabri  , Alain Durmus, Alexey Naumov, Eric Moulines
- [The computational and learning benefits of Daleian neural networks](#) Adam Haber, Elad Schneidman
- [CLIPDraw: Exploring Text-to-Drawing Synthesis through Language-Image Encoders](#) Kevin Frans, Lisa Soros, Olaf Witkowski
- [Chaotic Dynamics are Intrinsic to Neural Network Training with SGD](#) Luis Herrmann, Maximilian Granz, Tim Landgraf
- [Efficient and Modular Implicit Differentiation](#) Mathieu Blondel, Quentin Berthet, Marco Cuturi, Roy Frostig, Stephan Hoyer, Felipe Llinares-Lopez, Fabian Pedregosa, Jean-Philippe Vert
- [Effective Dimension in Bandit Problems under Censorship](#) Gauthier Guinet, Saurabh Amin, Patrick Jaillet
- [Which Explanation Should I Choose? A Function Approximation Perspective to Characterizing Post Hoc Explanations](#) Tessa Han, Suraj Srinivas, Himabindu Lakkaraju
- [Accelerating Certified Robustness Training via Knowledge Transfer](#) Pratik Vaishnavi, Kevin Eykholt, Amir Rahmati
- [MOMA-LRG: Language-Refined Graphs for Multi-Object Multi-Actor Activity Parsing](#) Zelun Luo, Zane Durante, Linden Li, Wanze Xie, Ruochen Liu, Emily Jin, Zhuoyi Huang, Lun Yu Li, Jiajun Wu, Juan Carlos Niebles, Ehsan Adeli, Fei-Fei Li
- [MTNeuro: A Benchmark for Evaluating Representations of Brain Structure Across Multiple Levels of Abstraction](#) Jorge Quesada, Lakshmi Sathidevi, Ran Liu, Nauman Ahad, Joy Jackson, Mehdi Azabou, Jingyun Xiao, Christopher Liding, Matthew Jin, Carolina Urzay, William Gray-Roncal, Erik Johnson, Eva Dyer
- [FLamby: Datasets and Benchmarks for Cross-Silo Federated Learning in Realistic Healthcare Settings](#) Jean Ogier du Terrail, Samy-Safwan Ayed, Edwige Cyffers, Felix Grimberg, Chaoyang He, Regis Loeb, Paul Mangold, Tanguy Marchand, Othmane Marfoq, Erum Mushtaq, Boris Muzellec, Constantin Philippenko, Santiago Silva, Maria Tele  czuk, Shadi Albarqouni, Salman Avestimehr, Aur  lien Bellet, Aymeric Dieuleveut, Martin Jaggi, Sai Praneeth Karimireddy, Marco Lorenzi, Giovanni Neglia, Marc Tommasi, Mathieu Andreux
- [Foundation Posteriors for Approximate Probabilistic Inference](#) Mike Wu, Noah Goodman
- [Blackbox Attacks via Surrogate Ensemble Search](#) Zikui Cai, Chengyu Song, Srikanth Krishnamurthy, Amit Roy-Chowdhury, Salman Asif
- [Transition to Linearity of General Neural Networks with Directed Acyclic Graph Architecture](#) Libin Zhu, Chaoyue Liu, Misha Belkin
- [A Comprehensive Study on Large-Scale Graph Training: Benchmarking and Rethinking](#) Keyu Duan, Zirui Liu, Peihao Wang, Wenqing Zheng, Kaiyuan Zhou, Tianlang Chen, Xiaoli Li, Zhaoyang Wang

[Report an Issue](#) | [Name Change Policy](#)

- [Scale-invariant Learning by Physics Inversion](#) Philipp Holl, Vladlen Koltun, Nils Thuerey
- [Regret Bounds for Multilabel Classification in Sparse Label Regimes](#) Róbert Busa-Fekete, Heejin Choi, Krzysztof Dembczynski, Claudio Gentile, Henry Reeve, Balazs Szorenyi
- [Sample Constrained Treatment Effect Estimation](#) Raghavendra Addanki, David Arbour, Tung Mai, Cameron Musco, Anup Rao
- [CATER: Intellectual Property Protection on Text Generation APIs via Conditional Watermarks](#) Xuanli He, Qionghai Xu, Yi Zeng, Lingjuan Lyu, Fangzhao Wu, Jiwei Li, Ruoxi Jia
- [On the symmetries of the synchronization problem in Cryo-EM: Multi-Frequency Vector Diffusion Maps on the Projective Plane](#) Gabriele Cesa, Arash Behboodi, Taco S. Cohen, Max Welling
- [Sketching based Representations for Robust Image Classification with Provable Guarantees](#) Nishanth Dikkala, Sankeerth Rao Karingula, Raghu Meka, Jelani Nelson, Rina Panigrahy, Xin Wang
- [ResQ: A Residual Q Function-based Approach for Multi-Agent Reinforcement Learning Value Factorization](#) Siqi Shen, Mengwei Qiu, Jun Liu, Weiquan Liu, Yongquan Fu, Xinwang Liu, Cheng Wang
- [Non-identifiability and the Blessings of Misspecification in Models of Molecular Fitness](#) Eli Weinstein, Alan Amin, Jonathan Frazer, Debora Marks
- [Conditional Diffusion Process for Inverse Halftoning](#) Hao Jiang, Yadong Mu
- [Towards a Standardised Performance Evaluation Protocol for Cooperative MARL](#) Rihab Gorsane, Omayma Mahjoub, Ruan John de Kock, Roland Dubb, Siddarth Singh, Arnu Pretorius
- [Size and depth of monotone neural networks: interpolation and approximation](#) Dan Mikulincer, Daniel Reichman
- [Action-modulated midbrain dopamine activity arises from distributed control policies](#) Jack Lindsey, Ashok Litwin-Kumar
- [Latent Hierarchical Causal Structure Discovery with Rank Constraints](#) Biwei Huang, Charles Jia Han Low, Feng Xie, Clark Glymour, Kun Zhang
- [Beyond Real-world Benchmark Datasets: An Empirical Study of Node Classification with GNNs](#) Seiji Maekawa, Koki Noda, Yuya Sasaki, makoto onizuka
- [Doubly-Asynchronous Value Iteration: Making Value Iteration Asynchronous in Actions](#) Tian Tian, Kenny Young, Richard S. Sutton
- [Turbocharging Solution Concepts: Solving NEs, CE and CCEs with Neural Equilibrium Solvers](#) Luke Marris, Ian Gemp, Thomas Anthony, Andrea Tacchetti, Siqi Liu, Karl Tuyls
- [NOMAD: Nonlinear Manifold Decoders for Operator Learning](#) Jacob Seidman, Georgios Kissas, Paris Perdikaris, George J. Pappas
- [Semi-Supervised Video Salient Object Detection Based on Uncertainty-Guided Pseudo Labels](#) Yongri Piao, Chenyang Lu, Miao Zhang, Huchuan Lu
- [Debiased Causal Tree: Heterogeneous Treatment Effects Estimation with Unmeasured Confounding](#) Caizhi Tang, Huiyuan Wang, Xinyu Li, Qing Cui, Ya-Lin Zhang, Feng Zhu, Longfei Li, Jun Zhou, Linbo Jiang
- [Asymptotic Properties for Bayesian Neural Network in Besov Space](#) Kyeongwon Lee, Jaeyong Lee
- [A PAC-Bayesian Generalization Bound for Equivariant Networks](#) Arash Behboodi, Gabriele Cesa, Taco S. Cohen
- [Neural Network Architecture Beyond Width and Depth](#) Shijun Zhang, Zuowei Shen, Haizhao Yang
- [S-Prompts Learning with Pre-trained Transformers: An Occam's Razor for Domain Incremental Learning](#) Yabin Wang, Zhiwu Huang, Xiaopeng Hong
- [OmniVL: One Foundation Model for Image-Language and Video-Language Tasks](#) Junke Wang, Dongdong Chen, Zuxuan Wu, Chong Luo, Luowei Zhou, Yucheng Zhao, Yujia Xie, Ce Liu, Yu-Gang Jiang, Lu Yuan
- [Self-Organized Group for Cooperative Multi-agent Reinforcement Learning](#) Jianzhun Shao, Zhiqiang Lou, Hongchang Zhang, Yuhang Jiang, Shuncheng He, Xiangyang Ji
- [\(Optimal\) Online Bipartite Matching with Degree Information](#) Anders Aamand, Justin Chen, Piotr Indyk
- [Fairness in Federated Learning via Core-Stability](#) Bhaskar Ray Chaudhury, Linyi Li, Mintong Kang, Bo Li, Ruta Mehta
- [Log-Polar Space Convolution Layers](#) Bing Su, Ji-Rong Wen
- [CoNSoLe: Convex Neural Symbolic Learning](#) Haoran Li, Yang Weng, Hanghang Tong
- [DPM-Solver: A Fast ODE Solver for Diffusion Probabilistic Model Sampling in Around 10 Steps](#) Cheng Lu, Yuhao Zhou, Fan Bao, Jianfei Chen, Chongxuan Li, Jun Zhu
- [Hierarchical Channel-spatial Encoding for Communication-efficient Collaborative Learning](#) Qihua ZHOU, Song Guo, YI LIU, Jie ZHANG, Jiewei Zhang, Tao GUO, Zhenda XU, Xun Liu, Zhihao Qu
- [Attracting and Dispersing: A Simple Approach for Source-free Domain Adaptation](#) Shiqi Yang, yaxing wang, kai wang, Shangling Jui, Joost van de Weijer
- [SCINet: Time Series Modeling and Forecasting with Sample Convolution and Interaction](#) Minhao LIU, Ailing Zeng, Muxi Chen, Zhijian Xu, Qiuxia LAI, Lingna Ma, Qiang Xu
- [Exploration-Guided Reward Shaping for Reinforcement Learning under Sparse Rewards](#) Rati Devidze, Parameswaran Kamalaruban, Adish Singla
- [Active Exploration for Inverse Reinforcement Learning](#) David Lindner, Andreas Krause, Giorgia Ramponi
- [Subspace Recovery from Heterogeneous Data with Non-isotropic Noise](#) John C. Duchi, Vitaly Feldman, Lunjia Hu, Kunal Talwar
- [Lifelong Neural Predictive Coding: Learning Cumulatively Online without Forgetting](#) Alex Ororbia, Ankur Mali, C Lee Giles, Daniel Kifer
- [Align then Fusion: Generalized Large-scale Multi-view Clustering with Anchor Matching Correspondences](#) Siwei Wang, Xinwang Liu,

- [Geodesic Graph Neural Network for Efficient Graph Representation Learning](#) Lecheng Kong, Yixin Chen, Muhan Zhang
- [Improved Differential Privacy for SGD via Optimal Private Linear Operators on Adaptive Streams](#) Sergey Denisov, H. Brendan McMahan, John Rush, Adam Smith, Abhradeep Guha Thakurta
- [On Privacy and Personalization in Cross-Silo Federated Learning](#) Ken Liu, Shengyuan Hu, Steven Z. Wu, Virginia Smith
- [SHAQ: Incorporating Shapley Value Theory into Multi-Agent Q-Learning](#) Jianhong Wang, Yuan Zhang, Yunjie Gu, Tae-Kyun Kim
- [Trajectory balance: Improved credit assignment in GFlowNets](#) Nikolay Malkin, Moksh Jain, Emmanuel Bengio, Chen Sun, Yoshua Bengio
- [Instance-Dependent Near-Optimal Policy Identification in Linear MDPs via Online Experiment Design](#) Andrew Wagenmaker, Kevin G. Jamieson
-  [ProcTHOR: Large-Scale Embodied AI Using Procedural Generation](#) Matt Deitke, Eli VanderBilt, Alvaro Herrasti, Luca Weihs, Kiana Ehsani, Jordi Salvador, Winson Han, Eric Kolve, Aniruddha Kembhavi, Roozbeh Mottaghi
- [Efficient Sampling on Riemannian Manifolds via Langevin MCMC](#) Xiang Cheng, Jingzhao Zhang, Suvrit Sra
- [OpenFWL: Large-scale Multi-structural Benchmark Datasets for Full Waveform Inversion](#) Chengyuan Deng, Shihang Feng, Hanchen Wang, Xitong Zhang, Peng Jin, Yinan Feng, Qili Zeng, Yinpeng Chen, Youzuo Lin
- [High-Order Pooling for Graph Neural Networks with Tensor Decomposition](#) Chenqing Hua, Guillaume Rabusseau, Jian Tang
- [Natural image synthesis for the retina with variational information bottleneck representation](#) Babak Rahmani, Demetri Psaltis, Christophe Moser
- [Imitating Past Successes can be Very Suboptimal](#) Benjamin Eysenbach, Soumith Udatha, Russ R. Salakhutdinov, Sergey Levine
- [A Communication-efficient Algorithm with Linear Convergence for Federated Minimax Learning](#) Zhenyu Sun, Ermin Wei
- [Dynamic Graph Neural Networks Under Spatio-Temporal Distribution Shift](#) Zeyang Zhang, Xin Wang, Ziwei Zhang, Haoyang Li, Zhou Qin, Wenwu Zhu
- [Biologically Inspired Dynamic Thresholds for Spiking Neural Networks](#) Jianchuan Ding, Bo Dong, Felix Heide, Yufei Ding, Yunduo Zhou, Baocai Yin, Xin Yang
- [GraphQNTK: Quantum Neural Tangent Kernel for Graph Data](#) Yehui Tang, Junchi Yan
- [Trajectory-guided Control Prediction for End-to-end Autonomous Driving: A Simple yet Strong Baseline](#) Penghao Wu, Xiaosong Jia, Li Chen, Junchi Yan, Hongyang Li, Yu Qiao
- [360-MLC: Multi-view Layout Consistency for Self-training and Hyper-parameter Tuning](#) Bolivar Solarte, Chin-Hsuan Wu, Yueh-Cheng Liu, Yi-Hsuan Tsai, Min Sun
- [Distributed Optimization for Overparameterized Problems: Achieving Optimal Dimension Independent Communication Complexity](#) Bingqing Song, Ioannis Tsaknakis, Chung-Yiu Yau, Hoi-To Wai, Mingyi Hong
- [Falsification before Extrapolation in Causal Effect Estimation](#) Zeshan M Hussain, Michael Oberst, Ming-Chieh Shih, David Sontag
- [Surprise Minimizing Multi-Agent Learning with Energy-based Models](#) Karush Suri, Xiao Qi Shi, Konstantinos N Plataniotis, Yuri Lawryshyn
- [Geodesic Self-Attention for 3D Point Clouds](#) Zhengyu Li, XUAN TANG, Zihao Xu, Xihao Wang, Hui Yu, Mingsong Chen, xian wei
- [Test Time Adaptation via Conjugate Pseudo-labels](#) Sachin Goyal, Mingjie Sun, Aditi Raghunathan, J. Zico Kolter
- [SPD domain-specific batch normalization to crack interpretable unsupervised domain adaptation in EEG](#) Reinmar Kobler, Jun-ichiro Hirayama, Qibin Zhao, Motoaki Kawanabe
- [AVLEN: Audio-Visual-Language Embodied Navigation in 3D Environments](#) Sudipta Paul, Amit Roy-Chowdhury, Anoop Cherian
- [Semantic uncertainty intervals for disentangled latent spaces](#) Swami Sankaranarayanan, Anastasios Angelopoulos, Stephen Bates, Yaniv Romano, Phillip Isola
- [A Fast Scale-Invariant Algorithm for Non-negative Least Squares with Non-negative Data](#) Jelena Diakonikolas, Chenghui Li, Swati Padmanabhan, Chaobing Song
- [Self-supervised Amodal Video Object Segmentation](#) Jian Yao, Yuxin Hong, Chiyu Wang, Tianjun Xiao, Tong He, Francesco Locatello, David P Wipf, Yanwei Fu, Zheng Zhang
- [Beyond the Return: Off-policy Function Estimation under User-specified Error-measuring Distributions](#) Audrey Huang, Nan Jiang
- [Disentangling Transfer in Continual Reinforcement Learning](#) Maciej Wolczyk, Michał Zając, Razvan Pascanu, Łukasz Kuciński, Piotr Miłoś
- [Meta-Learning with Self-Improving Momentum Target](#) Jihoon Tack, Jongjin Park, Hankook Lee, Jaeho Lee, Jinwoo Shin
- [Online Allocation and Learning in the Presence of Strategic Agents](#) Steven Yin, Shipra Agrawal, Assaf Zeevi
- [Learning Contrastive Embedding in Low-Dimensional Space](#) Shuo Chen, Chen Gong, Jun Li, Jian Yang, Gang Niu, Masashi Sugiyama
- [Near-Optimal Randomized Exploration for Tabular Markov Decision Processes](#) Zhihan Xiong, Ruoqi Shen, Qiwen Cui, Maryam Fazel, Simon S. Du
- [Defending Against Adversarial Attacks via Neural Dynamic System](#) Xiyuan Li, Zou Xin, Weiwei Liu
- [On the Robustness of Graph Neural Diffusion to Topology Perturbations](#) Yang Song, Qiyu Kang, Sijie Wang, Kai Zhao, Wee Peng Tay
- [Few-shot Relational Reasoning via Connection Subgraph Pretraining](#) Qian Huang, Hongyu Ren, Jure Leskovec
- [Equivariant Networks for Zero-Shot Coordination](#) Darius Muglich, Christian Schroeder de Witt, Elise van der Pol, Shimon Whiteson, Jakob Foerster
- [Randomized Sketches for Clustering: Fast and Optimal Kernel \$k\$ -Means](#) Rong Yin, Yong Liu, Weiping Wang, Dan Meng
- [Quantile Constrained Reinforcement Learning: A Reinforcement Learning Framework Constraining Outage Probability](#) Whiyoung Jung, Myungsik Cho, Jongeui Park, Youngchul Sung

- [Procedural Image Programs for Representation Learning](#) Manel Baradad, Richard Chen, Jonas Wulff, Tongzhou Wang, Rogerio Feris, Antonio Torralba, Phillip Isola
- [RényiCL: Contrastive Representation Learning with Skew Rényi Divergence](#) Kyungmin Lee, Jinwoo Shin
- [On the Efficient Implementation of High Accuracy Optimality of Profile Maximum Likelihood](#) Moses Charikar, Zhihao Jiang, Kirankumar Shiragur, Aaron Sidford
- [Paraphrasing Is All You Need for Novel Object Captioning](#) Cheng-Fu Yang, Yao-Hung Hubert Tsai, Wan-Cyuan Fan, Russ R. Salakhutdinov, Louis-Philippe Morency, Frank Wang
- [Adversarial Robustness is at Odds with Lazy Training](#) Yunjuan Wang, Enayat Ullah, Poorya Mianjy, Raman Arora
- [Uncertainty Estimation for Multi-view Data: The Power of Seeing the Whole Picture](#) Myong Chol Jung, He Zhao, Joanna Dipnall, Belinda Gabbe, Lan Du
- [Motion Transformer with Global Intention Localization and Local Movement Refinement](#) Shaoshuai Shi, Li Jiang, Dengxin Dai, Bernt Schiele
- [No-regret learning in games with noisy feedback: Faster rates and adaptivity via learning rate separation](#) Yu-Guan Hsieh, Kimon Antonakopoulos, Volkan Cevher, Panayotis Mertikopoulos
- [First-Order Algorithms for Min-Max Optimization in Geodesic Metric Spaces](#) Michael Jordan, Tianyi Lin, Emmanouil-Vasileios Vlatakis-Gkaragkounis
- [Feature-Proxy Transformer for Few-Shot Segmentation](#) Jian-Wei Zhang, Yifan Sun, Yi Yang, Wei Chen
- [Conformal Frequency Estimation with Sketched Data](#) Matteo Sesia, Stefano Favaro
- [Revisiting Active Sets for Gaussian Process Decoders](#) Pablo Moreno-Muñoz, Cilie Feldager, Søren Hauberg
- [Exact learning dynamics of deep linear networks with prior knowledge](#) Lukas Braun, Clémentine Dominé, James Fitzgerald, Andrew Saxe
- [AMP: Automatically Finding Model Parallel Strategies with Heterogeneity Awareness](#) Dacheng Li, Hongyi Wang, Eric Xing, Hao Zhang
- [A consistently adaptive trust-region method](#) Fadi Hamad, Oliver Hinder
- [Online Neural Sequence Detection with Hierarchical Dirichlet Point Process](#) Weihan Li, Yu Qi, Gang Pan
- [Leveraging Inter-Layer Dependency for Post -Training Quantization](#) changbao wang, DanDan Zheng, Yuanliu Liu, Liang Li
- [Emergence of Hierarchical Layers in a Single Sheet of Self-Organizing Spiking Neurons](#) Paul Bertens, Seong-Whan Lee
- [A gradient sampling method with complexity guarantees for Lipschitz functions in high and low dimensions](#) Damek Davis, Dmitriy Drusvyatskiy, Yin Tat Lee, Swati Padmanabhan, Guanghao Ye
- [CyCLIP: Cyclic Contrastive Language-Image Pretraining](#) Shashank Goel, Hritik Bansal, Sumit Bhatia, Ryan Rossi, Vishwa Vinay, Aditya Grover
- [When does dough become a bagel? Analyzing the remaining mistakes on ImageNet](#) Vijay Vasudevan, Benjamin Caine, Raphael Gontijo Lopes, Sara Fridovich-Keil, Rebecca Roelofs
- [Spatial Pruned Sparse Convolution for Efficient 3D Object Detection](#) Jianhui Liu, Yukang Chen, Xiaoqing Ye, Zhuotao Tian, Xiao Tan, Xiaojuan Qi
- [Sampling without Replacement Leads to Faster Rates in Finite-Sum Minimax Optimization](#) Aniket Das, Bernhard Schölkopf, Michael Muehlebach
- [Feature Learning in \$L_2\$ -regularized DNNs: Attraction/Repulsion and Sparsity](#) Arthur Jacot, Eugene Golikov, Clement Hongler, Franck Gabriel
- [Set-based Meta-Interpolation for Few-Task Meta-Learning](#) Seanie Lee, Bruno Andreis, Kenji Kawaguchi, Juho Lee, Sung Ju Hwang
- [Non-deep Networks](#) Ankit Goyal, Alexey Bochkovskiy, Jia Deng, Vladlen Koltun
- [Low-rank Optimal Transport: Approximation, Statistics and Debiasing](#) Meyer Scetbon, Marco Cuturi
- [Embodied Scene-aware Human Pose Estimation](#) Zhengyi Luo, Shun Iwase, Ye Yuan, Kris Kitani
- [Privacy Induces Robustness: Information-Computation Gaps and Sparse Mean Estimation](#) Kristian Georgiev, Samuel Hopkins
- [Batch Bayesian Optimization on Permutations using the Acquisition Weighted Kernel](#) Changyong Oh, Roberto Bondesan, Efstratios Gavves, Max Welling
- [Supervised Training of Conditional Monge Maps](#) Charlotte Bunne, Andreas Krause, Marco Cuturi
- [Formalizing Consistency and Coherence of Representation Learning](#) Harald Strömfelt, Luke Dickens, Artur Garcez, Alessandra Russo
- [Positively Weighted Kernel Quadrature via Subsampling](#) Satoshi Hayakawa, Harald Oberhauser, Terry Lyons
- [Guaranteed Conservation of Momentum for Learning Particle-based Fluid Dynamics](#) Lukas Prantl, Benjamin Ummenhofer, Vladlen Koltun, Nils Thuerey
- [M4Singer: A Multi-Style, Multi-Singer and Musical Score Provided Mandarin Singing Corpus](#) Lichao Zhang, Ruiqi Li, Shoutong Wang, Liqun Deng, Jinglin Liu, Yi Ren, Jinzheng He, Rongjie Huang, Jieming Zhu, Xiao Chen, Zhou Zhao
- [Bayesian Spline Learning for Equation Discovery of Nonlinear Dynamics with Quantified Uncertainty](#) Luning Sun, Daniel Huang, Hao Sun, Jian-Xun Wang
- [Convexity Certificates from Hessians](#) Julien Klaus, Niklas Merk, Konstantin Wiedom, Sören Laue, Joachim Giesen
- [FR: Folded Rationalization with a Unified Encoder](#) Wei Liu, Haozhao Wang, Jun Wang, Ruixuan Li, Chao Yue, YuanKai Zhang
- [Max-Min Off-Policy Actor-Critic Method Focusing on Worst-Case Robustness to Model Misspecification](#) Takumi Tanabe, Rei Sato, Kazuto Fukuchi, Jun Sakuma, Youhei Akimoto
- [One-Inlier is First: Towards Efficient Position Encoding for Point Cloud Registration](#) Fan Yang, Lin Guo, Zhi Chen, Wenbing Tao

- [Adaptive Distribution Calibration for Few-Shot Learning with Hierarchical Optimal Transport](#) Dandan Guo, Long Tian, He Zhao, Mingyuan Zhou, Hongyuan Zha
- [InsNet: An Efficient, Flexible, and Performant Insertion-based Text Generation Model](#) Sidi Lu, Tao Meng, Nanyun Peng
- [Weighted Distillation with Unlabeled Examples](#) Fotis Iliopoulos, Vasilis Kontonis, Cenk Baykal, Gaurav Menghani, Khoa Trinh, Erik Vee
- [When Does Group Invariant Learning Survive Spurious Correlations?](#) Yimeng Chen, Ruibin Xiong, Zhi-Ming Ma, Yanyan Lan
- [SNAKE: Shape-aware Neural 3D Keypoint Field](#) Chengliang Zhong, Peixing You, Xiaoxue Chen, Hao Zhao, Fuchun Sun, Guyue Zhou, Xiaodong Mu, Chuang Gan, Wenbing Huang
- [Single Loop Gaussian Homotopy Method for Non-convex Optimization](#) Hidenori Iwakiri, Yuhang Wang, Shinji Ito, Akiko Takeda
- [Minimax Optimal Online Imitation Learning via Replay Estimation](#) Gokul Swamy, Nived Rajaraman, Matt Peng, Sanjiban Choudhury, J. Bagnell, Steven Z. Wu, Jiantao Jiao, Kannan Ramchandran
- [Multi-layer State Evolution Under Random Convolutional Design](#) Max Daniels, Cedric Gerbelot, Florent Krzakala, Lenka Zdeborová
- [Mixture-of-Experts with Expert Choice Routing](#) Yanqi Zhou, Tao Lei, Hanxiao Liu, Nan Du, Yanping Huang, Vincent Zhao, Andrew M. Dai, zhifeng Chen, Quoc V Le, James Laudon
- [GULP: a prediction-based metric between representations](#) Enric Boix-Adsera, Hannah Lawrence, George Stepaniants, Philippe Rigollet
- [On the convergence of policy gradient methods to Nash equilibria in general stochastic games](#) Angeliki Giannou, Kyriakos Lotidis, Panayotis Mertikopoulos, Emmanouil-Vasileios Vlatakis-Gkaragkounis
- [SAGDA: Achieving \$\mathcal{O}\(\epsilon^{-2}\)\$ Communication Complexity in Federated Min-Max Learning](#) Haibo Yang, Zhuqing Liu, Xin Zhang, Jia Liu
- [ALMA: Hierarchical Learning for Composite Multi-Agent Tasks](#) Shariq Iqbal, Robby Costales, Fei Sha
- [Improved Bounds on Neural Complexity for Representing Piecewise Linear Functions](#) Kuan-Lin Chen, Harinath Garudadri, Bhaskar D Rao
- [Assaying Out-Of-Distribution Generalization in Transfer Learning](#) Florian Wenzel, Andrea Dittadi, Peter Gehler, Carl-Johann Simon-Gabriel, Max Horn, Dominik Zietlow, David Kernert, Chris Russell, Thomas Brox, Bernt Schiele, Bernhard Schölkopf, Francesco Locatello
- [M2N: Mesh Movement Networks for PDE Solvers](#) Wenbin Song, Mingrui Zhang, Joseph G Wallwork, Junpeng Gao, Zheng Tian, Fanglei Sun, Matthew Piggott, Junqing Chen, Zuoqiang Shi, Xiang Chen, Jun Wang
- [Physics-Informed Implicit Representations of Equilibrium Network Flows](#) Kevin D. Smith, Francesco Seccamonte, Ananthram Swami, Francesco Bullo
- [Learning Interface Conditions in Domain Decomposition Solvers](#) Ali Taghibakhshi, Nicolas Nytko, Tareq Uz Zaman, Scott MacLachlan, Luke Olson, Matthew West
- [TANKBind: Trigonometry-Aware Neural Networks for Drug-Protein Binding Structure Prediction](#) Wei Lu, Qifeng Wu, Jixian Zhang, Jiahua Rao, Chengtao Li, Shuangjia Zheng
- [Hamiltonian Latent Operators for content and motion disentanglement in image sequences](#) Asif Khan, Amos J. Storkey
- [Convolutional Neural Networks on Graphs with Chebyshev Approximation, Revisited](#) Mingguo He, Zhewei Wei, Ji-Rong Wen
- [A Kernelised Stein Statistic for Assessing Implicit Generative Models](#) Wenkai Xu, Gesine D Reinert
- [Fine-Grained Semantically Aligned Vision-Language Pre-Training](#) Juncheng Li, XIN HE, Longhui Wei, Long Qian, Linchao Zhu, Lingxi Xie, Yueting Zhuang, Qi Tian, Siliang Tang
- [Fast Distance Oracles for Any Symmetric Norm](#) Yichuan Deng, Zhao Song, OMRI WEINSTEIN, Ruizhe Zhang
- [Outlier-Robust Sparse Estimation via Non-Convex Optimization](#) Yu Cheng, Ilias Diakonikolas, Rong Ge, Shivam Gupta, Daniel Kane, Mahdi Soltanolkotabi
- [Recurrent Convolutional Neural Networks Learn Succinct Learning Algorithms](#) Surbhi Goel, Sham Kakade, Adam Kalai, Cyril Zhang
- [Navigating Memory Construction by Global Pseudo-Task Simulation for Continual Learning](#) Yejia Liu, Wang Zhu, Shaolei Ren
- [Contact-aware Human Motion Forecasting](#) Wei Mao, miaomiao Liu, Richard I Hartley, Mathieu Salzmann
- [Spectral Bias in Practice: The Role of Function Frequency in Generalization](#) Sara Fridovich-Keil, Raphael Gontijo Lopes, Rebecca Roelofs
- [Out-of-Distribution Detection with An Adaptive Likelihood Ratio on Informative Hierarchical VAE](#) Yewen Li, Chaojie Wang, Xiaobo Xia, Tongliang Liu, xin miao, Bo An
- [FOF: Learning Fourier Occupancy Field for Monocular Real-time Human Reconstruction](#) Qiao Feng, Yebin Liu, Yu-Kun Lai, Jingyu Yang, Kun Li
- [Rate-Optimal Online Convex Optimization in Adaptive Linear Control](#) Asaf Benjamin Cassel, Alon Peled-Cohen, Tomer Koren
- [RTFormer: Efficient Design for Real-Time Semantic Segmentation with Transformer](#) Jian Wang, Chenhui Gou, Qiman Wu, Haocheng Feng, Junyu Han, Errui Ding, Jingdong Wang
- [Staggered Rollout Designs Enable Causal Inference Under Interference Without Network Knowledge](#) Mayleen Cortez, Matthew Eichhorn, Christina Yu
- [Censored Quantile Regression Neural Networks for Distribution-Free Survival Analysis](#) Tim Pearce, Jong-Hyeon Jeong, yichen jia, Jun Zhu
- [Learning Graph-embedded Key-event Back-tracing for Object Tracking in Event Clouds](#) Zhiyu Zhu, Junhui Hou, Xianqiang Lyu
- [Efficient and Near-Optimal Smoothed Online Learning for Generalized Linear Functions](#) Adam Block, Max Simchowitz
- [Systematic improvement of neural network quantum states using Lanczos](#) Hongwei Chen, Douglas Hendry, Phillip Weinberg, Adrian Feiguin
- [Grounded Reinforcement Learning: Learning to Win the Game under Human Commands](#) Shusheng Xu, Huaijie Wang, YI WU
- [Enhance the Visual Representation via Discrete Adversarial Training](#) Xiaofeng Mao, YueFeng Chen, Ranjie Duan, Yao Zhu, Gege Qi, shaokai ye, Xiaodan Li, Rong Zhang, Hui Xue'

- [Natural Color Fool: Towards Boosting Black-box Unrestricted Attacks](#) Shengming Yuan, Qilong Zhang, Lianli Gao, Yaya Cheng, Jingkuan Song
- [Old can be Gold: Better Gradient Flow can Make Vanilla-GCNs Great Again](#) AJAY JAISWAL, Peihao Wang, Tianlong Chen, Justin Rousseau, Ying Ding, Zhangyang Wang
- [Egocentric Video-Language Pretraining](#) Kevin Qinghong Lin, Jinpeng Wang, Mattia Soldan, Michael Wray, Rui Yan, Eric Z. XU, Difei Gao, Rong-Cheng Tu, Wenzhe Zhao, Weijie Kong, Chengfei Cai, WANG HongFa, Dima Damen, Bernard Ghanem, Wei Liu, Mike Zheng Shou
- [Society of Agents: Regret Bounds of Concurrent Thompson Sampling](#) Yan Chen, Perry Dong, Qinxun Bai, Maria Dimakopoulou, Wei Xu, Zhengyuan Zhou
- [A Non-asymptotic Analysis of Non-parametric Temporal-Difference Learning](#) Eloise Berthier, Ziad Kobeissi, Francis Bach
- [CEIP: Combining Explicit and Implicit Priors for Reinforcement Learning with Demonstrations](#) Kai Yan, Alex Schwing, Yu-Xiong Wang
- [Memorization and Optimization in Deep Neural Networks with Minimum Over-parameterization](#) Simone Bombari, Mohammad Hossein Amani, Marco Mondelli
- [LAMP: Extracting Text from Gradients with Language Model Priors](#) Mislav Balunovic, Dimitar Dimitrov, Nikola Jovanović, Martin Vechev
- [Explainability Via Causal Self-Talk](#) Nicholas A. Roy, Junkyung Kim, Neil Rabinowitz
- [Continual Learning with Evolving Class Ontologies](#) Zhiqiu Lin, Deepak Pathak, Yu-Xiong Wang, Deva Ramanan, Shu Kong
- [A gradient estimator via L1-randomization for online zero-order optimization with two point feedback](#) Arya Akhavan, Evgenii Chzhen, Massimiliano Pontil, Alexandre Tsybakov
- [On the SDEs and Scaling Rules for Adaptive Gradient Algorithms](#) Sathika Malladi, Kaifeng Lyu, Abhishek Panigrahi, Sanjeev Arora
- [Improving Neural Ordinary Differential Equations with Nesterov's Accelerated Gradient Method](#) Ho Huu Nghia Nguyen, Tan Nguyen, Huyen Vo, Stanley Osher, Thieu Vo
- [VER: Scaling On-Policy RL Leads to the Emergence of Navigation in Embodied Rearrangement](#) Erik Wijmans, Irfan Essa, Dhruv Batra
- [Neural Attentive Circuits](#) Martin Weiss, Nasim Rahaman, Francesco Locatello, Chris Pal, Yoshua Bengio, Bernhard Schölkopf, Erran Li Li, Nicolas Ballas
- [CLEVRER-Humans: Describing Physical and Causal Events the Human Way](#) Jiayuan Mao, Xuelin Yang, Xikun Zhang, Noah Goodman, Jiajun Wu
- [3D Concept Grounding on Neural Fields](#) Yining Hong, Yilun Du, Chunru Lin, Josh Tenenbaum, Chuang Gan
- [Evaluating Graph Generative Models with Contrastively Learned Features](#) Hamed Shirzad, Kaveh Hassani, Danica J. Sutherland
- [Path Independent Equilibrium Models Can Better Exploit Test-Time Computation](#) Cem Anil, Ashwini Pople, Kaiqu Liang, Johannes Treutlein, Yuhuai Wu, Shaojie Bai, J. Zico Kolter, Roger B. Grosse
- [Implications of Model Indeterminacy for Explanations of Automated Decisions](#) Marc-Étienne Brunet, Ashton Anderson, Richard Zemel
- [Debugging and Explaining Metric Learning Approaches: An Influence Function Based Perspective](#) Ruofan Liu, Yun Lin, XIANGLIN YANG, Jin Song Dong
- [Meta-Reinforcement Learning with Self-Modifying Networks](#) Mathieu Chalvidal, Thomas Serre, Rufin VanRullen
- [FairVFL: A Fair Vertical Federated Learning Framework with Contrastive Adversarial Learning](#) Tao Qi, Fangzhao Wu, Chuhan Wu, Lingjuan Lyu, Tong Xu, Hao Liao, Zhongliang Yang, Yongfeng Huang, Xing Xie
- [Incorporating Bias-aware Margins into Contrastive Loss for Collaborative Filtering](#) An Zhang, Wenchang Ma, Xiang Wang, Tat-Seng Chua
- [Multiview Human Body Reconstruction from Uncalibrated Cameras](#) Zhixuan Yu, Linguang Zhang, Yuanlu Xu, Chengcheng Tang, LUAN TRAN, Cem Keskin, Hyun Soo Park
- [MACK: Multimodal Aligned Conceptual Knowledge for Unpaired Image-text Matching](#) Yan Huang, Yuming Wang, Yunan Zeng, Liang Wang
- [Learning Optical Flow from Continuous Spike Streams](#) Rui Zhao, Ruiqin Xiong, Jing Zhao, Zhaoqi Yu, Xiaopeng Fan, Tiejun Huang
- [Improving 3D-aware Image Synthesis with A Geometry-aware Discriminator](#) Zifan Shi, Yinghao Xu, Yujun Shen, Deli Zhao, Qifeng Chen, Dit-Yan Yeung
- [A Consistent and Differentiable Lp Canonical Calibration Error Estimator](#) Teodora Popordanoska, Raphael Sayer, Matthew Blaschko
- [Transform Once: Efficient Operator Learning in Frequency Domain](#) Michael Poli, Stefano Massaroli, Federico Berto, Jinkyoo Park, Tri Dao, Christopher Ré, Stefano Ermon
- [Fused Orthogonal Alternating Least Squares for Tensor Clustering](#) Jiacheng Wang, Dan Nicolae
- [A Solver-free Framework for Scalable Learning in Neural ILP Architectures](#) Yatin Nandwani, Rishabh Ranjan, - Mausam, Parag Singla
- [Distributed Distributionally Robust Optimization with Non-Convex Objectives](#) Yang Jiao, Kai Yang, Dongjin Song
- [Parameter-free Regret in High Probability with Heavy Tails](#) Jiujia Zhang, Ashok Cutkosky
- [Non-Convex Bilevel Games with Critical Point Selection Maps](#) Michael Arbel, Julien Mairal
- [Consistent Sufficient Explanations and Minimal Local Rules for explaining the decision of any classifier or regressor](#) Salim I. Amoukou, Nicolas Brunel
- [High-dimensional Additive Gaussian Processes under Monotonicity Constraints](#) Andrés López-Lopera, Francois Bachoc, Olivier Roustant
- [Spherical Channels for Modeling Atomic Interactions](#) Larry Zitnick, Abhishek Das, Adeesh Kolluru, Janice Lan, Muhammed Shuaibi, Anuroop Sriram, Zachary Ulissi, Brandon Wood
- [Handcrafted Backdoors in Deep Neural Networks](#) Sanghyun Hong, Nicholas Carlini, Alexey Kurakin
- [Touch and Go: Learning from Human-Collected Vision and Touch](#) Fengyu Yang, Chenyang Ma, Jiacheng Zhang, Jing Zhu, Wenzhen Yuan,

- [SoLAR: Sinkhorn Label Refinery for Imbalanced Partial-Label Learning](#) Haobo Wang, Mingxuan Xia, Yixuan Li, Yuren Mao, Lei Feng, Gang Chen, Junbo Zhao
- [Log-Linear-Time Gaussian Processes Using Binary Tree Kernels](#) Michael K. Cohen, Samuel Daulton, Michael A Osborne
- [Recovering Private Text in Federated Learning of Language Models](#) Samyak Gupta, Yangsibo Huang, Zexuan Zhong, Tianyu Gao, Kai Li, Danqi Chen
- [INRAS: Implicit Neural Representation for Audio Scenes](#) Kun Su, Mingfei Chen, Eli Shlizerman
- [Non-Monotonic Latent Alignments for CTC-Based Non-Autoregressive Machine Translation](#) Chenze Shao, Yang Feng
- [Attention-based Neural Cellular Automata](#) Mattie Tesfaldet, Derek Nowrouzezahrai, Chris Pal
- [Learning Deep Input-Output Stable Dynamics](#) Ryosuke Kojima, Yuji Okamoto
- [Grounded Video Situation Recognition](#) Zeeshan Khan, C.V. Jawahar, Makarand Tapaswi
- [Policy Optimization with Advantage Regularization for Long-Term Fairness in Decision Systems](#) Eric Yu, Zhizhen Qin, Min Kyung Lee, Sicun Gao
- [Gradient Descent: The Ultimate Optimizer](#) Kartik Chandra, Audrey Xie, Jonathan Ragan-Kelley, ERIK MEIJER
- [DAGMA: Learning DAGs via M-matrices and a Log-Determinant Acyclicity Characterization](#) Kevin Bello, Bryon Aragam, Pradeep Ravikumar
- [Forward-Backward Latent State Inference for Hidden Continuous-Time semi-Markov Chains](#) Nicolai Engelmann, Heinz Koeppl
- [LobsDICE: Offline Learning from Observation via Stationary Distribution Correction Estimation](#) Geon-Hyeong Kim, Jongmin Lee, Youngsoo Jang, Hongseok Yang, Kee-Eung Kim
- [Ensemble of Averages: Improving Model Selection and Boosting Performance in Domain Generalization](#) Devansh Arpit, Huan Wang, Yingbo Zhou, Caiming Xiong
- [Is \$L^2\$ Physics Informed Loss Always Suitable for Training Physics Informed Neural Network?](#) Chuwei Wang, Shanda Li, Di He, Liwei Wang
- [Vision GNN: An Image is Worth Graph of Nodes](#) Kai Han, Yunhe Wang, Jianyuan Guo, Yehui Tang, Enhua Wu
- [ELIGN: Expectation Alignment as a Multi-Agent Intrinsic Reward](#) Zixian Ma, Rose Wang, Fei-Fei Li, Michael Bernstein, Ranjay Krishna
- [Factored DRO: Factored Distributionally Robust Policies for Contextual Bandits](#) Tong Mu, Yash Chandak, Tatsunori B. Hashimoto, Emma Brunskill
- [Could Giant Pre-trained Image Models Extract Universal Representations?](#) Yutong Lin, Ze Liu, Zheng Zhang, Han Hu, Nanning Zheng, Stephen Lin, Yue Cao
- [A Scalable Deterministic Global Optimization Algorithm for Training Optimal Decision Tree](#) Kaixun Hua, Jiayang Ren, Yankai Cao
- [Thor: Wielding Hammers to Integrate Language Models and Automated Theorem Provers](#) Albert Qiaochu Jiang, Wenda Li, Szymon Tworkowski, Konrad Czechowski, Tomasz Odrzygóźdź, Piotr Miłoś, Yuhuai Wu, Mateja Jamnik
- [Rethinking Knowledge Graph Evaluation Under the Open-World Assumption](#) Haotong Yang, Zhouchen Lin, Muhan Zhang
- [KERPLE: Kernelized Relative Positional Embedding for Length Extrapolation](#) Ta-Chung Chi, Ting-Han Fan, Peter J Ramadge, Alexander Rudnicky
- [Zonotope Domains for Lagrangian Neural Network Verification](#) Matt Jordan, Jonathan Hayase, Alex Dimakis, Sewoong Oh
- [Neural Basis Models for Interpretability](#) Filip Radenovic, Abhimanyu Dubey, Dhruv Mahajan
- [RecursiveMix: Mixed Learning with History](#) Lingfeng Yang, Xiang Li, Borui Zhao, Renjie Song, Jian Yang
- [GAR: Generalized Autoregression for Multi-Fidelity Fusion](#) Yuxin Wang, Zheng Xing, WEI XING
- [Anonymized Histograms in Intermediate Privacy Models](#) Badih Ghazi, Pritish Kamath, Ravi Kumar, Pasin Manurangsi
- [Truly Deterministic Policy Optimization](#) Ehsan Saleh, Saba Ghaffari, Tim Bretl, Matthew West
- [Language Models with Image Descriptors are Strong Few-Shot Video-Language Learners](#) Zhenhailong Wang, Manling Li, Ruochen Xu, Luowei Zhou, Jie Lei, Xudong Lin, Shuohang Wang, Ziyi Yang, Chenguang Zhu, Derek Hoiem, Shih-Fu Chang, Mohit Bansal, Heng Ji
- [3DB: A Framework for Debugging Computer Vision Models](#) Guillaume Leclerc, Hadi Salman, Andrew Ilyas, Sai Vemprala, Logan Engstrom, Vibhav Vineet, Kai Xiao, Pengchuan Zhang, Shibani Santurkar, Greg Yang, Ashish Kapoor, Aleksander Madry
- [Empirical Gateaux Derivatives for Causal Inference](#) Michael Jordan, Yixin Wang, Angela Zhou
- [Calibrated Data-Dependent Constraints with Exact Satisfaction Guarantees](#) Songkai Xue, Yuekai Sun, Mikhail Yurochkin
- [Towards Consistency in Adversarial Classification](#) Laurent Meunier, Raphael Ettegui, Rafael Pinot, Yann Chevaleyre, Jamal Atif
- [Residual Multiplicative Filter Networks for Multiscale Reconstruction](#) Shayan Shekarfroush, David Lindell, David J. Fleet, Marcus A. Brubaker
- [WT-MVSNet: Window-based Transformers for Multi-view Stereo](#) Jinli Liao, Yikang Ding, Yoli Shavit, Dihe Huang, Shihao Ren, Jia Guo, Wensen Feng, Kai Zhang
- [Private Isotonic Regression](#) Badih Ghazi, Pritish Kamath, Ravi Kumar, Pasin Manurangsi
- [A Closer Look at Offline RL Agents](#) Yuwei Fu, Di Wu, Benoit Boulet
- [Generalization Bounds for Estimating Causal Effects of Continuous Treatments](#) Xin Wang, Shengfei Lyu, Xingyu Wu, Tianhao Wu, Huanhuan Chen
- [Better Uncertainty Calibration via Proper Scores for Classification and Beyond](#) Sebastian Gruber, Florian Buettner
- [Video Diffusion Models](#) Jonathan Ho, Tim Salimans, Alexey Gritsenko, William Chan, Mohammad Norouzi, David J. Fleet
- [Formulating Robustness Against Unforeseen Attacks](#) Sihui Dai, Saeed Mahloujifar, Prateek Mittal
- [Single Model Uncertainty Estimation via Stochastic Data Centering](#) Jayaraman Thiagarajan, Rushil Anirudh, Vivek Sivaraman

[Report an Issue](#) | [Name Change Policy](#)

- [Autoinverse: Uncertainty Aware Inversion of Neural Networks](#) Navid Ansari, Hans-peter Seidel, Nima Vahidi Ferdowsi, Vahid Babaei
- [FedPop: A Bayesian Approach for Personalised Federated Learning](#) Nikita Kotelevskii, Maxime Vono, Alain Durmus, Eric Moulines
- [Semantic Diffusion Network for Semantic Segmentation](#) Haoru Tan, Sitong Wu, Jimin Pi
- [Iterative Structural Inference of Directed Graphs](#) Aoran Wang, Jun Pang
- [An efficient graph generative model for navigating ultra-large combinatorial synthesis libraries](#) Aryan Pedawi, Pawel Gniewek, Chaoyi Chang, Brandon Anderson, Henry van den Bedem
- [Learning to Discover and Detect Objects](#) Vladimir Fomenko, Ismail Elezi, Deva Ramanan, Laura Leal-Taixé, Aljosa Osep
- [Simulation-guided Beam Search for Neural Combinatorial Optimization](#) Jinho Choo, Yeong-Dae Kwon, Jihoon Kim, Jeongwoo Jae, André Hottung, Kevin Tierney, Youngjune Gwon
- [FlowHMM: Flow-based continuous hidden Markov models](#) Pawel Lorek, Rafal Nowak, Tomasz Trzcinski, Maciej Zieba
- [Near Instance-Optimal PAC Reinforcement Learning for Deterministic MDPs](#) Andrea Tirinzoni, Aymen Al Marjani, Emilie Kaufmann
- [VICRegL: Self-Supervised Learning of Local Visual Features](#) Adrien Bardes, Jean Ponce, Yann LeCun
- [Alleviating Adversarial Attacks on Variational Autoencoders with MCMC](#) Anna Kuzina, Max Welling, Jakub Tomczak
- [Policy Gradient With Serial Markov Chain Reasoning](#) Edoardo Cetin, Oya Celiktutan
- [DTG-SSOD: Dense Teacher Guidance for Semi-Supervised Object Detection](#) Gang Li, Xiang Li, Yujie Wang, Wu Yichao, Ding Liang, Shanshan Zhang
- [Human-AI Shared Control via Policy Dissection](#) Quanyi Li, Zhenghao Peng, Haibin Wu, Lan Feng, Bolei Zhou
- [Improving Generative Adversarial Networks via Adversarial Learning in Latent Space](#) Yang Li, Yichuan Mo, Liangliang Shi, Junchi Yan
- [ShapeCrafter: A Recursive Text-Conditioned 3D Shape Generation Model](#) Rao Fu, Xiao Zhan, YIWEN CHEN, Daniel Ritchie, Srinath Sridhar
- [SoundSpaces 2.0: A Simulation Platform for Visual-Acoustic Learning](#) Changan Chen, Carl Schissler, Sanchit Garg, Philip Kobernik, Alexander Clegg, Paul Calamia, Dhruv Batra, Philip Robinson, Kristen Grauman
- [AutoWS-Bench-101: Benchmarking Automated Weak Supervision with 100 Labels](#) Nicholas Roberts, Xintong Li, Tzu-Heng Huang, Dyah Adila, Spencer Schoenberg, Cheng-Yu Liu, Lauren Pick, Haotian Ma, Aws Albarghouthi, Frederic Sala
- [Towards Disentangling Information Paths with Coded ResNeXt](#) Apostolos Avranas, Marios Kountouris
- [Time-Conditioned Dances with Simplicial Complexes: Zigzag Filtration Curve based Supra-Hodge Convolution Networks for Time-series Forecasting](#) Yuzhou Chen, Yulia Gel, H. Vincent Poor
- [Are Defenses for Graph Neural Networks Robust?](#) Felix Mujkanovic, Simon Geisler, Stephan Günnemann, Aleksandar Bojchevski
- [GraB: Finding Provably Better Data Permutations than Random Reshuffling](#) Yucheng Lu, Wentao Guo, Christopher M. De Sa
- [Amortized Proximal Optimization](#) Juhan Bae, Paul Vicol, Jeff Z. HaoChen, Roger B. Grosse
- [A Statistical Online Inference Approach in Averaged Stochastic Approximation](#) Chuhan Xie, Zhihua Zhang
- [SAViT: Structure-Aware Vision Transformer Pruning via Collaborative Optimization](#) Chuanyang Zheng, zheyang li, Kai Zhang, Zhi Yang, Wenming Tan, Jun Xiao, Ye Ren, Shiliang Pu
- [Deciding What to Model: Value-Equivalent Sampling for Reinforcement Learning](#) Dilip Arumugam, Benjamin Van Roy
- [Giga-scale Kernel Matrix-Vector Multiplication on GPU](#) Robert Hu, Siu Lun Chau, Dino Sejdicinovic, Joan Glaunès
- [MultiScan: Scalable RGBD scanning for 3D environments with articulated objects](#) Yongsan Mao, Yiming Zhang, Hanxiao Jiang, Angel Chang, Manolis Savva
- [Are GANs overkill for NLP?](#) David Alvarez-Melis, Vikas Garg, Adam Kalai
- [“Why Not Other Classes?”: Towards Class-Contrastive Back-Propagation Explanations](#) Yipei Wang, Xiaoqian Wang
- [Neural Stochastic Control](#) Jingdong Zhang, Qunxi Zhu, Wei LIN
- [Constrained Update Projection Approach to Safe Policy Optimization](#) Long Yang, Jiaming Ji, Juntao Dai, Linrui Zhang, Binbin Zhou, Pengfei Li, Yaodong Yang, Gang Pan
- [DetCLIP: Dictionary-Enriched Visual-Concept Paralleled Pre-training for Open-world Detection](#) Lewei Yao, Jianhua Han, Youpeng Wen, Xiaodan Liang, Dan Xu, Wei Zhang, Zhenguo Li, Chunjing XU, Hang Xu
- [The Sample Complexity of One-Hidden-Layer Neural Networks](#) Gal Vardi, Ohad Shamir, Nati Srebro
- [Learning from Few Samples: Transformation-Invariant SVMs with Composition and Locality at Multiple Scales](#) Tao Liu, P. R. Kumar, Ruida Zhou, Xi Liu
- [Learning Efficient Vision Transformers via Fine-Grained Manifold Distillation](#) Zhiwei Hao, Jianyuan Guo, Ding Jia, Kai Han, Yehui Tang, Chao Zhang, Han Hu, Yunhe Wang
- [Equivariant Graph Hierarchy-Based Neural Networks](#) Jiaqi Han, Wenbing Huang, Tingyang Xu, Yu Rong
- [Learning interacting dynamical systems with latent Gaussian process ODEs](#) Çağatay Yıldız, Melih Kandemir, Barbara Rakitsch
- [OLIVES Dataset: Ophthalmic Labels for Investigating Visual Eye Semantics](#) Mohit Prabhushankar, Kiran Kokilepersaud, Yash-ye Logan, Stephanie Trejo Corona, Ghassan AlRegib, Charles Wykoff
- [Off-Policy Evaluation for Action-Dependent Non-stationary Environments](#) Yash Chandak, Shiv Shankar, Nathaniel Bastian, Bruno da Silva, Emma Brunskill, Philip S. Thomas
- [Fast Mixing of Stochastic Gradient Descent with Normalization and Weight Decay](#) Zhiyuan Li, Tianhao Wang, Dingli Yu
- [OPEN: Orthogonal Propagation with Ego-Network Modeling](#) Liang Yang, Lina Kang, Qiuliang Zhang, Mengzhe Li, bingxin niu, Dongxiao He, Zhen Wang, Chuan Wang, Xiaochun Cao, Yuanfang Guo

- [Beyond the Best: Distribution Functional Estimation in Infinite-Armed Bandits](#) Yifei Wang, Tavor Baharav, Yanjun Han, Jiantao Jiao, David Tse
- [Adversarial training for high-stakes reliability](#) Daniel Ziegler, Seraphina Nix, Lawrence Chan, Tim Bauman, Peter Schmidt-Nielsen, Tao Lin, Adam Scherlis, Noa Nabeshima, Benjamin Weinstein-Raun, Daniel de Haas, Buck Shlegeris, Nate Thomas
- [ELEVATER: A Benchmark and Toolkit for Evaluating Language-Augmented Visual Models](#) Chunyuan Li, Haotian Liu, Liunian Li, Pengchuan Zhang, Jyoti Aneja, Jianwei Yang, Ping Jin, Houdong Hu, Zicheng Liu, Yong Jae Lee, Jianfeng Gao
- [Augmented RBMLE-UCB Approach for Adaptive Control of Linear Quadratic Systems](#) Akshay Mete, Rahul Singh, P. R. Kumar
- [Stochastic Window Transformer for Image Restoration](#) Jie Xiao, Xueyang Fu, Feng Wu, Zheng-Jun Zha
- [Phase Transition from Clean Training to Adversarial Training](#) Yue Xing, Qifan Song, Guang Cheng
- [pFL-Bench: A Comprehensive Benchmark for Personalized Federated Learning](#) Daoyuan Chen, Dawei Gao, Weirui Kuang, Yaliang Li, Bolin Ding
- [Squeezeformer: An Efficient Transformer for Automatic Speech Recognition](#) Sehoon Kim, Amir Gholami, Albert Shaw, Nicholas Lee, Karttikeya Mangalam, Jitendra Malik, Michael W. Mahoney, Kurt Keutzer
- [Deep Generalized Schrödinger Bridge](#) Guan-Horng Liu, Tianrong Chen, Oswin So, Evangelos Theodorou
- [Characterizing the Ventral Visual Stream with Response-Optimized Neural Encoding Models](#) Meenakshi Khosla, Keith Jamison, Amy Kuceyeski, Mert Sabuncu
- [Learning sparse features can lead to overfitting in neural networks](#) Leonardo Petrini, Francesco Cagnetta, Eric Vanden-Eijnden, Matthieu Wyart
- [Public Wisdom Matters! Discourse-Aware Hyperbolic Fourier Co-Attention for Social Text Classification](#) Karish Grover, S M Phaneendra Angara, Md Shad Akhtar, Tanmoy Chakraborty
- [Harmonizing the object recognition strategies of deep neural networks with humans](#) Thomas FEL, Ivan F Rodriguez Rodriguez, Drew Linsley, Thomas Serre
- [Learning Mixed Multinomial Logits with Provable Guarantees](#) Yiqun Hu, David Simchi-Levi, Zhenzhen Yan
- [Defining and Characterizing Reward Gaming](#) Joar Skalse, Nikolaus Howe, Dmitrii Krasheninnikov, David Krueger
- [Learning Distinct and Representative Modes for Image Captioning](#) Qi Chen, Chaorui Deng, Qi Wu
- [Lifting the Information Ratio: An Information-Theoretic Analysis of Thompson Sampling for Contextual Bandits](#) Gergely Neu, Iuliia Olkhovskaia, Matteo Papini, Ludovic Schwartz
- [Is this the Right Neighborhood? Accurate and Query Efficient Model Agnostic Explanations](#) Amit Dhurandhar, Karthikeyan Natesan Ramamurthy, Karthikeyan Shanmugam
- [Object Scene Representation Transformer](#) Mehdi S. M. Sajjadi, Daniel Duckworth, Aravindh Mahendran, Sjoerd van Steenkiste, Filip Pavetic, Mario Lucic, Leonidas J. Guibas, Klaus Greff, Thomas Kipf
- [Data-Driven Conditional Robust Optimization](#) Abhilash Reddy Chenreddy, Nymisha Bandi, Erick Delage
- [COLD Decoding: Energy-based Constrained Text Generation with Langevin Dynamics](#) Lianhui Qin, Sean Welleck, Daniel Khashabi, Yejin Choi
- [Matching in Multi-arm Bandit with Collision](#) YiRui Zhang, Siwei Wang, Zhixuan Fang
- [Multimodal Contrastive Learning with LIMoE: the Language-Image Mixture of Experts](#) Basil Mustafa, Carlos Riquelme, Joan Puigcerver, Rodolphe Jenatton, Neil Houlsby
- [Few-shot Learning for Feature Selection with Hilbert-Schmidt Independence Criterion](#) Atsutoshi Kumagai, Tomoharu Iwata, Yasutoshi Ida, Yasuhiro Fujiwara
- [Statistical Learning and Inverse Problems: A Stochastic Gradient Approach](#) Yuri Fonseca, Yuri Saporito
- [Hyperparameter Sensitivity in Deep Outlier Detection: Analysis and a Scalable Hyper-Ensemble Solution](#) Xueying Ding, Lingxiao Zhao, Leman Akoglu
- [TVLT: Textless Vision-Language Transformer](#) Zineng Tang, Jaemin Cho, Yixin Nie, Mohit Bansal
- [A Deep Reinforcement Learning Framework for Column Generation](#) Cheng Chi, Amine Aboussalah, Elias Khalil, Juyoung Wang, Zoha Sherkat-Masoumi
- [Learning to Drop Out: An Adversarial Approach to Training Sequence VAEs](#) Djordje Miladinovic, Kumar Shridhar, Kushal Jain, Max Paulus, Joachim M Buhmann, Carl Allen
- [On the Adversarial Robustness of Mixture of Experts](#) Joan Puigcerver, Rodolphe Jenatton, Carlos Riquelme, Pranjal Awasthi, Srinadh Bhojanapalli
- [Safety Guarantees for Neural Network Dynamic Systems via Stochastic Barrier Functions](#) Rayan Mazouz, Karan Muvvala, Akash Ratheesh Babu, Luca Laurenti, Morteza Lahijanian
- [Simple and Optimal Greedy Online Contention Resolution Schemes](#) Vasilis Livanos
- [Composition Theorems for Interactive Differential Privacy](#) Xin Lyu
- [On Margins and Generalisation for Voting Classifiers](#) Felix Biggs, Valentina Zantedeschi, Benjamin Guedj
- [Effective Backdoor Defense by Exploiting Sensitivity of Poisoned Samples](#) Weixin Chen, Baoyuan Wu, Haoqian Wang
- [Rethinking the Reverse-engineering of Trojan Triggers](#) Zhenting Wang, Kai Mei, Hailun Ding, Juan Zhai, Shiqing Ma
- [Antigen-Specific Antibody Design and Optimization with Diffusion-Based Generative Models for Protein Structures](#) Shitong Luo, Yufeng Su, Xingang Peng, Sheng Wang, Jian Peng, Jianzhu Ma

- [Information bottleneck theory of high-dimensional regression: relevancy, efficiency and optimality](#) Vudtiwat Ngampruetikorn, David J. Schwab
- [RainNet: A Large-Scale Imagery Dataset and Benchmark for Spatial Precipitation Downscaling](#) Xuanhong Chen, Kairui Feng, Naiyuan Liu, Bingbing Ni, Yifan Lu, Zhengyan Tong, Ziang Liu
- [Dataset Distillation using Neural Feature Regression](#) Yongchao Zhou, Ehsan Nezhadarya, Jimmy Ba
- [ZeroC: A Neuro-Symbolic Model for Zero-shot Concept Recognition and Acquisition at Inference Time](#) Tailin Wu, Megan Tjandrasuwita, Zhengxuan Wu, Xuelin Yang, Kevin Liu, Rok Sasic, Jure Leskovec
- [Theoretically Better and Numerically Faster Distributed Optimization with Smoothness-Aware Quantization Techniques](#) Bokun Wang, Mher Safaryan, Peter Richtarik
- [Sparse Structure Search for Delta Tuning](#) Shengding Hu, Zhen Zhang, Ning Ding, Yadao Wang, Yasheng Wang, Zhiyuan Liu, Maosong Sun
- [On the Safety of Interpretable Machine Learning: A Maximum Deviation Approach](#) Dennis Wei, Rahul Nair, Amit Dhurandhar, Kush R. Varshney, Elizabeth Daly, Moninder Singh
- [Non-stationary Transformers: Exploring the Stationarity in Time Series Forecasting](#) Yong Liu, Haixu Wu, Jianmin Wang, Mingsheng Long
- [Risk-Driven Design of Perception Systems](#) Anthony Corso, Sydney Katz, Craig Innes, Xin Du, Subramanian Ramamoorthy, Mykel J Kochenderfer
- [A Simple Approach to Automated Spectral Clustering](#) Jicong Fan, Yiheng Tu, Zhao Zhang, Mingbo Zhao, Haijun Zhang
- [Joint Entropy Search for Multi-Objective Bayesian Optimization](#) Ben Tu, Axel Gandy, Nikolas Kantas, Behrang Shafei
- [Subgroup Robustness Grows On Trees: An Empirical Baseline Investigation](#) Josh Gardner, Zoran Popovic, Ludwig Schmidt
- [Robustness to Unbounded Smoothness of Generalized SignSGD](#) Michael Crawshaw, Mingrui Liu, Francesco Orabona, Wei Zhang, Zhenxun Zhuang
- [GhostNetV2: Enhance Cheap Operation with Long-Range Attention](#) Yehui Tang, Kai Han, Jianyuan Guo, Chang Xu, Chao Xu, Yunhe Wang
- [Analyzing Sharpness along GD Trajectory: Progressive Sharpening and Edge of Stability](#) Zixuan Wang, Zhouzi Li, Jian Li
- [You Never Stop Dancing: Non-freezing Dance Generation via Bank-constrained Manifold Projection](#) Jiangxin Sun, Chunyu Wang, Huang Hu, Hanjiang Lai, Zhi Jin, Jian-Fang Hu
- [Nonnegative Tensor Completion via Integer Optimization](#) Caleb Bugg, Chen Chen, Anil Aswani
- [LION: Latent Point Diffusion Models for 3D Shape Generation](#) xiaohui zeng, Arash Vahdat, Francis Williams, Zan Gojcic, Or Litany, Sanja Fidler, Karsten Kreis
- [Distribution-Informed Neural Networks for Domain Adaptation Regression](#) Jun Wu, Jingrui He, Sheng Wang, Kaiyu Guan, Elizabeth Ainsworth
- [Momentum Adversarial Distillation: Handling Large Distribution Shifts in Data-Free Knowledge Distillation](#) Kien Do, Thai Hung Le, Dung Nguyen, Dang Nguyen, HARIPRIYA HARIKUMAR, Truyen Tran, Santu Rana, Svetha Venkatesh
- [Hard ImageNet: Segmentations for Objects with Strong Spurious Cues](#) Mazda Moayeri, Sahil Singla, Soheil Feizi
- [VideoMAE: Masked Autoencoders are Data-Efficient Learners for Self-Supervised Video Pre-Training](#) Zhan Tong, Yibing Song, Jue Wang, Limin Wang
- [Scalable design of Error-Correcting Output Codes using Discrete Optimization with Graph Coloring](#) Samarth Gupta, Saurabh Amin
- [A New Family of Generalization Bounds Using Samplewise Evaluated CMI](#) Fredrik Hellström, Giuseppe Durisi
- [Maximum-Likelihood Inverse Reinforcement Learning with Finite-Time Guarantees](#) Siliang Zeng, Chenliang Li, Alfredo Garcia, Mingyi Hong
- [Lost in Latent Space: Examining failures of disentangled models at combinatorial generalisation](#) Milton Montero, Jeffrey Bowers, Rui Ponte Costa, Casimir Ludwig, Gaurav Malhotra
- [MultiGuard: Provably Robust Multi-label Classification against Adversarial Examples](#) Jinyuan Jia, Wenjie Qu, Neil Gong
- [On Measuring Excess Capacity in Neural Networks](#) Florian Graf, Sebastian Zeng, Bastian Rieck, Marc Niethammer, Roland Kwitt
- [On Leave-One-Out Conditional Mutual Information For Generalization](#) Mohamad Rida Rammal, Alessandro Achille, Aditya Golatkar, Suhas Diggavi, Stefano Soatto
- [Near-Isometric Properties of Kronecker-Structured Random Tensor Embeddings](#) Qijia Jiang
- [MExMI: Pool-based Active Model Extraction Crossover Membership Inference](#) Yaxin Xiao, Qingqing Ye, Haibo Hu, Huadi Zheng, Chengfang Fang, Jie Shi
- [Parameter-Efficient Masking Networks](#) Yue Bai, Huan Wang, Xu Ma, Yitian Zhang, Zhiqiang Tao, Yun Fu
- [How Sampling Impacts the Robustness of Stochastic Neural Networks](#) Sina Däubener, Asja Fischer
- [Toward Robust Spiking Neural Network Against Adversarial Perturbation](#) LING LIANG, Kaidi Xu, Xing Hu, Lei Deng, Yuan Xie
- [GRASP: Navigating Retrosynthetic Planning with Goal-driven Policy](#) Yemin Yu, Ying Wei, Kun Kuang, Zhengxing Huang, Huaxiu Yao, Fei Wu
- [End-to-end Symbolic Regression with Transformers](#) Pierre-alexandre Kamienny, Stéphane d'Ascoli, Guillaume Lample, Francois Charton
- [Gold-standard solutions to the Schrödinger equation using deep learning: How much physics do we need?](#) Leon Gerard, Michael Scherbela, Philipp Marquetand, Philipp Grohs
- [EcoFormer: Energy-Saving Attention with Linear Complexity](#) Jing Liu, Zizheng Pan, Haoyu He, Jianfei Cai, Bohan Zhuang
- [Wild-Time: A Benchmark of in-the-Wild Distribution Shift over Time](#) Huaxiu Yao, Caroline Choi, Bochuan Cao, Yoonho Lee, Pang Wei W. Koh, Chelsea Finn

- [Sound and Complete Causal Identification with Latent Variables Given Local Background Knowledge](#) Tian-Zuo Wang, Tian Qin, Zhi-Hua Zhou
- [A Unified Diversity Measure for Multiagent Reinforcement Learning](#) Zongkai Liu, Chao Yu, Yaodong Yang, peng sun, Zifan Wu, Yuan Li
- [HorNet: Efficient High-Order Spatial Interactions with Recursive Gated Convolutions](#) Yongming Rao, Wenliang Zhao, Yansong Tang, Jie Zhou, Ser Nam Lim, Jiwen Lu
- [Learning Tractable Probabilistic Models from Inconsistent Local Estimates](#) Shasha Jin, Vasundhara Komaragiri, Tahrima Rahman, Vibhav Gogate
- [Decomposed Knowledge Distillation for Class-Incremental Semantic Segmentation](#) Donghyeon Baek, Youngmin Oh, Sanghoon Lee, Junghyup Lee, Bumsub Ham
- [Minimax Optimal Algorithms for Fixed-Budget Best Arm Identification](#) Junpei Komiyama, Taira Tsuchiya, Junya Honda
- [Reduced Representation of Deformation Fields for Effective Non-rigid Shape Matching](#) Ramana Subramanyam Sundararaman, Riccardo Marin, Emanuele Rodolà, Maks Ovsjanikov
- [BEVFusion: A Simple and Robust LiDAR-Camera Fusion Framework](#) Tingting Liang, Hongwei Xie, Kaicheng Yu, Zhongyu Xia, Zhiwei Lin, Yongtao Wang, Tao Tang, Bing Wang, Zhi Tang
- [Automatic Differentiation of Programs with Discrete Randomness](#) Gaurav Arya, Moritz Schauer, Frank Schäfer, Christopher Rackauckas
- [A Closer Look at the Adversarial Robustness of Deep Equilibrium Models](#) Zonghan Yang, Tianyu Pang, Yang Liu
- [Near-Optimal Private and Scalable \$k\$ -Clustering](#) Vincent Cohen-Addad, Alessandro Epasto, Vahab Mirrokni, Shyam Narayanan, Peilin Zhong
- [NS3: Neuro-symbolic Semantic Code Search](#) Shushan Arakelyan, Anna Hakhverdyan, Miltiadis Allamanis, Luis Garcia, Christophe Hauser, Xiang Ren
- [Revisiting Sparse Convolutional Model for Visual Recognition](#) xili dai, Mingyang Li, Pengyuan Zhai, Shengbang Tong, Xingjian Gao, Shao-Lun Huang, Zhihui Zhu, Chong You, Yi Ma
- [Temporal Latent Bottleneck: Synthesis of Fast and Slow Processing Mechanisms in Sequence Learning](#) Aniket Didolkar, Kshitij Gupta, Anirudh Goyal, Nitesh Bharadwaj Gundavarapu, Alex M. Lamb, Nan Rosemary Ke, Yoshua Bengio
- [A Near-Optimal Primal-Dual Method for Off-Policy Learning in CMDP](#) Fan Chen, Junyu Zhang, Zaiwen Wen
- [DReS-FL: Dropout-Resilient Secure Federated Learning for Non-IID Clients via Secret Data Sharing](#) Jiawei Shao, Yuchang Sun, Songze Li, Jun Zhang
- [BackdoorBench: A Comprehensive Benchmark of Backdoor Learning](#) Baoyuan Wu, Hongrui Chen, Mingda Zhang, Zihao Zhu, Shaokui Wei, Danni Yuan, Chao Shen
- [REVIVE: Regional Visual Representation Matters in Knowledge-Based Visual Question Answering](#) Yuanze Lin, Yujia Xie, Dongdong Chen, Yichong Xu, Chenguang Zhu, Lu Yuan
- [FedAvg with Fine Tuning: Local Updates Lead to Representation Learning](#) Liam Collins, Hamed Hassani, Aryan Mokhtari, Sanjay Shakkottai
- [Modeling Transitivity and Cyclicity in Directed Graphs via Binary Code Box Embeddings](#) Dongxu Zhang, Michael Boratko, Cameron Musco, Andrew McCallum
- [Generalization Bounds for Gradient Methods via Discrete and Continuous Prior](#) Xuanyuan Luo, Bei Luo, Jian Li
- [On Deep Generative Models for Approximation and Estimation of Distributions on Manifolds](#) Biraj Dahal, Alexander Havrilla, Minshuo Chen, Tuo Zhao, Wenjing Liao
- [Memory Efficient Continual Learning with Transformers](#) Beyza Ermis, Giovanni Zappella, Martin Wistuba, Aditya Rawal, Cedric Archambeau
- [Counterfactual Neural Temporal Point Process for Estimating Causal Influence of Misinformation on Social Media](#) Yizhou Zhang, Defu Cao, Yan Liu
- [Curriculum Reinforcement Learning using Optimal Transport via Gradual Domain Adaptation](#) Peide Huang, Mengdi Xu, Jiacheng Zhu, Laixi Shi, Fei Fang, DING ZHAO
- [Temporally-Consistent Survival Analysis](#) Lucas Maystre, Daniel Russo
- [Symbolic Distillation for Learned TCP Congestion Control](#) S P Sharan, Wenqing Zheng, Kuo-Feng Hsu, Jiarong Xing, Ang Chen, Zhangyang Wang
- [A Probabilistic Graph Coupling View of Dimension Reduction](#) Hugues Van Assel, Thibault Espinasse, Julien Chiquet, Franck Picard
- [Hardness of Noise-Free Learning for Two-Hidden-Layer Neural Networks](#) Sitan Chen, Aravind Gollakota, Adam Klivans, Raghu Meka
- [Convergence beyond the over-parameterized regime using Rayleigh quotients](#) David A. R. Robin, Kevin Scaman, marc lelarge
- [Optimistic Posterior Sampling for Reinforcement Learning with Few Samples and Tight Guarantees](#) Daniil Tiapkin, Denis Belomestny, Daniele Calandriello, Eric Moulines, Remi Munos, Alexey Naumov, Mark Rowland, Michal Valko, Pierre Ménard
- [On Convergence of FedProx: Local Dissimilarity Invariant Bounds, Non-smoothness and Beyond](#) Xiaotong Yuan, Ping Li
- [Can Push-forward Generative Models Fit Multimodal Distributions?](#) Antoine Salmon, Valentin De Bortoli, Julie Delon, Agnes Desolneux
- [Improved Imaging by Invex Regularizers with Global Optima Guarantees](#) Samuel Pinilla, Tingting Mu, Neil Bourne, Jeyan Thiyagalingam
- [The Neural Covariance SDE: Shaped Infinite Depth-and-Width Networks at Initialization](#) Mufan Li, Mihai Nica, Dan Roy
- [Rethinking and Scaling Up Graph Contrastive Learning: An Extremely Efficient Approach with Group Discrimination](#) YIZHEN ZHENG, Shirui Pan, Vincent CS Lee, Yu Zheng, Philip S Yu

- [Diverse Weight Averaging for Out-of-Distribution Generalization](#) Alexandre Rame, Matthieu Kirchmeyer, Thibaud Rahier, Alain Rakotomamonjy, Patrick Gallinari, Matthieu Cord
- [Bivariate Causal Discovery for Categorical Data via Classification with Optimal Label Permutation](#) Yang Ni
- [Phase transitions in when feedback is useful](#) Lokesh Boominathan, Xaq Pitkow
- [Stochastic Second-Order Methods Improve Best-Known Sample Complexity of SGD for Gradient-Dominated Functions](#) Saeed Masiha, Saber Salehkaleybar, Niao He, Negar Kiyavash, Patrick Thiran
- [Posterior and Computational Uncertainty in Gaussian Processes](#) Jonathan Wenger, Geoff Pleiss, Marvin Pförtner, Philipp Hennig, John P. Cunningham
- [Understanding the Evolution of Linear Regions in Deep Reinforcement Learning](#) Setareh Cohan, Nam Hee Kim, David Rolnick, Michiel van de Panne
- [Causal Discovery in Heterogeneous Environments Under the Sparse Mechanism Shift Hypothesis](#) Ronan Perry, Julius von Kügelgen, Bernhard Schölkopf
- [Towards Practical Control of Singular Values of Convolutional Layers](#) Alexandra Senderovich, Ekaterina Bulatova, Anton Obukhov, Maxim Rakhuba
- [Deep Multi-Modal Structural Equations For Causal Effect Estimation With Unstructured Proxies](#) Shachi Deshpande, Kaiwen Wang, Dhruv Sreenivas, Zheng Li, Volodymyr Kuleshov
- [Generic bounds on the approximation error for physics-informed \(and\) operator learning](#) Tim De Ryck, Siddhartha Mishra
- [Learning Distributions Generated by Single-Layer ReLU Networks in the Presence of Arbitrary Outliers](#) Saikiran Bulusu, Geethu Joseph, M. Cenk Gursoy, Pramod Varshney
- [GenerSpeech: Towards Style Transfer for Generalizable Out-Of-Domain Text-to-Speech](#) Rongjie Huang, Yi Ren, Jinglin Liu, Chenye Cui, Zhou Zhao
- [On Non-Linear operators for Geometric Deep Learning](#) Grégoire Sergeant-Perthuis, Jakob Maier, Joan Bruna, Edouard Oyallon
- [Momentum Aggregation for Private Non-convex ERM](#) Hoang Tran, Ashok Cutkosky
- [Learning in Congestion Games with Bandit Feedback](#) Qiwen Cui, Zhihan Xiong, Maryam Fazel, Simon S. Du
- [LOG: Active Model Adaptation for Label-Efficient OOD Generalization](#) Jie-Jing Shao, Lan-Zhe Guo, Xiao-wen Yang, Yu-Feng Li
- [PolarMix: A General Data Augmentation Technique for LiDAR Point Clouds](#) Aoran Xiao, Jiaying Huang, Dayan Guan, Kaiwen Cui, Shijian Lu, Ling Shao
- [Learning from Distributed Users in Contextual Linear Bandits Without Sharing the Context](#) Osama Hanna, Lin Yang, Christina Fragouli
- [Provably Feedback-Efficient Reinforcement Learning via Active Reward Learning](#) Dingwen Kong, Lin Yang
- [Recurrent Memory Transformer](#) Aydar Bulatov, Yury Kuratov, Mikhail Burtsev
- [Hierarchical Lattice Layer for Partially Monotone Neural Networks](#) Hiroki Yanagisawa, Kohei Miyaguchi, Takayuki Katsuki
- [Class-Dependent Label-Noise Learning with Cycle-Consistency Regularization](#) De Cheng, Yixiong Ning, Nannan Wang, Xinbo Gao, Heng Yang, Yuxuan Du, Bo Han, Tongliang Liu
- [SeqPATE: Differentially Private Text Generation via Knowledge Distillation](#) Zhiliang Tian, Yingxiu Zhao, Ziyue Huang, Yu-Xiang Wang, Nevin L. Zhang, He He
- [Bayesian Persuasion for Algorithmic Recourse](#) Keegan Harris, Valerie Chen, Joon Kim, Ameet Talwalkar, Hoda Heidari, Steven Z. Wu
- [Instance-Based Uncertainty Estimation for Gradient-Boosted Regression Trees](#) Jonathan Brophy, Daniel Lowd
- [Transcormer: Transformer for Sentence Scoring with Sliding Language Modeling](#) Kaitao Song, Yichong Leng, Xu Tan, Yicheng Zou, Tao Qin, Dongsheng Li
- [Models Out of Line: A Fourier Lens on Distribution Shift Robustness](#) Sara Fridovich-Keil, Brian Bartoldson, James Diffenderfer, Bhavya Kailkhura, Timo Bremer
- [Deep Learning Methods for Proximal Inference via Maximum Moment Restriction](#) Benjamin Kompa, David Bellamy, Tom Kolokotronis, James M. Robins, Andrew Beam
- [Nest Your Adaptive Algorithm for Parameter-Agnostic Nonconvex Minimax Optimization](#) Junchi Yang, Xiang Li, Niao He
- [Brownian Noise Reduction: Maximizing Privacy Subject to Accuracy Constraints](#) Justin Whitehouse, Aaditya Ramdas, Steven Z. Wu, Ryan M. Rogers
- [Are AlphaZero-like Agents Robust to Adversarial Perturbations?](#) Li-Cheng Lan, Huan Zhang, Ti-Rong Wu, Meng-Yu Tsai, I-Chen Wu, Chao-Jui Hsieh
- [AnimeSR: Learning Real-World Super-Resolution Models for Animation Videos](#) Yanze Wu, Xintao Wang, GEN LI, Ying Shan
- [On the Spectral Bias of Convolutional Neural Tangent and Gaussian Process Kernels](#) Amnon Geifman, Meirav Galun, David Jacobs, Basri Ronen
- [Fairness Transferability Subject to Bounded Distribution Shift](#) Yatong Chen, Reilly Raab, Jialu Wang, Yang Liu
- [Improving Self-Supervised Learning by Characterizing Idealized Representations](#) Yann Dubois, Stefano Ermon, Tatsunori B. Hashimoto, Percy S. Liang
- [On the difficulty of learning chaotic dynamics with RNNs](#) Jonas Mikhail, Zahra Monfared, Daniel Durstewitz
- [SKFlow: Learning Optical Flow with Super Kernels](#) SHANGKUN SUN, Yuanqi Chen, Yu Zhu, Guodong Guo, Ge Li
- [Mingling Foresight with Imagination: Model-Based Cooperative Multi-Agent Reinforcement Learning](#) Zhiwei Xu, Dapeng Li, Bin Zhang, Yuan Zhan, Yunpeng Bai, Guoliang Fan

- [End-to-end Stochastic Optimization with Energy-based Model](#) Lingkai Kong, Jiaming Cui, Yuchen Zhuang, Rui Feng, B. Aditya Prakash, Chao Zhang
- [Learning low-dimensional generalizable natural features from retina using a U-net](#) Siwei Wang, Benjamin Hoshal, Elizabeth de Laittre, Thierry Mora, Michael Berry, Stephanie Palmer
- [Split-kl and PAC-Bayes-split-kl Inequalities for Ternary Random Variables](#) Yi-Shan Wu, Yevgeny Seldin
- [Wasserstein \$K\$ -means for clustering probability distributions](#) Yubo Zhuang, Xiaohui Chen, Yun Yang
- [Learning Long-Term Crop Management Strategies with CyclesGym](#) Matteo Turchetta, Luca Corinzia, Scott Sussex, Amanda Burton, Juan Herrera, Ioannis Athanasiadis, Joachim M Buhmann, Andreas Krause
- [Identification, Amplification and Measurement: A bridge to Gaussian Differential Privacy](#) Yi Liu, Ke Sun, Bei Jiang, Linglong Kong
- [MACE: Higher Order Equivariant Message Passing Neural Networks for Fast and Accurate Force Fields](#) Ilyes Batatia, David P Kovacs, Gregor Simm, Christoph Ortner, Gabor Csanyi
- [Follow-the-Perturbed-Leader for Adversarial Markov Decision Processes with Bandit Feedback](#) Yan Dai, Haipeng Luo, Liyu Chen
- [Improving Multi-Task Generalization via Regularizing Spurious Correlation](#) Ziniu Hu, Zhe Zhao, Xinyang Yi, Tiansheng Yao, Lichan Hong, Yizhou Sun, Ed Chi
- [Estimating and Explaining Model Performance When Both Covariates and Labels Shift](#) Lingjiao Chen, Matei Zaharia, James Y. Zou
- [Tenrec: A Large-scale Multipurpose Benchmark Dataset for Recommender Systems](#) Guanghu Yuan, Fajie Yuan, Yudong Li, Beibei Kong, Shujie Li, Lei Chen, Min Yang, Chenyun YU, Bo Hu, Zang Li, Yu Xu, Xiaohu Qie
- [Joint Entropy Search For Maximally-Informed Bayesian Optimization](#) Carl Hvarfner, Frank Hutter, Luigi Nardi
- [Provable General Function Class Representation Learning in Multitask Bandits and MDP](#) Rui Lu, Andrew Zhao, Simon S. Du, Gao Huang
- [Weighted Mutual Learning with Diversity-Driven Model Compression](#) Miao Zhang, Li Wang, David Campos, Wei Huang, Chenjuan Guo, Bin Yang
- [S2P: State-conditioned Image Synthesis for Data Augmentation in Offline Reinforcement Learning](#) Daesol Cho, Dongseok Shim, H. Jin Kim
- [Neural Collapse with Normalized Features: A Geometric Analysis over the Riemannian Manifold](#) Can Yaras, Peng Wang, Zhihui Zhu, Laura Balzano, Qing Qu
- [Conformalized Fairness via Quantile Regression](#) Meichen Liu, Lei Ding, Dengdeng Yu, Wulong Liu, Linglong Kong, Bei Jiang
- [Efficient Methods for Non-stationary Online Learning](#) Peng Zhao, Yan-Feng Xie, Lijun Zhang, Zhi-Hua Zhou
- [Increasing the Scope as You Learn: Adaptive Bayesian Optimization in Nested Subspaces](#) Leonard Papenmeier, Luigi Nardi, Matthias Poloczek
- [ACIL: Analytic Class-Incremental Learning with Absolute Memorization and Privacy Protection](#) HUIPING ZHUANG, Zhenyu Weng, Hongxin Wei, RENCHUNZI XIE, Kar-Ann Toh, Zhiping Lin
- [MsSVT: Mixed-scale Sparse Voxel Transformer for 3D Object Detection on Point Clouds](#) Shaocong Dong, Lihe Ding, Haiyang Wang, Tingfa Xu, Xinli Xu, Jie Wang, Ziyang Bian, Ying Wang, Jianan Li
- [Deterministic Langevin Monte Carlo with Normalizing Flows for Bayesian Inference](#) Richard Grumitt, Biwei Dai, Uros Seljak
- [Evolution of Neural Tangent Kernels under Benign and Adversarial Training](#) Noel Loo, Ramin Hasani, Alexander Amini, Daniela Rus
- [Zero-Sum Stochastic Stackelberg Games](#) Denizalp Goktas, Sadie Zhao, Amy Greenwald
- [Evaluating Out-of-Distribution Performance on Document Image Classifiers](#) Stefan Larson, Yi Yang Gordon Lim, Yutong Ai, David Kuang, Kevin Leach
- [Distributed Learning of Conditional Quantiles in the Reproducing Kernel Hilbert Space](#) Heng Lian
- [Spatial Mixture-of-Experts](#) Nikoli Dryden, Torsten Hoeffler
- [Amortized Mixing Coupling Processes for Clustering](#) Huafeng Liu, Liping Jing
- [Hilbert Distillation for Cross-Dimensionality Networks](#) Dian Qin, Haishuai Wang, Zhe Liu, HONGJIA XU, Sheng Zhou, Jiajun Bu
- [Provably Efficient Offline Multi-agent Reinforcement Learning via Strategy-wise Bonus](#) Qiwen Cui, Simon S. Du
- [A Best-of-Both-Worlds Algorithm for Bandits with Delayed Feedback](#) Saeed Masoudian, Julian Zimmert, Yevgeny Seldin
- [LIFT: Language-Interfaced Fine-Tuning for Non-language Machine Learning Tasks](#) Tuan Dinh, Yuchen Zeng, Ruisu Zhang, Ziqian Lin, Michael Gira, Shashank Rajput, Jy-yong Sohn, Dimitris Papailiopoulos, Kangwook Lee
- [Capturing Failures of Large Language Models via Human Cognitive Biases](#) Erik Jones, Jacob Steinhardt
- [Template based Graph Neural Network with Optimal Transport Distances](#) Cédric Vincent-Cuaz, Rémi Flamary, Marco Corneli, Titouan Vayer, Nicolas Courty
- [Knowledge Distillation Improves Graph Structure Augmentation for Graph Neural Networks](#) Lirong Wu, Haitao Lin, Yufei Huang, Stan Z. Li
- [Learning Invariant Graph Representations for Out-of-Distribution Generalization](#) Haoyang Li, Ziwei Zhang, Xin Wang, Wenwu Zhu
- [A Quadrature Rule combining Control Variates and Adaptive Importance Sampling](#) Rémi Leluc, François Portier, Johan Segers, Aigerim Zhuman
- [GAL: Gradient Assisted Learning for Decentralized Multi-Organization Collaborations](#) Enmao Diao, Jie Ding, Vahid Tarokh
- [Direct Advantage Estimation](#) Hsiao-Ru Pan, Nico Gürtler, Alexander Neitz, Bernhard Schölkopf
- [Honor of Kings Arena: an Environment for Generalization in Competitive Reinforcement Learning](#) Hua Wei, Jingxiao Chen, Xiyang Ji, Hongyang Qin, Minwen Deng, Siqin Li, Liang Wang, Weinan Zhang, Yong Yu, Liu Linc, Lanxiao Huang, Deheng Ye, Qiang Fu, Wei Yang

- [On the Symmetries of Deep Learning Models and their Internal Representations](#) Charles Godfrey, Davis Brown, Tegan Emerson, Henry Kvinge
- [TabNAS: Rejection Sampling for Neural Architecture Search on Tabular Datasets](#) Chengrun Yang, Gabriel Bender, Hanxiao Liu, Pieter-Jan Kindermans, Madeleine Udell, Yifeng Lu, Quoc V Le, Da Huang
- [Embed and Emulate: Learning to estimate parameters of dynamical systems with uncertainty quantification](#) Ruoxi Jiang, Rebecca Willett
- [A Differentiable Semantic Metric Approximation in Probabilistic Embedding for Cross-Modal Retrieval](#) Hao Li, Jingkuan Song, Lianli Gao, Pengpeng Zeng, Haonan Zhang, Gongfu Li
- [Friendly Noise against Adversarial Noise: A Powerful Defense against Data Poisoning Attack](#) Tian Yu Liu, Yu Yang, Baharan Mirzasoleiman
- [Dict-TTS: Learning to Pronounce with Prior Dictionary Knowledge for Text-to-Speech](#) Ziyue Jiang, Zhe Su, Zhou Zhao, Qian Yang, Yi Ren, Jinglin Liu, 振辉 叶
- [AZ-whiteness test: a test for signal uncorrelation on spatio-temporal graphs](#) Daniele Zambon, Cesare Alippi
- [ViSioNS: Visual Search in Natural Scenes Benchmark](#) Fermín Travi, Gonzalo Ruarte, Gaston Bujia, Juan Esteban Kamienkowski
- [Value Function Decomposition for Iterative Design of Reinforcement Learning Agents](#) James MacGlashan, Evan Archer, Alisa Devlic, Takuma Seno, Craig Sherstan, Peter Wurman, Peter Stone
- [HandMeThat: Human-Robot Communication in Physical and Social Environments](#) Yanming Wan, Jiayuan Mao, Josh Tenenbaum
- [Task-Agnostic Graph Explanations](#) Yaochen Xie, Sumeet Katariya, Xianfeng Tang, Edward Huang, Nikhil Rao, Karthik Subbian, Shuiwang Ji
- [Embrace the Gap: VAEs Perform Independent Mechanism Analysis](#) Patrik Reizinger, Luigi Gresele, Jack Brady, Julius von Kügelgen, Dominik Zietlow, Bernhard Schölkopf, Georg Martius, Wieland Brendel, Michel Besserve
- [Dataset Inference for Self-Supervised Models](#) Adam Dziedzic, Haonan Duan, Muhammad Ahmad Kaleem, Nikita Dhawan, Jonas Guan, Yannis Cattan, Franziska Boenisch, Nicolas Papernot
- [Statistically Meaningful Approximation: a Case Study on Approximating Turing Machines with Transformers](#) Colin Wei, Yining Chen, Tengyu Ma
- [Improved Feature Distillation via Projector Ensemble](#) Yudong Chen, Sen Wang, Jiajun Liu, Xuwei Xu, Frank de Hoog, Zi Huang
- [Does GNN Pretraining Help Molecular Representation?](#) Ruoxi Sun, Hanjun Dai, Adams Wei Yu
- [ULNeF: Untangled Layered Neural Fields for Mix-and-Match Virtual Try-On](#) Igor Santesteban, Miguel Otaduy, Nils Thuerey, Dan Casas
- [Introspective Learning: A Two-Stage approach for Inference in Neural Networks](#) Mohit Prabhushankar, Ghassan AlRegib
- [Bayesian Active Learning with Fully Bayesian Gaussian Processes](#) Christoffer Riis, Francisco Antunes, Frederik Hüttel, Carlos Lima Azevedo, Francisco Pereira
- [E-MAPP: Efficient Multi-Agent Reinforcement Learning with Parallel Program Guidance](#) Can Chang, Ni Mu, Jiajun Wu, Ling Pan, Huazhe Xu
- [In Defense of the Unitary Scalarization for Deep Multi-Task Learning](#) Vitaly Kurin, Alessandro De Palma, Ilya Kostrikov, Shimon Whiteson, Pawan K Mudigonda
- [Flexible Neural Image Compression via Code Editing](#) Chenjian Gao, Tongda Xu, Dailan He, Yan Wang, Hongwei Qin
- [Homomorphic Matrix Completion](#) Xiao-Yang Liu, Zechu (Steven) Li, Xiaodong Wang
- [Spending Thinking Time Wisely: Accelerating MCTS with Virtual Expansions](#) Weirui Ye, Pieter Abbeel, Yang Gao
- [Intrinsic dimensionality estimation using Normalizing Flows](#) Christian Horvat, Jean-Pascal Pfister
- [Moment Distributionally Robust Tree Structured Prediction](#) Yeshe Li, Danyal Saeed, Xinhua Zhang, Brian Ziebart, Kevin Gimpel
- [Minimax Optimal Fixed-Budget Best Arm Identification in Linear Bandits](#) Junwen Yang, Vincent Tan
- [Tempo: Accelerating Transformer-Based Model Training through Memory Footprint Reduction](#) Muralidhar Andoorveedu, Zhanda Zhu, Bojian Zheng, Gennady Pekhimenko
- [I2DFormer: Learning Image to Document Attention for Zero-Shot Image Classification](#) Muhammad Ferjad Naeem, Yongqin Xian, Luc V Gool, Federico Tombari
- [Benchmarking Heterogeneous Treatment Effect Models through the Lens of Interpretability](#) Jonathan Crabbé, Alicia Curth, Ioana Bica, Mihaela van der Schaar
- [AD-DROP: Attribution-Driven Dropout for Robust Language Model Fine-Tuning](#) Tao Yang, Jinghao Deng, Xiaojun Quan, Qifan Wang, Shaoliang Nie
- [Reinforced Genetic Algorithm for Structure-based Drug Design](#) Tianfan Fu, Wenhao Gao, Connor Coley, Jimeng Sun
- [A Variational Edge Partition Model for Supervised Graph Representation Learning](#) Yilin He, Chaojie Wang, Hao Zhang, Bo Chen, Mingyuan Zhou
- [Learning Optimal Flows for Non-Equilibrium Importance Sampling](#) Yu Cao, Eric Vanden-Eijnden
- [Decision Trees with Short Explainable Rules](#) Victor Feitosa Souza, Ferdinando Cicalese, Eduardo Laber, Marco Molinaro
- [NAS-Bench-360: Benchmarking Neural Architecture Search on Diverse Tasks](#) Renbo Tu, Nicholas Roberts, Misha Khodak, Junhong Shen, Frederic Sala, Ameet Talwalkar
- [Sparse Fourier Backpropagation in Cryo-EM Reconstruction](#) Dari Kimanius, Kiarash Jamali, Sjors Scheres
- [Modular Flows: Differential Molecular Generation](#) Yogesh Verma, Samuel Kaski, Markus Heinonen, Vikas Garg
- [Anonymous Bandits for Multi-User Systems](#) Hossein Esfandiari, Vahab Mirrokni, Jon Schneider
- [Unsupervised Object Detection Pretraining with Joint Object Priors Generation and Detector Learning](#) Yizhou Wang, Meilin Chen,

- [Invariance Learning in Deep Neural Networks with Differentiable Laplace Approximations](#) Alexander Immer, Tycho van der Ouderaa, Gunnar Rätsch, Vincent Fortuin, Mark van der Wilk
- [QueryPose: Sparse Multi-Person Pose Regression via Spatial-Aware Part-Level Query](#) Yabo Xiao, Kai Su, Xiaojuan Wang, Dongdong Yu, Lei Jin, Mingshu He, Zehuan Yuan
- [EAGER: Asking and Answering Questions for Automatic Reward Shaping in Language-guided RL](#) Thomas Carta, Pierre-Yves Oudeyer, Olivier Sigaud, Sylvain Lamprier
- [OpenFilter: A Framework to Democratize Research Access to Social Media AR Filters](#) Piera Riccio, Bill Psomas, Francesco Galati, Francisco Escolano, Thomas Hofmann, Nuria Oliver
- [Improving Policy Learning via Language Dynamics Distillation](#) Victor Zhong, Jesse Mu, Luke Zettlemoyer, Edward Grefenstette, Tim Rocktäschel
- [Learning Energy Networks with Generalized Fenchel-Young Losses](#) Mathieu Blondel, Felipe Llinares-Lopez, Robert Dadashi, Leonard Hussenot, Matthieu Geist
- [Interaction-Grounded Learning with Action-Inclusive Feedback](#) Tengyang Xie, Akanksha Saran, Dylan J Foster, Lekan Molu, Ida Momennejad, Nan Jiang, Paul Mineiro, John Langford
- [Making Look-Ahead Active Learning Strategies Feasible with Neural Tangent Kernels](#) Mohamad Amin Mohamadi, Wonho Bae, Danica J. Sutherland
- [The Neural Testbed: Evaluating Joint Predictions](#) Ian Osband, Zheng Wen, Seyed Mohammad Asghari, Vikranth Dwaracherla, Xiuyuan Lu, MORTEZA IBRAHIMI, Dieterich Lawson, Botao Hao, Brendan O'Donoghue, Benjamin Van Roy
- [Learning Dense Object Descriptors from Multiple Views for Low-shot Category Generalization](#) Stefan Stojanov, Anh Thai, Zixuan Huang, James M. Rehg
- [Monte Carlo Tree Descent for Black-Box Optimization](#) Yaoguang Zhai, Sicun Gao
- [Teacher Forcing Recovers Reward Functions for Text Generation](#) Yongchang Hao, Yuxin Liu, Lili Mou
- [Masked Autoencoding for Scalable and Generalizable Decision Making](#) Fangchen Liu, Hao Liu, Aditya Grover, Pieter Abbeel
- [Distributional Reward Estimation for Effective Multi-agent Deep Reinforcement Learning](#) Jifeng Hu, Yanchao Sun, Hechang Chen, Sili Huang, haiyin piao, Yi Chang, Lichao Sun
- [Sampling from Log-Concave Distributions with Infinity-Distance Guarantees](#) Oren Mangoubi, Nisheeth Vishnoi
- [ELASTIC: Numerical Reasoning with Adaptive Symbolic Compiler](#) Jiaxin Zhang, Yashar Moshfeghi
- [Training Spiking Neural Networks with Local Tandem Learning](#) Qu Yang, Jibin Wu, Malu Zhang, Yansong Chua, Xinchao Wang, Haizhou Li
- [FiLM: Frequency improved Legendre Memory Model for Long-term Time Series Forecasting](#) Tian Zhou, Ziqing MA, xue wang, Qingsong Wen, Liang Sun, Tao Yao, Wotao Yin, Rong Jin
- [Differentially Private Linear Sketches: Efficient Implementations and Applications](#) Fuheng Zhao, Dan Qiao, Rachel Redberg, Divyakant Agrawal, Amr El Abbadi, Yu-Xiang Wang
- [Coordinates Are NOT Lonely - Codebook Prior Helps Implicit Neural 3D representations](#) Fukun Yin, Wen Liu, Zilong Huang, Pei Cheng, Tao Chen, Gang Yu
- [Scalable Neural Video Representations with Learnable Positional Features](#) Subin Kim, Sihyun Yu, Jaeho Lee, Jinwoo Shin
- [Data Augmentation MCMC for Bayesian Inference from Privatized Data](#) Nianqiao Ju, Jordan Awan, Ruobin Gong, Vinayak Rao
- [Neural Circuit Architectural Priors for Embodied Control](#) Nikhil Bhattasali, Anthony M Zador, Tatiana Engel
- [Bayesian Optimization over Discrete and Mixed Spaces via Probabilistic Reparameterization](#) Samuel Daulton, Xingchen Wan, David Eriksson, Maximilian Balandat, Michael A Osborne, Eytan Bakshy
- [Learning to Constrain Policy Optimization with Virtual Trust Region](#) Thai Hung Le, Thommen Karimpanal George, Majid Abdolshah, Dung Nguyen, Kien Do, Sunil Gupta, Svetha Venkatesh
- [Verification and search algorithms for causal DAGs](#) Davin Choo, Kirankumar Shiragur, Arnab Bhattacharyya
- [Causally motivated multi-shortcut identification and removal](#) Jiayun Zheng, Maggie Makar
- [Avalon: A Benchmark for RL Generalization Using Procedurally Generated Worlds](#) Joshua Albrecht, Abraham Fetterman, Bryden Fogelman, Ellie Kitanidis, Bartosz Wróblewski, Nicole Seo, Michael Rosenthal, Maksis Knutins, Zack Polizzi, James Simon, Kanjun Qiu
- [Learning Equivariant Segmentation with Instance-Unique Querying](#) Wenguan Wang, James Liang, Dongfang Liu
- [NeuPhysics: Editable Neural Geometry and Physics from Monocular Videos](#) Yi-Ling Qiao, Alexander Gao, Ming Lin
- [Measuring Data Reconstruction Defenses in Collaborative Inference Systems](#) Mengda Yang, Ziang Li, Juan Wang, Hongxin Hu, Ao Ren, Xiaoyang Xu, Wenzhe Yi
- [ZARTS: On Zero-order Optimization for Neural Architecture Search](#) Xiaoxing Wang, Wenxuan Guo, Jianlin Su, Xiaokang Yang, Junchi Yan
- [Make Some Noise: Reliable and Efficient Single-Step Adversarial Training](#) Pau de Jorge Aranda, Adel Bibi, Riccardo Volpi, Amartya Sanyal, Philip Torr, Gregory Rogez, Puneet Dokania
- [Structural Pruning via Latency-Saliency Knapsack](#) Maying Shen, Hongxu Yin, Pavlo Molchanov, Lei Mao, Jianna Liu, Jose M. Alvarez
- [Risk Bounds of Multi-Pass SGD for Least Squares in the Interpolation Regime](#) Difan Zou, Jingfeng Wu, Vladimir Braverman, Quanquan Gu, Sham Kakade
- [Decision-based Black-box Attack Against Vision Transformers via Patch-wise Adversarial Removal](#) Yucheng Shi, Yahong Han, Yu-an Tan, Xiaohui Kuang
- [EfficientFormer: Vision Transformers at MobileNet Speed](#) Yanyu Li, Geng Yuan, Yang Wen, Ju Hu, Georgios Evangelidis, Sergey Tulyakov, Yanqi Wang, Jian Ren

- [Holomorphic Equilibrium Propagation Computes Exact Gradients Through Finite Size Oscillations](#) Axel Laborieux, Friedemann Zenke
- [Learning Substructure Invariance for Out-of-Distribution Molecular Representations](#) Nianzu Yang, Kaipeng Zeng, Qitian Wu, Xiaosong Jia, Junchi Yan
- [The Dollar Street Dataset: Images Representing the Geographic and Socioeconomic Diversity of the World](#) William Gaviria Rojas, Sudnya Diamos, Keertan Kini, David Kanter, Vijay Janapa Reddi, Cody Coleman
- [LST: Ladder Side-Tuning for Parameter and Memory Efficient Transfer Learning](#) Yi-Lin Sung, Jaemin Cho, Mohit Bansal
- [CGLB: Benchmark Tasks for Continual Graph Learning](#) Xikun Zhang, Dongjin Song, Dacheng Tao
- [Pre-Trained Image Encoder for Generalizable Visual Reinforcement Learning](#) Zhecheng Yuan, Zhengrong Xue, Bo Yuan, Xueqian Wang, Yi WU, Yang Gao, Huazhe Xu
- [Amortized Inference for Heterogeneous Reconstruction in Cryo-EM](#) Axel Levy, Gordon Wetzstein, Julien N.P Martel, Frederic Poitevin, Ellen Zhong
- [RKHS-SHAP: Shapley Values for Kernel Methods](#) Siu Lun Chau, Robert Hu, Javier González, Dino Sejdinovic
- [Single-Stage Visual Relationship Learning using Conditional Queries](#) Alakh Desai, Tz-Ying Wu, Subarna Tripathi, Nuno Vasconcelos
- [Unlabelled Sample Compression Schemes for Intersection-Closed Classes and Extremal Classes](#) Joachim Rubinstein, Benjamin Rubinstein
- [VCT: A Video Compression Transformer](#) Fabian Mentzer, George D Toderici, David Minnen, Sergi Caelles, Sung Jin Hwang, Mario Lucic, Eirikur Agustsson
- [Amortized Inference for Causal Structure Learning](#) Lars Lorch, Scott Sussex, Jonas Rothfuss, Andreas Krause, Bernhard Schölkopf
- [Emergent Graphical Conventions in a Visual Communication Game](#) Shuwen Qiu, Sirui Xie, Lifeng Fan, Tao Gao, Jungseock Joo, Song-Chun Zhu, Yixin Zhu
- [Repairing Neural Networks by Leaving the Right Past Behind](#) Ryutaro Tanno, Melanie F. Pradier, Aditya Nori, Yingzhen Li
- [Selective compression learning of latent representations for variable-rate image compression](#) Jooyoung Lee, Seyoon Jeong, Munchurl Kim
- [Multi-LexSum: Real-world Summaries of Civil Rights Lawsuits at Multiple Granularities](#) Zejiang Shen, Kyle Lo, Lauren Yu, Nathan Dahlberg, Margo Schlanger, Doug Downey
- [Increasing Confidence in Adversarial Robustness Evaluations](#) Roland S. Zimmermann, Wieland Brendel, Florian Tramer, Nicholas Carlini
- [Local Bayesian optimization via maximizing probability of descent](#) Quan Nguyen, Kaiwen Wu, Jacob Gardner, Roman Garnett
- [Staircase Attention for Recurrent Processing of Sequences](#) Da JU, Stephen Roller, Sainbayar Sukhbaatar, Jason E Weston
- [Causal Inference with Non-IID Data using Linear Graphical Models](#) Chi Zhang, Karthika Mohan, Judea Pearl
- [Data-Driven Offline Decision-Making via Invariant Representation Learning](#) Han Qi, Yi Su, Aviral Kumar, Sergey Levine
- [Untargeted Backdoor Watermark: Towards Harmless and Stealthy Dataset Copyright Protection](#) Yiming Li, Yang Bai, Yong Jiang, Yong Yang, Shu-Tao Xia, Bo Li
- [Learning Symmetric Rules with SATNet](#) Sangho Lim, Eun-Gyeol Oh, Hongseok Yang
- [The Privacy Onion Effect: Memorization is Relative](#) Nicholas Carlini, Matthew Jagielski, Chiyuan Zhang, Nicolas Papernot, Andreas Terzis, Florian Tramer
- [Langevin Autoencoders for Learning Deep Latent Variable Models](#) Shohei Taniguchi, Yusuke Iwasawa, Wataru Kumagai, Yutaka Matsuo
- [Envy-free Policy Teaching to Multiple Agents](#) Jiarui Gan, R Majumdar, Adish Singla, Goran Radanovic
- [Provably Efficient Model-Free Constrained RL with Linear Function Approximation](#) Arnob Ghosh, Xingyu Zhou, Ness Shroff
- [Towards Optimal Communication Complexity in Distributed Non-Convex Optimization](#) Kumar Kshitij Patel, Lingxiao Wang, Blake E. Woodworth, Brian Bullins, Nati Srebro
- [Eliciting Thinking Hierarchy without a Prior](#) Yuqing Kong, Yunqi Li, Yubo Zhang, Zhihuan Huang, Jinzhao Wu
- [Fair and Efficient Allocations Without Obvious Manipulations](#) Alexandros Psomas, Paritosh Verma
- [Bridge the Gap Between Architecture Spaces via A Cross-Domain Predictor](#) Yuqiao Liu, Yehui Tang, Zeqiong Lv, Yunhe Wang, Yanan Sun
- [Shield Decentralization for Safe Multi-Agent Reinforcement Learning](#) Daniel Melcer, Christopher Amato, Stavros Tripakis
- [Domain Generalization without Excess Empirical Risk](#) Ozan Sener, Vladlen Koltun
- [Star Temporal Classification: Sequence Modeling with Partially Labeled Data](#) Vineel Pratap, Awni Hannun, Gabriel Synnaeve, Ronan Collobert
- [Signal Processing for Implicit Neural Representations](#) Dejia Xu, Peihao Wang, Yifan Jiang, Zhiwen Fan, Zhangyang Wang
- [Fault-Aware Neural Code Rankers](#) Jeevana Priya Inala, Chenglong Wang, Mei Yang, Andres Coda, Mark Encarnación, Shuvendu Lahiri, Madanlal Musuvathi, Jianfeng Gao
- [Prompt Certified Machine Unlearning with Randomized Gradient Smoothing and Quantization](#) Zijie Zhang, Yang Zhou, Xin Zhao, Tianshi Che, Lingjuan Lyu
- [What Makes a "Good" Data Augmentation in Knowledge Distillation - A Statistical Perspective](#) Huan Wang, Suhas Lohit, Michael N. Jones, Yun Fu
- [Supervising the Multi-Fidelity Race of Hyperparameter Configurations](#) Martin Wistuba, Arind Kadra, Josif Grabocka
- [Streaming Radiance Fields for 3D Video Synthesis](#) Lingzhi LI, Zhen Shen, Zhongshu Wang, Li Shen, Ping Tan
- [Byzantine-tolerant federated Gaussian process regression for streaming data](#) Xu Zhang, Zhenyuan Yuan, Minghui Zhu
- [Neural Matching Fields: Implicit Representation of Matching Fields for Visual Correspondence](#) Sunghwan Hong, Jisu Nam, Seokju Cho, Susung Hong, Sangryul Jeon, Dongbo Min, Seungryong Kim

- [Local Spatiotemporal Representation Learning for Longitudinally-consistent Neuroimage Analysis](#) Mengwei Ren, Neel Dey, Martin Styner, Kelly Botteron, Guido Gerig
- [Overparameterization from Computational Constraints](#) Sanjam Garg, Somesh Jha, Saeed Mahloujifar, Mohammad Mahmood, Mingyuan Wang
- [A Unifying Framework of Off-Policy General Value Function Evaluation](#) Tengyu Xu, Zhuoran Yang, Zhaoran Wang, Yingbin Liang
- [Anchor-Changing Regularized Natural Policy Gradient for Multi-Objective Reinforcement Learning](#) Ruida Zhou, Tao Liu, Dileep Kalathil, P. R. Kumar, Chao Tian
- [Do Current Multi-Task Optimization Methods in Deep Learning Even Help?](#) Derrick Xin, Behrooz Ghorbani, Justin Gilmer, Ankush Garg, Orhan Firat
- [TAP-Vid: A Benchmark for Tracking Any Point in a Video](#) Carl Doersch, Ankush Gupta, Larisa Markeeva, Adria Recasens, Lucas Smaira, Yusuf Aytar, Joao Carreira, Andrew Zisserman, Yi Yang
- [Alignment-guided Temporal Attention for Video Action Recognition](#) Yizhou Zhao, Zhenyang Li, Xun Guo, Yan Lu
- [Meta Reinforcement Learning with Finite Training Tasks - a Density Estimation Approach](#) Zohar Rimmon, Aviv Tamar, Gilad Adler
- [TotalSelfScan: Learning Full-body Avatars from Self-Portrait Videos of Faces, Hands, and Bodies](#) Junting Dong, Qi Fang, Yudong Guo, Sida Peng, Qing Shuai, Xiaowei Zhou, Hujun Bao
- [DHRL: A Graph-Based Approach for Long-Horizon and Sparse Hierarchical Reinforcement Learning](#) Seungjae Lee, Jigang Kim, Inkyu Jang, H. Jin Kim
- [A Classification of \$G\$ -invariant Shallow Neural Networks](#) Devanshu Agrawal, James Ostrowski
- [A Conditional Randomization Test for Sparse Logistic Regression in High-Dimension](#) Binh T. Nguyen, Bertrand Thirion, Sylvain Arlot
- [Biologically-Plausible Determinant Maximization Neural Networks for Blind Separation of Correlated Sources](#) Bariscan Bozkurt, Cengiz Pehlevan, Alper Erdogan
- [Generalized One-shot Domain Adaptation of Generative Adversarial Networks](#) Zicheng Zhang, Yinglu Liu, Congying Han, Tiande Guo, Ting Yao, Tao Mei
- [Representing Spatial Trajectories as Distributions](#) Didac Suris Coll-Vinent, Carl Vondrick
- [EPIC-KITCHENS VISOR Benchmark: Video Segmentations and Object Relations](#) Ahmad Darkhalil, Dandan Shan, Bin Zhu, Jian Ma, Amlan Kar, Richard Higgins, Sanja Fidler, David Fouhey, Dima Damen
- [Inverse Design for Fluid-Structure Interactions using Graph Network Simulators](#) Kelsey Allen, Tatiana Lopez-Guevara, Kimberly L. Stachenfeld, Alvaro Sanchez Gonzalez, Peter Battaglia, Jessica B. Hamrick, Tobias Pfaff
- [What Makes Graph Neural Networks Miscalibrated?](#) Hans Hao-Hsun Hsu, Yuesong Shen, Christian Tomani, Daniel Cremers
- [Stochastic Adaptive Activation Function](#) Kyungsu Lee, Jaeseung Yang, Haeyun Lee, Jae Youn Hwang
- [ActionSense: A Multimodal Dataset and Recording Framework for Human Activities Using Wearable Sensors in a Kitchen Environment](#) Joseph DelPreto, Chao Liu, Yiyue Luo, Michael Foshey, Yunzhu Li, Antonio Torralba, Wojciech Matusik, Daniela Rus
- [Video compression dataset and benchmark of learning-based video-quality metrics](#) Anastasia Antsiferova, Sergey Lavrushkin, Maksim Smirnov, Aleksandr Gushchin, Dmitriy Vatolin, Dmitriy Kulikov
- [Scalable Distributional Robustness in a Class of Non-Convex Optimization with Guarantees](#) Avinandan Bose, Arunesh Sinha, Tien Mai
- [Prototypical VoteNet for Few-Shot 3D Point Cloud Object Detection](#) Shizhen Zhao, Xiaojuan Qi
- [Learning from a Sample in Online Algorithms](#) C.J. Argue, Alan Frieze, Anupam Gupta, Christopher Seiler
- [Robust Rent Division](#) Dominik Peters, Ariel D. Procaccia, David Zhu
- [Efficient Dataset Distillation using Random Feature Approximation](#) Noel Loo, Ramin Hasani, Alexander Amini, Daniela Rus
- [Scalable Sensitivity and Uncertainty Analyses for Causal-Effect Estimates of Continuous-Valued Interventions](#) Andrew Jesson, Alyson Douglas, Peter Manshausen, Maëlys Solal, Nicolai Meinshausen, Philip Stier, Yarin Gal, Uri Shalit
- [Sparse Interaction Additive Networks via Feature Interaction Detection and Sparse Selection](#) James Enouen, Yan Liu
- [Faster Stochastic Algorithms for Minimax Optimization under Polyak-Łojasiewicz Condition](#) Lesi Chen, Boyuan Yao, Luo Luo
- [Kantorovich Strikes Back! Wasserstein GANs are not Optimal Transport?](#) Alexander Korotin, Alexander Kolesov, Evgeny Burnaev
- [List-Decodable Sparse Mean Estimation via Difference-of-Pairs Filtering](#) Ilias Diakonikolas, Daniel Kane, Sushrut Karmalkar, Ankit Pensia, Thanasis Pittas
- [Distributional Convergence of the Sliced Wasserstein Process](#) Jiaqi Xi, Jonathan Niles-Weed
- [VTC-LFC: Vision Transformer Compression with Low-Frequency Components](#) Zhenyu Wang, Hao Luo, Pichao WANG, Feng Ding, Fan Wang, Hao Li
- [A Theory of PAC Learnability under Transformation Invariances](#) Han Shao, Omar Montasser, Avrim Blum
- [PALBERT: Teaching ALBERT to Ponder](#) Nikita Balagansky, Daniil Gavrilov
- [Distributed Methods with Compressed Communication for Solving Variational Inequalities, with Theoretical Guarantees](#) Aleksandr Beznosikov, Peter Richtarik, Michael Diskin, Max Ryabinin, Alexander Gasnikov
- [Analyzing Data-Centric Properties for Graph Contrastive Learning](#) Puja Trivedi, Ekdeep S Lubana, Mark Heimann, Danai Koutra, Jayaraman Thiagarajan
- [Learning Bipartite Graphs: Heavy Tails and Multiple Components](#) José Vinícius de Miranda Cardoso, Jiaxi Ying, Daniel Palomar
- [Training Scale-Invariant Neural Networks on the Sphere Can Happen in Three Regimes](#) Maxim Kodryan, Ekaterina Lobacheva, Maksim Nakhodnov, Dmitry P. Vetrov

- [Exploring the Whole Rashomon Set of Sparse Decision Trees](#) Rui Xin, Chudi Zhong, Zhi Chen, Takuya Takagi, Margo Seltzer, Cynthia Rudin
- [Graph Self-supervised Learning with Accurate Discrepancy Learning](#) Dongki Kim, Jinheon Baek, Sung Ju Hwang
- [Multi-Scale Adaptive Network for Single Image Denoising](#) Yuanbiao Gou, Peng Hu, Jiancheng Lv, Joey Tianyi Zhou, Xi Peng
- [A theory of weight distribution-constrained learning](#) Weishun Zhong, Ben Sorscher, Daniel Lee, Haim Sompolsky
- [Hierarchical Normalization for Robust Monocular Depth Estimation](#) Chi Zhang, Wei Yin, Billz Wang, Gang Yu, BIN FU, Chunhua Shen
- [Deep Compression of Pre-trained Transformer Models](#) Naigang Wang, Chi-Chun (Charlie) Liu, Swagath Venkataramani, Sanchari Sen, Chia-Yu Chen, Kaoutar El Maghraoui, Vijayalakshmi (Viji) Srinivasan, Leland Chang
- [Constrained Predictive Coding as a Biologically Plausible Model of the Cortical Hierarchy](#) Siavash Golkar, Tiberiu Tesileanu, Yanis Bahroun, Anirvan Sengupta, Dmitri Chklovskii
- [Washing The Unwashable : On The \(Im\)possibility of Fairwashing Detection](#) Ali Shahin Shamsabadi, Mohammad Yaghini, Natalie Dullerud, Sierra Wyllie, Ulrich Aivodji, Aisha Alaagib, Sébastien Gambs, Nicolas Papernot
- [Near-Optimal Collaborative Learning in Bandits](#) Clémence Réda, Sattar Vakili, Emilie Kaufmann
- [Hypothesis Testing for Differentially Private Linear Regression](#) Daniel Alabi, Salil Vadhan
- [Bayesian Optimistic Optimization: Optimistic Exploration for Model-based Reinforcement Learning](#) Chenyang Wu, Tianci Li, Zongzhang Zhang, Yang Yu
- [TokenMixup: Efficient Attention-guided Token-level Data Augmentation for Transformers](#) Hyeong Kyu Choi, Joonmyung Choi, Hyunwoo J. Kim
- [Error Analysis of Tensor-Train Cross Approximation](#) Zhen Qin, Alexander Lidiak, Zhexuan Gong, Gongguo Tang, Michael B Wakin, Zhihui Zhu
- [Approximate Secular Equations for the Cubic Regularization Subproblem](#) Yihang Gao, Man-Chung Yue, Michael Ng
- [On the Theoretical Properties of Noise Correlation in Stochastic Optimization](#) Aurelien Lucchi, Frank Proske, Antonio Orvieto, Francis Bach, Hans Kersting
- [Test-Time Prompt Tuning for Zero-Shot Generalization in Vision-Language Models](#) Manli Shu, Weili Nie, De-An Huang, Zhiding Yu, Tom Goldstein, Anima Anandkumar, Chaowei Xiao
- [SemMAE: Semantic-Guided Masking for Learning Masked Autoencoders](#) Gang Li, Heliang Zheng, Daqing Liu, Chaoyue Wang, Bing Su, Changwen Zheng
- [BiT: Robustly Binarized Multi-distilled Transformer](#) Zechun Liu, Barlas Oguz, Aasish Pappu, Lin Xiao, Scott Yih, Meng Li, Raghuraman Krishnamoorthi, Yashar Mehdad
- [Perturbation Learning Based Anomaly Detection](#) Jinyu Cai, Jicong Fan
- [Knowledge-Aware Bayesian Deep Topic Model](#) Dongsheng Wang, Yi.shi Xu, Miaoge Li, Zhibin Duan, Chaojie Wang, Bo Chen, Mingyuan Zhou
- [SelecMix: Debaised Learning by Contradicting-pair Sampling](#) Inwoo Hwang, Sangjun Lee, Yunhyeok Kwak, Seong Joon Oh, Damien Teney, Jin-Hwa Kim, Byoung-Tak Zhang
- [Are all Frames Equal? Active Sparse Labeling for Video Action Detection](#) Aayush Rana, Yogesh Rawat
- [Unsupervised Domain Adaptation for Semantic Segmentation using Depth Distribution](#) Quanliang Wu, Huajun Liu
- [P2P: Tuning Pre-trained Image Models for Point Cloud Analysis with Point-to-Pixel Prompting](#) Ziyi Wang, Xumin Yu, Yongming Rao, Jie Zhou, Jiwen Lu
- [Finding Differences Between Transformers and ConvNets Using Counterfactual Simulation Testing](#) Nataniel Ruiz, Sarah Bargal, Cihang Xie, Kate Saenko, Stan Sclaroff
- [In the Eye of the Beholder: Robust Prediction with Causal User Modeling](#) Amir Feder, Guy Horowitz, Yoav Wald, Roi Reichart, Nir Rosenfeld
- [Variational inference via Wasserstein gradient flows](#) Marc Lambert, Sinho Chewi, Francis Bach, Silvére Bonnabel, Philippe Rigollet
- [projUNN: efficient method for training deep networks with unitary matrices](#) Bobak Kiani, Randall Balestriero, Yann LeCun, Seth Lloyd
- [Reinforcement Learning in a Birth and Death Process: Breaking the Dependence on the State Space](#) Jonatha Anselmi, Bruno Gaujal, Louis-Sébastien Rebuffi
- [Multi-dataset Training of Transformers for Robust Action Recognition](#) Junwei Liang, Enwei Zhang, Jun Zhang, Chunhua Shen
- [An Embarrassingly Simple Approach to Semi-Supervised Few-Shot Learning](#) Xiu-Shen Wei, H.-Y. Xu, Faen Zhang, Yuxin Peng, Wei Zhou
- [Recipe for a General, Powerful, Scalable Graph Transformer](#) Ladislav Rampásek, Michael Galkin, Vijay Prakash Dwivedi, Anh Tuan Luu, Guy Wolf, Dominique Beaini
- [ALIFE: Adaptive Logit Regularizer and Feature Replay for Incremental Semantic Segmentation](#) Youngmin Oh, Donghyeon Baek, Bumsub Ham
- [Rare Gems: Finding Lottery Tickets at Initialization](#) Kartik Sreenivasan, Jy-yong Sohn, Liu Yang, Matthew Grinde, Alliot Nagle, Hongyi Wang, Eric Xing, Kangwook Lee, Dimitris Papailiopoulos
- [Fast Vision Transformers with HiLo Attention](#) Zizheng Pan, Jianfei Cai, Bohan Zhuang
- [Online Bipartite Matching with Advice: Tight Robustness-Consistency Tradeoffs for the Two-Stage Model](#) Billy Jin, Will Ma
- [Identifying good directions to escape the NTK regime and efficiently learn low-degree plus sparse polynomials](#) Eshaan Nichani, Yu Bai, Jason D. Lee

- [Pure Transformers are Powerful Graph Learners](#) Jinwoo Kim, Dat Nguyen, Seonwoo Min, Sungjun Cho, Moontae Lee, Honglak Lee, Seunghoon Hong
- [Orthogonal Transformer: An Efficient Vision Transformer Backbone with Token Orthogonalization](#) Huaibo Huang, Xiaoqiang Zhou, Ran He
- [Using Mixup as a Regularizer Can Surprisingly Improve Accuracy & Out-of-Distribution Robustness](#) Francesco Pinto, Harry Yang, Ser Nam Lim, Philip Torr, Puneet Dokania
- [NeurOLight: A Physics-Agnostic Neural Operator Enabling Parametric Photonic Device Simulation](#) Jiaqi Gu, Zhengqi Gao, Chenghao Feng, Hanqing Zhu, Ray Chen, Duane Boning, David Pan
- [Learning the Structure of Large Networked Systems Obeying Conservation Laws](#) Anirudh Rayas, Rajasekhar Anguluri, Gautam Dasarathy
- [FP8 Quantization: The Power of the Exponent](#) Andrey Kuzmin, Mart van Baalen, Yuwei Ren, Markus Nagel, Jorn Peters, Tijmen Blankevoort
- [Bridging the Gap Between Vision Transformers and Convolutional Neural Networks on Small Datasets](#) Zhiying Lu, Hongtao Xie, Chuanbin Liu, Yongdong Zhang
- [Private Set Generation with Discriminative Information](#) Dingfan Chen, Raouf Kerkouche, Mario Fritz
- [The First Optimal Algorithm for Smooth and Strongly-Concave Strongly-Concave Minimax Optimization](#) Dmitry Kovalev, Alexander Gasnikov
- [Provable Defense against Backdoor Policies in Reinforcement Learning](#) Shubham Bharti, Xuezhou Zhang, Adish Singla, Jerry Zhu
- [Diffusion Models as Plug-and-Play Priors](#) Alexandros Graikos, Nikolay Malkin, Nebojsa Jojic, Dimitris Samaras
- [CoupAlign: Coupling Word-Pixel with Sentence-Mask Alignments for Referring Image Segmentation](#) Zicheng Zhang, Yi Zhu, Jianzhuang Liu, Xiaodan Liang, Wei Ke
- [Saliency-Aware Neural Architecture Search](#) Ramtin Hosseini, Pengtao Xie
- [VaiPhy: a Variational Inference Based Algorithm for Phylogeny](#) Hazal Koptagel, Oskar Kviman, Harald Melin, Negar Safinianaini, Jens Lagergren
- [A simple but strong baseline for online continual learning: Repeated Augmented Rehearsal](#) Yaqian Zhang, Bernhard Pfahringer, Eibe Frank, Albert Bifet, Nick Jin Sean Lim, Yunzhe Jia
- [You Only Live Once: Single-Life Reinforcement Learning](#) Annie Chen, Archit Sharma, Sergey Levine, Chelsea Finn
- [Compressible-composable NeRF via Rank-residual Decomposition](#) Jiaxiang Tang, Xiaokang Chen, Jingbo Wang, Gang Zeng
- [Data-Efficient Pipeline for Offline Reinforcement Learning with Limited Data](#) Allen Nie, Yannis Flet-Berliac, Deon Jordan, William Steenbergen, Emma Brunskill
- [Hardness in Markov Decision Processes: Theory and Practice](#) Michelangelo Conserva, Paulo Rauber
- [Injecting Domain Knowledge from Empirical Interatomic Potentials to Neural Networks for Predicting Material Properties](#) Zeren Shui, Daniel Karls, Mingjian Wen, Ilya Nikiforov, Ellad Tadmor, George Karypis
- [Learning Modular Simulations for Homogeneous Systems](#) Jayesh Gupta, Sai Vemprala, Ashish Kapoor
- [On Gap-dependent Bounds for Offline Reinforcement Learning](#) Xinqi Wang, Qiwen Cui, Simon S. Du
- [Semi-Discrete Normalizing Flows through Differentiable Tessellation](#) Ricky T. Q. Chen, Brandon Amos, Maximilian Nickel
- [Don't Pour Cereal into Coffee: Differentiable Temporal Logic for Temporal Action Segmentation](#) Ziwei Xu, Yogesh Rawat, Yongkang Wong, Mohan S. Kankanhalli, Mubarak Shah
- [Batch-Size Independent Regret Bounds for Combinatorial Semi-Bandits with Probabilistically Triggered Arms or Independent Arms](#) Xutong Liu, Jinhang Zuo, Siwei Wang, Carlee Joe-Wong, John C.S. Lui, Wei Chen
- [Less-forgetting Multi-lingual Fine-tuning](#) Yuren Mao, Yaobo Liang, Nan Duan, Haobo Wang, Kai Wang, Lu Chen, Yunjun Gao
- [Adversarial Attack on Attackers: Post-Process to Mitigate Black-Box Score-Based Query Attacks](#) Sizhe Chen, Zhehao Huang, Qinghua Tao, Yingwen Wu, Cihang Xie, Xiaolin Huang
- [Retaining Knowledge for Learning with Dynamic Definition](#) Zichang Liu, Benjamin Coleman, Tianyi Zhang, Anshumali Shrivastava
- [HUMANISE: Language-conditioned Human Motion Generation in 3D Scenes](#) Zan Wang, Yixin Chen, Tengyu Liu, Yixin Zhu, Wei Liang, Siyuan Huang
- [Keypoint-Guided Optimal Transport with Applications in Heterogeneous Domain Adaptation](#) Xiang Gu, Yucheng Yang, Wei Zeng, Jian Sun, Zongben Xu
- [Sequence-to-Set Generative Models](#) Longtao Tang, Ying Zhou, Yu Yang
- [Near-Optimal Multi-Agent Learning for Safe Coverage Control](#) Manish Prajapat, Matteo Turchetta, Melanie Zeilinger, Andreas Krause
- [Network change point localisation under local differential privacy](#) Mengchu Li, Tom Berrett, Yi Yu
- [Independence Testing for Bounded Degree Bayesian Networks](#) Arnab Bhattacharyya, Clément L Canonne, Qiping Yang
- [Beyond spectral gap: the role of the topology in decentralized learning](#) Thijs Vogels, Hadrien Hendrikx, Martin Jaggi
- [GriddlyJS: A Web IDE for Reinforcement Learning](#) Christopher Bamford, Minqi Jiang, Mikayel Samvelyan, Tim Rocktäschel
- [Periodic Graph Transformers for Crystal Material Property Prediction](#) Keqiang Yan, Yi Liu, Yuchao Lin, Shuiwang Ji
- [Decentralized, Communication- and Coordination-free Learning in Structured Matching Markets](#) Chinmay Maheshwari, Shankar Sastry, Eric Mazumdar
- [Improving GANs with A Dynamic Discriminator](#) Ceyuan Yang, Yujun Shen, Yinghao Xu, Deli Zhao, Bo Dai, Bolei Zhou
- [Deliberated Domain Bridging for Domain Adaptive Semantic Segmentation](#) Lin Chen, Zhixiang Wei, Xin Jin, Huaian Chen, Miao Zheng, Kai Chen, Yi Jin

- [Scalable and Efficient Non-adaptive Deterministic Group Testing](#) Dariusz Kowalski, Dominik Pajak
- [Structure-Preserving 3D Garment Modeling with Neural Sewing Machines](#) Xipeng Chen, Guangrun Wang, Dizhong Zhu, Xiaodan Liang, Philip Torr, Liang Lin
- [Rethinking Variational Inference for Probabilistic Programs with Stochastic Support](#) Tim Reichelt, Luke Ong, Thomas Rainforth
- [Variance Reduced ProxSkip: Algorithm, Theory and Application to Federated Learning](#) Grigory Malinovsky, Kai Yi, Peter Richtarik
- [DreamShard: Generalizable Embedding Table Placement for Recommender Systems](#) Daochen Zha, Louis Feng, Qiaoyu Tan, Zirui Liu, Kwei-Herng Lai, Bhargav Bhushanam, Yuandong Tian, Arun Kejariwal, Xia Hu
- [Order-Invariant Cardinality Estimators Are Differentially Private](#) Charlie Dickens, Justin Thaler, Daniel Ting
- [NeuForm: Adaptive Overfitting for Neural Shape Editing](#) Connor Lin, Niloy Mitra, Gordon Wetzstein, Leonidas J. Guibas, Paul Guerrero
- [Inductive Logical Query Answering in Knowledge Graphs](#) Michael Galkin, Zhaocheng Zhu, Hongyu Ren, Jian Tang
- [Drawing out of Distribution with Neuro-Symbolic Generative Models](#) Yichao Liang, Josh Tenenbaum, Tuan Anh Le, Siddharth N
- [Unsupervised Object Representation Learning using Translation and Rotation Group Equivariant VAE](#) Alireza Nasiri, Tristan Bepler
- [PointTAD: Multi-Label Temporal Action Detection with Learnable Query Points](#) Jing Tan, Xiaotong Zhao, Xintian Shi, Bin Kang, Limin Wang
- [Unpacking Reward Shaping: Understanding the Benefits of Reward Engineering on Sample Complexity](#) Abhishek Gupta, Aldo Pacchiano, Yuexiang Zhai, Sham Kakade, Sergey Levine
- [LASSIE: Learning Articulated Shapes from Sparse Image Ensemble via 3D Part Discovery](#) Chun-Han Yao, Wei-Chih Hung, Yuanzhen Li, Michael Rubinstein, Ming-Hsuan Yang, Varun Jampani
- [Retrieval-Augmented Diffusion Models](#) Andreas Blattmann, Robin Rombach, Kaan Oktay, Jonas Müller, Björn Ommer
- [Logical Credal Networks](#) Radu Marinescu, Haifeng Qian, Alexander Gray, Debarun Bhattacharjya, Francisco Barahona, Tian Gao, Ryan Riegel, Pravinda Sahu
- [Neural Set Function Extensions: Learning with Discrete Functions in High Dimensions](#) Nikolaos Karalias, Joshua Robinson, Andreas Loukas, Stefanie Jegelka
- [Minimax-Optimal Multi-Agent RL in Markov Games With a Generative Model](#) Gen Li, Yuejie Chi, Yuting Wei, Yuxin Chen
- [Bi-directional Weakly Supervised Knowledge Distillation for Whole Slide Image Classification](#) Linhao Qu, xiaoyuan luo, Manning Wang, Zhijian Song
- [Private and Communication-Efficient Algorithms for Entropy Estimation](#) Gecia Bravo-Hermesdorff, Róbert Busa-Fekete, Mohammad Ghavamzadeh, Andres Munoz Medina, Umar Syed
- [PKD: General Distillation Framework for Object Detectors via Pearson Correlation Coefficient](#) Weihan Cao, Yifan Zhang, Jianfei Gao, Anda Cheng, Ke Cheng, Jian Cheng
- [Off-Team Learning](#) Brandon Cui, Hengyuan Hu, Andrei Lupu, Samuel Sokota, Jakob Foerster
- [NUWA-Infinity: Autoregressive over Autoregressive Generation for Infinite Visual Synthesis](#) Jian Liang, Chenfei Wu, Xiaowei Hu, Zhe Gan, Jianfeng Wang, Lijuan Wang, Zicheng Liu, Yuejian Fang, Nan Duan
- [SoftPatch: Unsupervised Anomaly Detection with Noisy Data](#) Xi Jiang, Jianlin Liu, Jinbao Wang, Qiang Nie, Kai WU, Yong Liu, Chengjie Wang, Feng Zheng
- [Stability Analysis and Generalization Bounds of Adversarial Training](#) Jiancong Xiao, Yanbo Fan, Ruoyu Sun, Jue Wang, Zhi-Quan Luo
- [LasUIE: Unifying Information Extraction with Latent Adaptive Structure-aware Generative Language Model](#) Hao Fei, Shengqiong Wu, Jingye Li, Bobo Li, Fei Li, Libo Qin, Meishan Zhang, Min Zhang, Tat-Seng Chua
- [STaR: Bootstrapping Reasoning With Reasoning](#) Eric Zelikman, Yuhuai Wu, Jesse Mu, Noah Goodman
- [TPU-KNN: K Nearest Neighbor Search at Peak FLOP/s](#) Felix Chern, Blake Hechtman, Andy Davis, Ruiqi Guo, David Majnemer, Sanjiv Kumar
- [Efficient Meta Reinforcement Learning for Preference-based Fast Adaptation](#) Zhizhou Ren, Anji Liu, Yitao Liang, Jian Peng, Jianzhu Ma
- [Weakly Supervised Representation Learning with Sparse Perturbations](#) Kartik Ahuja, Jason S. Hartford, Yoshua Bengio
- [A Multi-Resolution Framework for U-Nets with Applications to Hierarchical VAEs](#) Fabian Falck, Christopher Williams, Dominic Danks, George Deligiannidis, Christopher Yau, Chris C Holmes, Arnaud Doucet, Matthew Willetts
- [Watermarking for Out-of-distribution Detection](#) Qizhou Wang, Feng Liu, Yonggang Zhang, Jing Zhang, Chen Gong, Tongliang Liu, Bo Han
- [K-LITE: Learning Transferable Visual Models with External Knowledge](#) Sheng Shen, Chunyuan Li, Xiaowei Hu, Yujia Xie, Jianwei Yang, Pengchuan Zhang, Zhe Gan, Lijuan Wang, Lu Yuan, Ce Liu, Kurt Keutzer, Trevor Darrell, Anna Rohrbach, Jianfeng Gao
- [VeriDark: A Large-Scale Benchmark for Authorship Verification on the Dark Web](#) Andrei Manolache, Florin Brad, Antonio Barbalau, Radu Tudor Ionescu, Marius Popescu
- [EHRSQL: A Practical Text-to-SQL Benchmark for Electronic Health Records](#) Gyubok Lee, Hyeonji Hwang, Seongsu Bae, Yeonsu Kwon, Woncheol Shin, Seongjun Yang, Minjoon Seo, Jong-Yeup Kim, Edward Choi
- [Where do Models go Wrong? Parameter-Space Saliency Maps for Explainability](#) Roman Levin, Manli Shu, Eitan Borgnia, Furong Huang, Micah Goldblum, Tom Goldstein
- [Using Embeddings for Causal Estimation of Peer Influence in Social Networks](#) Irina Cristali, Victor Veitch
- [Collaborative Learning by Detecting Collaboration Partners](#) Shu Ding, Wei Wang
- [Linear Label Ranking with Bounded Noise](#) Dimitris Fotakis, Alkis Kalavasis, Vasilis Kontonis, Christos Tzamos
- [Quo Vadis: Is Trajectory Forecasting the Key Towards Long-Term Multi-Object Tracking?](#) Patrick Dendorfer, Vladimir Yugay, Aljosa Osep,

- [Wasserstein Iterative Networks for Barycenter Estimation](#) Alexander Korotin, Vage Egiazarian, Lingxiao Li, Evgeny Burnaev
- [Identifiability of deep generative models without auxiliary information](#) Bohdan Kivva, Goutham Rajendran, Pradeep Ravikumar, Bryon Aragam
- [Task Discovery: Finding the Tasks that Neural Networks Generalize on](#) Andrei Atanov, Andrei Filatov, Teresa Yeo, Ajay Sohmshetty, Amir Zamir
- [Iron: Private Inference on Transformers](#) Meng Hao, Hongwei Li, Hanxiao Chen, Pengzhi Xing, Guowen Xu, Tianwei Zhang
- [On Batch Teaching with Sample Complexity Bounded by VCD](#) Farnam Mansouri, Hans Simon, Adish Singla, Sandra Zilles
- [Approaching Quartic Convergence Rates for Quasi-Stochastic Approximation with Application to Gradient-Free Optimization](#) Caio Kalil Lauand, Sean Meyn
- [PAC: Assisted Value Factorization with Counterfactual Predictions in Multi-Agent Reinforcement Learning](#) Hanhan Zhou, Tian Lan, Vaneet Aggarwal
- [Translation-equivariant Representation in Recurrent Networks with a Continuous Manifold of Attractors](#) Wenhao Zhang, Ying Nian Wu, Si Wu
- [OpenXAI: Towards a Transparent Evaluation of Model Explanations](#) Chirag Agarwal, Satyapriya Krishna, Eshika Saxena, Martin Pawelczyk, Nari Johnson, Isha Puri, Marinka Zitnik, Himabindu Lakkaraju
- [Better Best of Both Worlds Bounds for Bandits with Switching Costs](#) Idan Amir, Guy Azov, Tomer Koren, Roi Livni
- [Uncertainty-Aware Hierarchical Refinement for Incremental Implicitly-Refined Classification](#) Jian Yang, Kai Zhu, Kecheng Zheng, Yang Cao
- [The Hessian Screening Rule](#) Johan Larsson, Jonas Wallin
- [Sharp Analysis of Stochastic Optimization under Global Kurdyka-Lojasiewicz Inequality](#) Ilyas Fatkhullin, Jalal Etesami, Niao He, Negar Kiyavash
- [Plan To Predict: Learning an Uncertainty-Foreseeing Model For Model-Based Reinforcement Learning](#) Zifan Wu, Chao Yu, Chen Chen, Jianye Hao, Hankz Hankui Zhuo
- [Co-Modality Graph Contrastive Learning for Imbalanced Node Classification](#) Yiyue Qian, Chunhui Zhang, Yiming Zhang, Qianlong Wen, Yanfang Ye, Chuxu Zhang
- [Efficiency Ordering of Stochastic Gradient Descent](#) Jie Hu, Vishwaraj Doshi, Do-Young Eun
- [Muffliato: Peer-to-Peer Privacy Amplification for Decentralized Optimization and Averaging](#) Edwige Cyffers, Mathieu Even, Aurélien Bellet, Laurent Massoulié
- [Mining Multi-Label Samples from Single Positive Labels](#) Youngin Cho, Daejin Kim, MOHAMMAD AZAM KHAN, Jaegul Choo
- [Sequential Information Design: Learning to Persuade in the Dark](#) Martino Bernasconi, Matteo Castiglioni, Alberto Marchesi, Nicola Gatti, Francesco Trovò
- [AutoML Two-Sample Test](#) Jonas M. Kübler, Vincent Stimper, Simon Buchholz, Krikamol Muandet, Bernhard Schölkopf
- [Cross-Linked Unified Embedding for cross-modality representation learning](#) Xinming Tu, Zhi-Jie Cao, xia chenrui, Sara Mostafavi, Ge Gao
- [Learning Individualized Treatment Rules with Many Treatments: A Supervised Clustering Approach Using Adaptive Fusion](#) Haixu Ma, Donglin Zeng, Yufeng Liu
- [Robust Generalized Method of Moments: A Finite Sample Viewpoint](#) Dhruv Rohatgi, Vasilis Syrgkanis
- [k-Sliced Mutual Information: A Quantitative Study of Scalability with Dimension](#) Ziv Goldfeld, Kristjan Greenewald, Theshani Nuradha, Galen Reeves
- [What's the Harm? Sharp Bounds on the Fraction Negatively Affected by Treatment](#) Nathan Kallus
- [Bessel Equivariant Networks for Inversion of Transmission Effects in Multi-Mode Optical Fibres](#) Joshua Mitton, Simon Mekhail, Miles Padgett, Daniele Faccio, Marco Aversa, Roderick Murray-Smith
- [Training Subset Selection for Weak Supervision](#) Hunter Lang, Aravindan Vijayaraghavan, David Sontag
- [Expansion and Shrinkage of Localization for Weakly-Supervised Semantic Segmentation](#) Jinlong Li, Zequn Jie, Xu Wang, Xiaolin Wei, Lin Ma
- [No Free Lunch from Deep Learning in Neuroscience: A Case Study through Models of the Entorhinal-Hippocampal Circuit](#) Rylan Schaeffer, Mikail Khona, Ila Fiete
- [On the Importance of Gradient Norm in PAC-Bayesian Bounds](#) Itai Gat, Yossi Adi, Alex Schwing, Tamir Hazan
- [RAMBO-RL: Robust Adversarial Model-Based Offline Reinforcement Learning](#) Marc Rigter, Bruno Lacerda, Nick Hawes
- [Learning NP-Hard Multi-Agent Assignment Planning using GNN: Inference on a Random Graph and Provable Auction-Fitted Q-learning](#) HYUNWOOK KANG, Taehwan Kwon, Jinkyoo Park, James R. Morrison
- [Improved techniques for deterministic l2 robustness](#) Sahil Singla, Soheil Feizi
- [Normalizing Flows for Knockoff-free Controlled Feature Selection](#) Derek Hansen, Brian Manzo, Jeffrey Regier
- [ReFactor GNNs: Revisiting Factorisation-based Models from a Message-Passing Perspective](#) Yihong Chen, Pushkar Mishra, Luca Franceschi, Pasquale Minervini, Pontus Lars Erik Saito Stenertorp, Sebastian Riedel
- [Efficient Architecture Search for Diverse Tasks](#) Junhong Shen, Misha Khodak, Ameet Talwalkar
- [Beyond Not-Forgetting: Continual Learning with Backward Knowledge Transfer](#) Sen Lin, Li Yang, Deliang Fan, Junshan Zhang
- [Inherently Explainable Reinforcement Learning in Natural Language](#) Xiangyu Peng, Mark Riedl, Prithviraj Ammanabrolu
- [Conditional Independence Testing with Heteroskedastic Data and Applications to Causal Discovery](#) Wiebke Günther, Urmi Ninad, Jonas Wahl, Jakob Runge

- [On Reinforcement Learning and Distribution Matching for Fine-Tuning Language Models with no Catastrophic Forgetting](#) Tomasz Korbak, Hady Elsahar, Germán Kruszewski, Marc Dymetman
- [Ask4Help: Learning to Leverage an Expert for Embodied Tasks](#) Kunal Pratap Singh, Luca Weihs, Alvaro Herrasti, Jonghyun Choi, Aniruddha Kembhavi, Roozbeh Mottaghi
- [Human-AI Collaborative Bayesian Optimisation](#) Arun Kumar A V, Santu Rana, Alistair Shilton, Svetha Venkatesh
- [Robust Learning against Relational Adversaries](#) Yizhen Wang, Mohannad Alhanahnah, Xiaozhu Meng, Ke Wang, Mihai Christodorescu, Somesh Jha
- [Active Bayesian Causal Inference](#) Christian Toth, Lars Lorch, Christian Knoll, Andreas Krause, Franz Pernkopf, Robert Peharz, Julius von Kügelgen
- [Understanding and Improving Robustness of Vision Transformers through Patch-based Negative Augmentation](#) Yao Qin, Chiyuan Zhang, Ting Chen, Balaji Lakshminarayanan, Alex Beutel, Xuezhi Wang
- [LogiGAN: Learning Logical Reasoning via Adversarial Pre-training](#) Xinyu Pi, Wanjun Zhong, Yan Gao, Nan Duan, Jian-Guang Lou
- [Revisiting Non-Parametric Matching Cost Volumes for Robust and Generalizable Stereo Matching](#) Kelvin Cheng, Tianfu Wu, Christopher Healey
- [On the Interpretability of Regularisation for Neural Networks Through Model Gradient Similarity](#) Vincent Szolnoky, Viktor Andersson, Balazs Kulcsar, Rebecka Jörnsten
- [Controllable 3D Face Synthesis with Conditional Generative Occupancy Fields](#) Keqiang Sun, Shangzhe Wu, Zhaoyang Huang, Ning Zhang, Quan Wang, Hongsheng Li
- [FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness](#) Tri Dao, Dan Fu, Stefano Ermon, Atri Rudra, Christopher Ré
- [Locally Hierarchical Auto-Regressive Modeling for Image Generation](#) Tackgeun You, Saehoon Kim, Chiheon Kim, Doyup Lee, Bohyung Han
- [Learning Enhanced Representation for Tabular Data via Neighborhood Propagation](#) Kounianhua Du, Weinan Zhang, Ruiwen Zhou, Yangkun Wang, Xilong Zhao, Jiarui Jin, Quan Gan, Zheng Zhang, David P Wipf
- [On global convergence of ResNets: From finite to infinite width using linear parameterization](#) Raphaël Barboni, Gabriel Peyré, Francois-Xavier Vialard
- [Spherization Layer: Representation Using Only Angles](#) Hoyong Kim, kangil kim
- [On the Identifiability of Nonlinear ICA: Sparsity and Beyond](#) Yujia Zheng, Ignavier Ng, Kun Zhang
- [Self-Supervised Visual Representation Learning with Semantic Grouping](#) Xin Wen, Bingchen Zhao, Anlin Zheng, Xiangyu Zhang, Xiaojuan Qi
- [Optimistic Mirror Descent Either Converges to Nash or to Strong Coarse Correlated Equilibria in Bimatrix Games](#) Ioannis Anagnostides, Gabriele Farina, Ioannis Panageas, Tuomas Sandholm
- [Discovered Policy Optimisation](#) Chris Lu, Jakub Kuba, Alistair Letcher, Luke Metz, Christian Schroeder de Witt, Jakob Foerster
- [Robust Testing in High-Dimensional Sparse Models](#) Anand Jerry George, Clément L Canonne
- [Learning Consistency-Aware Unsigned Distance Functions Progressively from Raw Point Clouds](#) Junsheng Zhou, Baorui Ma, Yu-Shen Liu, Yi Fang, Zhizhong Han
- [Understanding Square Loss in Training Overparametrized Neural Network Classifiers](#) Tianyang Hu, Jun WANG, Wenjia Wang, Zhenguo Li
- [Multi-Agent Reinforcement Learning is a Sequence Modeling Problem](#) Muning Wen, Jakub Kuba, Runji Lin, Weinan Zhang, Ying Wen, Jun Wang, Yaodong Yang
- [Non-stationary Bandits with Knapsacks](#) Shang Liu, Jiashuo Jiang, Xiaocheng Li
- [Fast Bayesian Inference with Batch Bayesian Quadrature via Kernel Recombination](#) Masaki Adachi, Satoshi Hayakawa, Martin Jørgensen, Harald Oberhauser, Michael A Osborne
- [Are Two Heads the Same as One? Identifying Disparate Treatment in Fair Neural Networks](#) Michael Lohaus, Matthäus Kleindessner, Krishnaram Kenthapadi, Francesco Locatello, Chris Russell
- [Provable Generalization of Overparameterized Meta-learning Trained with SGD](#) Yu Huang, Yingbin Liang, Longbo Huang
- [When Do Flat Minima Optimizers Work?](#) Jean Kaddour, Linqing Liu, Ricardo Silva, Matt J. Kusner
- [Fast Instrument Learning with Faster Rates](#) Ziyu Wang, Yuhao Zhou, Jun Zhu
- [Sparse Gaussian Process Hyperparameters: Optimize or Integrate?](#) Vidhi Lalchand, Wessel Bruinsma, David Burt, Carl Edward Rasmussen
- [HierSpeech: Bridging the Gap between Text and Speech by Hierarchical Variational Inference using Self-supervised Representations for Speech Synthesis](#) Sang-Hoon Lee, Seung-Bin Kim, Ji-Hyun Lee, Eunwoo Song, Min-Jae Hwang, Seong-Whan Lee
- [Sharing Knowledge for Meta-learning with Feature Descriptions](#) Tomoharu Iwata, Atsutoshi Kumagai
- [A Variant of Anderson Mixing with Minimal Memory Size](#) Fuchao Wei, Chenglong Bao, Yang Liu, Guangwen Yang
- [AdaptFormer: Adapting Vision Transformers for Scalable Visual Recognition](#) Shoufa Chen, Chongjian GE, Zhan Tong, Jiangliu Wang, Yibing Song, Jue Wang, Ping Luo
- [Symmetry Teleportation for Accelerated Optimization](#) Bo Zhao, Nima Dehmamy, Robin Walters, Rose Yu
- [Wasserstein Logistic Regression with Mixed Features](#) Aras Selvi, Mohammad Reza Belbasi, Martin Haugh, Wolfram Wiesemann
- [TaiSu: A 166M Large-scale High-Quality Dataset for Chinese Vision-Language Pre-training](#) Yulong Liu, Guibo Zhu, Bin Zhu, Qi Song, Guojing Ge, Haoran Chen, GuanHui Qiao, Ru Peng, Lingxiang Wu, Jinqiao Wang
- [Robust Bayesian Regression via Hard Thresholding](#) Zheyi Fan, Zhaohui Li, Qingpei Hu

- [SwinTrack: A Simple and Strong Baseline for Transformer Tracking](#) Liting Lin, Heng Fan, Zhipeng Zhang, Yong Xu, Haibin Ling
- [Unknown-Aware Domain Adversarial Learning for Open-Set Domain Adaptation](#) JoonHo Jang, Byeonghu Na, Dong Hyeok Shin, Mingi Ji, Kyungwoo Song, Il-chul Moon
- [Learning Chaotic Dynamics in Dissipative Systems](#) Zongyi Li, Miguel Liu-Schiaffini, Nikola Kovachki, Kamyar Aizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew Stuart, Anima Anandkumar
- [Poisson Flow Generative Models](#) Yilun Xu, Ziming Liu, Max Tegmark, Tommi Jaakkola
- [Bounding and Approximating Intersectional Fairness through Marginal Fairness](#) Mathieu Molina, Patrick Loiseau
- [Semi-infinitely Constrained Markov Decision Processes](#) Liangyu Zhang, Yang Peng, Wenhao Yang, Zhihua Zhang
- [Latent Planning via Expansive Tree Search](#) Robert Giesemann, Florian T. Pokorny
- [Invertible Monotone Operators for Normalizing Flows](#) Byeongkeun Ahn, Chiyoon Kim, Youngjoon Hong, Hyunwoo J. Kim
- [Dynamic Sparse Network for Time Series Classification: Learning What to “See”](#) Qiao Xiao, Boqian Wu, Yu Zhang, Shiwei Liu, Mykola Pechenizkiy, Elena Mocanu, Decebal Constantin Mocanu
- [Learning to Sample and Aggregate: Few-shot Reasoning over Temporal Knowledge Graphs](#) Ruijie Wang, Zheng Li, Dachun Sun, Shengzhong Liu, Jinning Li, Bing Yin, Tarek Abdelzaher
- [Evaluating Robustness to Dataset Shift via Parametric Robustness Sets](#) Nikolaj Thams, Michael Oberst, David Sontag
- [CogView2: Faster and Better Text-to-Image Generation via Hierarchical Transformers](#) Ming Ding, Wendi Zheng, Wenyi Hong, Jie Tang
- [Recursive Reasoning in Minimax Games: A Level \$k\$ Gradient Play Method](#) Zichu Liu, Lacra Pavel
- [When to Ask for Help: Proactive Interventions in Autonomous Reinforcement Learning](#) Annie Xie, Fahim Tajwar, Archit Sharma, Chelsea Finn
- [Reinforcement Learning with Neural Radiance Fields](#) Danny Driess, Ingmar Schubert, Pete Florence, Yunzhu Li, Marc Toussaint
- [Function Classes for Identifiable Nonlinear Independent Component Analysis](#) Simon Buchholz, Michel Besserve, Bernhard Schölkopf
- [Self-supervised Heterogeneous Graph Pre-training Based on Structural Clustering](#) Yaming Yang, Ziyu Guan, Zhe Wang, Wei Zhao, Cai Xu, Weigang Lu, Jianbin Huang
- [Learning Audio-Visual Dynamics Using Scene Graphs for Audio Source Separation](#) Moitreyia Chatterjee, Narendra Ahuja, Anoop Cherian
- [On the Sample Complexity of Stabilizing LTI Systems on a Single Trajectory](#) Yang Hu, Adam Wierman, Guannan Qu
- [coVariance Neural Networks](#) Saurabh Sihag, Gonzalo Mateos, Corey McMillan, Alejandro Ribeiro
- [Taming Fat-Tailed \(“Heavier-Tailed” with Potentially Infinite Variance\) Noise in Federated Learning](#) Haibo Yang, Peiwen Qiu, Jia Liu
- [Exploring Figure-Ground Assignment Mechanism in Perceptual Organization](#) Wei Zhai, Yang Cao, Jing Zhang, Zheng-Jun Zha
- [Two-Stream Network for Sign Language Recognition and Translation](#) Yutong Chen, Ronglai Zuo, Fangyun Wei, Yu Wu, Shujie LIU, Brian Mak
- [VisFIS: Visual Feature Importance Supervision with Right-for-the-Right-Reason Objectives](#) Zhuofan Ying, Peter Hase, Mohit Bansal
- [Operative dimensions in unconstrained connectivity of recurrent neural networks](#) Renate Krause, Matthew Cook, Sepp Kollmorgen, Valerio Mante, Giacomo Indiveri
- [Batch size-invariance for policy optimization](#) Jacob Hilton, Karl Cobbe, John Schulman
- [Inference and Sampling for Archimax Copulas](#) Yuting Ng, Ali Hasan, Vahid Tarokh
- [Neural Shape Deformation Priors](#) Jiapeng Tang, Lev Markhasin, Bi Wang, Justus Thies, Matthias Niessner
- [The Curse of Unrolling: Rate of Differentiating Through Optimization](#) Damien Scieur, Gauthier Gidel, Quentin Bertrand, Fabian Pedregosa
- [A Deep Learning Dataloader with Shared Data Preparation](#) jian xie, Jingwei Xu, Guochang Wang, Yuan Yao, Zenan Li, Chun Cao, Hanghang Tong
- [Learning Physics Constrained Dynamics Using Autoencoders](#) Tsung-Yen Yang, Justinian Rosca, Karthik Narasimhan, Peter J Ramadge
- [Variational Model Perturbation for Source-Free Domain Adaptation](#) Mengmeng Jing, Xiantong Zhen, Jingjing Li, Cees Snoek
- [On the non-universality of deep learning: quantifying the cost of symmetry](#) Emmanuel Abbe, Enric Boix-Adsera
- [Sharper Convergence Guarantees for Asynchronous SGD for Distributed and Federated Learning](#) Anastasiia Koloskova, Sebastian U. Stich, Martin Jaggi
- [A Unified Framework for Alternating Offline Model Training and Policy Learning](#) Shentao Yang, Shujian Zhang, Yihao Feng, Mingyuan Zhou
- [Adjoint-aided inference of Gaussian process driven differential equations](#) Paterne GAHUNGU, Christopher Lanyon, Mauricio A Álvarez, Engineer Bainomugisha, Michael T Smith, Richard Wilkinson
- [BOME! Bilevel Optimization Made Easy: A Simple First-Order Approach](#) Bo Liu, Mao Ye, Stephen Wright, Peter Stone, Qiang Liu
- [Mirror Descent with Relative Smoothness in Measure Spaces, with application to Sinkhorn and EM](#) Pierre-Cyril Aubin-Frankowski, Anna Korba, Flavien Léger
- [Peer Prediction for Learning Agents](#) Shi Feng, Fang-Yi Yu, Yiling Chen
- [Visual Clues: Bridging Vision and Language Foundations for Image Paragraph Captioning](#) Yujia Xie, Luowei Zhou, Xiyang Dai, Lu Yuan, Nguyen Bach, Ce Liu, Michael Zeng
- [APT-36K: A Large-scale Benchmark for Animal Pose Estimation and Tracking](#) Yuxiang Yang, Junjie Yang, Yufei Xu, Jing Zhang, Long Lan, Dacheng Tao
- [ShuffleMixer: An Efficient ConvNet for Image Super-Resolution](#) Long Sun, Jinshan Pan, Jinhui Tang
- [Neural Topological Ordering for Computation Graphs](#) Mukul Gagrani, Corrado Rainone, Yang Yang, Harris Teague, Wonseok Jeon, Roberto

- [Probable Domain Generalization via Quantile Risk Minimization](#) Cian Eastwood, Alexander Robey, Shashank Singh, Julius von Kügelgen, Hamed Hassani, George J. Pappas, Bernhard Schölkopf
- [Locating and Editing Factual Associations in GPT](#) Kevin Meng, David Bau, Alex Andonian, Yonatan Belinkov
- [S-PIFu: Integrating Parametric Human Models with PIFu for Single-view Clothed Human Reconstruction](#) Kennard Chan, Guosheng Lin, Haiyu Zhao, Weisi Lin
- [A composable machine-learning approach for steady-state simulations on high-resolution grids](#) Rishikesh Ranade, Chris Hill, Lalit Ghule, Jay Pathak
- [Outlier Suppression: Pushing the Limit of Low-bit Transformer Language Models](#) Xiuying Wei, Yunchen Zhang, Xiangguo Zhang, Ruihao Gong, Shanghang Zhang, Qi Zhang, Fengwei Yu, Xianglong Liu
- [A Single-timescale Analysis for Stochastic Approximation with Multiple Coupled Sequences](#) Han Shen, Tianyi Chen
- [Bayesian Risk Markov Decision Processes](#) Yifan Lin, Yuxuan Ren, Enlu Zhou
- [All Politics is Local: Redistricting via Local Fairness](#) Shao-Heng Ko, Erin Taylor, Pankaj Agarwal, Kamesh Munagala
- [Confident Adaptive Language Modeling](#) Tal Schuster, Adam Fisch, Jai Gupta, Mostafa Dehghani, Dara Bahri, Vinh Tran, Yi Tay, Donald Metzler
- [Submodular Maximization in Clean Linear Time](#) Wenxin Li, Moran Feldman, Ehsan Kazemi, Amin Karbasi
- [Parameters or Privacy: A Provable Tradeoff Between Overparameterization and Membership Inference](#) Jasper Tan, Blake Mason, Hamid Javadi, Richard Baraniuk
- [EF-BV: A Unified Theory of Error Feedback and Variance Reduction Mechanisms for Biased and Unbiased Compression in Distributed Optimization](#) Laurent Condat, Kai Yi, Peter Richtarik
- [DataMUX: Data Multiplexing for Neural Networks](#) Vishvak Murahari, Carlos Jimenez, Runzhe Yang, Karthik Narasimhan
- [Biologically-plausible backpropagation through arbitrary timespans via local neuromodulators](#) Yuhan Helena Liu, Stephen Smith, Stefan Mihalas, Eric Shea-Brown, Uygur Sümbül
- [Revisiting Realistic Test-Time Training: Sequential Inference and Adaptation by Anchored Clustering](#) Yongyi Su, Xun Xu, Kui Jia
- [Label Noise in Adversarial Training: A Novel Perspective to Study Robust Overfitting](#) Chengyu Dong, Liyuan Liu, Jingbo Shang
- [Active Labeling: Streaming Stochastic Gradients](#) Vivien Cabannes, Francis Bach, Vianney Perchet, Alessandro Rudi
- [CEBaB: Estimating the Causal Effects of Real-World Concepts on NLP Model Behavior](#) Eldar D Abraham, Karel D'Oosterlinck, Amir Feder, Yair Gat, Atticus Geiger, Christopher Potts, Roi Reichart, Zhengxuan Wu
- [TOIST: Task Oriented Instance Segmentation Transformer with Noun-Pronoun Distillation](#) Pengfei Li, Beiwen Tian, Yongliang Shi, Xiaoxue Chen, Hao Zhao, Guyue Zhou, Ya-Qin Zhang
- [Mind the Gap: Understanding the Modality Gap in Multi-modal Contrastive Representation Learning](#) Victor Weixin Liang, Yuhui Zhang, Yongchan Kwon, Serena Yeung, James Y. Zou
- [\$\alpha\$ -ReQ : Assessing Representation Quality in Self-Supervised Learning by measuring eigenspectrum decay](#) Kumar K Agrawal, Arnab Kumar Mondal, Arna Ghosh, Blake Richards
- [When Combinatorial Thompson Sampling meets Approximation Regret](#) Pierre Perrault
- [Pruning has a disparate impact on model accuracy](#) Cuong Tran, Ferdinando Fioretto, Jung-Eun Kim, Rakshit Naidu
- [Sequence Model Imitation Learning with Unobserved Contexts](#) Gokul Swamy, Sanjiban Choudhury, J. Bagnell, Steven Z. Wu
- [Tikhonov Regularization is Optimal Transport Robust under Martingale Constraints](#) Jiajin Li, Sirui Lin, Jose Blanchet, Viet Anh Nguyen
- [Policy Optimization with Linear Temporal Logic Constraints](#) Cameron Voloshin, Hoang Le, Swarat Chaudhuri, Yisong Yue
- [Merging Models with Fisher-Weighted Averaging](#) Michael S Matena, Colin A. Raffel
- [Learning to Generate Inversion-Resistant Model Explanations](#) Hoyong Jeong, Suyoung Lee, Sung Ju Hwang, Sooel Son
- [Unsupervised Multi-View Object Segmentation Using Radiance Field Propagation](#) Xinhang Liu, Jiaben Chen, Huai Yu, Yu-Wing Tai, Chi-Keung Tang
- [Beyond Mahalanobis Distance for Textual OOD Detection](#) Pierre Colombo, Eduardo Dadalto, Guillaume Staerman, Nathan Noiry, Pablo Piantanida
- [FasterRisk: Fast and Accurate Interpretable Risk Scores](#) Jiachang Liu, Chudi Zhong, Boxuan Li, Margo Seltzer, Cynthia Rudin
- [Graph Learning Assisted Multi-Objective Integer Programming](#) Yaixin Wu, Wen Song, Zhiguang Cao, Jie Zhang, Abhishek Gupta, Mingyan Lin
- [Revisiting Sliced Wasserstein on Images: From Vectorization to Convolution](#) Khai Nguyen, Nhat Ho
- [SignRFF: Sign Random Fourier Features](#) Xiaoyun Li, Ping Li
- [The Role of Baselines in Policy Gradient Optimization](#) Jincheng Mei, Wesley Chung, Valentin Thomas, Bo Dai, Csaba Szepesvari, Dale Schuurmans
- [A Rotated Hyperbolic Wrapped Normal Distribution for Hierarchical Representation Learning](#) Seunghyuk Cho, Juyong Lee, Jaesik Park, Dongwoo Kim
- [Private Graph All-Pairwise-Shortest-Path Distance Release with Improved Error Rate](#) Chenglin Fan, Ping Li, Xiaoyun Li
- [Multi-Fidelity Best-Arm Identification](#) Riccardo Poiani, Alberto Maria Metelli, Marcello Restelli
- [SemiFL: Semi-Supervised Federated Learning for Unlabeled Clients with Alternate Training](#) Enmao Diao, Jie Ding, Vahid Tarokh
- [RankFeat: Rank-1 Feature Removal for Out-of-distribution Detection](#) Yue Song, Nicu Sebe, Wei Wang
- [Jump Self-attention: Capturing High-order Statistics in Transformers](#) Haoyi Zhou, Siyang Xiao, Shanghang Zhang, Jieqi Peng, Shuai

- [Learning Infinite-Horizon Average-Reward Restless Multi-Action Bandits via Index Awareness](#) GUOJUN XIONG, Shufan Wang, Jian Li
- [STNDT: Modeling Neural Population Activity with Spatiotemporal Transformers](#) Trung Le, Eli Shlizerman
- [ProtoVAE: A Trustworthy Self-Explainable Prototypical Variational Model](#) Srishti Gautam, AHCÈNE BOUBEKKI, Stine Hansen, Suaiba Salahuddin, Robert Jenssen, Marina Höhne, Michael Kampffmeyer
- [If Influence Functions are the Answer, Then What is the Question?](#) Juhan Bae, Nathan Ng, Alston Lo, Marzyeh Ghassemi, Roger B. Grosse
- [FACT: Learning Governing Abstractions Behind Integer Sequences](#) Peter Belcak, Ard Kastrati, Flavio Schenker, Roger Wattenhofer
- [SAPipe: Staleness-Aware Pipeline for Data Parallel DNN Training](#) Yangrui Chen, Cong Xie, Meng Ma, Juncheng Gu, Yanghua Peng, Haibin Lin, Chuan Wu, Yibo Zhu
- [Probing Classifiers are Unreliable for Concept Removal and Detection](#) Abhinav Kumar, Chenhao Tan, Amit Sharma
- [SCL-WC: Cross-Slide Contrastive Learning for Weakly-Supervised Whole-Slide Image Classification](#) Xiyue Wang, Jinxi Xiang, Jun Zhang, Sen Yang, Zhongyi Yang, Ming-Hui Wang, Jing Zhang, Wei Yang, Junzhou Huang, Xiao Han
- [Reproducibility in Optimization: Theoretical Framework and Limits](#) Kwangjun Ahn, Prateek Jain, Ziwei Ji, Satyen Kale, Praneeth Netrapalli, Gil I Shamir
- [Hierarchical classification at multiple operating points](#) Jack Valmadre
- [Kernel Multimodal Continuous Attention](#) Alexander Moreno, Zhenke Wu, Supriya Nagesh, Walter Dempsey, James M. Rehg
- [VisCo Grids: Surface Reconstruction with Viscosity and Coarea Grids](#) Albert Pumarola, Artsiom Sanakoyeu, Lior Yariv, Ali Thabet, Yaron Lipman
- [Exact Shape Correspondence via 2D graph convolution](#) Barakeel Fanseu Kamhoua, Lin Zhang, Yongqiang Chen, Han Yang, MA Kaili, Bo Han, Bo Li, James Cheng
- [Posterior Matching for Arbitrary Conditioning](#) Ryan Strauss, Junier B Oliva
- [CARD: Classification and Regression Diffusion Models](#) Xizewen Han, Huangjie Zheng, Mingyuan Zhou
- [What Can the Neural Tangent Kernel Tell Us About Adversarial Robustness?](#) Nikolaos Tsilivis, Julia Kempe
- [Benefits of Permutation-Equivariance in Auction Mechanisms](#) Tian Qin, Fengxiang He, Dingfeng Shi, Wenbing Huang, Dacheng Tao
- [MoCoDA: Model-based Counterfactual Data Augmentation](#) Silviu Pitis, Elliot Creager, Ajay Mandlekar, Animesh Garg
- [SkinCon: A skin disease dataset densely annotated by domain experts for fine-grained debugging and analysis](#) Roxana Daneshjou, Mert Yuksekgonul, Zhuo Ran Cai, Roberto Novoa, James Y. Zou
- [A permutation-free kernel two-sample test](#) Shubhanshu Shekhar, Ilmun Kim, Aaditya Ramdas
- [Simple Unsupervised Object-Centric Learning for Complex and Naturalistic Videos](#) Gautam Singh, Yi-Fu Wu, Sungjin Ahn
- [Communication-efficient distributed eigenspace estimation with arbitrary node failures](#) Vasileios Charisopoulos, Anil Damle
- [On Uncertainty, Tempering, and Data Augmentation in Bayesian Classification](#) Sanyam Kapoor, Wesley J. Maddox, Pavel Izmailov, Andrew G. Wilson
- [Optimal Binary Classification Beyond Accuracy](#) Shashank Singh, Justin T. Khim
- [Information-Theoretic GAN Compression with Variational Energy-based Model](#) Minsoo Kang, Hyewon Yoo, Eunhee Kang, Sehwan Ki, Hyong Euk Lee, Bohyung Han
- [On Learning and Refutation in Noninteractive Local Differential Privacy](#) Alexander Edmonds, Aleksandar Nikolov, Toniann Pitassi
- [Why So Pessimistic? Estimating Uncertainties for Offline RL through Ensembles, and Why Their Independence Matters](#) Kamyar Ghasemipour, Shixiang (Shane) Gu, Ofir Nachum
- [Private Synthetic Data for Multitask Learning and Marginal Queries](#) Giuseppe Vietri, Cedric Archambeau, Sergul Aydore, William Brown, Michael Kearns, Aaron Roth, Ankit Siva, Shuai Tang, Steven Z. Wu
- [Sample-Efficient Reinforcement Learning of Partially Observable Markov Games](#) Qinghua Liu, Csaba Szepesvari, Chi Jin
- [Advancing Model Pruning via Bi-level Optimization](#) Yihua Zhang, Yuguang Yao, Parikshit Ram, Pu Zhao, Tianlong Chen, Mingyi Hong, Yanzhi Wang, Sijia Liu
- [Optimal Positive Generation via Latent Transformation for Contrastive Learning](#) Yinqi Li, Hong Chang, Bingpeng MA, Shiguang Shan, Xilin Chen
- [MineDojo: Building Open-Ended Embodied Agents with Internet-Scale Knowledge](#) Linxi Fan, Guanzhi Wang, Yunfan Jiang, Ajay Mandlekar, Yuncong Yang, Haoyi Zhu, Andrew Tang, De-An Huang, Yuke Zhu, Anima Anandkumar
- [The price of unfairness in linear bandits with biased feedback](#) Solenne Gaucher, Alexandra Carpentier, Christophe Giraud
- [DeepFoids: Adaptive Bio-Inspired Fish Simulation with Deep Reinforcement Learning](#) Yuko Ishiwaka, Xiao Zeng, Shun Ogawa, Donovan Westwater, Tadayuki Tone, Masaki Nakada
- [Truncated Matrix Power Iteration for Differentiable DAG Learning](#) Zhen Zhang, Ignavier Ng, Dong Gong, Yuhang Liu, Ehsan Abbasnejad, Mingming Gong, Kun Zhang, Javen Qinfeng Shi
- [Learning Debiased Classifier with Biased Committee](#) Nayeong Kim, SEHYUN HWANG, Sungsoo Ahn, Jaesik Park, Suha Kwak
- [Modeling the Machine Learning Multiverse](#) Samuel J. Bell, Onno Kampman, Jesse Dodge, Neil Lawrence
- [Label-Aware Global Consistency for Multi-Label Learning with Single Positive Labels](#) Ming-Kun Xie, Jiahao Xiao, Sheng-Jun Huang
- [Unifying Voxel-based Representation with Transformer for 3D Object Detection](#) Yanwei Li, Yilun Chen, Xiaojuan Qi, Zeming Li, Jian Sun, Jiaya Jia
- [UDC: Unified DNAS for Compressible TinyML Models for Neural Processing Units](#) Igor Fedorov, Ramon Matas, Hokchhay Tann, Chuteng Zhou, Matthew Mattina, Paul Whatmough

- [Fine-Grained Analysis of Stability and Generalization for Modern Meta Learning Algorithms](#) Jiechao Guan, Yong Liu, Zhiwu Lu
- [On Scrambling Phenomena for Randomly Initialized Recurrent Networks](#) Vaggos Chatziafratis, Ioannis Panageas, Clayton Sanford, Stelios Stavroulakis
- [Learning to Branch with Tree MDPs](#) Lara Scavuzzo, Feng Chen, Didier Chetelat, Maxime Gasse, Andrea Lodi, Neil Yorke-Smith, Karen Aardal
- [Neural Sheaf Diffusion: A Topological Perspective on Heterophily and Oversmoothing in GNNs](#) Cristian Bodnar, Francesco Di Giovanni, Benjamin Chamberlain, Pietro Lió, Michael Bronstein
- [pyKT: A Python Library to Benchmark Deep Learning based Knowledge Tracing Models](#) Zitao Liu, Qiongqiong Liu, Jiahao Chen, Shuyan Huang, Jiliang Tang, Weiqi Luo
- [Adversarial Unlearning: Reducing Confidence Along Adversarial Directions](#) Amrith Setlur, Benjamin Eysenbach, Virginia Smith, Sergey Levine
- [How Would The Viewer Feel? Estimating Wellbeing From Video Scenarios](#) Mantas Mazeika, Eric Tang, Andy Zou, Steven Basart, Jun Shern Chan, Dawn Song, David Forsyth, Jacob Steinhardt, Dan Hendrycks
- [On Elimination Strategies for Bandit Fixed-Confidence Identification](#) Andrea Tirinzoni, Rémy Degenne
- [When Adversarial Training Meets Vision Transformers: Recipes from Training to Architecture](#) Yichuan Mo, Dongxian Wu, Yifei Wang, Yiwen Guo, Yisen Wang
- [Escaping from the Barren Plateau via Gaussian Initializations in Deep Variational Quantum Circuits](#) Kaining Zhang, Liu Liu, Min-Hsiu Hsieh, Dacheng Tao
- [A sharp NMF result with applications in network modeling](#) Jiashun Jin
- [Decoupling Classifier for Boosting Few-shot Object Detection and Instance Segmentation](#) Bin-Bin Gao, Xiaochen Chen, Zhongyi Huang, Congchong Nie, Jun Liu, Jinxiang Lai, GUANNAN JIANG, Xi Wang, Chengjie Wang
- [Private Estimation with Public Data](#) Alex Bie, Gautam Kamath, Vikrant Singhal
- [Pre-activation Distributions Expose Backdoor Neurons](#) Runkai Zheng, Rongjun Tang, Jianze Li, Li Liu
- [DaDA: Distortion-aware Domain Adaptation for Unsupervised Semantic Segmentation](#) Sujin Jang, Joohan Na, Dokwan Oh
- [Large-batch Optimization for Dense Visual Predictions: Training Faster R-CNN in 4.2 Minutes](#) Zeyue Xue, Jianming Liang, Guanglu Song, Zhuofan Zong, Liang Chen, Yu Liu, Ping Luo
- [Most Activation Functions Can Win the Lottery Without Excessive Depth](#) Rebekka Burkholz
- [Rapid Model Architecture Adaption for Meta-Learning](#) Yiren Zhao, Xitong Gao, I Shumailov, Nicolo Fusi, Robert Mullins
- [Exploring through Random Curiosity with General Value Functions](#) Aditya Ramesh, Louis Kirsch, Sjoerd van Steenkiste, Jürgen Schmidhuber
- [Optimal Transport-based Identity Matching for Identity-invariant Facial Expression Recognition](#) Daeha Kim, Byung Cheol Song
- [Unsupervised Learning under Latent Label Shift](#) Manley Roberts, Pranav Mani, Saurabh Garg, Zachary Lipton
- [SHINE: SubHypergraph Inductive Neural nEtwork](#) Yuan Luo
- [Efficient Aggregated Kernel Tests using Incomplete \$U\$ -statistics](#) Antonin Schrab, Ilmun Kim, Benjamin Guedj, Arthur Gretton
- [Influencing Long-Term Behavior in Multiagent Reinforcement Learning](#) Dong-Ki Kim, Matthew Riemer, Miao Liu, Jakob Foerster, Michael Everett, Chuangchuang Sun, Gerald Tesauero, Jonathan P. How
- [Quantized Training of Gradient Boosting Decision Trees](#) Yu Shi, Guolin Ke, Zhuoming Chen, Shuxin Zheng, Tie-Yan Liu
- [Convergent Representations of Computer Programs in Human and Artificial Neural Networks](#) Shashank Srikant, Ben Lipkin, Anna Ivanova, Evelina Fedorenko, Una-May O'Reilly
- [QC-StyleGAN - Quality Controllable Image Generation and Manipulation](#) Dat Viet Thanh Nguyen, Phong Tran The, Tan M. Dinh, Cuong Pham, Anh Tran
- [Retrieve, Reason, and Refine: Generating Accurate and Faithful Patient Instructions](#) Fenglin Liu, Bang Yang, Chenyu You, Xian Wu, Shen Ge, Zhangdaihong Liu, Xu Sun, Yang Yang, David Clifton
- [Data Distributional Properties Drive Emergent In-Context Learning in Transformers](#) Stephanie Chan, Adam Santoro, Andrew Lampinen, Jane Wang, Aaditya Singh, Pierre Richemond, James McClelland, Felix Hill
- [Algorithms that Approximate Data Removal: New Results and Limitations](#) Vinith Suriyakumar, Ashia C. Wilson
- [LOT: Layer-wise Orthogonal Training on Improving l2 Certified Robustness](#) Xiaojun Xu, Linyi Li, Bo Li
- [Lottery Tickets on a Data Diet: Finding Initializations with Sparse Trainable Networks](#) Mansheej Paul, Brett Larsen, Surya Ganguli, Jonathan Frankle, Gintare Karolina Dziugaite
- [Regularized Molecular Conformation Fields](#) Lihao Wang, Yi Zhou, Yiqun Wang, Xiaoqing Zheng, Xuanjing Huang, Hao Zhou
- [Reconstruction on Trees and Low-Degree Polynomials](#) Frederic Koehler, Elchanan Mossel
- [Lower Bounds and Nearly Optimal Algorithms in Distributed Learning with Communication Compression](#) Xinmeng Huang, Yiming Chen, Wotao Yin, Kun Yuan
- [Memory safe computations with XLA compiler](#) Artem Artemev, Yuze An, Tilman Roeder, Mark van der Wilk
- [Towards Theoretically Inspired Neural Initialization Optimization](#) Yibo Yang, Hong Wang, Haobo Yuan, Zhouchen Lin
- [AnimeRun: 2D Animation Visual Correspondence from Open Source 3D Movies](#) Li Siyao, Yuhang Li, Bo Li, Chao Dong, Ziwei Liu, Chen Change Loy
- [Uncertainty Estimation Using Riemannian Model Dynamics for Offline Reinforcement Learning](#) Guy Tennenholtz, Shie Mannor

- [Improved Convergence Rate of Stochastic Gradient Langevin Dynamics with Variance Reduction and its Application to Optimization](#) Yuri Kinoshita, Taiji Suzuki
- [Practical Adversarial Attacks on Spatiotemporal Traffic Forecasting Models](#) Fan LIU, Hao Liu, Wenzhao Jiang
- [Efficient learning of nonlinear prediction models with time-series privileged information](#) Bastian Jung, Fredrik D. Johansson
- [Layer Freezing & Data Sieving: Missing Pieces of a Generic Framework for Sparse Training](#) Geng Yuan, Yanyu Li, Sheng Li, Zhenglun Kong, Sergey Tulyakov, Xulong Tang, Yanzhi Wang, Jian Ren
- [ETAB: A Benchmark Suite for Visual Representation Learning in Echocardiography](#) Ahmed M. Alaa, Anthony Philippakis, David Sontag
- [Trustworthy Monte Carlo](#) Juha Harviainen, Mikko Koivisto, Petteri Kaski
- [Double Bubble, Toil and Trouble: Enhancing Certified Robustness through Transitivity](#) Andrew Cullen, Paul Montague, Shijie Liu, Sarah Erfani, Benjamin Rubinstein
- [An \$\alpha\$ -No-Regret Algorithm For Graphical Bilinear Bandits](#) Geovani Rizk, Igor Colin, Albert Thomas, Rida Laraki, Yann Chevaleyre
- [A Unified Analysis of Federated Learning with Arbitrary Client Participation](#) Shiqiang Wang, Mingyue Ji
- [Deconfounded Representation Similarity for Comparison of Neural Networks](#) Tianyu Cui, Yogesh Kumar, Pekka Marttinen, Samuel Kaski
- [Fairness without Demographics through Knowledge Distillation](#) Junyi Chai, Taeuk Jang, Xiaoqian Wang
- [Sleeping Agent: Scalable Hidden Trigger Backdoors for Neural Networks Trained from Scratch](#) Hossein Souri, Liam Fowl, Rama Chellappa, Micah Goldblum, Tom Goldstein
- [Dynamic Pricing with Monotonicity Constraint under Unknown Parametric Demand Model](#) Su Jia, Andrew Li, R Ravi
- [A Win-win Deal: Towards Sparse and Robust Pre-trained Language Models](#) Yuanxin Liu, Fandong Meng, Zheng Lin, Jiangnan Li, Peng Fu, Yanan Cao, Weiping Wang, Jie Zhou
- [Asymptotics of smoothed Wasserstein distances in the small noise regime](#) Yunzi Ding, Jonathan Niles-Weed
- [Fine-tuning Language Models over Slow Networks using Activation Quantization with Guarantees](#) Jue WANG, Binhang Yuan, Luka Rimanic, Yongjun He, Tri Dao, Beidi Chen, Christopher Ré, Ce Zhang
- [Pareto Set Learning for Expensive Multi-Objective Optimization](#) Xi Lin, Zhiyuan Yang, Xiaoyuan Zhang, Qingfu Zhang
- [Non-monotonic Resource Utilization in the Bandits with Knapsacks Problem](#) Raunak Kumar, Robert Kleinberg
- [Efficient identification of informative features in simulation-based inference](#) Jonas Beck, Michael Deistler, Yves Bernaerts, Jakob H Macke, Philipp Berens
- [Agreement-on-the-line: Predicting the Performance of Neural Networks under Distribution Shift](#) Christina Baek, Yiding Jiang, Aditi Raghunathan, J. Zico Kolter
- [Large-Scale Differentiable Causal Discovery of Factor Graphs](#) Romain Lopez, Jan-Christian Huetter, Jonathan Pritchard, Aviv Regev
- [Diagnosing failures of fairness transfer across distribution shift in real-world medical settings](#) Jessica Schrouff, Natalie Harris, Sanmi Koyejo, Ibrahim M. Alabdulmohsin, Eva Schnider, Krista Opsahl-Ong, Alexander Brown, Subhrajit Roy, Diana Mincu, Christina Chen, Awa Dieng, Yuan Liu, Vivek Natarajan, Alan Karthikesalingam, Katherine A. Heller, Silvia Chiappa, Alexander D'Amour
- [Towards Lightweight Black-Box Attack Against Deep Neural Networks](#) Chenghao Sun, Yonggang Zhang, Wan Chaoqun, Qizhou Wang, Ya Li, Tongliang Liu, Bo Han, Xinmei Tian
- [Federated Learning from Pre-Trained Models: A Contrastive Learning Approach](#) Yue Tan, Guodong Long, Jie Ma, LU LIU, Tianyi Zhou, Jing Jiang
- [Theoretically Provable Spiking Neural Networks](#) Shao-Qun Zhang, Zhi-Hua Zhou
- [Approximate Euclidean lengths and distances beyond Johnson-Lindenstrauss](#) Aleksandros Sobczyk, Mathieu Luisier
- [InsPro: Propagating Instance Query and Proposal for Online Video Instance Segmentation](#) Fei He, Haoyang Zhang, Naiyu Gao, Jian Jia, Yanhu Shan, Xin Zhao, Kaiqi Huang
- [Measuring and Reducing Model Update Regression in Structured Prediction for NLP](#) Deng Cai, Elman Mansimov, Yi-An Lai, Yixuan Su, Lei Shu, Yi Zhang
- [Rethinking Lipschitz Neural Networks and Certified Robustness: A Boolean Function Perspective](#) Bohang Zhang, Du Jiang, Di He, Liwei Wang
- [Multivariate Time-Series Forecasting with Temporal Polynomial Graph Neural Networks](#) Yijing Liu, Qinxian Liu, Jian-Wei Zhang, Haozhe Feng, Zhongwei Wang, Zihan Zhou, Wei Chen
- [Few-shot Image Generation via Adaptation-Aware Kernel Modulation](#) Yunqing Zhao, Keshigeyan Chandrasegaran, Milad Abdollahzadeh, Ngai-Man (Man) Cheung
- [Learning to Follow Instructions in Text-Based Games](#) Mathieu Tuli, Andrew Li, Pashootan Vaezipoor, Toryn Klassen, Scott Sanner, Sheila McIlraith
- [Concentration of Data Encoding in Parameterized Quantum Circuits](#) Guangxi Li, Ruilin Ye, Xuanqiang Zhao, Xin Wang
- [Improving Variational Autoencoders with Density Gap-based Regularization](#) Jianfei Zhang, Jun Bai, Chenghua Lin, Yanmeng Wang, Wenge Rong
- [RISE: Robust Individualized Decision Learning with Sensitive Variables](#) Xiaoqing Tan, Zhengling Qi, Christopher Seymour, Lu Tang
- [Bayesian Clustering of Neural Spiking Activity Using a Mixture of Dynamic Poisson Factor Analyzers](#) Ganchao Wei, Ian H Stevenson, Xiaojing Wang
- [Understanding Deep Contrastive Learning via Coordinate-wise Optimization](#) Yuandong Tian
- [Beyond neural scaling laws: beating power law scaling via data pruning](#) Ben Sorscher, Robert Geirhos, Shashank Shekhar, Surya Ganguli,

- [Neural Estimation of Submodular Functions with Applications to Differentiable Subset Selection](#) Abir De, Soumen Chakrabarti
- [Maximum Class Separation as Inductive Bias in One Matrix](#) Tejaswi Kasarla, Gertjan Burghouts, Max van Spengler, Elise van der Pol, Rita Cucchiara, Pascal Mettes
- [Surprising Instabilities in Training Deep Networks and a Theoretical Analysis](#) Yuxin Sun, DONG LAO, Ganesh Sundaramoorthi, Anthony Yezzi
- [WaveBound: Dynamic Error Bounds for Stable Time Series Forecasting](#) Youngin Cho, Daejin Kim, DONGMIN KIM, MOHAMMAD AZAM KHAN, Jaegul Choo
- [Finite-Time Analysis of Adaptive Temporal Difference Learning with Deep Neural Networks](#) Tao Sun, Dongsheng Li, Bao Wang
- [Benign Underfitting of Stochastic Gradient Descent](#) Tomer Koren, Roi Livni, Yishay Mansour, Uri Sherman
- [A Geometric Perspective on Variational Autoencoders](#) Clément Chadebec, Stephanie Allassonniere
- [A Data-Augmentation Is Worth A Thousand Samples: Analytical Moments And Sampling-Free Training](#) Randall Balestriero, Ishan Misra, Yann LeCun
- [Effective Adaptation in Multi-Task Co-Training for Unified Autonomous Driving](#) Xiwen Liang, Yangxin Wu, Jianhua Han, Hang Xu, Chunjing XU, Xiaodan Liang
- [Redundant representations help generalization in wide neural networks](#) Diego Doimo, Aldo Glielmo, Sebastian Goldt, Alessandro Laio
- [Mean Estimation in High-Dimensional Binary Markov Gaussian Mixture Models](#) Yihan Zhang, Nir Weinberger
- [Rate-Distortion Theoretic Bounds on Generalization Error for Distributed Learning](#) Milad Sefidgaran, Romain Chor, Abdellatif Zaidi
- [Online Frank-Wolfe with Arbitrary Delays](#) Yuanyu Wan, Wei-Wei Tu, Lijun Zhang
- [Isometric 3D Adversarial Examples in the Physical World](#) yibo miao, Yinpeng Dong, Jun Zhu, Xiao-Shan Gao
- [Tight Lower Bounds on Worst-Case Guarantees for Zero-Shot Learning with Attributes](#) Alessio Mazzetto, Cristina Menghini, Andrew Yuan, Eli Upfal, Stephen Bach
- [Graph Neural Networks with Adaptive Readouts](#) David Buterez, Jon Paul Janet, Steven J Kiddle, Dino Oglic, Pietro Liò
- [Robust Imitation of a Few Demonstrations with a Backwards Model](#) Jung Yeon Park, Lawson Wong
- [Communication Efficient Distributed Learning for Kernelized Contextual Bandits](#) Chuanhao Li, Huazheng Wang, Mengdi Wang, Hongning Wang
- [Trust Region Policy Optimization with Optimal Transport Discrepancies: Duality and Algorithm for Continuous Actions](#) Antonio Terpin, Nicolas Lanzetti, Batuhan Yardim, Florian Dorfler, Giorgia Ramponi
- [ELIAS: End-to-End Learning to Index and Search in Large Output Spaces](#) Nilesh Gupta, Patrick Chen, Hsiang-Fu Yu, Cho-Jui Hsieh, Inderjit Dhillon
- [GStarX: Explaining Graph Neural Networks with Structure-Aware Cooperative Games](#) Shichang Zhang, Yozen Liu, Neil Shah, Yizhou Sun
- [Bridging the Gap from Asymmetry Tricks to Decorrelation Principles in Non-contrastive Self-supervised Learning](#) Kang-Jun Liu, Masanori Suganuma, Takayuki Okatani
- [Lipschitz Bandits with Batched Feedback](#) Yasong Feng, zengfeng Huang, Tianyu Wang
- [Improving Barely Supervised Learning by Discriminating Unlabeled Samples with Super-Class](#) Guan Gui, Zhen Zhao, Lei Qi, Luping Zhou, Lei Wang, Yinghuan Shi
- [Low-Rank Modular Reinforcement Learning via Muscle Synergy](#) Heng Dong, Tonghan Wang, Jiayuan Liu, Chongjie Zhang
- [Neural Temporal Walks: Motif-Aware Representation Learning on Continuous-Time Dynamic Graphs](#) Ming Jin, Yuan-Fang Li, Shirui Pan
- [Understanding Benign Overfitting in Gradient-Based Meta Learning](#) Lisha Chen, Songtao Lu, Tianyi Chen
- [Generative Neural Articulated Radiance Fields](#) Alexander Bergman, Petr Kellnhofer, Wang Yifan, Eric Chan, David Lindell, Gordon Wetzstein
- [AutoMS: Automatic Model Selection for Novelty Detection with Error Rate Control](#) Yifan Zhang, Haiyan Jiang, Haojie Ren, Changliang Zou, Dejing Dou
- [GALOIS: Boosting Deep Reinforcement Learning via Generalizable Logic Synthesis](#) Yushi Cao, Zhiming Li, Tianpei Yang, Hao Zhang, YAN ZHENG, Yi Li, Jianye Hao, Yang Liu
- [Faster Deep Reinforcement Learning with Slower Online Network](#) Kavosh Asadi, Rasool Fakoor, Omer Gottesman, Taesup Kim, Michael Littman, Alexander J. Smola
- [Learn to Match with No Regret: Reinforcement Learning in Markov Matching Markets](#) Yifei Min, Tianhao Wang, Ruitu Xu, Zhaoran Wang, Michael Jordan, Zhuoran Yang
- [Efficient Frameworks for Generalized Low-Rank Matrix Bandit Problems](#) Yue Kang, Cho-Jui Hsieh, Thomas Chun Man Lee
- [A Projection-free Algorithm for Constrained Stochastic Multi-level Composition Optimization](#) Tesi Xiao, Krishnakumar Balasubramanian, Saeed Ghadimi
- [Green Hierarchical Vision Transformer for Masked Image Modeling](#) Lang Huang, Shan You, Mingkai Zheng, Fei Wang, Chen Qian, Toshihiko Yamasaki
- [Simultaneous Missing Value Imputation and Structure Learning with Groups](#) Pablo Morales-Alvarez, Wenbo Gong, Angus Lamb, Simon Woodhead, Simon Peyton Jones, Nick Pawlowski, Miltiadis Allamanis, Cheng Zhang
- [Towards Reliable Simulation-Based Inference with Balanced Neural Ratio Estimation](#) Arnaud Delaunoy, Joeri Hermans, François Rozet, Antoine Wehenkel, Gilles Louppe
- [Interaction Modeling with Multiplex Attention](#) Fan-Yun Sun, Isaac Kauvar, Ruohan Zhang, Jiachen Li, Mykel J Kochenderfer, Jiajun Wu, ...

- [Low-rank lottery tickets: finding efficient low-rank neural networks via matrix differential equations](#) Steffen Schotthöfer, Emanuele Zangrando, Jonas Kusch, Gianluca Ceruti, Francesco Tudisco
- [Causality-driven Hierarchical Structure Discovery for Reinforcement Learning](#) shaohui peng, Xing Hu, Rui Zhang, Ke Tang, Jiaming Guo, Qi Yi, Ruizhi Chen, xishan zhang, Zidong Du, Ling Li, Qi Guo, Yunji Chen
- [Pay attention to your loss : understanding misconceptions about Lipschitz neural networks](#) Louis Béthune, Thibaut Boissin, Mathieu Serrurier, Franck Mamalet, Corentin Friedrich, Alberto Gonzalez Sanz
- [Large-Scale Retrieval for Reinforcement Learning](#) Peter Humphreys, Arthur Guez, Olivier Tieleman, Laurent Sifre, Theophane Weber, Timothy Lillicrap
- [Gradient flow dynamics of shallow ReLU networks for square loss and orthogonal inputs](#) Etienne Boursier, Loucas PILLAUD-VIVIEN, Nicolas Flammarion
- [Weisfeiler and Leman Go Walking: Random Walk Kernels Revisited](#) Nils M. Kriege
- [Outsourcing Training without Uploading Data via Efficient Collaborative Open-Source Sampling](#) Junyuan Hong, Lingjuan Lyu, Jiayu Zhou, Michael Spranger
- [Multi-agent Dynamic Algorithm Configuration](#) Ke Xue, Jiacheng Xu, Lei Yuan, Miqing Li, Chao Qian, Zongzhang Zhang, Yang Yu
- [TaSIL: Taylor Series Imitation Learning](#) Daniel Pfrommer, Thomas Zhang, Stephen Tu, Nikolai Matni
- [Experimental Design for Linear Functionals in Reproducing Kernel Hilbert Spaces](#) Mojmir Mutny, Andreas Krause
- [Continuous MDP Homomorphisms and Homomorphic Policy Gradient](#) Sahand Rezaei-Shoshtari, Rosie Zhao, Prakash Panangaden, David Meger, Doina Precup
- [Score-based Generative Modeling Secretly Minimizes the Wasserstein Distance](#) Dohyun Kwon, Ying Fan, Kangwook Lee
- [Uncertainty-Aware Reinforcement Learning for Risk-Sensitive Player Evaluation in Sports Game](#) Guiliang Liu, Yudong Luo, Oliver Schulte, Pascal Poupart
- [End-to-end Algorithm Synthesis with Recurrent Networks: Extrapolation without Overthinking](#) Arpit Bansal, Avi Schwarzschild, Eitan Borgnia, Zeyad Emam, Furong Huang, Micah Goldblum, Tom Goldstein
- [Smooth Fictitious Play in Stochastic Games with Perturbed Payoffs and Unknown Transitions](#) Lucas Baudin, Rida Laraki
- [OOD Link Prediction Generalization Capabilities of Message-Passing GNNs in Larger Test Graphs](#) Yangze Zhou, Gitta Kutyniok, Bruno Ribeiro
- [Algorithms with Prediction Portfolios](#) Michael Dinitz, Sungjin Im, Thomas Lavastida, Benjamin Moseley, Sergei Vassilvitskii
- [A Unified Hard-Constraint Framework for Solving Geometrically Complex PDEs](#) Songming Liu, Hao Zhongkai, Chengyang Ying, Hang Su, Jun Zhu, Ze Cheng
- [Instability and Local Minima in GAN Training with Kernel Discriminators](#) Evan Becker, Parthe Pandit, Sundeep Rangan, Alyson K. Fletcher
- [Versatile Multi-stage Graph Neural Network for Circuit Representation](#) Shuwen Yang, Zhihao Yang, Dong Li, Yingxueff Zhang, Zhanguang Zhang, Guojie Song, Jianye Hao
- [Non-Stationary Bandits under Recharging Payoffs: Improved Planning with Sublinear Regret](#) Orestis Papadigenopoulos, Constantine Caramanis, Sanjay Shakkottai
- [Optimal and Adaptive Monteiro-Svaiter Acceleration](#) Yair Carmon, Danielle Hausler, Arun Jambulapati, Yujia Jin, Aaron Sidford
- [Chroma-VAE: Mitigating Shortcut Learning with Generative Classifiers](#) Wanqian Yang, Polina Kirichenko, Micah Goldblum, Andrew G. Wilson
- [SparCL: Sparse Continual Learning on the Edge](#) Zifeng Wang, Zheng Zhan, Yifan Gong, Geng Yuan, Wei Niu, Tong Jian, Bin Ren, Stratis Ioannidis, Yanzhi Wang, Jennifer Dy
- [Adaptively Exploiting d-Separators with Causal Bandits](#) Blair Bilodeau, Linbo Wang, Dan Roy
- [Spectrum Random Masking for Generalization in Image-based Reinforcement Learning](#) Yangru Huang, Peixi Peng, Yifan Zhao, Guangyao Chen, Yonghong Tian
- [Contrastive Graph Structure Learning via Information Bottleneck for Recommendation](#) Chunyu Wei, Jian Liang, Di Liu, Fei Wang
- [Pyramid Attention For Source Code Summarization](#) Lei Chai, Ming Li
- [SIREN: Shaping Representations for Detecting Out-of-Distribution Objects](#) Xuefeng Du, Gabriel Gozum, Yifei Ming, Yixuan Li
- [CLOOB: Modern Hopfield Networks with InfoLOOB Outperform CLIP](#) Andreas Fürst, Elisabeth Rumetshofer, Johannes Lehner, Viet T. Tran, Fei Tang, Hubert Ramsauer, David Kreil, Michael Kopp, Günter Klambauer, Angela Bitto, Sepp Hochreiter
- [I2Q: A Fully Decentralized Q-Learning Algorithm](#) Jiechuan Jiang, Zongqing Lu
- [Planning to the Information Horizon of BAMDPs via Epistemic State Abstraction](#) Dilip Arumugam, Satinder Singh
- [AutoST: Towards the Universal Modeling of Spatio-temporal Sequences](#) Jianxin Li, Shuai Zhang, Hui Xiong, Haoyi Zhou
- [Constrained Langevin Algorithms with L-mixing External Random Variables](#) Yuping Zheng, Andrew Lamperski
- [Language Conditioned Spatial Relation Reasoning for 3D Object Grounding](#) Shizhe Chen, Pierre-Louis Guhur, Makarand Tapaswi, Cordelia Schmid, Ivan Laptev
- [MVP-N: A Dataset and Benchmark for Real-World Multi-View Object Classification](#) REN WANG, Jiayue Wang, Tae Sung Kim, JINSUNG KIM, Hyuk-Jae Lee
- [Data Augmentation for Compositional Data: Advancing Predictive Models of the Microbiome](#) Elliott Gordon-Rodriguez, Thomas Quinn, John P. Cunningham
- [On the Robustness of Deep Clustering Models: Adversarial Attacks and Defenses](#) Anshuman Chhabra, Ashwin Sekhari, Prasant

- [Exploiting the Relationship Between Kendall's Rank Correlation and Cosine Similarity for Attribution Protection](#) Fan Wang, Adams Wai Kin Kong
- [Wavelet Feature Maps Compression for Image-to-Image CNNs](#) Shahaf E. Finder, Yair Zohav, Maor Ashkenazi, Eran Treister
- [Biologically plausible solutions for spiking networks with efficient coding](#) Veronika Koren, Stefano Panzeri
- [Finding Optimal Arms in Non-stochastic Combinatorial Bandits with Semi-bandit Feedback and Finite Budget](#) Jasmin Brandt, Viktor Bengs, Björn Haddendorst, Eyke Hüllermeier
- [Graph Neural Networks are Dynamic Programmers](#) Andrew J Dudzik, Petar Veličković
- [Evaluated CMI Bounds for Meta Learning: Tightness and Expressiveness](#) Fredrik Hellström, Giuseppe Durisi
- [SPD: Synergy Pattern Diversifying Oriented Unsupervised Multi-agent Reinforcement Learning](#) Yuhang Jiang, Jianzhun Shao, Shuncheng He, Hongchang Zhang, Xiangyang Ji
- [Self-Aware Personalized Federated Learning](#) Huili Chen, Jie Ding, Eric W. Tramel, Shuang Wu, Anit Kumar Sahu, Salman Avestimehr, Tao Zhang
- [Additive MIL: Intrinsically Interpretable Multiple Instance Learning for Pathology](#) Syed Ashar Javed, Dinkar Juyal, Harshith Padigela, Amaro Taylor-Weiner, Limin Yu, Aaditya Prakash
- [Model-Based Imitation Learning for Urban Driving](#) Anthony Hu, Gianluca Corrado, Nicolas Griffiths, Zachary Murez, Corina Gurau, Hudson Yeo, Alex Kendall, Roberto Cipolla, Jamie Shotton
- [Online Training Through Time for Spiking Neural Networks](#) Mingqing Xiao, Qingyan Meng, Zongpeng Zhang, Di He, Zhouchen Lin
- [SCONE: Surface Coverage Optimization in Unknown Environments by Volumetric Integration](#) Antoine Guedon, Pascal Monasse, Vincent Lepetit
- [WebShop: Towards Scalable Real-World Web Interaction with Grounded Language Agents](#) Shunyu Yao, Howard Chen, John Yang, Karthik Narasimhan
- [Boosting Out-of-distribution Detection with Typical Features](#) Yao Zhu, YueFeng Chen, Chuanlong Xie, Xiaodan Li, Rong Zhang, Hui Xue, Xiang Tian, bolun zheng, Yaowu Chen
- [Dynamic Learning in Large Matching Markets](#) Anand Kalvit, Assaf Zeevi
- [Invariant and Transportable Representations for Anti-Causal Domain Shifts](#) Yibo Jiang, Victor Veitch
- [Symplectic Spectrum Gaussian Processes: Learning Hamiltonians from Noisy and Sparse Data](#) Yusuke Tanaka, Tomoharu Iwata, naonori ueda
- [Multiclass Learnability Beyond the PAC Framework: Universal Rates and Partial Concept Classes](#) Alkis Kalavasis, Grigoris Velegkas, Amin Karbasi
- [MetaTeacher: Coordinating Multi-Model Domain Adaptation for Medical Image Classification](#) Zhenbin Wang, Mao Ye, Xiatian Zhu, Liuhan Peng, Liang Tian, Yingying Zhu
- [Biological Learning of Irreducible Representations of Commuting Transformations](#) Alexander Genkin, David Lipshutz, Siavash Golkar, Tiberiu Tesileanu, Dmitri Chklovskii
- [A Fourier Approach to Mixture Learning](#) Mingda Qiao, Guru Guruganesh, Ankit Rawat, Kumar Avinava Dubey, Manzil Zaheer
- [Structured Energy Network As a Loss](#) Jay Yoon Lee, Dhruvesh Patel, Purujit Goyal, Wenlong Zhao, Zhiyang Xu, Andrew McCallum
- [Bayesian inference via sparse Hamiltonian flows](#) Naitong Chen, Zuheng Xu, Trevor Campbell
- [SAPA: Similarity-Aware Point Affiliation for Feature Upsampling](#) Hao Lu, Wenze Liu, Zixuan Ye, Hongtao Fu, Yuliang Liu, Zhiguo Cao
- [Losses Can Be Blessings: Routing Self-Supervised Speech Representations Towards Efficient Multilingual and Multitask Speech Processing](#) Yonggan Fu, Yang Zhang, Kaizhi Qian, Zhifan Ye, Zhongzhi Yu, Cheng-I Jeff Lai, Celine Lin
- [Gradient Methods Provably Converge to Non-Robust Networks](#) Gal Vardi, Gilad Yehudai, Ohad Shamir
- [Diversity vs. Recognizability: Human-like generalization in one-shot generative models](#) Victor Boutin, Lakshya Singhal, Xavier Thomas, Thomas Serre
- [Task-level Differentially Private Meta Learning](#) Xinyu Zhou, Raef Bassily
- [On the Statistical Efficiency of Reward-Free Exploration in Non-Linear RL](#) Jinglin Chen, Aditya Modi, Akshay Krishnamurthy, Nan Jiang, Alekh Agarwal
- [Pessimism for Offline Linear Contextual Bandits using \$\ell_p\$ Confidence Sets](#) Gene Li, Cong Ma, Nati Srebro
- [Discrete-Convex-Analysis-Based Framework for Warm-Starting Algorithms with Predictions](#) Shinsaku Sakaue, Taihei Oki
- [CAESAR: An Embodied Simulator for Generating Multimodal Referring Expression Datasets](#) Md Mofijul Islam, Reza Mirzaiee, Alexi Gladstone, Haley Green, Tariq Iqbal
- [Generalizing Bayesian Optimization with Decision-theoretic Entropies](#) Willie Neiswanger, Lantao Yu, Shengjia Zhao, Chenlin Meng, Stefano Ermon
- [Ordered Subgraph Aggregation Networks](#) Chendi Qian, Gaurav Rattan, Floris Geerts, Mathias Niepert, Christopher Morris
- [Neuron with Steady Response Leads to Better Generalization](#) Qiang Fu, Lun Du, Haitao Mao, Xu Chen, Wei Fang, Shi Han, Dongmei Zhang
- [Laplacian Autoencoders for Learning Stochastic Representations](#) Marco Miani, Frederik Warburg, Pablo Moreno-Muñoz, Nicki Skafté, Søren Hauberg
- [Alleviating the Sample Selection Bias in Few-shot Learning by Removing Projection to the Centroid](#) Jing Xu, Xu Luo, Xinglin Pan, Yanan Li, Wenjie Pei, Zenglin Xu

- [Cooperative Distribution Alignment via JSD Upper Bound](#) Wonwoong Cho, ZIYU GONG, David I. Inouye
- [An Analytical Theory of Curriculum Learning in Teacher-Student Networks](#) Luca Saglietti, Stefano Mannelli, Andrew Saxe
- [RSA: Reducing Semantic Shift from Aggressive Augmentations for Self-supervised Learning](#) Yingbin Bai, Erkun Yang, Zhaoqing Wang, Yuxuan Du, Bo Han, Cheng Deng, Dadong Wang, Tongliang Liu
- [Evaluating Latent Space Robustness and Uncertainty of EEG-ML Models under Realistic Distribution Shifts](#) Neeraj Wagh, Jionghao Wei, Samarth Rawal, Brent M Berry, Yogatheesan Varatharajah
- [u-HuBERT: Unified Mixed-Modal Speech Pretraining And Zero-Shot Transfer to Unlabeled Modality](#) Wei-Ning Hsu, Bowen Shi
- [Hierarchical Graph Transformer with Adaptive Node Sampling](#) ZAIXI ZHANG, Qi Liu, Qingyong Hu, Chee-Kong Lee
- [Learning Options via Compression](#) Yiding Jiang, Evan Liu, Benjamin Eysenbach, J. Zico Kolter, Chelsea Finn
- [Fixed-Distance Hamiltonian Monte Carlo](#) Hadi Mohasel Afshar, Sally Cripps
- [GlanceNets: Interpretable, Leak-proof Concept-based Models](#) Emanuele Marconato, Andrea Passerini, Stefano Teso
- [3DOS: Towards 3D Open Set Learning - Benchmarking and Understanding Semantic Novelty Detection on Point Clouds](#) Antonio Alliegro, Francesco Cappio Borlino, Tatiana Tommasi
- [Masked Prediction: A Parameter Identifiability View](#) Bingbin Liu, Daniel J. Hsu, Pradeep Ravikumar, Andrej Risteski
- [Self-Supervised Learning of Brain Dynamics from Broad Neuroimaging Data](#) Armin Thomas, Christopher Ré, Russell Poldrack
- [Characterization of Excess Risk for Locally Strongly Convex Population Risk](#) Mingyang Yi, Ruoyu Wang, Zhi-Ming Ma
- [A Non-Asymptotic Moreau Envelope Theory for High-Dimensional Generalized Linear Models](#) Lijia Zhou, Frederic Koehler, Pragya Sur, Danica J. Sutherland, Nati Srebro
- [Towards Efficient 3D Object Detection with Knowledge Distillation](#) Jihan Yang, Shaoshuai Shi, Runyu Ding, Zhe Wang, Xiaojuan Qi
- [CodeRL: Mastering Code Generation through Pretrained Models and Deep Reinforcement Learning](#) Hung Le, Yue Wang, Akhilesh Deepak Gotmare, Silvio Savarese, Steven Chu Hong Hoi
- [Leveraging the Hints: Adaptive Bidding in Repeated First-Price Auctions](#) Wei Zhang, Yanjun Han, Zhengyuan Zhou, Aaron Flores, Tsachy Weissman
- [Sample Efficiency Matters: A Benchmark for Practical Molecular Optimization](#) Wenhao Gao, Tianfan Fu, Jimeng Sun, Connor Coley
- [MGNNI: Multiscale Graph Neural Networks with Implicit Layers](#) Juncheng Liu, Bryan Hooi, Kenji Kawaguchi, Xiaokui Xiao
- [UQGAN: A Unified Model for Uncertainty Quantification of Deep Classifiers trained via Conditional GANs](#) Philipp Oberdiek, Gernot Fink, Matthias Rottmann
- [Audio-Driven Co-Speech Gesture Video Generation](#) Xian Liu, Qianyi Wu, Hang Zhou, Yuanqi Du, Wayne Wu, Dahua Lin, Ziwei Liu
- [Concept Embedding Models: Beyond the Accuracy-Explainability Trade-Off](#) Mateo Espinosa Zarlenga, Pietro Barbiero, Gabriele Ciravegna, Giuseppe Marra, Francesco Giannini, Michelangelo Diligenti, Zohreh Shams, Frederic Precioso, Stefano Melacci, Adrian Weller, Pietro Lió, Mateja Jamnik
- [DENSE: Data-Free One-Shot Federated Learning](#) Jie Zhang, Chen Chen, Bo Li, Lingjuan Lyu, Shuang Wu, Shouhong Ding, Chunhua Shen, Chao Wu
- [Searching for Better Spatio-temporal Alignment in Few-Shot Action Recognition](#) Yichao Cao, Xiu Su, Qingfei Tang, Shan You, Xiaobo Lu, Chang Xu
- [Fine-Tuning Pre-Trained Language Models Effectively by Optimizing Subnetworks Adaptively](#) Haojie Zhang, Ge Li, Jia Li, Zhongjin Zhang, YUQI ZHU, Zhi Jin
- [Quality Not Quantity: On the Interaction between Dataset Design and Robustness of CLIP](#) Thao Nguyen, Gabriel Ilharco, Mitchell Wortsman, Sewoong Oh, Ludwig Schmidt
- [Stars: Tera-Scale Graph Building for Clustering and Learning](#) CJ Carey, Jonathan Halcrow, Rajesh Jayaram, Vahab Mirrokni, Warren Schudy, Peilin Zhong
- [Score-Based Diffusion meets Annealed Importance Sampling](#) Arnaud Doucet, Will Grathwohl, Alexander G. Matthews, Heiko Strathmann
- [PaCo: Parameter-Compositional Multi-task Reinforcement Learning](#) Lingfeng Sun, Haichao Zhang, Wei Xu, Masayoshi TOMIZUKA
- [Entropy-Driven Mixed-Precision Quantization for Deep Network Design](#) Zhenhong Sun, Ce Ge, Junyan Wang, Ming Lin, Hesun Chen, Hao Li, Xiuyu Sun
- [Non-Gaussian Tensor Programs](#) Eugene Golikov, Greg Yang
- [Adaptation Accelerating Sampling-based Bayesian Inference in Attractor Neural Networks](#) Xingsi Dong, Zilong Ji, Tianhao Chu, Tiejun Huang, Wenhao Zhang, Si Wu
- [A Contrastive Framework for Neural Text Generation](#) Yixuan Su, Tian Lan, Yan Wang, Dani Yogatama, Lingpeng Kong, Nigel Collier
- [Exploring the Latent Space of Autoencoders with Interventional Assays](#) Felix Leeb, Stefan Bauer, Michel Besserve, Bernhard Schölkopf
- [Pythae: Unifying Generative Autoencoders in Python - A Benchmarking Use Case](#) Clément Chadebec, Louis Vincent, Stephanie Allasoiniere
- [Adaptive Multi-stage Density Ratio Estimation for Learning Latent Space Energy-based Model](#) Zhisheng Xiao, Tian Han
- [An Analysis of Ensemble Sampling](#) Chao Qin, Zheng Wen, Xiuyuan Lu, Benjamin Van Roy
- [Fair Wrapping for Black-box Predictions](#) Alexander Soen, Ibrahim M. Alabdulmohsin, Sanmi Koyejo, Yishay Mansour, Nyalleng Moorosi, Richard Nock, Ke Sun, Lexing Xie
- [Bridging Central and Local Differential Privacy in Data Acquisition Mechanisms](#) Alireza Fallah, Ali Makhdoumi, azarakhsh malekian, Asuman Ozdaglar

- [Multi-fidelity Monte Carlo: a pseudo-marginal approach](#) Diana Cai, Ryan P. Adams
- [SAPD+: An Accelerated Stochastic Method for Nonconvex-Concave Minimax Problems](#) Xuan Zhang, Necdet Serhat Aybat, Mert Gurbuzbalaban
- [Contrastive Adapters for Foundation Model Group Robustness](#) Michael Zhang, Christopher Ré
- [Decoupled Context Processing for Context Augmented Language Modeling](#) Zonglin Li, Ruiqi Guo, Sanjiv Kumar
- [LISA: Learning Interpretable Skill Abstractions from Language](#) Divyansh Garg, Skanda Vaidyanath, Kuno Kim, Jiaming Song, Stefano Ermon
- [Accelerated Primal-Dual Gradient Method for Smooth and Convex-Concave Saddle-Point Problems with Bilinear Coupling](#) Dmitry Kovalev, Alexander Gasnikov, Peter Richtarik
- [Diffusion Curvature for Estimating Local Curvature in High Dimensional Data](#) Dhananjay Bhaskar, Kincaid MacDonald, Oluwadamilola Fasina, Dawson Thomas, Bastian Rieck, Ian Adelstein, Smita Krishnaswamy
- [Hidden Progress in Deep Learning: SGD Learns Parities Near the Computational Limit](#) Boaz Barak, Benjamin Edelman, Surbhi Goel, Sham Kakade, Eran Malach, Cyril Zhang
- [One Positive Label is Sufficient: Single-Positive Multi-Label Learning with Label Enhancement](#) Ning Xu, Congyu Qiao, Jiaqi Lv, Xin Geng, Min-Ling Zhang
- [Communication Acceleration of Local Gradient Methods via an Accelerated Primal-Dual Algorithm with an Inexact Prox](#) Abdurakhmon Sadiev, Dmitry Kovalev, Peter Richtarik
- [Aligning individual brains with fused unbalanced Gromov Wasserstein](#) Alexis Thual, Quang Huy TRAN, Tatiana Zemskova, Nicolas Courty, Rémi Flamary, Stanislas Dehaene, Bertrand Thirion
- [This is the way: designing and compiling LEPISZCZE, a comprehensive NLP benchmark for Polish](#) Lukasz Augustyniak, Kamil Tagowski, Albert Sawczyn, Denis Janiak, Roman Bartusiak, Adrian Szymczak, Arkadiusz Janz, Piotr Szymański, Marcin Wątroba, Mikołaj Morzy, Tomasz Kajdanowicz, Maciej Piasecki
- [An In-depth Study of Stochastic Backpropagation](#) Jun Fang, Mingze Xu, Hao Chen, Bing Shuai, Zhuowen Tu, Joseph Tighe
- [Transformer Memory as a Differentiable Search Index](#) Yi Tay, Vinh Tran, Mostafa Dehghani, Jianmo Ni, Dara Bahri, Harsh Mehta, Zhen Qin, Kai Hui, Zhe Zhao, Jai Gupta, Tal Schuster, William W. Cohen, Donald Metzler
- [Insights into Pre-training via Simpler Synthetic Tasks](#) Yuhuai Wu, Felix Li, Percy S. Liang
- [Last-Iterate Convergence of Optimistic Gradient Method for Monotone Variational Inequalities](#) Eduard Gorbunov, Adrien Taylor, Gauthier Gidel
- [3DILG: Irregular Latent Grids for 3D Generative Modeling](#) Biao Zhang, Matthias Niessner, Peter Wonka
- [Policy Optimization for Markov Games: Unified Framework and Faster Convergence](#) Runyu Zhang, Qinghua Liu, Huan Wang, Caiming Xiong, Na Li, Yu Bai
- [C2FAR: Coarse-to-Fine Autoregressive Networks for Precise Probabilistic Forecasting](#) Shane Bergsma, Tim Zeyl, Javad Rahimipour Anaraki, Lei Guo
- [METS-CoV: A Dataset of Medical Entity and Targeted Sentiment on COVID-19 Related Tweets](#) Peilin Zhou, Zeqiang Wang, Dading Chong, Zhijiang Guo, Yining Hua, Zichang Su, Zhiyang Teng, Jiageng Wu, Jie Yang
- [Respecting Transfer Gap in Knowledge Distillation](#) Yulei Niu, Long Chen, Chang Zhou, Hanwang Zhang
- [Target alignment in truncated kernel ridge regression](#) Arash Amini, Richard Baumgartner, Dai Feng
- [Continual Learning In Environments With Polynomial Mixing Times](#) Matthew Riemer, Sharath Chandra Raparthy, Ignacio Cases, Gopeshh Subbaraj, Maximilian Puelma Touzel, Irina Rish
- [ENS-10: A Dataset For Post-Processing Ensemble Weather Forecasts](#) Saleh Ashkboos, Langwen Huang, Nikoli Dryden, Tal Ben-Nun, Peter Dueben, Lukas Gianinazzi, Luca Kummer, Torsten Hoefler
- [Panchromatic and Multispectral Image Fusion via Alternating Reverse Filtering Network](#) Keyu Yan, Man Zhou, Jie Huang, Feng Zhao, Chengjun Xie, Chongyi Li, Danfeng Hong
- [Unsupervised Cross-Task Generalization via Retrieval Augmentation](#) Bill Yuchen Lin, Kangmin Tan, Chris Miller, Beiwen Tian, Xiang Ren
- [Natural gradient enables fast sampling in spiking neural networks](#) Paul Masset, Jacob Zavatone-Veth, J. Patrick Connor, Venkatesh Murthy, Cengiz Pehlevan
- [Adversarial Auto-Augment with Label Preservation: A Representation Learning Principle Guided Approach](#) Kaiwen Yang, Yanchao Sun, Jiahao Su, Fengxiang He, Xinmei Tian, Furong Huang, Tianyi Zhou, Dacheng Tao
- [Coordinate Linear Variance Reduction for Generalized Linear Programming](#) Chaobing Song, Cheuk Yin Lin, Stephen Wright, Jelena Diakonikolas
- [Probabilistic Missing Value Imputation for Mixed Categorical and Ordered Data](#) Yuxuan Zhao, Alex Townsend, Madeleine Udell
- [Gaussian Copula Embeddings](#) Chien Lu, Jaakko Peltonen
- [SGAM: Building a Virtual 3D World through Simultaneous Generation and Mapping](#) Yuan Shen, Wei-Chiu Ma, Shenlong Wang
- [Finding Naturally Occurring Physical Backdoors in Image Datasets](#) Emily Wenger, Roma Bhattacharjee, Arjun Nitin Bhagoji, Josephine Passananti, Emilio Andere, Heather Zheng, Ben Zhao
- [Unsupervised Representation Learning from Pre-trained Diffusion Probabilistic Models](#) Zijian Zhang, Zhou Zhao, Zhijie Lin
- [Learning Causally Invariant Representations for Out-of-Distribution Generalization on Graphs](#) Yongqiang Chen, Yonggang Zhang, Yatao Bian, Han Yang, MA Kaili, Binghui Xie, Tongliang Liu, Bo Han, James Cheng
- [To update or not to update? Neurons at equilibrium in deep models](#) Andrea Bruggenolte, Enzo Tartaalione, Marco Grangetto

[Report an Issue](#) | [Name Change Policy](#)

- [Non-convex online learning via algorithmic equivalence](#) Udaya Ghai, Zhou Lu, Elad Hazan
- [TA-MoE: Topology-Aware Large Scale Mixture-of-Expert Training](#) Chang Chen, Min Li, Zhihua Wu, Dianhai Yu, Chao Yang
- [A Combinatorial Perspective on the Optimization of Shallow ReLU Networks](#) Michael S Matena, Colin A. Raffel
- [Large Language Models are Zero-Shot Reasoners](#) Takeshi Kojima, Shixiang (Shane) Gu, Machel Reid, Yutaka Matsuo, Yusuke Iwasawa
- [Trading off Utility, Informativeness, and Complexity in Emergent Communication](#) Mycal Tucker, Roger Levy, Julie A. Shah, Noga Zaslavsky
- [FiLM-Ensemble: Probabilistic Deep Learning via Feature-wise Linear Modulation](#) Mehmet Ozgur Turkoglu, Alexander Becker, Hüseyin Anil Gündüz, Mina Rezaei, Bernd Bischl, Rodrigo Caye Daudt, Stefano D'Aronco, Jan Wegner, Konrad Schindler
- [Meta-DMoE: Adapting to Domain Shift by Meta-Distillation from Mixture-of-Experts](#) Tao Zhong, Zhixiang Chi, Li Gu, Yang Wang, Yuanhao Yu, Jin Tang
- [Beyond Time-Average Convergence: Near-Optimal Uncoupled Online Learning via Clairvoyant Multiplicative Weights Update](#) Georgios Piliouras, Ryann Sim, Stratis Skoulakis
- [Meta-Reward-Net: Implicitly Differentiable Reward Learning for Preference-based Reinforcement Learning](#) Runze Liu, Fengshuo Bai, Yali Du, Yaodong Yang
- [One-shot Neural Backdoor Erasing via Adversarial Weight Masking](#) Shuwen Chai, Jinghui Chen
- [Revisiting Neural Scaling Laws in Language and Vision](#) Ibrahim M. Alabdulmohsin, Behnam Neyshabur, Xiaohua Zhai
- [Efficient Phi-Regret Minimization in Extensive-Form Games via Online Mirror Descent](#) Yu Bai, Chi Jin, Song Mei, Ziang Song, Tiancheng Yu
- [Long Range Graph Benchmark](#) Vijay Prakash Dwivedi, Ladislav Rampásek, Michael Galkin, Ali Parviz, Guy Wolf, Anh Tuan Luu, Dominique Beaini
- [On the inability of Gaussian process regression to optimally learn compositional functions](#) Matteo Giordano, Kolyan Ray, Johannes Schmidt-Hieber
- [Active Learning Through a Covering Lens](#) Ofer Yehuda, Avihu Dekel, Guy Hacohen, Daphna Weinshall
- [Uplifting Bandits](#) Yu-Guan Hsieh, Shiva Kasiviswanathan, Branislav Kveton
- [Training Uncertainty-Aware Classifiers with Conformalized Deep Learning](#) Bat-Sheva Einbinder, Yaniv Romano, Matteo Sesia, Yanfei Zhou
- [A general approximation lower bound in \$L^p\$ norm, with applications to feed-forward neural networks](#) El Mehdi Achour, Armand Foucault, Sébastien Gerchinovitz, François Malgouyres
- [EnvPool: A Highly Parallel Reinforcement Learning Environment Execution Engine](#) Jiayi Weng, Min Lin, Shengyi Huang, Bo Liu, Denys Makoviichuk, Viktor Makovychuk, Zichen Liu, Yufan Song, Ting Luo, Yukun Jiang, Zhongwen Xu, Shuicheng Yan
- [Generative Visual Prompt: Unifying Distributional Control of Pre-Trained Generative Models](#) Chen Henry Wu, Saman Motamed, Shaunak Srivastava, Fernando D De la Torre
- [Implicit Warping for Animation with Image Sets](#) Arun Mallya, Ting-Chun Wang, Ming-Yu Liu
- [FNeVR: Neural Volume Rendering for Face Animation](#) Bohan Zeng, Boyu Liu, Hong Li, Xuhui Liu, Jianzhuang Liu, Dapeng Chen, Wei Peng, Baochang Zhang
- [Ontologue: Declarative Benchmark Construction for Ontological Multi-Label Classification](#) Sean Yang, Bernease Herman, Bill Howe
- [Make an Omelette with Breaking Eggs: Zero-Shot Learning for Novel Attribute Synthesis](#) Yu-Hsuan Li, Tzu-Yin Chao, Ching-Chun Huang, Pin-Yu Chen, Wei-Chen Chiu
- [Exploiting Semantic Relations for Glass Surface Detection](#) Jiaying Lin, Yuen-Hei Yeung, Rynson Lau
- [Differentially Private Generalized Linear Models Revisited](#) Raman Arora, Raef Bassily, Cristóbal Guzmán, Michael Menart, Enayat Ullah
- [GREED: A Neural Framework for Learning Graph Distance Functions](#) Rishabh Ranjan, Siddharth Grover, Sourav Medya, Venkatesan Chakaravarthy, Yogish Sabharwal, Sayan Ranu
- [Domain Adaptation under Open Set Label Shift](#) Saurabh Garg, Sivaraman Balakrishnan, Zachary Lipton
- [Efficient Adversarial Training without Attacking: Worst-Case-Aware Robust Reinforcement Learning](#) Yongyuan Liang, Yanchao Sun, Ruijie Zheng, Furong Huang
- [Alleviating "Posterior Collapse" in Deep Topic Models via Policy Gradient](#) Yewen Li, Chaojie Wang, Zhibin Duan, Dongsheng Wang, Bo Chen, Bo An, Mingyuan Zhou
- [Tracking Functional Changes in Nonstationary Signals with Evolutionary Ensemble Bayesian Model for Robust Neural Decoding](#) Xinyun Zhu, Yu Qi, Gang Pan, Yueming Wang
- [Zeroth-Order Hard-Thresholding: Gradient Error vs. Expansivity](#) William de Vazelhes, Hualin Zhang, Huimin Wu, Xiaotong Yuan, Bin Gu
- [Collaborative Linear Bandits with Adversarial Agents: Near-Optimal Regret Bounds](#) Aritra Mitra, Arman Adibi, George J. Pappas, Hamed Hassani
- [Differentially Private Graph Learning via Sensitivity-Bounded Personalized PageRank](#) Alessandro Epasto, Vahab Mirrokni, Bryan Perozzi, Anton Tsitsulin, Peilin Zhong
- [How Well Do Unsupervised Learning Algorithms Model Human Real-time and Life-long Learning?](#) Chengxu Zhuang, Ziyu Xiang, Yoon Bai, Xiaoxuan Jia, Nicholas Turk-Browne, Kenneth Norman, James J DiCarlo, Dan Yamins
- [Stochastic Multiple Target Sampling Gradient Descent](#) Hoang Phan, Ngoc Tran, Trung Le, Toan Tran, Nhat Ho, Dinh Phung
- [Towards Out-of-Distribution Sequential Event Prediction: A Causal Treatment](#) Chenxiao Yang, Qitian Wu, Qingsong Wen, Zhiqiang Zhou, Liang Sun, Junchi Yan
- [Hybrid Neural Autoencoders for Stimulus Encoding in Visual and Other Sensory Neuroprostheses](#) Jacob Granley, Lucas Relic, Michael

- [Tractable Function-Space Variational Inference in Bayesian Neural Networks](#) Tim G. J. Rudner, Zonghao Chen, Yee Whye Teh, Yarin Gal
- [BMU-MoCo: Bidirectional Momentum Update for Continual Video-Language Modeling](#) Yizhao Gao, Nanyi Fei, Haoyu Lu, Zhiwu Lu, Hao Jiang, Yijie Li, Zhao Cao
- [Can Hybrid Geometric Scattering Networks Help Solve the Maximum Clique Problem?](#) Yimeng Min, Frederik Wenkel, Michael Perlmutter, Guy Wolf
- [Expected Improvement for Contextual Bandits](#) Hung Tran-The, Sunil Gupta, Santu Rana, Tuan Truong, Long Tran-Thanh, Svetha Venkatesh
- [Redistribution of Weights and Activations for AdderNet Quantization](#) Ying Nie, Kai Han, Haikang Diao, Chuanjian Liu, Enhua Wu, Yunhe Wang
- [Physically-Based Face Rendering for NIR-VIS Face Recognition](#) Yunqi Miao, Alexandros Lattas, Jiankang Deng, Jungong Han, Stefanos Zafeiriou
- [DGraph: A Large-Scale Financial Dataset for Graph Anomaly Detection](#) Xuanwen Huang, Yang Yang, Yang Wang, Chunping Wang, Zhisheng Zhang, Jiarong Xu, Lei Chen, Michalis Vazirgiannis
- [Globally Convergent Policy Search for Output Estimation](#) Jack Umenberger, Max Simchowitz, Juan Perdomo, Kaiqing Zhang, Russ Tedrake
- [Escaping Saddle Points for Effective Generalization on Class-Imbalanced Data](#) Harsh Rangwani, Sumukh K Aithal, Mayank Mishra, Venkatesh Babu R
- [Continuous Deep Q-Learning in Optimal Control Problems: Normalized Advantage Functions Analysis](#) Anton Plaksin, Stepan Martyanov
- [Partial Identification of Treatment Effects with Implicit Generative Models](#) Vahid Balazadeh Meresht, Vasilis Syrgkanis, Rahul G. Krishnan
- [Robust Model Selection and Nearly-Proper Learning for GMMs](#) Allen Liu, Jerry Li, Ankur Moitra
- [A2: Efficient Automated Attacker for Boosting Adversarial Training](#) Zhuoer Xu, Guanghui Zhu, Changhua Meng, shiwen cui, Zhenzhe Ying, Weiqiang Wang, Ming GU, Yihua Huang
- [Shape, Light, and Material Decomposition from Images using Monte Carlo Rendering and Denoising](#) Jon Hasselgren, Nikolai Hofmann, Jacob Munkberg
- [Convergence for score-based generative modeling with polynomial complexity](#) Holden Lee, Jianfeng Lu, Yixin Tan
- [Kernel Interpolation with Sparse Grids](#) Mohit Yadav, Daniel R. Sheldon, Cameron Musco
- [Contrastive Language-Image Pre-Training with Knowledge Graphs](#) Xuran Pan, Tianzhu Ye, Dongchen Han, Shiji Song, Gao Huang
- [Reconstructing Training Data From Trained Neural Networks](#) Niv Haim, Gal Vardi, Gilad Yehudai, Ohad Shamir, Michal Irani
- [Hierarchical Agglomerative Graph Clustering in Poly-Logarithmic Depth](#) Laxman Dhulipala, David Eisenstat, Jakub Lacki, Vahab Mirrokni, Jessica Shi
- [On-Device Training Under 256KB Memory](#) Ji Lin, Ligeng Zhu, Wei-Ming Chen, Wei-Chen Wang, Chuang Gan, Song Han
- [Behavior Transformers: Cloning \$k\$ modes with one stone](#) Nur Muhammad Shafiullah, Zichen Cui, Ariuntuya (Arty) Altanzaya, Lerrel Pinto
- [Performative Power](#) Moritz Hardt, Meena Jagadeesan, Celestine Mender-Dünner
- [Diagonal State Spaces are as Effective as Structured State Spaces](#) Ankit Gupta, Albert Gu, Jonathan Berant
- [Deep Fourier Up-Sampling](#) man zhou, Hu Yu, Jie Huang, Feng Zhao, Jinwei Gu, Chen Change Loy, Deyu Meng, Chongyi Li
- [Generative Time Series Forecasting with Diffusion, Denoise, and Disentanglement](#) Yan Li, Xinjiang Lu, Yaqing Wang, Dejing Dou
- [The Missing Invariance Principle found -- the Reciprocal Twin of Invariant Risk Minimization](#) Dongsung Huh, Avinash Baidya
- [Signal Recovery with Non-Expansive Generative Network Priors](#) Jorio Cocola
- [Towards Understanding the Mixture-of-Experts Layer in Deep Learning](#) Zixiang Chen, Yihe Deng, Yue Wu, Quanquan Gu, Yuanzhi Li
- [Indicators of Attack Failure: Debugging and Improving Optimization of Adversarial Examples](#) Maura Pintor, Luca Demetrio, Angelo Sotgiu, Ambra Demontis, Nicholas Carlini, Battista Biggio, Fabio Roli
- [Beyond accuracy: generalization properties of bio-plausible temporal credit assignment rules](#) Yuhan Helena Liu, Arna Ghosh, Blake Richards, Eric Shea-Brown, Guillaume Lajoie
- [Parameter tuning and model selection in Optimal Transport with semi-dual Brenier formulation](#) Adrien Vacher, Francois-Xavier Vialard
- [VITA: Video Instance Segmentation via Object Token Association](#) Miran Heo, Sukjun Hwang, Seoung Wug Oh, Joon-Young Lee, Seon Joo Kim
- [Bridging the Gap: Unifying the Training and Evaluation of Neural Network Binary Classifiers](#) Nathan Tsoi, Kate Candon, Deyuan Li, Yofti Milkessa, Marynel Vázquez
- [Truncated proposals for scalable and hassle-free simulation-based inference](#) Michael Deistler, Pedro J. Goncalves, Jakob H Macke
- [When to Update Your Model: Constrained Model-based Reinforcement Learning](#) Tianying Ji, Yu Luo, Fuchun Sun, Mingxuan Jing, Fengxiang He, Wenbing Huang
- [Noise Attention Learning: Enhancing Noise Robustness by Gradient Scaling](#) Yangdi Lu, Yang Bo, Wenbo He
- [Iso-Dream: Isolating and Leveraging Noncontrollable Visual Dynamics in World Models](#) Minting Pan, Xiangming Zhu, Yunbo Wang, Xiaokang Yang
- [PointNeXt: Revisiting PointNet++ with Improved Training and Scaling Strategies](#) Guocheng Qian, Yuchen Li, Houwen Peng, Jinjie Mai, Hasan Hammoud, Mohamed Elhoseiny, Bernard Ghanem
- [Quantum Algorithms for Sampling Log-Concave Distributions and Estimating Normalizing Constants](#) Andrew M. Childs, Tongyang Li, Jin-Peng Liu, Chunhao Wang, Ruizhe Zhang
- [Physics-Embedded Neural Networks: Graph Neural PDE Solvers with Mixed Boundary Conditions](#) Masanobu Horie, NAOTO MITSUME

- [Mismatched No More: Joint Model-Policy Optimization for Model-Based RL](#) Benjamin Eysenbach, Alexander Khazatsky, Sergey Levine, Russ R. Salakhutdinov
- [Phase diagram of Stochastic Gradient Descent in high-dimensional two-layer neural networks](#) Rodrigo Veiga, Ludovic Stephan, Bruno Loureiro, Florent Krzakala, Lenka Zdeborová
- [Constrained Stochastic Nonconvex Optimization with State-dependent Markov Data](#) Abhishek Roy, Krishnakumar Balasubramanian, Saeed Ghadimi
- [Adapting Self-Supervised Vision Transformers by Probing Attention-Conditioned Masking Consistency](#) Viraj Prabhu, Sriram Yenamandra, Aaditya Singh, Judy Hoffman
- [MaskTune: Mitigating Spurious Correlations by Forcing to Explore](#) Saeid Asgari, Aliasghar Khani, Fereshte Khani, Ali Gholami, Linh Tran, Ali Mahdavi Amiri, Ghassan Hamarneh
- [A Dataset for Efforts Towards Achieving the Sustainable Development Goal of Safe Working Environments](#) Eirik Lund Flogard, Ole Jakob Mengshoel
- [Decomposing NeRF for Editing via Feature Field Distillation](#) Sosuke Kobayashi, Eiichi Matsumoto, Vincent Sitzmann
- [Sample-Then-Optimize Batch Neural Thompson Sampling](#) Zhongxiang Dai, YAO SHU, Bryan Kian Hsiang Low, Patrick Jaillet
- [Video-based Human-Object Interaction Detection from Tubelet Tokens](#) Danyang Tu, Wei Sun, Xionghuo Min, Guangtao Zhai, Wei Shen
- [Efficient and Stable Fully Dynamic Facility Location](#) Sayan Bhattacharya, Silvio Lattanzi, Nikos Parotsidis
- [MCVD - Masked Conditional Video Diffusion for Prediction, Generation, and Interpolation](#) Vikram Voleti, Alexia Jolicoeur-Martineau, Chris Pal
- [Addressing Leakage in Concept Bottleneck Models](#) Marton Havasi, Sonali Parbhoo, Finale Doshi-Velez
- [Adaptive Oracle-Efficient Online Learning](#) Guanghui Wang, Zihao Hu, Vidya Muthukumar, Jacob D. Abernethy
- [MoVQ: Modulating Quantized Vectors for High-Fidelity Image Generation](#) Chuanxia Zheng, Tung-Long Vuong, Jianfei Cai, Dinh Phung
- [Meta-Auto-Decoder for Solving Parametric Partial Differential Equations](#) Xiang Huang, Zhanhong Ye, Hongsheng Liu, Shi Ji, Zidong Wang, Kang Yang, Yang Li, Min Wang, Haotian CHU, Fan Yu, Bei Hua, Lei Chen, Bin Dong
- [Sharpness-Aware Training for Free](#) JIAWEI DU, Daquan Zhou, Jiashi Feng, Vincent Tan, Joey Tianyi Zhou
- [Generalization Error Bounds on Deep Learning with Markov Datasets](#) Lan V. Truong
- [AirFRANS: High Fidelity Computational Fluid Dynamics Dataset for Approximating Reynolds-Averaged Navier–Stokes Solutions](#) Florent Bonnet, Jocelyn Mazari, Paola Cinnella, Patrick Gallinari
- [Generalization for multiclass classification with overparameterized linear models](#) Vignesh Subramanian, Rahul Arya, Anant Sahai
- [Inception Transformer](#) Chenyang Si, Weihao Yu, Pan Zhou, Yichen Zhou, Xinchao Wang, Shuicheng Yan
- [ElasticMVS: Learning elastic part representation for self-supervised multi-view stereopsis](#) Jinzhi Zhang, Ruofan Tang, Zheng Cao, Jing Xiao, Ruqi Huang, LU FANG
- [Adaptive Stochastic Variance Reduction for Non-convex Finite-Sum Minimization](#) Ali Kavis, Stratis Skoulakis, Kimon Antonakopoulos, Leello Tadesse Dadi, Volkan Cevher
- [Thinking Outside the Ball: Optimal Learning with Gradient Descent for Generalized Linear Stochastic Convex Optimization](#) Idan Amir, Roi Livni, Nati Srebro
- [Generalization Properties of NAS under Activation and Skip Connection Search](#) Zhenyu Zhu, Fanghui Liu, Grigorios Chrysos, Volkan Cevher
- [Mesoscopic modeling of hidden spiking neurons](#) Shuqi Wang, Valentin Schmutz, Guillaume Bellec, Wulfram Gerstner
- [SageMix: Saliency-Guided Mixup for Point Clouds](#) Sanghyeok Lee, Minkyu Jeon, Injae Kim, Yongyang Xiong, Hyunwoo J. Kim
- [Denoising Diffusion Restoration Models](#) Bahjat Kawar, Michael Elad, Stefano Ermon, Jiaming Song
- [On Translation and Reconstruction Guarantees of the Cycle-Consistent Generative Adversarial Networks](#) Anish Chakrabarty, Swagatam Das
- [Adversarial Training with Complementary Labels: On the Benefit of Gradually Informative Attacks](#) Jianan Zhou, Jianing Zhu, Jingfeng ZHANG, Tongliang Liu, Gang Niu, Bo Han, Masashi Sugiyama
- [Tractable Optimality in Episodic Latent MABs](#) Jeongyeol Kwon, Yonathan Efroni, Constantine Caramanis, Shie Mannor
- [A Characterization of Semi-Supervised Adversarially Robust PAC Learnability](#) Idan Attias, Steve Hanneke, Yishay Mansour
- [Data-IQ: Characterizing subgroups with heterogeneous outcomes in tabular data](#) Nabeel Seedat, Jonathan Crabbé, Ioana Bica, Mihaela van der Schaar
- [Task-Free Continual Learning via Online Discrepancy Distance Learning](#) Fei Ye, Adrian G. Bors
- [BinauralGrad: A Two-Stage Conditional Diffusion Probabilistic Model for Binaural Audio Synthesis](#) Yichong Leng, Zehua Chen, Junliang Guo, Haohe Liu, Jiawei Chen, Xu Tan, Danilo Mandic, Lei He, Xiangyang Li, Tao Qin, sheng zhao, Tie-Yan Liu
- [ConfLab: A Data Collection Concept, Dataset, and Benchmark for Machine Analysis of Free-Standing Social Interactions in the Wild](#) Chirag Raman, Jose Vargas Quiros, Stephanie Tan, Ashraf Islam, Ekin Gedik, Hayley Hung
- [Flamingo: a Visual Language Model for Few-Shot Learning](#) Jean-Baptiste Alayrac, Jeff Donahue, Pauline Luc, Antoine Miech, Iain Barr, Yana Hasson, Karel Lenc, Arthur Mensch, Katherine Millican, Malcolm Reynolds, Roman Ring, Eliza Rutherford, Serkan Cabi, Tengda Han, Zhitao Gong, Sina Samangooei, Marianne Monteiro, Jacob L Menick, Sebastian Borgeaud, Andy Brock, Aida Nematzadeh, Sahand Sharifzadeh, Mikołaj Bińkowski, Ricardo Barreira, Oriol Vinyals, Andrew Zisserman, Karén Simonyan
- [Understanding the Eluder Dimension](#) Gene Li, Pritish Kamath, Dylan J Foster, Nati Srebro

- [Learning State-Aware Visual Representations from Audible Interactions](#) Himangi Mittal, Pedro Morgado, Unnat Jain, Abhinav Gupta
- [Context-Based Dynamic Pricing with Partially Linear Demand Model](#) Jinzhi Bu, David Simchi-Levi, Chonghuan Wang
- [NeuroSchedule: A Novel Effective GNN-based Scheduling Method for High-level Synthesis](#) Jun Zeng, Mingyang Kou, Hailong Yao
- [Hand-Object Interaction Image Generation](#) Hezhen Hu, Weilun Wang, Wengang Zhou, Houqiang Li
- [DISCO: Adversarial Defense with Local Implicit Functions](#) Chih-Hui Ho, Nuno Vasconcelos
- [Distinguishing discrete and continuous behavioral variability using warped autoregressive HMMs](#) Julia Costacurta, Lea Duncker, Blue Sheffer, Winthrop Gillis, Caleb Weinreb, Jeffrey Markowitz, Sandeep R Datta, Alex Williams, Scott Linderman
- [RORL: Robust Offline Reinforcement Learning via Conservative Smoothing](#) Rui Yang, Chenjia Bai, Xiaoteng Ma, Zhaoran Wang, Chongjie Zhang, Lei Han
- [Optimal Scaling for Locally Balanced Proposals in Discrete Spaces](#) Haoran Sun, Hanjun Dai, Dale Schuurmans
- [The Impact of Task Underspecification in Evaluating Deep Reinforcement Learning](#) Vindula Jayawardana, Catherine Tang, Sirui Li, Dajiang Suo, Cathy Wu
- [Zero-Shot 3D Drug Design by Sketching and Generating](#) Siyu Long, Yi Zhou, Xinyu Dai, Hao Zhou
- [Decoupling Knowledge from Memorization: Retrieval-augmented Prompt Learning](#) Xiang Chen, Lei Li, Ningyu Zhang, Xiaozhuan Liang, Shumin Deng, Chuanqi Tan, Fei Huang, Luo Si, Huajun Chen
- [Polynomial time guarantees for the Burer-Monteiro method](#) Diego Cifuentes, Ankur Moitra
- [Optimal Comparator Adaptive Online Learning with Switching Cost](#) Zhiyu Zhang, Ashok Cutkosky, Yannis Paschalidis
- [A Robust Phased Elimination Algorithm for Corruption-Tolerant Gaussian Process Bandits](#) Ilija Bogunovic, Zihan Li, Andreas Krause, Jonathan Scarlett
- [Fair Rank Aggregation](#) Diptarka Chakraborty, Syamantak Das, Arindam Khan, Aditya Subramanian
- [Beyond IID: data-driven decision-making in heterogeneous environments](#) Omar Besbes, Will Ma, Omar Mouchtaki
- [Neur2SP: Neural Two-Stage Stochastic Programming](#) Rahul Mihir Patel, Justin Dumouchelle, Elias Khalil, Merve Bodur
- [SQ Lower Bounds for Learning Single Neurons with Massart Noise](#) Ilias Diakonikolas, Daniel Kane, Lisheng Ren, Yuxin Sun
- [MaskPlace: Fast Chip Placement via Reinforced Visual Representation Learning](#) Yao Lai, Yao Mu, Ping Luo
- [List-Decodable Sparse Mean Estimation](#) Shiwei Zeng, Jie Shen
- [Structure-Aware Image Segmentation with Homotopy Warping](#) Xiaoling Hu
- [Positive-Unlabeled Learning using Random Forests via Recursive Greedy Risk Minimization](#) Jonathan Wilton, Abigail Koay, Ryan Ko, Miao Xu, Nan Ye
- [Extracting computational mechanisms from neural data using low-rank RNNs](#) Adrian Valente, Jonathan W. Pillow, Srdjan Ostojic
- [Pluralistic Image Completion with Gaussian Mixture Models](#) Xiaobo Xia, Wenhao Yang, Jie Ren, Yewen Li, Yibing Zhan, Bo Han, Tongliang Liu
- [A Fast Post-Training Pruning Framework for Transformers](#) Woosuk Kwon, Sehoon Kim, Michael W. Mahoney, Joseph Hassoun, Kurt Keutzer, Amir Gholami
- [SurDis: A Surface Discontinuity Dataset for Wearable Technology to Assist Blind Navigation in Urban Environments](#) Kuan Yew Leong, Siew Mooi Lim
- [Interventions, Where and How? Experimental Design for Causal Models at Scale](#) Panagiotis Tigas, Yashas Annadani, Andrew Jesson, Bernhard Schölkopf, Yarin Gal, Stefan Bauer
- [Dual-discriminative Graph Neural Network for Imbalanced Graph-level Anomaly Detection](#) GE ZHANG, Zhenyu Yang, Jia Wu, Jian Yang, Shan Xue, Hao Peng, Jianlin Su, Chuan Zhou, Quan Z. Sheng, Leman Akoglu, Charu Aggarwal
- [Posted Pricing and Dynamic Prior-independent Mechanisms with Value Maximizers](#) Yuan Deng, Vahab Mirrokni, Hanrui Zhang
- [Curious Exploration via Structured World Models Yields Zero-Shot Object Manipulation](#) Cansu Sancaktar, Sebastian Blaes, Georg Martius
- [Estimating Noise Transition Matrix with Label Correlations for Noisy Multi-Label Learning](#) Shikun Li, Xiaobo Xia, Hansong Zhang, Yibing Zhan, Shiming Ge, Tongliang Liu
- [Active Learning Polynomial Threshold Functions](#) Omri Ben-Eliezer, Max Hopkins, Chutong Yang, Hantao Yu
- [Single-phase deep learning in cortico-cortical networks](#) Will Greedy, Heng Wei Zhu, Joseph Pemberton, Jack Mellor, Rui Ponte Costa
- [Domain Generalization by Learning and Removing Domain-specific Features](#) Yu Ding, Lei Wang, Bin Liang, Shuming Liang, Yang Wang, Fang Chen
- [Torsional Diffusion for Molecular Conformer Generation](#) Bowen Jing, Gabriele Corso, Jeffrey Chang, Regina Barzilay, Tommi Jaakkola
- [LiteTransformerSearch: Training-free Neural Architecture Search for Efficient Language Models](#) Mojan Javaheripi, Gustavo de Rosa, Subhabrata Mukherjee, Shital Shah, Tomasz Religa, Caio Cesar Teodoro Mendes, Sebastien Bubeck, Farinaz Koushanfar, Debadeepta Dey
- [AgraSSt: Approximate Graph Stein Statistics for Interpretable Assessment of Implicit Graph Generators](#) Wenkai Xu, Gesine D Reinert
- [On the Limitations of Stochastic Pre-processing Defenses](#) Yue Gao, I Shumailov, Kassem Fawaz, Nicolas Papernot
- [Iterative Scene Graph Generation](#) Siddhesh Khandelwal, Leonid Sigal
- [Proximal Point Imitation Learning](#) Luca Viano, Angeliki Kamoutsis, Gergely Neu, Igor Krawczuk, Volkan Cevher
- [Heatmap Distribution Matching for Human Pose Estimation](#) Haoxuan Qu, Li Xu, Yujun Cai, Lin Geng Foo, Jun Liu
- [Mining Unseen Classes via Regional Objectness: A Simple Baseline for Incremental Segmentation](#) Zekang Zhang, Guangyu Gao, Zhiyuan Fang, Jianbo Jiao, Yunchao Wei
- [Bezier Gaussian Processes for Tall and Wide Data](#) Martin Jørgensen, Michael A Osborne

[Report an Issue](#) | [Name Change Policy](#)

- [Smoothed Embeddings for Certified Few-Shot Learning](#) Mikhail Pautov, Olesya Kuznetsova, Nurislam Tursynbek, Aleksandr Petiushko, Ivan Oseledets
- [Finite Sample Analysis Of Dynamic Regression Parameter Learning](#) Mark Kozdoba, Edward Moroshko, Shie Mannor, Yacov Crammer
- [Group Meritocratic Fairness in Linear Contextual Bandits](#) Riccardo Grazi, Arya Akhavan, John IF Falk, Leonardo Cella, Massimiliano Pontil
- [New Lower Bounds for Private Estimation and a Generalized Fingerprinting Lemma](#) Gautam Kamath, Argyris Mouzakis, Vikrant Singhal
- [Lower Bounds on Randomly Preconditioned Lasso via Robust Sparse Designs](#) Jonathan Kelner, Frederic Koehler, Raghu Meka, Dhruv Rohatgi
- [Model-based Safe Deep Reinforcement Learning via a Constrained Proximal Policy Optimization Algorithm](#) Ashish K Jayant, Shalabh Bhatnagar
- [Exact Solutions of a Deep Linear Network](#) Liu Ziyin, Botao Li, Xiangming Meng
- [An Adaptive Kernel Approach to Federated Learning of Heterogeneous Causal Effects](#) Thanh Vinh Vo, Arnab Bhattacharyya, Young Lee, Tze-Yun Leong
- [FeLMi : Few shot Learning with hard Mixup](#) Aniket Roy, Anshul Shah, Ketul Shah, Prithviraj Dhar, Anoop Cherian, Rama Chellappa
- [EpiGRAF: Rethinking training of 3D GANs](#) Ivan Skorokhodov, Sergey Tulyakov, Yiqun Wang, Peter Wonka
- [Towards Trustworthy Automatic Diagnosis Systems by Emulating Doctors' Reasoning with Deep Reinforcement Learning](#) Arsene Fansi Tchango, Rishab Goel, Julien Martel, Zhi Wen, Gaetan Marceau Caron, Joumana Ghosn
- [Towards Improving Faithfulness in Abstractive Summarization](#) Xiuying Chen, Mingzhe Li, Xin Gao, Xiangliang Zhang
- [ZIN: When and How to Learn Invariance Without Environment Partition?](#) Yong Lin, Shengyu Zhu, Lu Tan, Peng Cui
- [Random Sharpness-Aware Minimization](#) Yong Liu, Siqi Mai, Minhao Cheng, Xiangning Chen, Cho-Jui Hsieh, Yang You
- [Active Surrogate Estimators: An Active Learning Approach to Label-Efficient Model Evaluation](#) Jannik Kossen, Sebastian Farquhar, Yarin Gal, Thomas Rainforth
- [HAPI: A Large-scale Longitudinal Dataset of Commercial ML API Predictions](#) Lingjiao Chen, Zhihua Jin, Evan Sabri Eyuboglu, Christopher Ré, Matei Zaharia, James Y. Zou
- [Near-Optimal Regret Bounds for Multi-batch Reinforcement Learning](#) Zihan Zhang, Yuhang Jiang, Yuan Zhou, Xiangyang Ji
- [OST: Improving Generalization of DeepFake Detection via One-Shot Test-Time Training](#) Liang Chen, Yong Zhang, Yibing Song, Jue Wang, Lingqiao Liu
- [The Surprising Effectiveness of PPO in Cooperative Multi-Agent Games](#) Chao Yu, Akash Velu, Eugene Vinitzky, Jiaxuan Gao, Yu Wang, Alexandre Bayen, YI WU
- [Resolving the data ambiguity for periodic crystals](#) Daniel Widdowson, Vitaliy Kurlin
- [Video PreTraining \(VPT\): Learning to Act by Watching Unlabeled Online Videos](#) Bowen Baker, Ilge Akkaya, Peter Zhokov, Joost Huizinga, Jie Tang, Adrien Ecoffet, Brandon Houghton, Raul Sampedro, Jeff Clune
- [GLOBEM Dataset: Multi-Year Datasets for Longitudinal Human Behavior Modeling Generalization](#) Xuhai Xu, Han Zhang, Yasaman Sefidgar, Yiyi Ren, Xin Liu, Woosuk Seo, Jennifer Brown, Kevin Kuehn, Mike Merrill, Paula Nurius, Shwetak Patel, Tim Althoff, Margaret Morris, Eve Riskin, Jennifer Mankoff, Anind Dey
- [Shape And Structure Preserving Differential Privacy](#) Carlos Soto, Karthik Bharath, Matthew Reimherr, Aleksandra Slavković
- [Transformers meet Stochastic Block Models: Attention with Data-Adaptive Sparsity and Cost](#) Sungjun Cho, Seonwoo Min, Jinwoo Kim, Moontae Lee, Honglak Lee, Seunghoon Hong
- [Characteristics of Harmful Text: Towards Rigorous Benchmarking of Language Models](#) Maribeth Rauh, John Mellor, Jonathan Uesato, Po-Sen Huang, Johannes Welbl, Laura Weidinger, Sumanth Dathathri, Amelia Glaese, Geoffrey Irving, Iason Gabriel, William Isaac, Lisa Anne Hendricks
- [APG: Adaptive Parameter Generation Network for Click-Through Rate Prediction](#) Bencheng Yan, Pengjie Wang, Kai Zhang, Feng Li, Hongbo Deng, Jian Xu, Bo Zheng
- [NeoRL: A Near Real-World Benchmark for Offline Reinforcement Learning](#) Rong-Jun Qin, Xingyuan Zhang, Songyi Gao, Xiong-Hui Chen, Zewen Li, Weinan Zhang, Yang Yu
- [Stochastic Halpern Iteration with Variance Reduction for Stochastic Monotone Inclusions](#) Xufeng Cai, Chaobing Song, Cristóbal Guzmán, Jelena Diakonikolas
- [SNN-RAT: Robustness-enhanced Spiking Neural Network through Regularized Adversarial Training](#) Jianhao Ding, Tong Bu, Zhao Fei Yu, Tiejun Huang, Jian Liu
- [The Mechanism of Prediction Head in Non-contrastive Self-supervised Learning](#) Zixin Wen, Yuanzhi Li
- [Counterfactual Temporal Point Processes](#) Kimia Noorbakhsh, Manuel Rodriguez
- [Chain-of-Thought Prompting Elicits Reasoning in Large Language Models](#) Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, brian ichter, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou
- [House of Cans: Covert Transmission of Internal Datasets via Capacity-Aware Neuron Steganography](#) Xudong Pan, Shengyao Zhang, Mi Zhang, Yifan Yan, Min Yang
- [Polynomial-Time Optimal Equilibria with a Mediator in Extensive-Form Games](#) Brian Zhang, Tuomas Sandholm
- [Dungeons and Data: A Large-Scale NetHack Dataset](#) Eric Hambro, Roberta Raileanu, Danielle Rothermel, Vegard Mella, Tim Rocktäschel, Heinrich Küttler, Naila Murray
- [Optimal Dynamic Regret in LQR Control](#) Dheeraj Baby, Yu-Xiang Wang

- [Gradient Descent Is Optimal Under Lower Restricted Secant Inequality And Upper Error Bound](#) Charles Guille-Escuret, Adam Ibrahim, Baptiste Goujaud, Ioannis Mitliagkas
- [GenSDF: Two-Stage Learning of Generalizable Signed Distance Functions](#) Gene Chou, Ilya Chugunov, Felix Heide
- [Forecasting Human Trajectory from Scene History](#) Mancheng Meng, Ziyang Wu, Terrence Chen, Xiran Cai, Xiang Zhou, Fan Yang, Dinggang Shen
- [Debiasing Graph Neural Networks via Learning Disentangled Causal Substructure](#) Shaohua Fan, Xiao Wang, Yanhu Mo, Chuan Shi, Jian Tang
- [Stochastic Online Learning with Feedback Graphs: Finite-Time and Asymptotic Optimality](#) Teodor Vanislavov Marinov, Mehryar Mohri, Julian Zimmert
- [Asymptotics of \$\ell_2\$ Regularized Network Embeddings](#) Andrew Davison
- [Differentiable hierarchical and surrogate gradient search for spiking neural networks](#) Kaiwei Che, Luziwei Leng, Kaixuan Zhang, Jianguo Zhang, Qinghu Meng, Jie Cheng, Qinghai Guo, Jianxing Liao
- [On Embeddings for Numerical Features in Tabular Deep Learning](#) Yury Gorishniy, Ivan Rubachev, Artem Babenko
- [Visual Prompting via Image Inpainting](#) Amir Bar, Yossi Gandelsman, Trevor Darrell, Amir Globerson, Alexei Efros
- [MonoSDF: Exploring Monocular Geometric Cues for Neural Implicit Surface Reconstruction](#) Zehao Yu, Songyou Peng, Michael Niemeyer, Torsten Sattler, Andreas Geiger
- [OpenAUC: Towards AUC-Oriented Open-Set Recognition](#) Zitai Wang, Qianqian Xu, Zhiyong Yang, Yuan He, Xiaochun Cao, Qingming Huang
- [Reduction Algorithms for Persistence Diagrams of Networks: CoralTDA and PrunIT](#) Cuneyt G Akcora, Murat Kantarcioglu, Yulia Gel, Baris Coskunuzer
- [TreeMoCo: Contrastive Neuron Morphology Representation Learning](#) Hanbo Chen, Jiawei Yang, Daniel Iascone, Lijuan Liu, Lei He, Hanchuan Peng, Jianhua Yao
- [Compositional Generalization in Unsupervised Compositional Representation Learning: A Study on Disentanglement and Emergent Language](#) Zhenlin Xu, Marc Niethammer, Colin A. Raffel
- [Improving Zero-Shot Generalization in Offline Reinforcement Learning using Generalized Similarity Functions](#) Bogdan Mazouze, Ilya Kostrikov, Ofir Nachum, Jonathan J. Tompson
- [GAUDI: A Neural Architect for Immersive 3D Scene Generation](#) Miguel Angel Bautista, Pengsheng Guo, Samira Abnar, Walter Talbott, Alexander Toshev, Zhuoyuan Chen, Laurent Dinh, Shuangfei Zhai, Hanlin Goh, Daniel Ulbricht, Afshin Dehghan, Joshua Susskind
- [Mask-based Latent Reconstruction for Reinforcement Learning](#) Tao Yu, Zhizheng Zhang, Cuiling Lan, Yan Lu, Zhibo Chen
- [Product Ranking for Revenue Maximization with Multiple Purchases](#) Renzhe Xu, Xingxuan Zhang, Bo Li, Yafeng Zhang, Xiaolong Chen, Peng Cui
- [One Model to Edit Them All: Free-Form Text-Driven Image Manipulation with Semantic Modulations](#) Yiming Zhu, Hongyu Liu, Yibing Song, Ziyang Yuan, Xintong Han, Chun Yuan, Qifeng Chen, Jue Wang
- [Luckiness in Multiscale Online Learning](#) Wouter M. Koolen, Muriel F. Pérez-Ortiz
- [Deep Active Learning by Leveraging Training Dynamics](#) Haonan Wang, Wei Huang, Ziwei Wu, Hanghang Tong, Andrew J Margenot, Jingrui He
- [Learning and Covering Sums of Independent Random Variables with Unbounded Support](#) Alkis Kalavasis, Konstantinos Stavropoulos, Emmanouil Zampetakis
- [LGDN: Language-Guided Denoising Network for Video-Language Modeling](#) Haoyu Lu, Mingyu Ding, Nanyi Fei, Yuqi Huo, Zhiwu Lu
- [LieGG: Studying Learned Lie Group Generators](#) Artem Moskalev, Anna Sepiarskaia, Ivan Sosnovik, Arnold Smeulders
- [FourierNets enable the design of highly non-local optical encoders for computational imaging](#) Diptodip Deb, Zhenfei Jiao, Ruth Sims, Alex Chen, Michael Broxton, Misha B Ahrens, Kaspar Podgorski, Srinivas C Turaga
- [Benign Overfitting in Two-layer Convolutional Neural Networks](#) Yuan Cao, Zixiang Chen, Misha Belkin, Quanquan Gu
- [Discovery of Single Independent Latent Variable](#) Uri Shaham, Jonathan Svirsky, Ori Katz, Ronen Talmon
- [Meta-ticket: Finding optimal subnetworks for few-shot learning within randomly initialized neural networks](#) Daiki Chijiwa, Shin'ya Yamaguchi, Atsutoshi Kumagai, Yasutoshi Ida
- [LAION-5B: An open large-scale dataset for training next generation image-text models](#) Christoph Schuhmann, Romain Beaumont, Richard Vencu, Cade Gordon, Ross Wightman, Mehdi Cherti, Theo Coombes, Aarush Katta, Clayton Mullis, Mitchell Wortsman, Patrick Schramowski, Srivatsa Kundurthy, Katherine Crowson, Ludwig Schmidt, Robert Kaczmarczyk, Jenia Jitsev
- [Constants of motion network](#) Muhammad Firmansyah Kasim, Yi Heng Lim
- [HUMUS-Net: Hybrid Unrolled Multi-scale Network Architecture for Accelerated MRI Reconstruction](#) Zalan Fabian, Berk Tinaz, Mahdi Soltanolkotabi
- [A Damped Newton Method Achieves Global \$\mathcal{O}\left\(\frac{1}{k^2}\right\)\$ and Local Quadratic Convergence Rate](#) Slavomír Hanzely, Dmitry Kamzolov, Dmitry Pasechnyuk, Alexander Gasnikov, Peter Richtarik, Martin Takac
- [Disentangling the Predictive Variance of Deep Ensembles through the Neural Tangent Kernel](#) Seijin Kobayashi, Pau Vilimelis Aceituno, Johannes von Oswald
- [High-dimensional limit theorems for SGD: Effective dynamics and critical scaling](#) Gerard Ben Arous, Reza Gheissari, Aukosh Jagannath
- [Online Deep Equilibrium Learning for Regularization by Denoising](#) Jiaming Liu, Xiaojian Xu, Weijie Gan, shirin shoushtari, Ulugbek

- [Semantic Exploration from Language Abstractions and Pretrained Representations](#) Allison Tam, Neil Rabinowitz, Andrew Lampinen, Nicholas A. Roy, Stephanie Chan, DJ Strouse, Jane Wang, Andrea Banino, Felix Hill
- [Earthformer: Exploring Space-Time Transformers for Earth System Forecasting](#) Zhihan Gao, Xingjian Shi, Hao Wang, Yi Zhu, Yuyang (Bernie) Wang, Mu Li, Dit-Yan Yeung
- [Benchopt: Reproducible, efficient and collaborative optimization benchmarks](#) Thomas Moreau, Mathurin Massias, Alexandre Gramfort, Pierre Ablin, Pierre-Antoine Bannier, Benjamin Charlier, Mathieu Dagr  ou, Tom Dupre la Tour, Ghislain DURIF, Cassio F. Dantas, Quentin Klopfenstein, Johan Larsson, En Lai, Tanguy Lefort, Beno  t Mal  zieux, Badr MOUFAD, Binh T. Nguyen, Alain Rakotomamonjy, Zaccharie Ramzi, Joseph Salmon, Samuel Vaiter
- [SketchBoost: Fast Gradient Boosted Decision Tree for Multioutput Problems](#) Leonid Iosipoi, Anton Vakhrushev
- [Deep Attentive Belief Propagation: Integrating Reasoning and Learning for Solving Constraint Optimization Problems](#) Yanchen Deng, Shufeng Kong, Caihua Liu, Bo An
- [Conservative Dual Policy Optimization for Efficient Model-Based Reinforcement Learning](#) Shenao Zhang
- [Decentralized Training of Foundation Models in Heterogeneous Environments](#) Binhang Yuan, Yongjun He, Jared Davis, Tianyi Zhang, Tri Dao, Beidi Chen, Percy S. Liang, Christopher R  , Ce Zhang
- [Cross Aggregation Transformer for Image Restoration](#) Zheng Chen, Yulun Zhang, Jinjin Gu, yongbing zhang, Linghe Kong, Xin Yuan
- [Neural Payoff Machines: Predicting Fair and Stable Payoff Allocations Among Team Members](#) Daphne Cornelisse, Thomas Rood, Yoram Bachrach, Mateusz Malinowski, Tal Kachman
- [Learning Recourse on Instance Environment to Enhance Prediction Accuracy](#) Lokesh N, Guntakanti Sai Koushik, Abir De, Sunita Sarawagi
- [Learning to Re-weight Examples with Optimal Transport for Imbalanced Classification](#) Dandan Guo, Zhuo Li, meixi zheng, He Zhao, Mingyuan Zhou, Hongyuan Zha
- [DIMES: A Differentiable Meta Solver for Combinatorial Optimization Problems](#) Ruizhong Qiu, Zhiqing Sun, Yiming Yang
- [Confident Approximate Policy Iteration for Efficient Local Planning in \$q^\pi\$ -realizable MDPs](#) Gell  rt Weisz, Andr  s Gy  rgy, Tadashi Kozuno, Csaba Szepesvari
- [Finding Correlated Equilibrium of Constrained Markov Game: A Primal-Dual Approach](#) Ziyi Chen, Shaocong Ma, Yi Zhou
- [NSNet: A General Neural Probabilistic Framework for Satisfiability Problems](#) Zhaoyu Li, Xujie Si
- [Brain Network Transformer](#) Xuan Kan, Wei Dai, Hejie Cui, Zilong Zhang, Ying Guo, Carl Yang
- [A Neural Corpus Indexer for Document Retrieval](#) Yujing Wang, Yingyan Hou, Haonan Wang, Ziming Miao, Shibin Wu, Qi Chen, Yuqing Xia, Chengmin Chi, Guoshuai Zhao, Zheng Liu, Xing Xie, Hao Sun, Weiwei Deng, Qi Zhang, Mao Yang
- [Perfect Sampling from Pairwise Comparisons](#) Dimitris Fotakis, Alkis Kalavasis, Christos Tzamos
- [Improved Utility Analysis of Private CountSketch](#) Rasmus Pagh, Mikkel Thorup
- [Optimal Efficiency-Envy Trade-Off via Optimal Transport](#) Steven Yin, Christian Kroer
- [Non-Linear Coordination Graphs](#) Yipeng Kang, Tonghan Wang, Qianlan Yang, Xiaoran Wu, Chongjie Zhang
- [SafeBench: A Benchmarking Platform for Safety Evaluation of Autonomous Vehicles](#) Chejian Xu, Wenhao Ding, Weijie Lyu, ZUXIN LIU, Shuai Wang, Yihan He, Hanjiang Hu, DING ZHAO, Bo Li
- [Improving Diffusion Models for Inverse Problems using Manifold Constraints](#) Hyungjin Chung, Byeongsu Sim, Dohoon Ryu, Jong Chul Ye
- [Semi-supervised Vision Transformers at Scale](#) Zhaowei Cai, Avinash Ravichandran, Paolo Favaro, Manchen Wang, Davide Modolo, Rahul Bhotika, Zhuowen Tu, Stefano Soatto
- [Fast Bayesian Estimation of Point Process Intensity as Function of Covariates](#) Hideaki Kim, Taichi Asami, Hiroyuki Toda
- [Online PAC-Bayes Learning](#) Maxime Haddouche, Benjamin Guedj
- [Deep Model Reassembly](#) Xingyi Yang, Daquan Zhou, Songhua Liu, Jingwen Ye, Xinchao Wang
- [Fast Algorithms for Packing Proportional Fairness and its Dual](#) Francisco Criado, David Martinez-Rubio, Sebastian Pokutta
- [A Closer Look at Prototype Classifier for Few-shot Image Classification](#) Mingcheng Hou, Issei Sato
- [When are Offline Two-Player Zero-Sum Markov Games Solvable?](#) Qiwen Cui, Simon S. Du
- [BigBio: A Framework for Data-Centric Biomedical Natural Language Processing](#) Jason Fries, Leon Weber, Natasha Seelam, Gabriel Altay, Debajyoti Datta, Samuele Garda, Sunny Kang, Rosaline Su, Wojciech Kusa, Samuel Cahyawijaya, Fabio Barth, Simon Ott, Matthias Samwald, Stephen Bach, Stella Biderman, Mario S  nger, Bo Wang, Alison Callahan, Daniel Le  n Peri    n, Th  o Gigant, Patrick Haller, Jenny Chim, Jose Posada, John Giorgi, Karthik Rangasai Sivaraman, Marc P  mies, Marianna Nezhurina, Robert Martin, Michael Cullan, Moritz Freidank, Nathan Dahlberg, Shubhanshu Mishra, Shamik Bose, Nicholas Broad, Yanis Labrak, Shlok Deshmukh, Sid Kiblawi, Ayush Singh, Minh Chien Vu, Trishala Neeraj, Jonas Golde, Albert Villanova del Moral, Benjamin Beilharz
- [Branch & Learn for Recursively and Iteratively Solvable Problems in Predict+Optimize](#) Xinyi Hu, Jasper Lee, Jimmy Lee, Allen Z. Zhong
- [Linear tree shap](#) peng yu, Albert Bifet, Jesse Read, Chao Xu
- [Gradient Estimation with Discrete Stein Operators](#) Jiaxin Shi, Yuhao Zhou, Jessica Hwang, Michalis Titsias, Lester Mackey
- [Rapidly Mixing Multiple-try Metropolis Algorithms for Model Selection Problems](#) Hyunwoong Chang, Changwoo Lee, Zhao Tang Luo, Huiyan Sang, Quan Zhou
- [Distributionally Adaptive Meta Reinforcement Learning](#) Anurag Ajay, Abhishek Gupta, Dibya Ghosh, Sergey Levine, Pulkit Agrawal
- [Toward Efficient Robust Training against Union of \$\ell_p\$ Threat Models](#) Gaurang Sriramanan, Maharshi Gor, Soheil Feizi
- [Multi-view Subspace Clustering on Topological Manifold](#) Shudong Huang, Hongjie Wu, Yazhou Ren, Ivor Tsang, Zenglin Xu, Wentao Feng, Jiancheng Lv

- [Online Agnostic Multiclass Boosting](#) Vinod Raman, Ambuj Tewari
- [A contrastive rule for meta-learning](#) Nicolas Zucchet, Simon Schug, Johannes von Oswald, Dominic Zhao, João Sacramento
- [Distinguishing Learning Rules with Brain Machine Interfaces](#) Jacob Portes, Christian Schmid, James M Murray
- [Efficient Training of Low-Curvature Neural Networks](#) Suraj Srinivas, Kyle Matoba, Himabindu Lakkaraju, François Fleuret
- [Infinite-Fidelity Coregionalization for Physical Simulation](#) Shibo Li, Zheng Wang, Robert Kirby, Shandian Zhe
- [Open High-Resolution Satellite Imagery: The WorldStrat Dataset – With Application to Super-Resolution](#) Julien Cornebise, Ivan Oršolić, Freddie Kalaitzis
- [Evaluation beyond Task Performance: Analyzing Concepts in AlphaZero in Hex](#) Charles Lovering, Jessica Forde, George Konidakis, Ellie Pavlick, Michael Littman
- [Semi-supervised Semantic Segmentation with Prototype-based Consistency Regularization](#) Haiming Xu, Lingqiao Liu, Qiuchen Bian, Zhen Yang
- [Empirical Phase Diagram for Three-layer Neural Networks with Infinite Width](#) Hanxu Zhou, Zhou Qixuan, Zhenyuan Jin, Tao Luo, Yaoyu Zhang, Zhi-Qin Xu
- [Benchmarking and Analyzing 3D Human Pose and Shape Estimation Beyond Algorithms](#) Hui En Pang, Zhongang Cai, Lei Yang, Tianwei Zhang, Ziwei Liu
- [TTOpt: A Maximum Volume Quantized Tensor Train-based Optimization and its Application to Reinforcement Learning](#) Konstantin Sozykin, Andrei Chertkov, Roman Schutski, Anh-Huy Phan, Andrzej S CICHOCKI, Ivan Oseledets
- [Diversified Recommendations for Agents with Adaptive Preferences](#) William Brown, Arpit Agarwal
- [A Mixture Of Surprises for Unsupervised Reinforcement Learning](#) Andrew Zhao, Matthieu Lin, Yangguang Li, Yong-jin Liu, Gao Huang
- [Deep Hierarchical Planning from Pixels](#) Danijar Hafner, Kuang-Huei Lee, Ian Fischer, Pieter Abbeel
- [PeRFception: Perception using Radiance Fields](#) Yoonwoo Jeong, Seungjoo Shin, Junha Lee, Chris Choy, Anima Anandkumar, Minsu Cho, Jaesik Park
- [Efficient Submodular Optimization under Noise: Local Search is Robust](#) Lingxiao Huang, Yuyi Wang, Chunxue Yang, Huanjian Zhou
- [On the Generalization Power of the Overfitted Three-Layer Neural Tangent Kernel Model](#) Peizhong Ju, Xiaojun Lin, Ness Shroff
- [Chromatic Correlation Clustering, Revisited](#) Qing Xiu, Kai Han, Jing Tang, Shuang Cui, He Huang
- [Gradient-Free Methods for Deterministic and Stochastic Nonsmooth Nonconvex Optimization](#) Tianyi Lin, Zeyu Zheng, Michael Jordan
- [Multilingual Abusive Comment Detection at Scale for Indic Languages](#) Vikram Gupta, Sumegh Roychowdhury, Mithun Das, Somnath Banerjee, Punyajoy Saha, Binny Mathew, hastagiri prakash vanchinathan, Animesh Mukherjee
- [Generalized Delayed Feedback Model with Post-Click Information in Recommender Systems](#) Jiaqi Yang, De-Chuan Zhan
- [A Communication-Efficient Distributed Gradient Clipping Algorithm for Training Deep Neural Networks](#) Mingrui Liu, Zhenxun Zhuang, Yunwen Lei, Chunyang Liao
- [On Analyzing Generative and Denoising Capabilities of Diffusion-based Deep Generative Models](#) Kamil Deja, Anna Kuzina, Tomasz Trzcinski, Jakub Tomczak
- [On the Tradeoff Between Robustness and Fairness](#) Xinsong Ma, Zekai Wang, Weiwei Liu
- [Learning Expressive Meta-Representations with Mixture of Expert Neural Processes](#) Qi Wang, Herke van Hoof
- [Learning Physical Dynamics with Subequivariant Graph Neural Networks](#) Jiaqi Han, Wenbing Huang, Hengbo Ma, Jiachen Li, Josh Tenenbaum, Chuang Gan
- [DiSC: Differential Spectral Clustering of Features](#) Ram Dyuthi Sristi, Gal Mishne, Ariel Jaffe
- [Learn what matters: cross-domain imitation learning with task-relevant embeddings](#) Tim Franzmeyer, Philip Torr, João F. Henriques
- [UViM: A Unified Modeling Approach for Vision with Learned Guiding Codes](#) Alexander Kolesnikov, André Susano Pinto, Lucas Beyer, Xiaohua Zhai, Jeremiah Harmsen, Neil Houlsby
- [Rotation-Equivariant Conditional Spherical Neural Fields for Learning a Natural Illumination Prior](#) James Gardner, Bernhard Egger, William Smith
- [Proximal Learning With Opponent-Learning Awareness](#) Stephen Zhao, Chris Lu, Roger B. Grosse, Jakob Foerster
- [HyperTree Proof Search for Neural Theorem Proving](#) Guillaume Lample, Timothee Lacroix, Marie-Anne Lachaux, Aurelien Rodriguez, Amaury Hayat, Thibaut Lavril, Gabriel Ebner, Xavier Martinet
- [The Policy-gradient Placement and Generative Routing Neural Networks for Chip Design](#) Ruoyu Cheng, Xianglong Lyu, Yang Li, Junjie Ye, Jianye Hao, Junchi Yan
- [IMED-RL: Regret optimal learning of ergodic Markov decision processes](#) Fabien Pesquerel, Odalric-Ambrym Maillard
- [Coresets for Wasserstein Distributionally Robust Optimization Problems](#) Ruomin Huang, Jiawei Huang, Wenjie Liu, Hu Ding
- [SAMURAI: Shape And Material from Unconstrained Real-world Arbitrary Image collections](#) Mark Boss, Andreas Engelhardt, Abhishek Kar, Yuanzhen Li, Deqing Sun, Jonathan Barron, Hendrik PA Lensch, Varun Jampani
- [Automatic differentiation of nonsmooth iterative algorithms](#) Jerome Bolte, Edouard Pauwels, Samuel Vaiter
- [Wukong: A 100 Million Large-scale Chinese Cross-modal Pre-training Benchmark](#) Jiaxi Gu, Xiaojun Meng, Guansong Lu, Lu Hou, Niu Minzhe, Xiaodan Liang, Lewei Yao, Runhui Huang, Wei Zhang, Xin Jiang, Chunjing XU, Hang Xu
- [Neural Abstractions](#) Alessandro Abate, Alec Edwards, Mirco Giacobbe
- [Few-Shot Non-Parametric Learning with Deep Latent Variable Model](#) Zhiying Jiang, Yiqin Dai, Ji Xin, Ming Li, Jimmy Lin
- [ST-Adapter: Parameter-Efficient Image-to-Video Transfer Learning](#) Juntong Pan, Ziyi Lin, Xiatian Zhu, Jing Shao, Hongsheng Li

- [Distilled Gradient Aggregation: Purify Features for Input Attribution in the Deep Neural Network](#) Giyoung Jeon, Haedong Jeong, Jaesik Choi
- [Temporally Disentangled Representation Learning](#) Weiran Yao, Guangyi Chen, Kun Zhang
- [Can Adversarial Training Be Manipulated By Non-Robust Features?](#) Lue Tao, Lei Feng, Hongxin Wei, Jinfeng Yi, Sheng-Jun Huang, Songcan Chen
- [On the Effectiveness of Fine-tuning Versus Meta-reinforcement Learning](#) Mandi Zhao, Pieter Abbeel, Stephen James
- [Generalizing Goal-Conditioned Reinforcement Learning with Variational Causal Reasoning](#) Wenhao Ding, Haohong Lin, Bo Li, DING ZHAO
- [WinoGAViL: Gamified Association Benchmark to Challenge Vision-and-Language Models](#) Yonatan Bitton, Nitzan Bitton Guetta, Ron Yosef, Yuval Elovici, Mohit Bansal, Gabriel Stanovsky, Roy Schwartz
- [Elucidating the Design Space of Diffusion-Based Generative Models](#) Tero Karras, Miika Aittala, Timo Aila, Samuli Laine
- [TREC: Transient Redundancy Elimination-based Convolution](#) Jiawei Guan, Feng Zhang, Jiesong Liu, Hsin-Hsuan Sung, Ruofan Wu, Xiaoyong Du, Xipeng Shen
- [Chaotic Regularization and Heavy-Tailed Limits for Deterministic Gradient Descent](#) Soon Hoe Lim, Yijun Wan, Umut Simsekli
- [DARE: Disentanglement-Augmented Rationale Extraction](#) Linan Yue, Qi Liu, Yichao Du, Yanqing An, Li Wang, Enhong Chen
- [The Gyro-Structure of Some Matrix Manifolds](#) Xuan Son Nguyen
- [SMPL: Simulated Industrial Manufacturing and Process Control Learning Environments](#) Mohan Zhang, Xiaozhou Wang, Benjamin Decardi-Nelson, Bo Song, An Zhang, Jinfeng Liu, Sile Tao, Jiayi Cheng, Xiaohong Liu, Dengdeng Yu, Matthew Poon, Animesh Garg
- [CoPur: Certifiably Robust Collaborative Inference via Feature Purification](#) Jing Liu, Chulin Xie, Sanmi Koyejo, Bo Li
- [Toward Understanding Privileged Features Distillation in Learning-to-Rank](#) Shuo Yang, Sujay Sanghavi, Holakou Rahmanian, Jan Bakus, Vishwanathan S. V. N.
- [Contrastive and Non-Contrastive Self-Supervised Learning Recover Global and Local Spectral Embedding Methods](#) Randall Balestrieri, Yann LeCun
- [Subgame Solving in Adversarial Team Games](#) Brian Zhang, Luca Carminati, Federico Cacciamani, Gabriele Farina, Pierricardo Olivieri, Nicola Gatti, Tuomas Sandholm
- [A framework for bilevel optimization that enables stochastic and global variance reduction algorithms](#) Mathieu Dagr  ou, Pierre Ablin, Samuel Vaiter, Thomas Moreau
- [Predicting Cellular Responses to Novel Drug Perturbations at a Single-Cell Resolution](#) Leon Hetzel, Simon Boehm, Niki Kilbertus, Stephan G  nnemann, mohammad lotfollahi, Fabian Theis
- [Accelerated Projected Gradient Algorithms for Sparsity Constrained Optimization Problems](#) Jan Harold Alcantara, Ching-pei Lee
- [The Stability-Efficiency Dilemma: Investigating Sequence Length Warmup for Training GPT Models](#) Conglong Li, Minjia Zhang, Yuxiong He
- [Trading off Image Quality for Robustness is not Necessary with Regularized Deterministic Autoencoders](#) Amrutha Saseendran, Kathrin Skubch, Stefan Falkner, Margret Keuper
- [Does Momentum Change the Implicit Regularization on Separable Data?](#) Bohan Wang, Qi Meng, Huishuai Zhang, Ruoyu Sun, Wei Chen, Zhi-Ming Ma, Tie-Yan Liu
- [Generalization Gap in Amortized Inference](#) Mingtian Zhang, Peter Hayes, David Barber
- [Top Two Algorithms Revisited](#) Marc Jourdan, R  my Degenne, Dorian Baudry, Rianne de Heide, Emilie Kaufmann
- [Relation-Constrained Decoding for Text Generation](#) Xiang Chen, Zhixian Yang, Xiaojun Wan
- [Learning General World Models in a Handful of Reward-Free Deployments](#) Yingchen Xu, Jack Parker-Holder, Aldo Pacchiano, Philip Ball, Oleh Rybkin, S Roberts, Tim Rockt  schel, Edward Grefenstette
- [Bringing Image Scene Structure to Video via Frame-Clip Consistency of Object Tokens](#) Elad Ben Avraham, Roei Herzig, Karttikeya Mangalam, Amir Bar, Anna Rohrbach, Leonid Karlinsky, Trevor Darrell, Amir Globerson
- [Probabilistic Transformer: Modelling Ambiguities and Distributions for RNA Folding and Molecule Design](#) J  rg Franke, Frederic Runge, Frank Hutter
- [PulseImpute: A Novel Benchmark Task for Pulsative Physiological Signal Imputation](#) Maxwell Xu, Alexander Moreno, Supriya Nagesh, Varol Aydemir, David Wetter, Santosh Kumar, James M. Rehg
- [Beyond Separability: Analyzing the Linear Transferability of Contrastive Representations to Related Subpopulations](#) Jeff Z. HaoChen, Colin Wei, Ananya Kumar, Tengyu Ma
- [Precise Regret Bounds for Log-loss via a Truncated Bayesian Algorithm](#) Changlong Wu, Mohsen Heidari, Ananth Grama, Wojciech Szpankowski
- [What are the best Systems? New Perspectives on NLP Benchmarking](#) Pierre Colombo, Nathan Noiry, Ekhine Irurozki, Stephan Cl  men  on
- [Learning from Label Proportions by Learning with Label Noise](#) Jianxin Zhang, Yutong Wang, Clay Scott
- [Dynamics of SGD with Stochastic Polyak Stepsizes: Truly Adaptive Variants and Convergence to Exact Solution](#) Antonio Orvieto, Simon Lacoste-Julien, Nicolas Loizou
- [MORA: Improving Ensemble Robustness Evaluation with Model Reweighting Attack](#) yunrui yu, Xitong Gao, Cheng-Zhong Xu
- [Faster and Scalable Algorithms for Densest Subgraph and Decomposition](#) Elfarouk Harb, Kent Quanrud, Chandra Chekuri
- [Extrapolation and Spectral Bias of Neural Nets with Hadamard Product: a Polynomial Net Study](#) Yongtao Wu, Zhenyu Zhu, Fanghui Liu, Grigoris Chrysos, Volkan Cevher

[Report an Issue](#) | [Name Change Policy](#)

- [Cost-Sensitive Self-Training for Optimizing Non-Decomposable Metrics](#) Harsh Rangwani, shrinivas ramasubramanian, Sho Takemori, Kato Takashi, Yuhei Umeda, Venkatesh Babu R
- [Tensor Wheel Decomposition and Its Tensor Completion Application](#) Zhong-Cheng Wu, Ting-Zhu Huang, Liang-Jian Deng, Hong-Xia Dou, Deyu Meng
- [BOND: Benchmarking Unsupervised Outlier Node Detection on Static Attributed Graphs](#) Kay Liu, Yingdong Dou, Yue Zhao, Xueying Ding, Xiyang Hu, Ruitong Zhang, Kaize Ding, Canyu Chen, Hao Peng, Kai Shu, Lichao Sun, Jundong Li, George H Chen, Zhihao Jia, Philip S Yu
- [Movement Penalized Bayesian Optimization with Application to Wind Energy Systems](#) Shyam Sundhar Ramesh, Pier Giuseppe Sessa, Andreas Krause, Ilija Bogunovic
- [Associating Objects and Their Effects in Video through Coordination Games](#) Erika Lu, Forrester Cole, Weidi Xie, Tali Dekel, Bill Freeman, Andrew Zisserman, Michael Rubinstein
- [Point-M2AE: Multi-scale Masked Autoencoders for Hierarchical Point Cloud Pre-training](#) Renrui Zhang, Ziyu Guo, Peng Gao, Rongyao Fang, Bin Zhao, Dong Wang, Yu Qiao, Hongsheng Li
- [Exploring Example Influence in Continual Learning](#) Qing Sun, Fan Lyu, Fanhua Shang, Wei Feng, Liang Wan
- [Subspace clustering in high-dimensions: Phase transitions & Statistical-to-Computational gap](#) Luca Pesce, Bruno Loureiro, Florent Krzakala, Lenka Zdeborová
- [Self-Supervised Fair Representation Learning without Demographics](#) Junyi Chai, Xiaoqian Wang
- [CEDe: A collection of expert-curated datasets with atom-level entity annotations for Optical Chemical Structure Recognition](#) Rodrigo Hormazabal, Changyoung Park, Soonyoung Lee, Sehui Han, Yeonsik Jo, Jaewan Lee, Ahra Jo, Seung Hwan Kim, Jaegul Choo, Moontae Lee, Honglak Lee
- [How Mask Matters: Towards Theoretical Understandings of Masked Autoencoders](#) Qi Zhang, Yifei Wang, Yisen Wang
- [Improving Out-of-Distribution Generalization by Adversarial Training with Structured Priors](#) Qixun Wang, Yifei Wang, Hong Zhu, Yisen Wang
- [Graphein - a Python Library for Geometric Deep Learning and Network Analysis on Biomolecular Structures and Interaction Networks](#) Arian Jamasb, Ramon Viñas Torné, Eric Ma, Yuanqi Du, Charles Harris, Kexin Huang, Dominic Hall, Pietro Lió, Tom Blundell
- [ZeroQuant: Efficient and Affordable Post-Training Quantization for Large-Scale Transformers](#) Zhewei Yao, Reza Yazdani Aminabadi, Minjia Zhang, Xiaoxia Wu, Conglong Li, Yuxiong He
- [Simplified Graph Convolution with Heterophily](#) Sudhanshu Chanpuriya, Cameron Musco
- [Signal Propagation in Transformers: Theoretical Perspectives and the Role of Rank Collapse](#) Lorenzo Noci, Sotiris Anagnostidis, Luca Biggio, Antonio Orvieto, Sidak Pal Singh, Aurelien Lucchi
- [Maximum a posteriori natural scene reconstruction from retinal ganglion cells with deep denoiser priors](#) Eric Wu, Nora Brackbill, Alexander Sher, Alan Litke, Eero Simoncelli, E.J. Chichilnisky
- [Imbalance Trouble: Revisiting Neural-Collapse Geometry](#) Christos Thrampoulidis, Ganesh Ramachandra Kini, Vala Vakilian, Tina Behnia
- [ProtoX: Explaining a Reinforcement Learning Agent via Prototyping](#) Ronilo Ragodos, Tong Wang, Qihang Lin, Xun Zhou
- [NOTE: Robust Continual Test-time Adaptation Against Temporal Correlation](#) Taesik Gong, Jongheon Jeong, Taewon Kim, Yewon Kim, Jinwoo Shin, Sung-Ju Lee
- [Margin-Based Few-Shot Class-Incremental Learning with Class-Level Overfitting Mitigation](#) Yixiong Zou, Shanghang Zhang, Yuhua Li, Ruixuan Li
- [First Hitting Diffusion Models for Generating Manifold, Graph and Categorical Data](#) Mao Ye, Lemeng Wu, Qiang Liu
- [Forecasting Future World Events With Neural Networks](#) Andy Zou, Tristan Xiao, Ryan Jia, Joe Kwon, Mantas Mazeika, Richard Li, Dawn Song, Jacob Steinhardt, Owain Evans, Dan Hendrycks
- [Human-Robotic Prosthesis as Collaborating Agents for Symmetrical Walking](#) Ruofan Wu, Junmin Zhong, Brent Wallace, Xiang Gao, He Huang, Jennie Si
- [Support Recovery in Sparse PCA with Incomplete Data](#) Hanbyul Lee, Qifan Song, Jean Honorio
- [Exponentially Improving the Complexity of Simulating the Weisfeiler-Lehman Test with Graph Neural Networks](#) Anders Aamand, Justin Chen, Piotr Indyk, Shyam Narayanan, Ronitt Rubinfeld, Nicholas Schiefer, Sandeep Silwal, Tal Wagner
- [Zero-shot Transfer Learning within a Heterogeneous Graph via Knowledge Transfer Networks](#) Minji Yoon, John Palowitch, Dustin Zelle, Ziniu Hu, Ruslan Salakhutdinov, Bryan Perozzi
- [Nonlinear Sufficient Dimension Reduction with a Stochastic Neural Network](#) SIQI LIANG, Yan Sun, Faming Liang
- [Autoregressive Perturbations for Data Poisoning](#) Pedro Sandoval-Segura, Vasu Singla, Jonas Geiping, Micah Goldblum, Tom Goldstein, David Jacobs
- [NodeFormer: A Scalable Graph Structure Learning Transformer for Node Classification](#) Qitian Wu, Wentao Zhao, Zenan Li, David P Wipf, Junchi Yan
- [Is Integer Arithmetic Enough for Deep Learning Training?](#) Alireza Ghaffari, Marzieh S. Tahaei, Mohammadreza Tayarani, Masoud Asgharian, Vahid Partovi Nia
- [mRI: Multi-modal 3D Human Pose Estimation Dataset using mmWave, RGB-D, and Inertial Sensors](#) Sizhe An, Yin Li, Umit Ogras
- [Effectiveness of Vision Transformer for Fast and Accurate Single-Stage Pedestrian Detection](#) Jing Yuan, Panagiotis Barmpoutis, Tania Stathaki
- [ESCADA: Efficient Safety and Context Aware Dose Allocation for Precision Medicine](#) Ilker Demirel, Ahmet Alparslan Celik, Cem Tekin
- [Grow and Merge: A Unified Framework for Continuous Categories Discovery](#) Xinwei Zhang, Jianwen Liang, Yutong Feng, Zhi-Fan Wu

- [Category-Level 6D Object Pose Estimation in the Wild: A Semi-Supervised Learning Approach and A New Dataset](#) Yanjie Ze, Xiaolong Wang
- [Change Event Dataset for Discovery from Spatio-temporal Remote Sensing Imagery](#) Utkarsh Mall, Bharath Hariharan, Kavita Bala
- [Improved Algorithms for Neural Active Learning](#) Yikun Ban, Yuheng Zhang, Hanghang Tong, Arindam Banerjee, Jingrui He
- [Robust Calibration with Multi-domain Temperature Scaling](#) Yaodong Yu, Stephen Bates, Yi Ma, Michael Jordan
- [Independence Testing-Based Approach to Causal Discovery under Measurement Error and Linear Non-Gaussian Models](#) Haoyue Dai, Peter Spirtes, Kun Zhang
- [CUP: Critic-Guided Policy Reuse](#) Jin Zhang, Siyuan Li, Chongjie Zhang
- [Local Identifiability of Deep ReLU Neural Networks: the Theory](#) Joachim Bona-Pellissier, François Malgouyres, Francois Bachoc
- [DOMINO: Decomposed Mutual Information Optimization for Generalized Context in Meta-Reinforcement Learning](#) Yao Mu, Yuzheng Zhuang, Fei Ni, Bin Wang, Jianyu Chen, Jianye Hao, Ping Luo
- [Generalization Bounds with Minimal Dependency on Hypothesis Class via Distributionally Robust Optimization](#) Yibo Zeng, Henry Lam
- [QUARK: Controllable Text Generation with Reinforced Unlearning](#) Ximing Lu, Sean Welleck, Jack Hessel, Liwei Jiang, Lianhui Qin, Peter West, Prithviraj Ammanabrolu, Yejin Choi
- [Safe Opponent-Exploitation Subgame Refinement](#) Mingyang Liu, Chengjie Wu, Qihan Liu, Yansen Jing, Jun Yang, Pingzhong Tang, Chongjie Zhang
- [Parameter-free Dynamic Graph Embedding for Link Prediction](#) Jiahao Liu, Dongsheng Li, Hansu Gu, Tun Lu, Peng Zhang, Ning Gu
- [Towards a Unified Framework for Uncertainty-aware Nonlinear Variable Selection with Theoretical Guarantees](#) Wenying Deng, Beau Coker, Rajarshi Mukherjee, Jeremiah Liu, Brent Coull
- [Non-Markovian Reward Modelling from Trajectory Labels via Interpretable Multiple Instance Learning](#) Joseph Early, Tom Bewley, Christine Evers, Sarvapali Ramchurn
- [Explaining Preferences with Shapley Values](#) Robert Hu, Siu Lun Chau, Jaime Ferrando Huertas, Dino Sejdinovic
- [Structuring Uncertainty for Fine-Grained Sampling in Stochastic Segmentation Networks](#) Frank Nussbaum, Jakob Gawlikowski, Julia Niebling
- [Fair Bayes-Optimal Classifiers Under Predictive Parity](#) Xianli Zeng, Edgar Dobriban, Guang Cheng
- [Pre-Train Your Loss: Easy Bayesian Transfer Learning with Informative Priors](#) Ravid Shwartz-Ziv, Micah Goldblum, Hossein Souri, Sanyam Kapoor, Chen Zhu, Yann LeCun, Andrew G. Wilson
- [Unsupervised Adaptation from Repeated Traversals for Autonomous Driving](#) Yurong You, Cheng Perng Phoo, Katie Luo, Travis Zhang, Wei-Lun Chao, Bharath Hariharan, Mark Campbell, Kilian Q. Weinberger
- [Training language models to follow instructions with human feedback](#) Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F. Christiano, Jan Leike, Ryan Lowe
- [Byzantine Spectral Ranking](#) Arnhav Datar, Arun Rajkumar, John Augustine
- [Non-rigid Point Cloud Registration with Neural Deformation Pyramid](#) YANG LI, Tatsuya Harada
- [Provably tuning the ElasticNet across instances](#) Maria-Florina F. Balcan, Misha Khodak, Dravyansh Sharma, Ameet Talwalkar
- [PlasticityNet: Learning to Simulate Metal, Sand, and Snow for Optimization Time Integration](#) Xuan Li, Yadi Cao, Minchen Li, Yin Yang, Craig Schroeder, Chenfanfu Jiang
- [Synthetic Model Combination: An Instance-wise Approach to Unsupervised Ensemble Learning](#) Alex Chan, Mihaela van der Schaar
- [Power and limitations of single-qubit native quantum neural networks](#) Zhan Yu, Hongshun Yao, Mujin Li, Xin Wang
- [Hedging as Reward Augmentation in Probabilistic Graphical Models](#) Debarun Bhattacharjya, Radu Marinescu
- [Debiased, Longitudinal and Coordinated Drug Recommendation through Multi-Visit Clinic Records](#) Hongda Sun, Shufang Xie, Shuqi Li, Yuhang Chen, Ji-Rong Wen, Rui Yan
- [Disentangling Causal Effects from Sets of Interventions in the Presence of Unobserved Confounders](#) Olivier Jeunen, Ciarán Gilligan-Lee, Rishabh Mehrotra, Mounia Lalmas
- [MATE: Benchmarking Multi-Agent Reinforcement Learning in Distributed Target Coverage Control](#) Xuehai Pan, Mickel Liu, Fangwei Zhong, Yaodong Yang, Song-Chun Zhu, Yizhou Wang
- [Receding Horizon Inverse Reinforcement Learning](#) Yiqing Xu, Wei Gao, David Hsu
- [A Reparametrization-Invariant Sharpness Measure Based on Information Geometry](#) Cheongjae Jang, Sungyoon Lee, Frank Park, Yung-Kyun Noh
- [Hyper-Representations as Generative Models: Sampling Unseen Neural Network Weights](#) Konstantin Schürholt, Boris Knyazev, Xavier Giró-i-Nieto, Damian Borth
- [Multi-Game Decision Transformers](#) Kuang-Huei Lee, Ofir Nachum, Mengjiao (Sherry) Yang, Lisa Lee, Daniel Freeman, Sergio Guadarrama, Ian Fischer, Winnie Xu, Eric Jang, Henryk Michalewski, Igor Mordatch
- [Improving Transformer with an Admixture of Attention Heads](#) Tan Nguyen, Tam Nguyen, Hai Do, Khai Nguyen, Vishwanath Saragadam, Minh Pham, Khuong Duy Nguyen, Nhat Ho, Stanley Osher
- [Flexible Diffusion Modeling of Long Videos](#) William Harvey, Saeid Naderiparizi, Vaden Masrani, Christian Weilbach, Frank Wood
- [Towards Reasonable Budget Allocation in Untargeted Graph Structure Attacks via Gradient Debias](#) Zihan Liu, Yun Luo, Lirong Wu, Zicheng Liu, Stan Z. Li

- [ToDD: Topological Compound Fingerprinting in Computer-Aided Drug Discovery](#) Andaç Demir, Baris Coskunuzer, Yulia Gel, Ignacio Segovia-Dominguez, Yuzhou Chen, Bulent Kiziltan
- [Geo-SIC: Learning Deformable Geometric Shapes in Deep Image Classifiers](#) Jian Wang, Miaomiao Zhang
- [A Survey and Datasheet Repository of Publicly Available US Criminal Justice Datasets](#) Miri Zilka, Bradley Butcher, Adrian Weller
- [Segmenting Moving Objects via an Object-Centric Layered Representation](#) Junyu Xie, Weidi Xie, Andrew Zisserman
- [NAS-Bench-Suite-Zero: Accelerating Research on Zero Cost Proxies](#) Arjun Krishnakumar, Colin White, Arber Zela, Renbo Tu, Mahmoud Safari, Frank Hutter
- [On the Strong Correlation Between Model Invariance and Generalization](#) Weijian Deng, Stephen Gould, Liang Zheng
- [On Scalable Testing of Samplers](#) Yash Pote, Kuldeep S Meel
- [How and Why to Manipulate Your Own Agent: On the Incentives of Users of Learning Agents](#) Yoav Kolumbus, Noam Nisan
- [LIPS - Learning Industrial Physical Simulation benchmark suite](#) Milad LEYLI ABADI, Antoine Marot, Jérôme Picault, David Danan, Mouadh Yagoubi, Benjamin Donnot, Seif Attoui, Pavel Dimitrov, Asma Farjallah, Clement Etienam
- [FlyView: a bio-informed optical flow truth dataset for visual navigation using panoramic stereo vision](#) Alix Leroy, Graham W. Taylor
- [Controllable Text Generation with Neurally-Decomposed Oracle](#) Tao Meng, Sidi Lu, Nanyun Peng, Kai-Wei Chang
- [Active Learning Helps Pretrained Models Learn the Intended Task](#) Alex Tamkin, Dat Nguyen, Salil Deshpande, Jesse Mu, Noah Goodman
- [Pragmatically Learning from Pedagogical Demonstrations in Multi-Goal Environments](#) Hugo Caselles-Dupré, Olivier Sigaud, Mohamed CHETOUANI
- [Beyond black box densities: Parameter learning for the deviated components](#) Dat Do, Nhat Ho, XuanLong Nguyen
- [Statistical, Robustness, and Computational Guarantees for Sliced Wasserstein Distances](#) Sloan Nietert, Ziv Goldfeld, Ritwik Sadhu, Kengo Kato
- [Oracle Inequalities for Model Selection in Offline Reinforcement Learning](#) Jonathan N Lee, George Tucker, Ofir Nachum, Bo Dai, Emma Brunskill
- [Model-Based Opponent Modeling](#) XiaoPeng Yu, Jiechuan Jiang, Wanpeng Zhang, Haobin Jiang, Zongqing Lu
- [What is Where by Looking: Weakly-Supervised Open-World Phrase-Grounding without Text Inputs](#) Tal Shaharabany, Yoad Tewel, Lior Wolf
- [DeepMed: Semiparametric Causal Mediation Analysis with Debiased Deep Learning](#) Siqi Xu, Lin Liu, Zhonghua Liu
- [On the Learning Mechanisms in Physical Reasoning](#) Shiqian Li, Kewen Wu, Chi Zhang, Yixin Zhu
- [A Continuous Time Framework for Discrete Denoising Models](#) Andrew Campbell, Joe Benton, Valentin De Bortoli, Thomas Rainforth, George Deligiannidis, Arnaud Doucet
- [Simple Mechanisms for Welfare Maximization in Rich Advertising Auctions](#) Gagan Aggarwal, Kshipra Bhawalkar, Aranyak Mehta, Divyarthi Mohan, Alexandros Psomas
- [Continuously Tempered PDMP samplers](#) Matthew Sutton, Robert Salomone, Augustin Chevallier, Paul Fearnhead
- [Distributed Influence-Augmented Local Simulators for Parallel MARL in Large Networked Systems](#) Miguel Suaui, Jinke He, Mustafa Mert Çelikok, Matthijs Spaan, Frans Oliehoek
- [Tight Mutual Information Estimation With Contrastive Fenchel-Legendre Optimization](#) Qing Guo, Junya Chen, Dong Wang, Yuewei Yang, Xinwei Deng, Jing Huang, Larry Carin, Fan Li, Chenyang Tao
- [Exploring the Algorithm-Dependent Generalization of AUPRC Optimization with List Stability](#) Peisong Wen, Qianqian Xu, Zhiyong Yang, Yuan He, Qingming Huang
- [LTMD: Learning Improvement of Spiking Neural Networks with Learnable Thresholding Neurons and Moderate Dropout](#) SIQI WANG, Tee Hiang Cheng, Meng-Hiot Lim
- [TransBoost: Improving the Best ImageNet Performance using Deep Transduction](#) Omer Belhasin, Guy Bar-Shalom, Ran El-Yaniv
- [Sparse Probabilistic Circuits via Pruning and Growing](#) Meihua Dang, Anji Liu, Guy Van den Broeck
- [Adam Can Converge Without Any Modification On Update Rules](#) Yushun Zhang, Congliang Chen, Naichen Shi, Ruoyu Sun, Zhi-Quan Luo
- [Nonlinear MCMC for Bayesian Machine Learning](#) James Vuckovic
- [Continual learning: a feature extraction formalization, an efficient algorithm, and fundamental obstructions](#) Binghui Peng, Andrej Risteski
- [IKEA-Manual: Seeing Shape Assembly Step by Step](#) Ruocheng Wang, Yunzhi Zhang, Jiayuan Mao, Ran Zhang, Chin-Yi Cheng, Jiajun Wu
- [M³ViT: Mixture-of-Experts Vision Transformer for Efficient Multi-task Learning with Model-Accelerator Co-design](#) hanxue liang, Zhiwen Fan, Rishov Sarkar, Ziyu Jiang, Tianlong Chen, Kai Zou, Yu Cheng, Cong Hao, Zhangyang Wang
- [When to Make Exceptions: Exploring Language Models as Accounts of Human Moral Judgment](#) Zhijing Jin, Sydney Levine, Fernando Gonzalez Adauto, Ojasv Kamal, Maarten Sap, Mrinmaya Sachan, Rada Mihalcea, Josh Tenenbaum, Bernhard Schölkopf
- [Exponential Family Model-Based Reinforcement Learning via Score Matching](#) Gene Li, Junbo Li, Anmol Kabra, Nati Srebro, Zhaoran Wang, Zhuoran Yang
- [Monte Carlo Tree Search based Variable Selection for High Dimensional Bayesian Optimization](#) Lei Song, Ke Xue, Xiaobin Huang, Chao Qian
- [OpenSRH: optimizing brain tumor surgery using intraoperative stimulated Raman histology](#) Cheng Jiang, Asadur Chowdury, Xinhai Hou, Akhil Kondepudi, Christian Freudiger, Kyle Conway, Sandra Camelo-Piragua, Daniel Orringer, Honglak Lee, Todd Hollon
- [Assistive Teaching of Motor Control Tasks to Humans](#) Megha Srivastava, Erdem Biyik, Suvir Mirchandani, Noah Goodman, Dorsa Sadigh
- [Retrospective Adversarial Replay for Continual Learning](#) Lilly Kumari, Shengjie Wang, Tianyi Zhou, Jeff A Bilmes

- [CASA: Category-agnostic Skeletal Animal Reconstruction](#) Yuefan Wu, Zeyuan Chen, Shaowei Liu, Zhongzheng Ren, Shenlong Wang
- [Regret Bounds for Information-Directed Reinforcement Learning](#) Botao Hao, Tor Lattimore
- [Adversarial Reprogramming Revisited](#) Matthias Englert, Ranko Lazic
- [Coded Residual Transform for Generalizable Deep Metric Learning](#) Shichao Kan, Yixiong Liang, Min Li, Yigang Cen, Jianxin Wang, Zhihai He
- [When Does Differentially Private Learning Not Suffer in High Dimensions?](#) Xuechen Li, Daogao Liu, Tatsunori B. Hashimoto, Huseyin A. Inan, Janardhan Kulkarni, Yin-Tat Lee, Abhradeep Guha Thakurta
- [Nearly Optimal Best-of-Both-Worlds Algorithms for Online Learning with Feedback Graphs](#) Shinji Ito, Taira Tsuchiya, Junya Honda
- [Few-shot Task-agnostic Neural Architecture Search for Distilling Large Language Models](#) Dongkuan (DK) Xu, Subhabrata Mukherjee, Xiaodong Liu, Debadeepta Dey, Wenhui Wang, Xiang Zhang, Ahmed Awadallah, Jianfeng Gao
- [Universal Rates for Interactive Learning](#) Steve Hanneke, Amin Karbasi, Shay Moran, Grigoris Velegkas
- [Learning Active Camera for Multi-Object Navigation](#) Peihao Chen, Dongyu Ji, Kunyang Lin, Weiwen Hu, Wenbing Huang, Thomas Li, Mingkui Tan, Chuang Gan
- [Knowledge Distillation: Bad Models Can Be Good Role Models](#) Gal Kaplun, Eran Malach, Preetum Nakkiran, Shai Shalev-Shwartz
- [On Computing Probabilistic Explanations for Decision Trees](#) Marcelo Arenas, Pablo Barceló, Miguel Romero Orth, Bernardo Subercaseaux
- [Masked Autoencoders that Listen](#) Po-Yao Huang, Hu Xu, Juncheng Li, Alexei Baevski, Michael Auli, Wojciech Galuba, Florian Metze, Christoph Feichtenhofer
- [AUTOMATA: Gradient Based Data Subset Selection for Compute-Efficient Hyper-parameter Tuning](#) Krishnateja Killamsetty, Guttu Sai Abhishek, Aakriti Lnu, Ganesh Ramakrishnan, Alexandre Evfimievski, Lucian Popa, Rishabh Iyer
- [DeepTOP: Deep Threshold-Optimal Policy for MDPs and RMABs](#) Khaled Nakhleh, I-Hong Hou
- [Is a Modular Architecture Enough?](#) Sarthak Mittal, Yoshua Bengio, Guillaume Lajoie
- [Exploration via Planning for Information about the Optimal Trajectory](#) Viraj Mehta, Ian Char, Joseph Abbate, Rory Conlin, Mark Boyer, Stefano Ermon, Jeff Schneider, Willie Neiswanger
- [Subquadratic Kronecker Regression with Applications to Tensor Decomposition](#) Matthew Fahrbach, Gang Fu, Mehrdad Ghadiri
- [Robust Anytime Learning of Markov Decision Processes](#) Marnix Suilen, Thiago D. Simão, David Parker, Nils Jansen
- [Discovering Design Concepts for CAD Sketches](#) Yuezhi Yang, Hao Pan
- [Learning Distributed and Fair Policies for Network Load Balancing as Markov Potential Game](#) Zhiyuan Yao, Zihan Ding
- [Para-CFlows: \$C^k\$ -universal diffeomorphism approximators as superior neural surrogates](#) Junlong Lyu, Zhitang Chen, Chang Feng, Wenjing Cun, Shengyu Zhu, Yanhui Geng, ZHIJIE XU, Chen Yongwei
- [NCP: Neural Correspondence Prior for Effective Unsupervised Shape Matching](#) Souhaib Attaiki, Maks Ovsjanikov
- [Efficient Spatially Sparse Inference for Conditional GANs and Diffusion Models](#) Muyang Li, Ji Lin, Chenlin Meng, Stefano Ermon, Song Han, Jun-Yan Zhu
- [A Greek Parliament Proceedings Dataset for Computational Linguistics and Political Analysis](#) Konstantina Dritsa, Aikaterini Thoma, Ioannis Pavlopoulos, Panos Louridas
- [Multi-objective Deep Data Generation with Correlated Property Control](#) Shiyu Wang, Xiaojie Guo, Xuanyang Lin, Bo Pan, Yuanqi Du, Yinkai Wang, Yanfang Ye, Ashley Petersen, Austin Leitgeb, Saleh Alkhalifa, Kevin Minbiole, William M. Wuest, Amarda Shehu, Liang Zhao
- [Domain Adaptation meets Individual Fairness. And they get along.](#) Debarghya Mukherjee, Felix Petersen, Mikhail Yurochkin, Yuekai Sun
- [An Asymptotically Optimal Batched Algorithm for the Dueling Bandit Problem](#) Arpit Agarwal, Rohan Ghuge, viswanath nagarajan
- [Enhanced Bilevel Optimization via Bregman Distance](#) Feihu Huang, Junyi Li, Shangqian Gao, Heng Huang
- [SAVi+ +: Towards End-to-End Object-Centric Learning from Real-World Videos](#) Gamaleldin Elsayed, Aravindh Mahendran, Sjoerd van Steenkiste, Klaus Greff, Michael C. Mozer, Thomas Kipf
- [Reincarnating Reinforcement Learning: Reusing Prior Computation to Accelerate Progress](#) Rishabh Agarwal, Max Schwarzer, Pablo Samuel Castro, Aaron C. Courville, Marc Bellemare
- [Global Linear and Local Superlinear Convergence of IRLS for Non-Smooth Robust Regression](#) Liangzu Peng, Christian Kümmerle, Rene Vidal
- [Rashomon Capacity: A Metric for Predictive Multiplicity in Classification](#) Hsiang Hsu, Flavio Calmon
- [Inducing Equilibria via Incentives: Simultaneous Design-and-Play Ensures Global Convergence](#) Boyi Liu, Jiayang Li, Zhuoran Yang, Hoi-To Wai, Mingyi Hong, Yu Nie, Zhaoran Wang
- [Learning with convolution and pooling operations in kernel methods](#) Theodor Misiakiewicz, Song Mei
- [In Differential Privacy, There is Truth: on Vote-Histogram Leakage in Ensemble Private Learning](#) JIAQI WANG, Roei Schuster, I Shumailov, David Lie, Nicolas Papernot
- [Estimating graphical models for count data with applications to single-cell gene network](#) Feiyi Xiao, Junjie Tang, Huaying Fang, Ruibin Xi
- [Online Minimax Multiobjective Optimization: Multicalibating and Other Applications](#) Daniel Lee, Georgy Noarov, Malleesh Pai, Aaron Roth
- [Roadblocks for Temporarily Disabling Shortcuts and Learning New Knowledge](#) Hongjing Niu, Hanting Li, Feng Zhao, Bin Li
- [Random Rank: The One and Only Strategyproof and Proportionally Fair Randomized Facility Location Mechanism](#) Haris Aziz, Alexander Lam, Mashbat Suzuki, Toby Walsh
- [Distributionally robust weighted k-nearest neighbors](#) Shixiang Zhu, Liyan Xie, Minghe Zhang, Rui Gao, Yao Xie

- [Neural Lyapunov Control of Unknown Nonlinear Systems with Stability Guarantees](#) Ruikun Zhou, Thanin Quartz, Hans De Sterck, Jun Liu
- [Minimax Regret for Cascading Bandits](#) Daniel Vial, Sujay Sanghavi, Sanjay Shakkottai, R. Srikant
- [Mean Estimation with User-level Privacy under Data Heterogeneity](#) Rachel Cummings, Vitaly Feldman, Audra McMillan, Kunal Talwar
- [Propo: a Message Passing Framework for Customizable and Composable Learning Algorithms](#) Paavo Parmas, Takuma Seno
- [A Lower Bound of Hash Codes' Performance](#) Xiaosu Zhu, Jingkuan Song, Yu Lei, Lianli Gao, Hengtao Shen
- [Self-Supervised Image Restoration with Blurry and Noisy Pairs](#) Zhilu Zhang, RongJian Xu, Ming Liu, Zifei Yan, Wangmeng Zuo
- [Embracing Consistency: A One-Stage Approach for Spatio-Temporal Video Grounding](#) Yang Jin, yongzhi li, Zehuan Yuan, Yadong Mu
- [Pitfalls of Epistemic Uncertainty Quantification through Loss Minimisation](#) Viktor Bengs, Eyke Hüllermeier, Willem Waegeman
- [Pile of Law: Learning Responsible Data Filtering from the Law and a 256GB Open-Source Legal Dataset](#) Peter Henderson, Mark Krass, Lucia Zheng, Neel Guha, Christopher D Manning, Dan Jurafsky, Daniel Ho
- [Active-Passive SimStereo - Benchmarking the Cross-Generalization Capabilities of Deep Learning-based Stereo Methods](#) Laurent Jospin, Allen Antony, Lian Xu, Hamid Laga, Farid Boussaid, Mohammed Bennamoun
- [Back Razor: Memory-Efficient Transfer Learning by Self-Sparsified Backpropagation](#) Ziyu Jiang, Xuxi Chen, Xueqin Huang, Xianzhi Du, Denny Zhou, Zhangyang Wang
- [Patching open-vocabulary models by interpolating weights](#) Gabriel Ilharco, Mitchell Wortsman, Samir Yitzhak Gadre, Shuran Song, Hannaneh Hajishirzi, Simon Kornblith, Ali Farhadi, Ludwig Schmidt
- [Training stochastic stabilized supralinear networks by dynamics-neutral growth](#) Wayne Soo, Mate Lengyel
- [Post-hoc estimators for learning to defer to an expert](#) Harikrishna Narasimhan, Wittawat Jitkrittum, Aditya K. Menon, Ankit Rawat, Sanjiv Kumar
- [Learning Generalizable Part-based Feature Representation for 3D Point Clouds](#) Xin Wei, Xiang Gu, Jian Sun
- [FourierFormer: Transformer Meets Generalized Fourier Integral Theorem](#) Tan Nguyen, Minh Pham, Tam Nguyen, Khai Nguyen, Stanley Osher, Nhat Ho
- [Learning Multi-resolution Functional Maps with Spectral Attention for Robust Shape Matching](#) Lei Li, Nicolas Donati, Maks Ovsjanikov
- [Label-invariant Augmentation for Semi-Supervised Graph Classification](#) Han Yue, Chunhui Zhang, Chuxu Zhang, Hongfu Liu
- [Practical Adversarial Multivalid Conformal Prediction](#) Osbert Bastani, Varun Gupta, Christopher Jung, Georgy Noarov, Ramya Ramalingam, Aaron Roth
- [Test-Time Training with Masked Autoencoders](#) Yossi Gandelsman, Yu Sun, Xinlei Chen, Alexei Efros
- [Constraining Gaussian Processes to Systems of Linear Ordinary Differential Equations](#) Andreas Besginow, Markus Lange-Hegermann
- [So3krates: Equivariant attention for interactions on arbitrary length-scales in molecular systems](#) Thorben Frank, Oliver Unke, Klaus-Robert Müller
- [HyperDomainNet: Universal Domain Adaptation for Generative Adversarial Networks](#) Aibek Alanov, Vadim Titov, Dmitry P. Vetrov
- [Multiagent Q-learning with Sub-Team Coordination](#) Wenhan Huang, Kai Li, Kun Shao, Tianze Zhou, Matthew Taylor, Jun Luo, Dongge Wang, Hangyu Mao, Jianye Hao, Jun Wang, Xiaotie Deng
- [CLiMB: A Continual Learning Benchmark for Vision-and-Language Tasks](#) Tejas Srinivasan, Ting-Yun Chang, Leticia Pinto Alva, Georgios Chochlakis, Mohammad Rostami, Jesse Thomason
- [Bidirectional Learning for Offline Infinite-width Model-based Optimization](#) Can Chen, Yingxueff Zhang, Jie Fu, Xue (Steve) Liu, Mark Coates
- [Differentially Private Model Compression](#) FatemehSadat Mireshghallah, Arturs Backurs, Huseyin A. Inan, Lukas Wutschitz, Janardhan Kulkarni
- [Dual-Curriculum Contrastive Multi-Instance Learning for Cancer Prognosis Analysis with Whole Slide Images](#) CHAO TU, YU ZHANG, Zhenyuan Ning
- [Meta-Complementing the Semantics of Short Texts in Neural Topic Models](#) Delvin Ce Zhang, Hady Lauw
- [Unified Optimal Transport Framework for Universal Domain Adaptation](#) Wanxing Chang, Ye Shi, Hoang Tuan, Jingya Wang
- [GBA: A Tuning-free Approach to Switch between Synchronous and Asynchronous Training for Recommendation Models](#) Wenbo Su, Yuanxing Zhang, Yufeng Cai, Kaixu Ren, Pengjie Wang, Huimin Yi, Yue Song, Jing Chen, Hongbo Deng, Jian Xu, Lin Qu, Bo Zheng
- [TUSK: Task-Agnostic Unsupervised Keypoints](#) Yuhe Jin, Weiwei Sun, Jan Hosang, Eduard Trulls, Kwang Moo Yi
- [Multi-block Min-max Bilevel Optimization with Applications in Multi-task Deep AUC Maximization](#) Quanqi Hu, YONGJIAN ZHONG, Tianbao Yang
- [Coresets for Vertical Federated Learning: Regularized Linear Regression and \$K\$ -Means Clustering](#) Lingxiao Huang, Zhize Li, Jialin Sun, Haoyu Zhao
- [Class-Aware Adversarial Transformers for Medical Image Segmentation](#) Chenyu You, Ruihan Zhao, Fenglin Liu, Siyuan Dong, Sandeep Chinchali, Ufuk Topcu, Lawrence Staib, James Duncan
- [MCL-GAN: Generative Adversarial Networks with Multiple Specialized Discriminators](#) Jinyoung Choi, Bohyung Han
- [Unsupervised Visual Representation Learning via Mutual Information Regularized Assignment](#) Dong Hoon Lee, Sungik Choi, Hyunwoo J. Kim, Sae-Young Chung
- [Mind Reader: Reconstructing complex images from brain activities](#) Sikun Lin, Thomas Sprague, Ambuj K Singh
- [A Few Expert Queries Suffices for Sample-Efficient RL with Resets and Linear Value Approximation](#) Philip Amortila, Nan Jiang, Dhruv Madeka, Dean P. Foster

- [Subsidiary Prototype Alignment for Universal Domain Adaptation](#) Jogendra Nath Kundu, Suvaansh Bhambri, Akshay R Kulkarni, Hiran Sarkar, Varun Jampani, Venkatesh Babu R
- [Dynamic Inverse Reinforcement Learning for Characterizing Animal Behavior](#) Zoe Ashwood, Aditi Jha, Jonathan W. Pillow
- [FedRolex: Model-Heterogeneous Federated Learning with Rolling Sub-Model Extraction](#) Samiul Alam, Luyang Liu, Ming Yan, Mi Zhang
- [Multi-Lingual Acquisition on Multimodal Pre-training for Cross-modal Retrieval](#) Liang Zhang, Anwen Hu, Qin Jin
- [Manifold Interpolating Optimal-Transport Flows for Trajectory Inference](#) Guillaume Huguet, Daniel Sumner Magruder, Alexander Tong, Oluwadamilola Fasina, Manik Kuchroo, Guy Wolf, Smita Krishnaswamy
- [Depth is More Powerful than Width with Prediction Concatenation in Deep Forest](#) Shen-Huan Lyu, Yi-Xiao He, Zhi-Hua Zhou
- [Logical Activation Functions: Logit-space equivalents of Probabilistic Boolean Operators](#) Scott Lowe, Robert Earle, Jason d'Eon, Thomas Trappenberg, Sageev Oore
- [An Investigation into Whitening Loss for Self-supervised Learning](#) Xi Weng, Lei Huang, Lei Zhao, Rao Anwer, Salman H. Khan, Fahad Shahbaz Khan
- [Geometric Knowledge Distillation: Topology Compression for Graph Neural Networks](#) Chenxiao Yang, Qitian Wu, Junchi Yan
- [A Benchmark for Compositional Visual Reasoning](#) Aimen Zerroug, Mohit Vaishnav, Julien Colin, Sebastian Musslick, Thomas Serre
- [Learning Articulated Rigid Body Dynamics with Lagrangian Graph Neural Network](#) Ravinder Bhattoo, Sayan Ranu, N M Anoop Krishnan
- [Myriad: a real-world testbed to bridge trajectory optimization and deep learning](#) Nikolaus Howe, Simon Dufort-Labbé, Nitarshan Rajkumar, Pierre-Luc Bacon
- [Batch Bayesian optimisation via density-ratio estimation with guarantees](#) Rafael Oliveira, Louis Tiao, Fabio T. Ramos
- [Amplifying Membership Exposure via Data Poisoning](#) Yufei Chen, Chao Shen, Yun Shen, Cong Wang, Yang Zhang
- [Boosting the Transferability of Adversarial Attacks with Reverse Adversarial Perturbation](#) Zeyu Qin, Yanbo Fan, Yi Liu, Li Shen, Yong Zhang, Jue Wang, Baoyuan Wu
- [Extra-Newton: A First Approach to Noise-Adaptive Accelerated Second-Order Methods](#) Kimon Antonakopoulos, Ali Kavis, Volkan Cevher
- [FETA: Towards Specializing Foundational Models for Expert Task Applications](#) Amit Alfassy, Assaf Arbelle, Oshri Halimi, Sivan Harary, Roei Herzig, Eli Schwartz, Rameswar Panda, Michele Dolfi, Christoph Auer, Peter Staar, Kate Saenko, Rogerio Feris, Leonid Karlinsky
- [Global Convergence of Federated Learning for Mixed Regression](#) Lili Su, Jiaming Xu, Pengkun Yang
- [BayesPCN: A Continually Learnable Predictive Coding Associative Memory](#) Jinsoo Yoo, Frank Wood
- [Optimizing Data Collection for Machine Learning](#) Rafid Mahmood, James Lucas, Jose M. Alvarez, Sanja Fidler, Marc Law
- [DP-PCA: Statistically Optimal and Differentially Private PCA](#) Xiyang Liu, Weihao Kong, Prateek Jain, Sewoong Oh
- [Semantic Probabilistic Layers for Neuro-Symbolic Learning](#) Kareem Ahmed, Stefano Teso, Kai-Wei Chang, Guy Van den Broeck, Antonio Vergari
- [Adapting to Online Label Shift with Provable Guarantees](#) Yong Bai, Yu-Jie Zhang, Peng Zhao, Masashi Sugiyama, Zhi-Hua Zhou
- [CAGroup3D: Class-Aware Grouping for 3D Object Detection on Point Clouds](#) Haiyang Wang, Lihe Ding, Shaocong Dong, Shaoshuai Shi, Aoxue Li, Jianan Li, Zhenguo Li, Liwei Wang
- [Revisiting Injective Attacks on Recommender Systems](#) Haoyang Li, Shimin Di, Lei Chen
- [Learning-based Motion Planning in Dynamic Environments Using GNNs and Temporal Encoding](#) Ruipeng Zhang, Chenning Yu, Jingkai Chen, Chuchu Fan, Sicun Gao
- [An empirical analysis of compute-optimal large language model training](#) Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, Thomas Hennigan, Eric Noland, Katherine Millican, George van den Driessche, Bogdan Damoc, Aurelia Guy, Simon Osindero, Karén Simonyan, Erich Elsen, Oriol Vinyals, Jack Rae, Laurent Sifre
- [Robust Imitation via Mirror Descent Inverse Reinforcement Learning](#) Dong-Sig Han, Hyunseo Kim, Hyundo Lee, JeHwan Ryu, Byoung-Tak Zhang
- [Characterizing Datapoints via Second-Split Forgetting](#) Pratyush Maini, Saurabh Garg, Zachary Lipton, J. Zico Kolter
- [VoiceBlock: Privacy through Real-Time Adversarial Attacks with Audio-to-Audio Models](#) Patrick O'Reilly, Andreas Bugler, Keshav Bhandari, Max Morrison, Bryan Pardo
- [A Theoretical View on Sparsely Activated Networks](#) Cenk Baykal, Nishanth Dikkala, Rina Panigrahy, Cyrus Rashtchian, Xin Wang
- [Invariance Learning based on Label Hierarchy](#) Shoji Toyota, Kenji Fukumizu
- [Effects of Data Geometry in Early Deep Learning](#) Saket Tiwari, George Konidakis
- [When to Intervene: Learning Optimal Intervention Policies for Critical Events](#) Niranjana Damera Venkata, Chiranjib Bhattacharyya
- [Draft-and-Revise: Effective Image Generation with Contextual RQ-Transformer](#) Doyup Lee, Chiheon Kim, Saehoon Kim, Minsu Cho, WOOK SHIN HAN
- [Finite-Sample Maximum Likelihood Estimation of Location](#) Shivam Gupta, Jasper Lee, Eric Price, Paul Valiant
- [GENIE: Higher-Order Denoising Diffusion Solvers](#) Tim Dockhorn, Arash Vahdat, Karsten Kreis
- [Provably Adversarially Robust Detection of Out-of-Distribution Data \(Almost\) for Free](#) Alexander Meinke, Julian Bitterwolf, Matthias Hein
- [Efficient Graph Similarity Computation with Alignment Regularization](#) Wei Zhuo, Guang Tan
- [Tsetlin Machine for Solving Contextual Bandit Problems](#) Raihan Seraj, Jivitesh Sharma, Ole-Christoffer Granmo
- [Prune and distill: similar reformatting of image information along rat visual cortex and deep neural networks](#) Paolo Muratore, Sina Tafazoli, Eugenio Piasini, Alessandro Laio, Davide Zoccolan

- [Matryoshka Representation Learning](#) Aditya Kusupati, Gantavya Bhatt, Aniket Rege, Matthew Wallingford, Aditya Sinha, Vivek Ramanujan, William Howard-Snyder, Kaifeng Chen, Sham Kakade, Prateek Jain, Ali Farhadi
- [Off-Policy Evaluation with Deficient Support Using Side Information](#) Nicolò Felicioni, Maurizio Ferrari Dacrema, Marcello Restelli, Paolo Cremonesi
- [The Nature of Temporal Difference Errors in Multi-step Distributional Reinforcement Learning](#) Yunhao Tang, Remi Munos, Mark Rowland, Bernardo Avila Pires, Will Dabney, Marc Bellemare
- [GraphDE: A Generative Framework for Debiased Learning and Out-of-Distribution Detection on Graphs](#) Zenan Li, Qitian Wu, Fan Nie, Junchi Yan
- [Expectation-Maximization Contrastive Learning for Compact Video-and-Language Representations](#) Peng Jin, Jinfa Huang, Fenglin Liu, Xian Wu, Shen Ge, Guoli Song, David Clifton, Jie Chen
- [Accelerating Sparse Convolution with Column Vector-Wise Sparsity](#) Yijun Tan, Kai Han, Kang Zhao, Xianzhi Yu, Zidong Du, Yunji Chen, Yunhe Wang, Jun Yao
- [GPT3.int8\(\): 8-bit Matrix Multiplication for Transformers at Scale](#) Tim Dettmers, Mike Lewis, Younes Belkada, Luke Zettlemoyer
- [Posterior Refinement Improves Sample Efficiency in Bayesian Neural Networks](#) Agustinus Kristiadi, Runa Eschenhagen, Philipp Hennig
- [Micro and Macro Level Graph Modeling for Graph Variational Auto-Encoders](#) Kiarash Zahirnia, Oliver Schulte, Parmis Naddaf, Ke Li
- [Spectral Bias Outside the Training Set for Deep Networks in the Kernel Regime](#) Benjamin Bowman, Guido F. Montufar
- [The Unreliability of Explanations in Few-shot Prompting for Textual Reasoning](#) Xi Ye, Greg Durrett
- [Self-Similarity Priors: Neural Collages as Differentiable Fractal Representations](#) Michael Poli, Winnie Xu, Stefano Massaroli, Chenlin Meng, Kuno Kim, Stefano Ermon
- [Scalable Multi-agent Covering Option Discovery based on Kronecker Graphs](#) Jiayu Chen, Jingdi Chen, Tian Lan, Vaneet Aggarwal
- [An Algorithm for Learning Switched Linear Dynamics from Data](#) Guillaume Berger, Monal Narasimhamurthy, Kandai Watanabe, Morteza Lahijanian, Sriram Sankaranarayanan
- [LECO: Learnable Episodic Count for Task-Specific Intrinsic Reward](#) Daejin Jo, Sungwoong Kim, Daniel Nam, Taehwan Kwon, Seungeun Rho, Jongmin Kim, Donghoon Lee
- [Generalised Implicit Neural Representations](#) Daniele Grattarola, Pierre Vanderghenst
- [EZNAS: Evolving Zero-Cost Proxies For Neural Architecture Scoring](#) Yash Akhauri, Juan Munoz, Nilesh Jain, Ravishankar Iyer
- [Personalized Federated Learning towards Communication Efficiency, Robustness and Fairness](#) Shiyun Lin, Yuze Han, Xiang Li, Zhihua Zhang
- [Local Linear Convergence of Gradient Methods for Subspace Optimization via Strict Complementarity](#) Ron Fisher, Dan Garber
- [Joint Learning of 2D-3D Weakly Supervised Semantic Segmentation](#) Hyeokjun Kweon, Kuk-Jin Yoon
- [RNNs of RNNs: Recursive Construction of Stable Assemblies of Recurrent Neural Networks](#) Leo Kozachkov, Michaela Ennis, Jean-Jacques Slotine
- [Training Spiking Neural Networks with Event-driven Backpropagation](#) Yaoyu Zhu, Zhaofei Yu, Wei Fang, Xiaodong Xie, Tiejun Huang, Timothée Masquelier
- [Trade-off between Payoff and Model Rewards in Shapley-Fair Collaborative Machine Learning](#) Quoc Phong Nguyen, Bryan Kian Hsiang Low, Patrick Jaillet
- [Efficient Non-Parametric Optimizer Search for Diverse Tasks](#) Ruochen Wang, Yuanhao Xiong, Minhao Cheng, Cho-Jui Hsieh
- [DualCoOp: Fast Adaptation to Multi-Label Recognition with Limited Annotations](#) Ximeng Sun, Ping Hu, Kate Saenko
- [What Can Transformers Learn In-Context? A Case Study of Simple Function Classes](#) Shivam Garg, Dimitris Tsipras, Percy S. Liang, Gregory Valiant
- [Towards Robust Blind Face Restoration with Codebook Lookup Transformer](#) Shangchen Zhou, Kelvin Chan, Chongyi Li, Chen Change Loy
- [Few-Shot Continual Active Learning by a Robot](#) Ali Ayub, Carter Fendley
- [Learning Generalized Policy Automata for Relational Stochastic Shortest Path Problems](#) Rushang Karia, Rashmeet Kaur Nayyar, Siddharth Srivastava
- [A Stochastic Linearized Augmented Lagrangian Method for Decentralized Bilevel Optimization](#) Songtao Lu, Siliang Zeng, Xiaodong Cui, Mark Squillante, Lior Horesh, Brian Kingsbury, Jia Liu, Mingyi Hong
- [Boosting the Performance of Generic Deep Neural Network Frameworks with Log-supermodular CRFs](#) Hao Xiong, Yangxiao Lu, Nicholas Ruozzi
- [Don't Roll the Dice, Ask Twice: The Two-Query Distortion of Matching Problems and Beyond](#) Georgios Amanatidis, Georgios Birmpas, Aris Filos-Ratsikas, Alexandros Voudouris
- [On the Effect of Pre-training for Transformer in Different Modality on Offline Reinforcement Learning](#) Shiro Takagi
- [Look where you look! Saliency-guided Q-networks for generalization in visual Reinforcement Learning](#) David Bertoin, Adil Zouitine, Mehdi Zouitine, Emmanuel Rachelson
- [Information-Theoretic Safe Exploration with Gaussian Processes](#) Alessandro Bottero, Carlos Luis, Julia Vinogradska, Felix Berkenkamp, Jan R. Peters
- [On the consistent estimation of optimal Receiver Operating Characteristic \(ROC\) curve](#) Renxiong Liu, Yunzhang Zhu
- [Online Algorithms for the Santa Claus Problem](#) Max Springer, MohammadTaghi Hajiaghayi, Debmalya Panigrahi, Mohammad Khani
- [Instance-based Learning for Knowledge Base Completion](#) Wanyun Cui, Xingran Chen

- [Decomposable Non-Smooth Convex Optimization with Nearly-Linear Gradient Oracle Complexity](#) Sally Dong, Haotian Jiang, Yin Tat Lee, Swati Padmanabhan, Guanghao Ye
- [On Optimal Learning Under Targeted Data Poisoning](#) Steve Hanneke, Amin Karbasi, Mohammad Mahmoody, Idan Mehalel, Shay Moran
- [Generalized Laplacian Eigenmaps](#) Hao Zhu, Piotr Koniusz
- [OGC: Unsupervised 3D Object Segmentation from Rigid Dynamics of Point Clouds](#) Ziyang Song, Bo Yang
- [Look More but Care Less in Video Recognition](#) Yitian Zhang, Yue Bai, Huan Wang, Yi Xu, Yun Fu
- [Transferring Pre-trained Multimodal Representations with Cross-modal Similarity Matching](#) Byoungjip Kim, Sungik Choi, Dasol Hwang, Moontae Lee, Honglak Lee
- [Get More at Once: Alternating Sparse Training with Gradient Correction](#) Li Yang, Jian Meng, Jae-sun Seo, Deliang Fan
- [BLOX: Macro Neural Architecture Search Benchmark and Algorithms](#) Thomas Chau, Łukasz Dudziak, Hongkai Wen, Nicholas Lane, Mohamed Abdelfattah
- [Does Self-supervised Learning Really Improve Reinforcement Learning from Pixels?](#) Xiang Li, Jinghuan Shang, Srijan Das, Michael Ryoo
- [TCT: Convexifying Federated Learning using Bootstrapped Neural Tangent Kernels](#) Yaodong Yu, Alexander Wei, Sai Praneeth Karimireddy, Yi Ma, Michael Jordan
- [Graph Coloring via Neural Networks for Haplotype Assembly and Viral Quasispecies Reconstruction](#) Hansheng Xue, Vaibhav Rajan, Yu Lin
- [Active Learning of Classifiers with Label and Seed Queries](#) Marco Bressan, Nicolò Cesa-Bianchi, Silvio Lattanzi, Andrea Paudice, Maximilian Thiessen
- [TANGO: Text-driven Photorealistic and Robust 3D Stylization via Lighting Decomposition](#) Yongwei Chen, Rui Chen, Jiabao Lei, Yabin Zhang, Kui Jia
- [Analyzing Lottery Ticket Hypothesis from PAC-Bayesian Theory Perspective](#) Keitaro Sakamoto, Issei Sato
- [Make Sharpness-Aware Minimization Stronger: A Sparsified Perturbation Approach](#) Peng Mi, Li Shen, Tianhe Ren, Yiyi Zhou, Xiaoshuai Sun, Rongrong Ji, Dacheng Tao
- [From Gradient Flow on Population Loss to Learning with Stochastic Gradient Descent](#) Christopher M. De Sa, Satyen Kale, Jason D. Lee, Ayush Sekhari, Karthik Sridharan
- [Distributional Reinforcement Learning for Risk-Sensitive Policies](#) Shiau Hong Lim, ILYAS MALIK
- [Text Classification with Born's Rule](#) Emanuele Guidotti, Alfio Ferrara
- [Learning a Condensed Frame for Memory-Efficient Video Class-Incremental Learning](#) Yixuan Pei, Zhiwu Qing, Jun CEN, Xiang Wang, Shiwei Zhang, Yaxiong Wang, Mingqian Tang, Nong Sang, Xueming Qian
- [Conformal Prediction with Temporal Quantile Adjustments](#) Zhen Lin, Shubhendu Trivedi, Jimeng Sun
- [Learning Viewpoint-Agnostic Visual Representations by Recovering Tokens in 3D Space](#) Jinghuan Shang, Srijan Das, Michael Ryoo
- [Certifying Some Distributional Fairness with Subpopulation Decomposition](#) Mintong Kang, Linyi Li, Maurice Weber, Yang Liu, Ce Zhang, Bo Li
- [A Theoretical Understanding of Gradient Bias in Meta-Reinforcement Learning](#) Bo Liu, Xidong Feng, Jie Ren, Luo Mai, Rui Zhu, Haifeng Zhang, Jun Wang, Yaodong Yang
- [Optimal Algorithms for Decentralized Stochastic Variational Inequalities](#) Dmitry Kovalev, Aleksandr Beznosikov, Abdurakhmon Sadiev, Michael Persiianov, Peter Richtarik, Alexander Gasnikov
- [Mirror Descent Maximizes Generalized Margin and Can Be Implemented Efficiently](#) Haoyuan Sun, Kwangjun Ahn, Christos Thrampoulidis, Navid Azizan
- [Adversarial Task Up-sampling for Meta-learning](#) Yichen WU, Long-Kai Huang, Ying Wei
- [MAtt: A Manifold Attention Network for EEG Decoding](#) Yue-Ting Pan, Jing-Lun Chou, Chun-Shu Wei
- [Geometry-aware Two-scale PIFu Representation for Human Reconstruction](#) Zheng Dong, Ke Xu, Ziheng Duan, Hujun Bao, Weiwei Xu, Rynson Lau
- [Relational Proxies: Emergent Relationships as Fine-Grained Discriminators](#) ABHRA CHAUDHURI, Massimiliano Mancini, Zeynep Akata, Anjan Dutta
- [Root Cause Analysis of Failures in Microservices through Causal Discovery](#) Azam Ikram, Sarthak Chakraborty, Subrata Mitra, Shiv Saini, Saurabh Bagchi, Murat Kocaoglu
- [Anticipating Performativity by Predicting from Predictions](#) Celestine Mendler-Dünnner, Frances Ding, Yixin Wang
- [Falconn+ +: A Locality-sensitive Filtering Approach for Approximate Nearest Neighbor Search](#) Ninh Pham, Tao Liu
- [Pre-Trained Language Models for Interactive Decision-Making](#) Shuang Li, Xavier Puig, Chris Paxton, Yilun Du, Clinton Wang, Linxi Fan, Tao Chen, De-An Huang, Ekin Akyürek, Anima Anandkumar, Jacob Andreas, Igor Mordatch, Antonio Torralba, Yuke Zhu
- [Tight Analysis of Extra-gradient and Optimistic Gradient Methods For Nonconvex Minimax Problems](#) Pouria Mahdavinia, Yuyang Deng, Haochuan Li, Mehrdad Mahdavi
- [Learning Generalizable Models for Vehicle Routing Problems via Knowledge Distillation](#) Jieyi Bi, Yining Ma, Jiahai Wang, Zhiguang Cao, Jinbiao Chen, Yuan Sun, Yeow Meng Chee
- [Large-scale Optimization of Partial AUC in a Range of False Positive Rates](#) Yao Yao, Qihang Lin, Tianbao Yang
- [The Burer-Monteiro SDP method can fail even above the Barvinok-Pataki bound](#) Liam O'Carroll, Vaidehi Srinivas, Aravindan Vijayaraghavan

- [MinVIS: A Minimal Video Instance Segmentation Framework without Video-based Training](#) De-An Huang, Zhiding Yu, Anima Anandkumar
- [Supported Policy Optimization for Offline Reinforcement Learning](#) Jialong Wu, Haixu Wu, Zihan Qiu, Jianmin Wang, Mingsheng Long
- [Infinite Recommendation Networks: A Data-Centric Approach](#) Naveen Sachdeva, Mehak Dhaliwal, Carole-Jean Wu, Julian McAuley
- [DDXPlus: A New Dataset For Automatic Medical Diagnosis](#) Arsene Fansi Tchango, Rishab Goel, Zhi Wen, Julien Martel, Joumana Ghosn
- [Clipped Stochastic Methods for Variational Inequalities with Heavy-Tailed Noise](#) Eduard Gorbunov, Marina Danilova, David Dobre, Pavel Dvurechenskii, Alexander Gasnikov, Gauthier Gidel
- [A Unified Sequence Interface for Vision Tasks](#) Ting Chen, Saurabh Saxena, Lala Li, Tsung-Yi Lin, David J. Fleet, Geoffrey E. Hinton
- [Expected Frequency Matrices of Elections: Computation, Geometry, and Preference Learning](#) Niclas Boehmer, Robert Bredereck, Edith Elkind, Piotr Faliszewski, Stanisław Szufa
- [GMMSeg: Gaussian Mixture based Generative Semantic Segmentation Models](#) Chen Liang, Wenguan Wang, Jiaxu Miao, Yi Yang
- [Understanding and Extending Subgraph GNNs by Rethinking Their Symmetries](#) Fabrizio Frasca, Beatrice Bevilacqua, Michael Bronstein, Haggai Maron
- [Parametrically Retargetable Decision-Makers Tend To Seek Power](#) Alex Turner, Prasad Tadepalli
- [CageNeRF: Cage-based Neural Radiance Field for Generalized 3D Deformation and Animation](#) Yicong Peng, Yichao Yan, Shengqi Liu, Yuhao Cheng, Shanyan Guan, Bowen Pan, Guangtao Zhai, Xiaokang Yang
- [Meta-Query-Net: Resolving Purity-Informativeness Dilemma in Open-set Active Learning](#) Dongmin Park, Yooju Shin, Jihwan Bang, Youngjun Lee, Hwanjun Song, Jae-Gil Lee
- [Dance of SNN and ANN: Solving binding problem by combining spike timing and reconstructive attention](#) Hao Zheng, Hui Lin, Rong Zhao, Luping Shi
- [Unsupervised Learning for Combinatorial Optimization with Principled Objective Relaxation](#) Haoyu Peter Wang, Nan Wu, Hang Yang, Cong Hao, Pan Li
- [PAC-Bayes Compression Bounds So Tight That They Can Explain Generalization](#) Sanae Lotfi, Marc Finzi, Sanyam Kapoor, Andres Potapczynski, Micah Goldblum, Andrew G. Wilson
- [BagFlip: A Certified Defense Against Data Poisoning](#) Yuhao Zhang, Aws Albarghouthi, Loris D'Antoni
- [Data augmentation for efficient learning from parametric experts](#) Alexandre Galashov, Josh S. Merel, Nicolas Heess
- [Symmetry-induced Disentanglement on Graphs](#) Giangiacomo Mercatali, Andre Freitas, Vikas Garg
- [Conformal Off-Policy Prediction in Contextual Bandits](#) Muhammad Faaiz Taufiq, Jean-Francois Ton, Rob Cornish, Yee Whye Teh, Arnaud Doucet
- [Black-Box Generalization: Stability of Zeroth-Order Learning](#) Konstantinos Nikolakakis, Farzin Haddadpour, Dionysis Kalogerias, Amin Karbasi
- [First Contact: Unsupervised Human-Machine Co-Adaptation via Mutual Information Maximization](#) Siddharth Reddy, Sergey Levine, Anca Dragan
- [HyperMiner: Topic Taxonomy Mining with Hyperbolic Embedding](#) Yi.shi Xu, Dongsheng Wang, Bo Chen, Ruiying Lu, Zhibin Duan, Mingyuan Zhou
- [Visual Concepts Tokenization](#) Tao Yang, Yuwang Wang, Yan Lu, Nanning Zheng
- [ZooD: Exploiting Model Zoo for Out-of-Distribution Generalization](#) Qishi Dong, Awais Muhammad, Fengwei Zhou, Chuanlong Xie, Tianyang Hu, Yongxin Yang, Sung-Ho Bae, Zhenguo Li
- [Nearly-Tight Bounds for Testing Histogram Distributions](#) Clément L. Canonne, Ilias Diakonikolas, Daniel Kane, Sihan Liu
- [TGEA 2.0: A Large-Scale Diagnostically Annotated Dataset with Benchmark Tasks for Text Generation of Pretrained Language Models](#) Huibin Ge, Xiaohu Zhao, Chuang Liu, Yulong Zeng, Qun Liu, Deyi Xiong
- [Quantifying Statistical Significance of Neural Network-based Image Segmentation by Selective Inference](#) Vo Nguyen Le Duy, Shogo Iwazaki, Ichiro Takeuchi
- [A Differentially Private Linear-Time fPTAS for the Minimum Enclosing Ball Problem](#) Bar Mahpud, Or Sheffet
- [BEER: Fast \$O\(1/T\)\$ Rate for Decentralized Nonconvex Optimization with Communication Compression](#) Haoyu Zhao, Boyue Li, Zhize Li, Peter Richtarik, Yuejie Chi
- [Autoregressive Search Engines: Generating Substrings as Document Identifiers](#) Michele Bevilacqua, Giuseppe Ottaviano, Patrick Lewis, Scott Yih, Sebastian Riedel, Fabio Petroni
- [Sampling with Riemannian Hamiltonian Monte Carlo in a Constrained Space](#) Yunbum Kook, Yin-Tat Lee, Ruoyi Shen, Santosh Vempala
- [Are All Losses Created Equal: A Neural Collapse Perspective](#) Jinxin Zhou, Chong You, Xiao Li, Kangning Liu, Sheng Liu, Qing Qu, Zhihui Zhu
- [Fair Ranking with Noisy Protected Attributes](#) Anay Mehrotra, Nisheeth Vishnoi
- [Social-Inverse: Inverse Decision-making of Social Contagion Management with Task Migrations](#) Guangmo Tong
- [Provable Benefit of Multitask Representation Learning in Reinforcement Learning](#) Yuan Cheng, Songtao Feng, Jing Yang, Hong Zhang, Yingbin Liang
- [Modeling Human Exploration Through Resource-Rational Reinforcement Learning](#) Marcel Binz, Eric Schulz
- [Generating Long Videos of Dynamic Scenes](#) Tim Brooks, Janne Hellsten, Miika Aittala, Ting-Chun Wang, Timo Aila, Jaakko Lehtinen, Ming-Yu Liu, Alexei Efros, Tero Karras

- [ORIENT: Submodular Mutual Information Measures for Data Subset Selection under Distribution Shift](#) Athresh Karanam, Krishnateja Killamsetty, Harsha Kokel, Rishabh Iyer
- [The BigScience ROOTS Corpus: A 1.6TB Composite Multilingual Dataset](#) Hugo Laurençon, Lucile Saulnier, Thomas Wang, Christopher Akiki, Albert Villanova del Moral, Teven Le Scao, Leandro Von Werra, Chenghao Mou, Eduardo González Ponferrada, Huu Nguyen, Jörg Froberg, Mario Šaško, Quentin Lhoest, Angelina McMillan-Major, Gerard Dupont, Stella Biderman, Anna Rogers, Loubna Ben allal, Francesco De Toni, Giada Pistilli, Olivier Nguyen, Somaieh Nikpoor, Maraim Masoud, Pierre Colombo, Javier de la Rosa, Paulo Villegas, Tristan Thrush, Shayne Longpre, Sebastian Nagel, Leon Weber, Manuel Muñoz, Jian Zhu, Daniel Van Strien, Zaid Alyafeai, Khalid Almubarak, Minh Chien Vu, Itziar Gonzalez-Dios, Aitor Soroa, Kyle Lo, Manan Dey, Pedro Ortiz Suarez, Aaron Gokaslan, Shamik Bose, David Adelani, Long Phan, Hieu Tran, Ian Yu, Suhas Pai, Jenny Chim, Violette Lepercq, Suzana Ilic, Margaret Mitchell, Sasha Alexandra Luccioni, Yacine Jernite
- [Module-Aware Optimization for Auxiliary Learning](#) Hong Chen, Xin Wang, Yue Liu, Yuwei Zhou, Chaoyu Guan, Wenwu Zhu
- [GET3D: A Generative Model of High Quality 3D Textured Shapes Learned from Images](#) Jun Gao, Tianchang Shen, Zian Wang, Wenzheng Chen, Kangxue Yin, Daiqing Li, Or Litany, Zan Gojcic, Sanja Fidler
- [BYOL-Explore: Exploration by Bootstrapped Prediction](#) Zhaohan Guo, Shantanu Thakoor, Miruna Pislari, Bernardo Avila Pires, Florent Althé, Corentin Tallec, Alaa Saade, Daniele Calandriello, Jean-Bastien Grill, Yunhao Tang, Michal Valko, Remi Munos, Mohammad Gheshlaghi Azar, Bilal Piot
- [SizeShiftReg: a Regularization Method for Improving Size-Generalization in Graph Neural Networks](#) Davide Buffelli, Pietro Lió, Fabio Vandin
- [On the Effectiveness of Lipschitz-Driven Rehearsal in Continual Learning](#) Lorenzo Bonicelli, Matteo Boschini, Angelo Porrello, Concetto Spampinato, SIMONE CALDERARA
- [Interpreting Operation Selection in Differentiable Architecture Search: A Perspective from Influence-Directed Explanations](#) Miao Zhang, Wei Huang, Bin Yang
- [Collaborative Learning of Discrete Distributions under Heterogeneity and Communication Constraints](#) Xinmeng Huang, Donghwan Lee, Edgar Dobriban, Hamed Hassani
- [On the Stability and Scalability of Node Perturbation Learning](#) Naoki Hiratani, Yash Mehta, Timothy Lillicrap, Peter E Latham
- [Unsupervised Learning of Group Invariant and Equivariant Representations](#) Robin Winter, Marco Bertolini, Tuan Le, Frank Noe, Djork-Arné Clevert
- [Factored Adaptation for Non-Stationary Reinforcement Learning](#) Fan Feng, Biwei Huang, Kun Zhang, Sara Magliacane
- [MAgNet: Mesh Agnostic Neural PDE Solver](#) Oussama Boussif, Yoshua Bengio, Loubna Benabbou, Dan Assouline
- [TarGF: Learning Target Gradient Field to Rearrange Objects without Explicit Goal Specification](#) Mingdong Wu, Fangwei Zhong, Yulong Xia, Hao Dong
- [Learning to Compare Nodes in Branch and Bound with Graph Neural Networks](#) Abdel Ghani Labassi, Didier Chetelat, Andrea Lodi
- [Teach Less, Learn More: On the Undistillable Classes in Knowledge Distillation](#) Yichen Zhu, Ning Liu, Zhiyuan Xu, Xin Liu, Weibin Meng, Louis Wang, Zhicai Ou, Jian Tang
- [GAGA: Deciphering Age-path of Generalized Self-paced Regularizer](#) Xingyu Qu, Diyang Li, Xiaohan Zhao, Bin Gu
- [ATD: Augmenting CP Tensor Decomposition by Self Supervision](#) Chaoqi Yang, Cheng Qian, Navjot Singh, Cao (Danica) Xiao, M Westover, Edgar Solomonik, Jimeng Sun
- [Towards Learning Universal Hyperparameter Optimizers with Transformers](#) Yutian Chen, Xingyou Song, Chansoo Lee, Zi Wang, Richard Zhang, David Dohan, Kazuya Kawakami, Greg Kochanski, Arnaud Doucet, Marc'Aurelio Ranzato, Sagi Perel, Nando de Freitas
- [Learning to Reconstruct Missing Data from Spatiotemporal Graphs with Sparse Observations](#) Ivan Marisca, Andrea Cini, Cesare Alippi
- [Robust Graph Structure Learning via Multiple Statistical Tests](#) Yaohua Wang, Fangyi Zhang, Ming Lin, Senzhang Wang, Xiuyu Sun, Rong Jin
- [Peripheral Vision Transformer](#) Juhong Min, Yucheng Zhao, Chong Luo, Minsu Cho
- [Maximum Common Subgraph Guided Graph Retrieval: Late and Early Interaction Networks](#) Indradyumna Roy, Soumen Chakrabarti, Abir De
- [Differentially Private Learning with Margin Guarantees](#) Raef Bassily, Mehryar Mohri, Ananda Theertha Suresh
- [ADBench: Anomaly Detection Benchmark](#) Songqiao Han, Xiyang Hu, Hailiang Huang, Minqi Jiang, Yue Zhao
- [GLIF: A Unified Gated Leaky Integrate-and-Fire Neuron for Spiking Neural Networks](#) Xingting Yao, Fanrong Li, Zitao Mo, Jian Cheng
- [HSDF: Hybrid Sign and Distance Field for Modeling Surfaces with Arbitrary Topologies](#) Li Wang, jie Yang, Weikai Chen, Xiaoxu Meng, Bo Yang, Jintao Li, Lin Gao
- [Inverse Game Theory for Stackelberg Games: the Blessing of Bounded Rationality](#) Jibang Wu, Weiran Shen, Fei Fang, Haifeng Xu
- [When Privacy Meets Partial Information: A Refined Analysis of Differentially Private Bandits](#) Achraf Azize, Debabrota Basu
- [Robust Reinforcement Learning using Offline Data](#) Kishan Panaganti, Zaiyan Xu, Dileep Kalathil, Mohammad Ghavamzadeh
- [Compositional generalization through abstract representations in human and artificial neural networks](#) Takuya Ito, Tim Klinger, Doug Schultz, John Murray, Michael Cole, Mattia Rigotti
- [Self-Consistent Dynamical Field Theory of Kernel Evolution in Wide Neural Networks](#) Blake Bordelon, Cengiz Pehlevan
- [Provably expressive temporal graph networks](#) Amauri Souza, Diego Mesquita, Samuel Kaski, Vikas Garg
- [Maximum Likelihood Training of Implicit Nonlinear Diffusion Model](#) Dongjun Kim, Byeonghu Na, Se Jung Kwon, Dongsoo Lee, Wanmo Kana. Il-chul Moon

- [Deep Combinatorial Aggregation](#) Yuesong Shen, Daniel Cremers
- [Efficient coding, channel capacity, and the emergence of retinal mosaics](#) Na Young Jun, Greg Field, John Pearson
- [TA-GATES: An Encoding Scheme for Neural Network Architectures](#) Xuefei Ning, Zixuan Zhou, Junbo Zhao, Tianchen Zhao, Yiping Deng, Changcheng Tang, Shuang Liang, Huazhong Yang, Yu Wang
- [ZSON: Zero-Shot Object-Goal Navigation using Multimodal Goal Embeddings](#) Arjun Majumdar, Gunjan Aggarwal, Bhavika Devnani, Judy Hoffman, Dhruv Batra
- [Autoformalization with Large Language Models](#) Yuhuai Wu, Albert Qiaochu Jiang, Wenda Li, Markus Rabe, Charles Staats, Mateja Jamnik, Christian Szegedy
- [Model-based Lifelong Reinforcement Learning with Bayesian Exploration](#) Haotian Fu, Shangqun Yu, Michael Littman, George Konidaris
- [On Efficient Online Imitation Learning via Classification](#) Yichen Li, Chicheng Zhang
- [Let Images Give You More: Point Cloud Cross-Modal Training for Shape Analysis](#) Xu Yan, Heshen Zhan, Chaoda Zheng, Jiantao Gao, Ruimao Zhang, Shuguang Cui, Zhen Li
- [On Robust Multiclass Learnability](#) Jingyuan Xu, Weiwei Liu
- [Debiased Self-Training for Semi-Supervised Learning](#) Baixu Chen, Junguang Jiang, Ximei Wang, Pengfei Wan, Jianmin Wang, Mingsheng Long
- [Rethinking Individual Global Max in Cooperative Multi-Agent Reinforcement Learning](#) Yitian Hong, Yaochu Jin, Yang Tang
- [Archimedes Meets Privacy: On Privately Estimating Quantiles in High Dimensions Under Minimal Assumptions](#) Omri Ben-Eliezer, Dan Mikulincer, Ilias Zadik
- [Uncovering the Structural Fairness in Graph Contrastive Learning](#) Ruijia Wang, Xiao Wang, Chuan Shi, Le Song
- [Estimation of Entropy in Constant Space with Improved Sample Complexity](#) Maryam Aliakbarpour, Andrew McGregor, Jelani Nelson, Erik Waingarten
- [Average Sensitivity of Euclidean k-Clustering](#) Yuichi Yoshida, Shinji Ito
- [Multi-block-Single-probe Variance Reduced Estimator for Coupled Compositional Optimization](#) Wei Jiang, Gang Li, Yibo Wang, Lijun Zhang, Tianbao Yang
- [Differentially Private Online-to-batch for Smooth Losses](#) Qinzi Zhang, Hoang Tran, Ashok Cutkosky
- [Weak-shot Semantic Segmentation via Dual Similarity Transfer](#) Junjie Chen, Li Niu, Siyuan Zhou, Jianlou Si, Chen Qian, Liqing Zhang
- [A Multi-Task Benchmark for Korean Legal Language Understanding and Judgement Prediction](#) Wonseok Hwang, Dongjun Lee, Kyoungyeon Cho, Hanuhl Lee, Minjoon Seo
- [Synergy-of-Experts: Collaborate to Improve Adversarial Robustness](#) Sen Cui, Jingfeng ZHANG, Jian Liang, Bo Han, Masashi Sugiyama, Changshui Zhang
- [Improved Fine-Tuning by Better Leveraging Pre-Training Data](#) Ziquan Liu, Yi Xu, Yuanhong Xu, Qi Qian, Hao Li, Xiangyang Ji, Antoni Chan, Rong Jin
- [Transferring Fairness under Distribution Shifts via Fair Consistency Regularization](#) Bang An, Zora Che, Mucong Ding, Furong Huang
- [OpenOOD: Benchmarking Generalized Out-of-Distribution Detection](#) Jingkang Yang, Pengyun Wang, Dejian Zou, Zitang Zhou, Kunyuan Ding, WENXUAN PENG, Haoqi Wang, Guangyao Chen, Bo Li, Yiyu Sun, Xuefeng Du, Kaiyang Zhou, Wayne Zhang, Dan Hendrycks, Yixuan Li, Ziwei Liu
- [Two-layer neural network on infinite dimensional data: global optimization guarantee in the mean-field regime](#) Naoki Nishikawa, Taiji Suzuki, Atsushi Nitanda, Denny Wu
- [KSD Aggregated Goodness-of-fit Test](#) Antonin Schrab, Benjamin Guedj, Arthur Gretton
- [Efficient Risk-Averse Reinforcement Learning](#) Ido Greenberg, Yinlam Chow, Mohammad Ghavamzadeh, Shie Mannor
- [D²NeRF: Self-Supervised Decoupling of Dynamic and Static Objects from a Monocular Video](#) Tianhao Wu, Fangcheng Zhong, Andrea Tagliasacchi, Forrester Cole, Cengiz Oztireli
- [On the Effective Number of Linear Regions in Shallow Univariate ReLU Networks: Convergence Guarantees and Implicit Bias](#) Itay Safran, Gal Vardi, Jason D. Lee
- [Robust \$\phi\$ -Divergence MDPs](#) Chin Pang Ho, Marek Petrik, Wolfram Wiesemann
- [Object Representations as Fixed Points: Training Iterative Refinement Algorithms with Implicit Differentiation](#) Michael Chang, Tom Griffiths, Sergey Levine
- [Benefits of Additive Noise in Composing Classes with Bounded Capacity](#) Alireza Fathollah Pour, Hassan Ashtiani
- [Multi-Objective Deep Learning with Adaptive Reference Vectors](#) Weiyu Chen, James Kwok
- [Towards Hard-pose Virtual Try-on via 3D-aware Global Correspondence Learning](#) Zaiyu Huang, Hanhui Li, Zhenyu Xie, Michael Kampffmeyer, qingling Cai, Xiaodan Liang
- [Pushing the limits of fairness impossibility: Who's the fairest of them all?](#) Brian Hsu, Rahul Mazumder, Preetam Nandy, Kinjal Basu
- [Recall Distortion in Neural Network Pruning and the Undecayed Pruning Algorithm](#) Aidan Good, Jiaqi Lin, Xin Yu, Hannah Sieg, Mikey Fergusson, Shandian Zhe, Jerzy Wiecezorek, Thiago Serra
- [Generating multivariate time series with Common Source Coordinated GAN \(COSCI-GAN\)](#) Ali Seyfi, Jean-Francois Rajotte, Raymond Ng
- [ConfounderGAN: Protecting Image Data Privacy with Causal Confounder](#) Qi Tian, Kun Kuang, Kelu Jiang, Furui Liu, Zhihua Wang, Fei Wu
- [Global Convergence of Direct Policy Search for State-Feedback \$\mathcal{H}_\infty\$ Robust Control: A Revisit of Nonsmooth Synthesis with Goldstein Subdifferential](#) Xingang Guo, Bin Hu

- [In What Ways Are Deep Neural Networks Invariant and How Should We Measure This?](#) Henry Kvinge, Tegan Emerson, Grayson Jorgenson, Scott Vasquez, Tim Doster, Jesse Lew
- [Optimal Weak to Strong Learning](#) Kasper Green Larsen, Martin Ritzert
- [A Neural Pre-Conditioning Active Learning Algorithm to Reduce Label Complexity](#) Seo Taek Kong, Soomin Jeon, Dongbin Na, Jaewon Lee, Hong-Seok Lee, Kyu-Hwan Jung
- [AnoShift: A Distribution Shift Benchmark for Unsupervised Anomaly Detection](#) Marius Dragoi, Elena Burceanu, Emanuela Haller, Andrei Manolache, Florin Brad
- [Grounding Aleatoric Uncertainty for Unsupervised Environment Design](#) Minqi Jiang, Michael Dennis, Jack Parker-Holder, Andrei Lupu, Heinrich Küttler, Edward Grefenstette, Tim Rocktäschel, Jakob Foerster
- [PROSPECT: Labeled Tandem Mass Spectrometry Dataset for Machine Learning in Proteomics](#) Omar Shouman, Wassim Gabriel, Victor-George Giurcoiu, Vitor Sternlicht, Mathias Wilhelm
- [VLMo: Unified Vision-Language Pre-Training with Mixture-of-Modality-Experts](#) Hangbo Bao, Wenhui Wang, Li Dong, Qiang Liu, Owais Khan Mohammed, Kriti Aggarwal, Subhojit Som, Songhao Piao, Furu Wei
- [Hub-Pathway: Transfer Learning from A Hub of Pre-trained Models](#) Yang Shu, Zhangjie Cao, Ziyang Zhang, Jianmin Wang, Mingsheng Long
- [Towards Better Evaluation for Dynamic Link Prediction](#) Farimah Poursafaei, Shenyang Huang, Kellin Pelrine, Reihaneh Rabbany
- [Coarse-to-Fine Vision-Language Pre-training with Fusion in the Backbone](#) Zi-Yi Dou, Aishwarya Kamath, Zhe Gan, Pengchuan Zhang, Jianfeng Wang, Linjie Li, Zicheng Liu, Ce Liu, Yann LeCun, Nanyun Peng, Jianfeng Gao, Lijuan Wang
- [On the Discrimination Risk of Mean Aggregation Feature Imputation in Graphs](#) Arjun Subramonian, Kai-Wei Chang, Yizhou Sun
- [VRL3: A Data-Driven Framework for Visual Deep Reinforcement Learning](#) Che Wang, Xufang Luo, Keith Ross, Dongsheng Li
- [Efficient and Effective Optimal Transport-Based Biclustering](#) Chakib Fettaf, lazhar labiod, Mohamed NADIF
- [Unifying and Boosting Gradient-Based Training-Free Neural Architecture Search](#) YAO SHU, Zhongxiang Dai, Zhaoxuan Wu, Bryan Kian Hsiang Low
- [Hyperbolic Embedding Inference for Structured Multi-Label Prediction](#) Bo Xiong, Michael Cochez, Mojtaba Nayyeri, Steffen Staab
- [Approximate Value Equivalence](#) Christopher Grimm, Andre Barreto, Satinder Singh
- [The Power and Limitation of Pretraining-Finetuning for Linear Regression under Covariate Shift](#) Jingfeng Wu, Difan Zou, Vladimir Braverman, Quanquan Gu, Sham Kakade
- [Rank Diminishing in Deep Neural Networks](#) Ruili Feng, Kecheng Zheng, Yukun Huang, Deli Zhao, Michael Jordan, Zheng-Jun Zha
- [Single-pass Streaming Lower Bounds for Multi-armed Bandits Exploration with Instance-sensitive Sample Complexity](#) Sepehr Assadi, Chen Wang
- [Optimistic Tree Searches for Combinatorial Black-Box Optimization](#) Cedric Malherbe, Antoine Grosnit, Rasul Tutunov, Haitham Bou Ammar, Jun Wang
- [Robust Feature-Level Adversaries are Interpretability Tools](#) Stephen Casper, Max Nadeau, Dylan Hadfield-Menell, Gabriel Kreiman
- [A Unified Convergence Theorem for Stochastic Optimization Methods](#) Xiao Li, Andre Milzarek
- [New Definitions and Evaluations for Saliency Methods: Staying Intrinsic, Complete and Sound](#) Arushi Gupta, Nikunj Saunshi, Dingli Yu, Kaifeng Lyu, Sanjeev Arora
- [Exponential Separations in Symmetric Neural Networks](#) Aaron Zweig, Joan Bruna
- [Randomized Message-Interception Smoothing: Gray-box Certificates for Graph Neural Networks](#) Yan Scholten, Jan Schuchardt, Simon Geisler, Aleksandar Bojchevski, Stephan Günnemann
- [Oscillatory Tracking of Continuous Attractor Neural Networks Account for Phase Precession and Procession of Hippocampal Place Cells](#) Tianhao Chu, Zilong Ji, Junfeng Zuo, Wenhao Zhang, Tiejun Huang, Yuanyuan Mi, Si Wu
- [Finding and Listing Front-door Adjustment Sets](#) Hyunchai Jeong, Jin Tian, Elias Bareinboim
- [MorphTE: Injecting Morphology in Tensorized Embeddings](#) Guobing Gan, Peng Zhang, Sunzhu Li, Xiuqing Lu, Benyou Wang
- [Active Learning with Safety Constraints](#) Romain Camilleri, Andrew Wagenmaker, Jamie H. Morgenstern, Lalit Jain, Kevin G. Jamieson
- [Is one annotation enough? - A data-centric image classification benchmark for noisy and ambiguous label estimation](#) Lars Schmarje, Vasco Grossmann, Claudius Zelenka, Sabine Dippel, Rainer Kiko, Mariusz Oszust, Matti Pastell, Jenny Stracke, Anna Valros, Nina Volkmann, Reinhard Koch
- [Sobolev Acceleration and Statistical Optimality for Learning Elliptic Equations via Gradient Descent](#) Yiping Lu, Jose Blanchet, Lexing Ying
- [Block-Recurrent Transformers](#) DeLesley Hutchins, Imanol Schlag, Yuhuai Wu, Ethan Dyer, Behnam Neyshabur
- [Learning Two-Player Markov Games: Neural Function Approximation and Correlated Equilibrium](#) Chris Junchi Li, Dongruo Zhou, Quanquan Gu, Michael Jordan
- [Dense Interspecies Face Embedding](#) Sejong Yang, Subin Jeon, Seonghyeon Nam, Seon Joo Kim
- [Efficient Scheduling of Data Augmentation for Deep Reinforcement Learning](#) Byungchan Ko, Jungseul Ok
- [Mix and Reason: Reasoning over Semantic Topology with Data Mixing for Domain Generalization](#) Chaoqi Chen, Luyao Tang, Feng Liu, Gangming Zhao, Yue Huang, Yizhou Yu
- [Personalized Online Federated Learning with Multiple Kernels](#) Pouya M. Ghari, Yanning Shen
- [Point Transformer V2: Grouped Vector Attention and Partition-based Pooling](#) Xiaoyang Wu, Yixing Lao, Li Jiang, Xihui Liu, Hengshuang Zhao

- [Neural Approximation of Graph Topological Features](#) Zuoyu Yan, Tengfei Ma, Liangcai Gao, Zhi Tang, Yusu Wang, Chao Chen
- [MOVE: Unsupervised Movable Object Segmentation and Detection](#) Adam Bielski, Paolo Favaro
- [FreGAN: Exploiting Frequency Components for Training GANs under Limited Data](#) mengping yang, Zhe Wang, Ziqiu Chi, Yanbing Zhang
- [LBD: Decouple Relevance and Observation for Individual-Level Unbiased Learning to Rank](#) Mouxian Chen, Chenghao Liu, Zemin Liu, Jianling Sun
- [MetricFormer: A Unified Perspective of Correlation Exploring in Similarity Learning](#) Jiexi Yan, Erkun Yang, Cheng Deng, Heng Huang
- [Toward a realistic model of speech processing in the brain with self-supervised learning](#) Juliette MILLET, Charlotte Caucheteux, pierre orhan, Yves Boubenec, Alexandre Gramfort, Ewan Dunbar, Christophe Pallier, Jean-Remi King
- [Distributed Inverse Constrained Reinforcement Learning for Multi-agent Systems](#) Shicheng Liu, Minghui Zhu
- [Collaborative Decision Making Using Action Suggestions](#) Dylan Asmar, Mykel J Kochenderfer
- [Near-Optimal Regret for Adversarial MDP with Delayed Bandit Feedback](#) Tiancheng Jin, Tal Lancewicki, Haipeng Luo, Yishay Mansour, Aviv Rosenberg
- [Provably sample-efficient RL with side information about latent dynamics](#) Yao Liu, Dipendra Misra, Miro Dudik, Robert E. Schapire
- [Optimal Gradient Sliding and its Application to Optimal Distributed Optimization Under Similarity](#) Dmitry Kovalev, Aleksandr Beznosikov, Ekaterina Borodich, Alexander Gasnikov, Gesualdo Scutari
- [Universally Expressive Communication in Multi-Agent Reinforcement Learning](#) Matthew Morris, Thomas D Barrett, Arnu Pretorius
- [Fast Stochastic Composite Minimization and an Accelerated Frank-Wolfe Algorithm under Parallelization](#) Benjamin Dubois-Taine, Francis Bach, Quentin Berthet, Adrien Taylor
- [Multi-Granularity Cross-modal Alignment for Generalized Medical Visual Representation Learning](#) Fuying Wang, Yuyin Zhou, Shujun WANG, Varut Vardhanabhuti, Lequan Yu
- [Learning Robust Rule Representations for Abstract Reasoning via Internal Inferences](#) Wenbo Zhang, likai tang, Site Mo, Xianggen Liu, Sen Song
- [Turning the Tables: Biased, Imbalanced, Dynamic Tabular Datasets for ML Evaluation](#) Sérgio Jesus, José Pombal, Duarte Alves, André Cruz, Pedro Saleiro, Rita Ribeiro, João Gama, Pedro Bizarro
- [DropCov: A Simple yet Effective Method for Improving Deep Architectures](#) Qilong Wang, Mingze Gao, Zhaolin Zhang, Jiangtao Xie, Peihua Li, Qinghua Hu
- [A Unifying Framework for Online Optimization with Long-Term Constraints](#) Matteo Castiglioni, Andrea Celli, Alberto Marchesi, Giulia Romano, Nicola Gatti
- [The least-control principle for local learning at equilibrium](#) Alexander Meulemans, Nicolas Zucchet, Seijin Kobayashi, Johannes von Oswald, João Sacramento
- [Optimizing Relevance Maps of Vision Transformers Improves Robustness](#) Hila Chefer, Idan Schwartz, Lior Wolf
- [Stability and Generalization of Kernel Clustering: from Single Kernel to Multiple Kernel](#) Weixuan Liang, Xinwang Liu, Yong Liu, sihang zhou, Jun-Jie Huang, Siwei Wang, Jiyuan Liu, Yi Zhang, En Zhu
- [Deep Ensembles Work, But Are They Necessary?](#) Taiga Abe, Estefany Kelly Buchanan, Geoff Pleiss, Richard Zemel, John P. Cunningham
- [VICE: Variational Interpretable Concept Embeddings](#) Lukas Muttenthaler, Charles Y. Zheng, Patrick McClure, Robert A. Vandermeulen, Martin N Hebart, Francisco Pereira
- [Random Normalization Aggregation for Adversarial Defense](#) Minjing Dong, Xinghao Chen, Yunhe Wang, Chang Xu
- [Distributionally Robust Optimization with Data Geometry](#) Jiashuo Liu, Jiayun Wu, Bo Li, Peng Cui
- [Near-Optimal Correlation Clustering with Privacy](#) Vincent Cohen-Addad, Chenglin Fan, Silvio Lattanzi, Slobodan Mitrovic, Ashkan Norouzi-Fard, Nikos Parotsidis, Jakub M. Tarnawski
- [Knowledge Distillation from A Stronger Teacher](#) Tao Huang, Shan You, Fei Wang, Chen Qian, Chang Xu
- [Optimal Transport of Classifiers to Fairness](#) Maarten Buyt, Tijl De Bie
- [Rethinking the compositionality of point clouds through regularization in the hyperbolic space](#) Antonio Montanaro, Diego Valsesia, Enrico Magli
- [ReCo: Retrieve and Co-segment for Zero-shot Transfer](#) Gyungin Shin, Weidi Xie, Samuel Albanie
- [Monocular Dynamic View Synthesis: A Reality Check](#) Hang Gao, Ruilong Li, Shubham Tulsiani, Bryan Russell, Angjoo Kanazawa
- [Bridging the Gap between Object and Image-level Representations for Open-Vocabulary Detection](#) Hanoona Bangalath, Muhammad Maaz, Muhammad Uzair Khattak, Salman H. Khan, Fahad Shahbaz Khan
- [A Simple and Optimal Policy Design for Online Learning with Safety against Heavy-tailed Risk](#) David Simchi-Levi, Zeyu Zheng, Feng Zhu
- [A Boosting Approach to Reinforcement Learning](#) Nataly Brukhim, Elad Hazan, Karan Singh
- [Relaxing Equivariance Constraints with Non-stationary Continuous Filters](#) Tycho van der Ouderaa, David W. Romero, Mark van der Wilk
- [The Franz-Parisi Criterion and Computational Trade-offs in High Dimensional Statistics](#) Afonso S Bandeira, Ahmed El Alaoui, Samuel Hopkins, Tselil Schramm, Alexander S Wein, Ilias Zadik
- [Robust Neural Posterior Estimation and Statistical Model Criticism](#) Daniel Ward, Patrick Cannon, Mark Beaumont, Matteo Fasiolo, Sebastian Schmon
- [Why do We Need Large Batchesizes in Contrastive Learning? A Gradient-Bias Perspective](#) Changyou Chen, Jianyi Zhang, Yi Xu, Liqun Chen, Jiali Duan, Yiran Chen, Son Tran, Belinda Zeng, Trishul Chilimbi
- [Randomized Channel Shuffling: Minimal-Overhead Backdoor Attack Detection without Clean Datasets](#) Ruisi Cai, Zhenyu Zhang, Tianlong

- [Structural Analysis of Branch-and-Cut and the Learnability of Gomory Mixed Integer Cuts](#) Maria-Florina F. Balcan, Siddharth Prasad, Tuomas Sandholm, Ellen Vitercik
- [Finite-Time Last-Iterate Convergence for Learning in Multi-Player Games](#) Yang Cai, Argyris Oikonomou, Weiqiang Zheng
- [When are Local Queries Useful for Robust Learning?](#) Pascale Gourdeau, Varun Kanade, Marta Kwiatkowska, James Worrell
- [SPoVT: Semantic-Prototype Variational Transformer for Dense Point Cloud Semantic Completion](#) Sheng Yu Huang, Hao-Yu Hsu, Frank Wang
- [Improving Intrinsic Exploration with Language Abstractions](#) Jesse Mu, Victor Zhong, Roberta Raileanu, Mingqi Jiang, Noah Goodman, Tim Rocktäschel, Edward Grefenstette
- [The trade-offs of model size in large recommendation models : 100GB to 10MB Criteo-tb DLRM model](#) Aditya Desai, Anshumali Shrivastava
- [Near-Optimal Goal-Oriented Reinforcement Learning in Non-Stationary Environments](#) Liyu Chen, Haipeng Luo
- [A Unified Framework for Deep Symbolic Regression](#) Mikel Landajuela, Chak Shing Lee, Jiachen Yang, Ruben Glatt, Claudio P Santiago, Ignacio Aravena, Terrell Mundhenk, Garrett Mulcahy, Brenden K Petersen
- [VoxGRAF: Fast 3D-Aware Image Synthesis with Sparse Voxel Grids](#) Katja Schwarz, Axel Sauer, Michael Niemeyer, Yiyi Liao, Andreas Geiger
- [Sparse Hypergraph Community Detection Thresholds in Stochastic Block Model](#) Erchuan Zhang, David Suter, Giang Truong, Syed Zulqarnain Gilani
- [Understanding Non-linearity in Graph Neural Networks from the Bayesian-Inference Perspective](#) Rongzhe Wei, Haoteng YIN, Junteng Jia, Austin R. Benson, Pan Li
- [Explainable Reinforcement Learning via Model Transforms](#) Mira Finkelstein, Nitsan Levy, Lucy Liu, Yoav Kolumbus, David C. Parkes, Jeffrey S Rosenschein, Sarah Keren
- [Real-Valued Backpropagation is Unsuitable for Complex-Valued Neural Networks](#) Zhi-Hao Tan, Yi Xie, Yuan Jiang, Zhi-Hua Zhou
- [Learning dynamics of deep linear networks with multiple pathways](#) Jianghong Shi, Eric Shea-Brown, Michael Buice
- [Self-Supervised Aggregation of Diverse Experts for Test-Agnostic Long-Tailed Recognition](#) Yifan Zhang, Bryan Hooi, Lanqing Hong, Jiashi Feng
- [Self-Supervised Learning via Maximum Entropy Coding](#) Xin Liu, Zhongdao Wang, Ya-Li Li, Shengjin Wang
- [A Practical, Progressively-Expressive GNN](#) Lingxiao Zhao, Neil Shah, Leman Akoglu
- [On Learning Fairness and Accuracy on Multiple Subgroups](#) Changjian Shui, Gezheng Xu, Qi CHEN, Jiaqi Li, Charles X. Ling, Tal Arbel, Boyu Wang, Christian Gagné
- [Adv-Attribute: Inconspicuous and Transferable Adversarial Attack on Face Recognition](#) Shuai Jia, Bangjie Yin, Taiping Yao, Shouhong Ding, Chunhua Shen, Xiaokang Yang, Chao Ma
- [Understanding Aesthetics with Language: A Photo Critique Dataset for Aesthetic Assessment](#) Daniel Vera Nieto, Luigi Celona, Clara Fernandez Labrador
- [Structuring Representations Using Group Invariants](#) Mehran Shakerinava, Arnab Kumar Mondal, Siamak Ravanbakhsh
- [Black-box coreset variational inference](#) Dionysis Manousakas, Hippolyt Ritter, Theofanis Karaletsos
- [Distilling Representations from GAN Generator via Squeeze and Span](#) Yu Yang, Xiaotian Cheng, Chang Liu, Hakan Bilen, Xiangyang Ji
- [A Quantitative Geometric Approach to Neural-Network Smoothness](#) Zi Wang, Gautam Prakriya, Somesh Jha
- [AutoMTL: A Programming Framework for Automating Efficient Multi-Task Learning](#) Lijun Zhang, Xiao Liu, Hui Guan
- [Log-Concave and Multivariate Canonical Noise Distributions for Differential Privacy](#) Jordan Awan, Jinshuo Dong
- [On the generalization of learning algorithms that do not converge](#) Nisha Chandramoorthy, Andreas Loukas, Khashayar Gatmiry, Stefanie Jegelka
- [PALMER: Perception - Action Loop with Memory for Long-Horizon Planning](#) Onur Beker, Mohammad Mohammadi, Amir Zamir
- [Leveraging Factored Action Spaces for Efficient Offline Reinforcement Learning in Healthcare](#) Shengpu Tang, Maggie Makar, Michael Sjoding, Finale Doshi-Velez, Jenna Wiens
- [Visual correspondence-based explanations improve AI robustness and human-AI team accuracy](#) Mohammad Reza Taesiri, Giang Nguyen, Anh Nguyen
- [Invariance-Aware Randomized Smoothing Certificates](#) Jan Schuchardt, Stephan Günnemann
- [On Image Segmentation With Noisy Labels: Characterization and Volume Properties of the Optimal Solutions to Accuracy and Dice](#) Marcus Nordstrom, Henrik Hult, Fredrik Löfman, Jonas Söderberg
- [Blessing of Depth in Linear Regression: Deeper Models Have Flatter Landscape Around the True Solution](#) Jianhao Ma, Salar Fattahi
- [Fairness Reprogramming](#) Guanhua Zhang, Yihua Zhang, Yang Zhang, Wenqi Fan, Qing Li, Sijia Liu, Shiyu Chang
- [WeightedSHAP: analyzing and improving Shapley based feature attributions](#) Yongchan Kwon, James Y. Zou
- [Temporal Effective Batch Normalization in Spiking Neural Networks](#) Chaoteng Duan, Jianhao Ding, Shiyang Chen, Zhaofei Yu, Tiejun Huang
- [Remember the Past: Distilling Datasets into Addressable Memories for Neural Networks](#) Zhiwei Deng, Olga Russakovsky
- [Robustness Analysis of Video-Language Models Against Visual and Language Perturbations](#) Madeline Schiappa, Shruti Vyas, Hamid Palangi, Yogesh Rawat, Vibhav Vineet
- [Hyperbolic Feature Augmentation via Distribution Estimation and Infinite Sampling on Manifolds](#) Zhi Gao, Yuwei Wu, Yunde Jia, Mehrtash Harandi

- [PatchComplete: Learning Multi-Resolution Patch Priors for 3D Shape Completion on Unseen Categories](#) Yuchen Rao, Yinyu Nie, Angela Dai
- [Q-ViT: Accurate and Fully Quantized Low-bit Vision Transformer](#) Yanjing Li, Sheng Xu, Baochang Zhang, Xianbin Cao, Peng Gao, Guodong Guo
- [Enhancing Safe Exploration Using Safety State Augmentation](#) Aivar Sootla, Alexander Cowen-Rivers, Jun Wang, Haitham Bou Ammar
- [Unsupervised Reinforcement Learning with Contrastive Intrinsic Control](#) Michael Laskin, Hao Liu, Xue Bin Peng, Denis Yarats, Aravind Rajeswaran, Pieter Abbeel
- [Distributed Online Convex Optimization with Compressed Communication](#) Zhipeng Tu, Xi Wang, Yiguang Hong, Lei Wang, Deming Yuan, Guodong Shi
- [Local Latent Space Bayesian Optimization over Structured Inputs](#) Natalie Maus, Haydn Jones, Juston Moore, Matt J. Kusner, John Bradshaw, Jacob Gardner
- [Graph Neural Network Bandits](#) Parnian Kassraie, Andreas Krause, Ilija Bogunovic
- [Concrete Score Matching: Generalized Score Matching for Discrete Data](#) Chenlin Meng, Kristy Choi, Jiaming Song, Stefano Ermon
- [Regularized Gradient Descent Ascent for Two-Player Zero-Sum Markov Games](#) Sihan Zeng, Thinh Doan, Justin Romberg
- [Chefs' Random Tables: Non-Trigonometric Random Features](#) Valerii Likhoshesterov, Krzysztof M Choromanski, Kumar Avinava Dubey, Frederick Liu, Tamas Sarlos, Adrian Weller
- [CS-Shapley: Class-wise Shapley Values for Data Valuation in Classification](#) Stephanie Schoch, Haifeng Xu, Yangfeng Ji
- [Factuality Enhanced Language Models for Open-Ended Text Generation](#) Nayeon Lee, Wei Ping, Peng Xu, Mostofa Patwary, Pascale N Fung, Mohammad Shoeybi, Bryan Catanzaro
- [On the Representation Collapse of Sparse Mixture of Experts](#) Zewen Chi, Li Dong, Shaohan Huang, Damai Dai, Shuming Ma, Barun Patra, Saksham Singhal, Payal Bajaj, XIA SONG, Xian-Ling Mao, Heyan Huang, Furu Wei
- [Nearly Optimal Algorithms for Linear Contextual Bandits with Adversarial Corruptions](#) Jiafan He, Dongruo Zhou, Tong Zhang, Quanquan Gu
- [Implicit Bias of Gradient Descent on Reparametrized Models: On Equivalence to Mirror Descent](#) Zhiyuan Li, Tianhao Wang, Jason D. Lee, Sanjeev Arora
- [HYPRO: A Hybridly Normalized Probabilistic Model for Long-Horizon Prediction of Event Sequences](#) Siqiao Xue, Xiaoming Shi, James Zhang, Hongyuan Mei
- [Towards Understanding Grokking: An Effective Theory of Representation Learning](#) Ziming Liu, Ouail Kitouni, Niklas S Nolte, Eric Michaud, Max Tegmark, Mike Williams
- [Asymptotic Behaviors of Projected Stochastic Approximation: A Jump Diffusion Perspective](#) Jiadong Liang, Yuze Han, Xiang Li, Zhihua Zhang
- [Towards Practical Few-shot Query Sets: Transductive Minimum Description Length Inference](#) Ségolène Martin, Malik Boudiaf, Emilie Chouzenoux, Jean-Christophe Pesquet, Ismail Ayed
- [Understanding the Generalization Benefit of Normalization Layers: Sharpness Reduction](#) Kaifeng Lyu, Zhiyuan Li, Sanjeev Arora
- [Learning Manifold Dimensions with Conditional Variational Autoencoders](#) Yijia Zheng, Tong He, Yixuan Qiu, David P Wipf
- [Multi-modal Grouping Network for Weakly-Supervised Audio-Visual Video Parsing](#) Shentong Mo, Yapeng Tian
- [Optimal-er Auctions through Attention](#) Dmitry Ivanov, Iskander Safiulin, Igor Filippov, Ksenia Balabaeva
- [Bootstrapped Transformer for Offline Reinforcement Learning](#) Kerong Wang, Hanye Zhao, Xufang Luo, Kan Ren, Weinan Zhang, Dongsheng Li
- [How to talk so AI will learn: Instructions, descriptions, and autonomy](#) Theodore Sumers, Robert Hawkins, Mark K. Ho, Tom Griffiths, Dylan Hadfield-Menell
- [Learning on the Edge: Online Learning with Stochastic Feedback Graphs](#) Emmanuel Esposito, Federico Fusco, Dirk van der Hoeven, Nicolò Cesa-Bianchi
- [Globally Gated Deep Linear Networks](#) Qianyi Li, Haim Sompolsky
- [Markov Chain Score Ascent: A Unifying Framework of Variational Inference with Markovian Gradients](#) Kyurae Kim, Jisu Oh, Jacob Gardner, Adji Bousso Dieng, Hongseok Kim
- [Cache-Augmented Inbatch Importance Resampling for Training Recommender Retriever](#) Jin Chen, Defu Lian, Yucheng Li, Baoyun Wang, Kai Zheng, Enhong Chen
- [Doubly Robust Counterfactual Classification](#) Kwangho Kim, Edward Kennedy, Jose Zubizarreta
- [Rethinking Value Function Learning for Generalization in Reinforcement Learning](#) Seungyong Moon, JunYeong Lee, Hyun Oh Song
- [Improving Certified Robustness via Statistical Learning with Logical Reasoning](#) Zhuolin Yang, Zhikuan Zhao, Boxin Wang, Jiawei Zhang, Linyi Li, Hengzhi Pei, Bojan Karlaš, Ji Liu, Heng Guo, Ce Zhang, Bo Li
- [Transformer-based Working Memory for Multiagent Reinforcement Learning with Action Parsing](#) Yaodong Yang, Guangyong Chen, Weixun Wang, Xiaotian Hao, Jianye Hao, Pheng-Ann Heng
- [Recruitment Strategies That Take a Chance](#) Gregory Kehne, Ariel D. Procaccia, Jingyan Wang
- [Fully Convolutional One-Stage 3D Object Detection on LiDAR Range Images](#) Zhi Tian, Xiangxiang Chu, Xiaoming Wang, Xiaolin Wei, Chunhua Shen
- [Understanding Robust Learning through the Lens of Representation Similarities](#) Christian Cianfarani, Arjun Nitin Bhagoji, Vikash Sehwal, Ren Zhao, Heather Zhang, Prateek Mittal

- [Chartalist: Labeled Graph Datasets for UTXO and Account-based Blockchains](#) Kiarash Shamsi, Friedhelm Victor, Murat Kantarcioglu, Yulia Gel, Cuneyt G Akcora
- [Multi-Instance Causal Representation Learning for Instance Label Prediction and Out-of-Distribution Generalization](#) Weijia Zhang, Xuanhui Zhang, hanwen deng, Min-Ling Zhang
- [PerfectDou: Dominating DouDizhu with Perfect Information Distillation](#) Guan Yang, Minghuan Liu, Weijun Hong, Weinan Zhang, Fei Fang, Guangjun Zeng, Yue Lin
- [On the Double Descent of Random Features Models Trained with SGD](#) Fanghui Liu, Johan Suykens, Volkan Cevher
- [SALSA: Attacking Lattice Cryptography with Transformers](#) Emily Wenger, Mingjie Chen, Francois Charton, Kristin E. Lauter
- [Variable-rate hierarchical CPC leads to acoustic unit discovery in speech](#) Santiago Cuervo, Adrian Lancucki, Ricard Marxer, Paweł Rychlikowski, Jan K. Chorowski
- [Learning to Attack Federated Learning: A Model-based Reinforcement Learning Attack Framework](#) Henger Li, Xiaolin Sun, Zizhan Zheng
- [Learning Neural Set Functions Under the Optimal Subset Oracle](#) Zijing Ou, Tingyang Xu, Qinliang Su, Yingzhen Li, Peilin Zhao, Yatao Bian
- [A Near-Optimal Best-of-Both-Worlds Algorithm for Online Learning with Feedback Graphs](#) Chloé Rouyer, Dirk van der Hoeven, Nicolò Cesa-Bianchi, Yevgeny Seldin
- [Optimal Query Complexities for Dynamic Trace Estimation](#) David Woodruff, Fred Zhang, Richard Zhang
- [Learning from Stochastically Revealed Preference](#) John Birge, Xiaocheng Li, Chunlin Sun
- [Mutual Information Divergence: A Unified Metric for Multimodal Generative Models](#) Jin-Hwa Kim, Yunji Kim, Jiyoung Lee, Kang Min Yoo, Sang-Woo Lee
- [Delving into Out-of-Distribution Detection with Vision-Language Representations](#) Yifei Ming, Ziyang Cai, Jiuxiang Gu, Yiyao Sun, Wei Li, Yixuan Li
- [OnePose+ +: Keypoint-Free One-Shot Object Pose Estimation without CAD Models](#) Xingyi He, Jiaming Sun, Yuang Wang, Di Huang, Hujun Bao, Xiaowei Zhou
- [Confidence-based Reliable Learning under Dual Noises](#) Peng Cui, Yang Yue, Zhijie Deng, Jun Zhu
- [Multitasking Models are Robust to Structural Failure: A Neural Model for Bilingual Cognitive Reserve](#) Giannis Daras, Negin Raoof, Zoi Gkalitsiou, Alex Dimakis
- [Learning \(Very\) Simple Generative Models Is Hard](#) Sitan Chen, Jerry Li, Yuanzhi Li
- [PEER: A Comprehensive and Multi-Task Benchmark for Protein Sequence Understanding](#) Minghao Xu, Zuobai Zhang, Jiarui Lu, Zhaocheng Zhu, Yangtian Zhang, Ma Chang, Runcheng Liu, Jian Tang
- [Bring Your Own Algorithm for Optimal Differentially Private Stochastic Minimax Optimization](#) Liang Zhang, Kiran K. Thekumparampil, Sewoong Oh, Niao He
- [A Regret-Variance Trade-Off in Online Learning](#) Dirk van der Hoeven, Nikita Zhivotovskiy, Nicolò Cesa-Bianchi
- [Alternating Mirror Descent for Constrained Min-Max Games](#) Andre Wibisono, Molei Tao, Georgios Piliouras
- [Deep Counterfactual Estimation with Categorical Background Variables](#) Edward De Brouwer
- [SnAKE: Bayesian Optimization with Pathwise Exploration](#) Jose Pablo Folch, Shiqiang Zhang, Robert Lee, Behrang Shafei, David Walz, Calvin Tsay, Mark van der Wilk, Ruth Misener
- [Self-Supervised Learning with an Information Maximization Criterion](#) Serdar Ozsoy, Shadi Hamdan, Sercan Arik, Deniz Yuret, Alper Erdogan
- [Twibot-22: Towards Graph-Based Twitter Bot Detection](#) Shangbin Feng, Zhaoxuan Tan, Herun Wan, Ningnan Wang, Zilong Chen, Binchi Zhang, Qinghua Zheng, Wenqian Zhang, Zhenyu Lei, Shujie Yang, Xinshun Feng, Qingyue Zhang, Hongrui Wang, Yuhao Liu, Yuyang Bai, Heng Wang, Zijian Cai, Yanbo Wang, Lijing Zheng, Zihan Ma, Jundong Li, Minnan Luo
- [Listen to Interpret: Post-hoc Interpretability for Audio Networks with NMF](#) Jayneel Parekh, Sanjeel Parekh, Pavlo Mozharovskiy, Florence d'Alché-Buc, Gaël Richard
- [Model-based RL with Optimistic Posterior Sampling: Structural Conditions and Sample Complexity](#) Alekh Agarwal, Tong Zhang
- [Deep Architecture Connectivity Matters for Its Convergence: A Fine-Grained Analysis](#) Wuyang Chen, Wei Huang, Xinyu Gong, Boris Hanin, Zhangyang Wang
- [OrdinalCLIP: Learning Rank Prompts for Language-Guided Ordinal Regression](#) Wanhua Li, Xiaoke Huang, Zheng Zhu, Yansong Tang, Xiu Li, Jie Zhou, Jiwen Lu
- [Kernel Memory Networks: A Unifying Framework for Memory Modeling](#) Georgios Iatropoulos, Johanni Brea, Wulfram Gerstner
- [The First Optimal Acceleration of High-Order Methods in Smooth Convex Optimization](#) Dmitry Kovalev, Alexander Gasnikov
- [CCCP is Frank-Wolfe in disguise](#) Alp Yurtsever, Suvrit Sra
- [Uni\[MASK\]: Unified Inference in Sequential Decision Problems](#) Micah Carroll, Orr Paradise, Jessy Lin, Raluca Georgescu, Mingfei Sun, David Bignell, Stephanie Milani, Katja Hofmann, Matthew Hausknecht, Anca Dragan, Sam Devlin
- [Efficient Active Learning with Abstention](#) Yinglun Zhu, Robert Nowak
- [SInGE: Sparsity via Integrated Gradients Estimation of Neuron Relevance](#) Edouard YVINEC, Arnaud Dapogny, Matthieu Cord, Kevin Bailly
- [On the Complexity of Adversarial Decision Making](#) Dylan J Foster, Alexander Rakhlin, Ayush Sekhari, Karthik Sridharan
- [MoCapAct: A Multi-Task Dataset for Simulated Humanoid Control](#) Nolan Wagener, Andrey Kolobov, Felipe Vieira Frujeri, Ricky Loynd, Ching-An Cheng, Matthew Hausknecht
- [On the Effectiveness of Persistent Homology](#) Renata Turkes, Guido F. Montufar, Nina Otter

- [An Adaptive Deep RL Method for Non-Stationary Environments with Piecewise Stable Context](#) Xiaoyu Chen, Xiangming Zhu, Yufeng Zheng, Pushi Zhang, Li Zhao, Wenxue Cheng, Peng CHENG, Yongqiang Xiong, Tao Qin, Jianyu Chen, Tie-Yan Liu
- [Expediting Large-Scale Vision Transformer for Dense Prediction without Fine-tuning](#) WEICONG LIANG, YUHUI YUAN, Henghui Ding, Xiao Luo, Weihong Lin, Ding Jia, Zheng Zhang, Chao Zhang, Han Hu
- [Flowification: Everything is a normalizing flow](#) Bálint Máté, Samuel Klein, Tobias Golling, François Fleuret
- [A time-resolved theory of information encoding in recurrent neural networks](#) Rainer Engelken, Sven Goedeke
- [A Unified Analysis of Mixed Sample Data Augmentation: A Loss Function Perspective](#) Chanwoo Park, Sangdoo Yun, Sanghyuk Chun
- [Recursive Reinforcement Learning](#) Ernst Moritz Hahn, Mateo Perez, Sven Schewe, Fabio Somenzi, Ashutosh Trivedi, Dominik Wojtczak
- [Non-Linguistic Supervision for Contrastive Learning of Sentence Embeddings](#) Yiren Jian, Chongyang Gao, Soroush Vosoughi
- [Towards Video Text Visual Question Answering: Benchmark and Baseline](#) Minyi Zhao, Bingjia Li, Jie Wang, Wanqing Li, Wenjing Zhou, Lan Zhang, Shijie Xuyang, Zhihang Yu, Xinkun Yu, Guangze Li, Aobotao Dai, Shuigeng Zhou
- [4D Unsupervised Object Discovery](#) Yuqi Wang, Yuntao Chen, ZHAO-XIANG ZHANG
- [Faster Linear Algebra for Distance Matrices](#) Piotr Indyk, Sandeep Silwal
- [Contextual Bandits with Knapsacks for a Conversion Model](#) Zhen LI, Gilles Stoltz
- [Contrastive Learning as Goal-Conditioned Reinforcement Learning](#) Benjamin Eysenbach, Tianjun Zhang, Sergey Levine, Russ R. Salakhutdinov
- [The Pitfalls of Regularization in Off-Policy TD Learning](#) Gaurav Manek, J. Zico Kolter
- [MCMAE: Masked Convolution Meets Masked Autoencoders](#) Peng Gao, Teli Ma, Hongsheng Li, Ziyi Lin, Jifeng Dai, Yu Qiao
- [Robustness to Label Noise Depends on the Shape of the Noise Distribution](#) Diane Oyen, Michal Kucer, Nicolas Hengartner, Har Simrat Singh
- [Fast Neural Kernel Embeddings for General Activations](#) Insu Han, Amir Zandieh, Jaehoon Lee, Roman Novak, Lechao Xiao, Amin Karbasi
- [Deep invariant networks with differentiable augmentation layers](#) Cédric ROMMEL, Thomas Moreau, Alexandre Gramfort
- [Factorized-FL: Personalized Federated Learning with Parameter Factorization & Similarity Matching](#) Wonyong Jeong, Sung Ju Hwang
- [Reinforcement Learning with a Terminator](#) Guy Tennenholtz, Nadav Merlis, Lior Shani, Shie Mannor, Uri Shalit, Gal Chechik, Assaf Hallak, Gal Dalal
- [How Transferable are Video Representations Based on Synthetic Data?](#) Yo-whan Kim, Samarth Mishra, SouYoung Jin, Rameswar Panda, Hilde Kuehne, Leonid Karlinsky, Venkatesh Saligrama, Kate Saenko, Aude Oliva, Rogerio Feris
- [Sub-exponential time Sum-of-Squares lower bounds for Principal Components Analysis](#) Aaron Potechin, Goutham Rajendran
- [A Multilabel Classification Framework for Approximate Nearest Neighbor Search](#) Ville Hyvönen, Elias Jääsaari, Teemu Roos
- [Learnable Polyphase Sampling for Shift Invariant and Equivariant Convolutional Networks](#) Renan A. Rojas-Gomez, Teck-Yian Lim, Alex Schwing, Minh Do, Raymond A. Yeh
- [Not All Bits have Equal Value: Heterogeneous Precisions via Trainable Noise](#) Pedro Savarese, Xin Yuan, Yanjing Li, Michael Maire
- [Tensor Program Optimization with Probabilistic Programs](#) Junru Shao, Xiyu Zhou, Siyuan Feng, Bohan Hou, Ruihang Lai, Hongyi Jin, Wuwei Lin, Masahiro Masuda, Cody Hao Yu, Tianqi Chen
- [Deep Generative Model for Periodic Graphs](#) Shiyu Wang, Xiaojie Guo, Liang Zhao
- [Exploring the Limits of Domain-Adaptive Training for Detoxifying Large-Scale Language Models](#) Boxin Wang, Wei Ping, Chaowei Xiao, Peng Xu, Mostofa Patwary, Mohammad Shoeybi, Bo Li, Anima Anandkumar, Bryan Catanzaro
- [Relational Reasoning via Set Transformers: Provable Efficiency and Applications to MARL](#) Fengzhuo Zhang, Boyi Liu, Kaixin Wang, Vincent Tan, Zhuoran Yang, Zhaoran Wang
- [Missing Data Imputation and Acquisition with Deep Hierarchical Models and Hamiltonian Monte Carlo](#) Ignacio Peis, Chao Ma, José Miguel Hernández-Lobato
- [Score-Based Generative Models Detect Manifolds](#) Jakiw Pidstrigach
- [Distributionally Robust Optimization via Ball Oracle Acceleration](#) Yair Carmon, Danielle Hausler
- [Palm up: Playing in the Latent Manifold for Unsupervised Pretraining](#) Hao Liu, Tom Zahavy, Volodymyr Mnih, Satinder Singh
- [Fair Infinitesimal Jackknife: Mitigating the Influence of Biased Training Data Points Without Refitting](#) Prasanna Sattigeri, Soumya Ghosh, Inkit Padhi, Pierre Dognin, Kush R. Varshney
- [JAWS: Auditing Predictive Uncertainty Under Covariate Shift](#) Drew Prinster, Anqi Liu, Suchi Saria
- [DNA: Proximal Policy Optimization with a Dual Network Architecture](#) Matthew Aitchison, Penny Sweetser
- [Triangulation candidates for Bayesian optimization](#) Robert B. Gramacy, Annie Sauer, Nathan Wycoff
- [Masked Autoencoders As Spatiotemporal Learners](#) Christoph Feichtenhofer, haoqi fan, Yanghao Li, Kaiming He
- [PyramidCLIP: Hierarchical Feature Alignment for Vision-language Model Pretraining](#) Yuting Gao, Jinfeng Liu, Zihan Xu, Jun Zhang, Ke Li, Rongrong Ji, Chunhua Shen
- [On the Parameterization and Initialization of Diagonal State Space Models](#) Albert Gu, Karan Goel, Ankit Gupta, Christopher Ré
- [Implicit Regularization or Implicit Conditioning? Exact Risk Trajectories of SGD in High Dimensions](#) Courtney Paquette, Elliot Paquette, Ben Adlam, Jeffrey Pennington
- [Extrapolative Continuous-time Bayesian Neural Network for Fast Training-free Test-time Adaptation](#) Hengguan Huang, Xiangming Gu, Hao Wang, Chang Xiao, Hongfu Liu, Ye Wang
- [Global Convergence and Stability of Stochastic Gradient Descent](#) Vivak Patel, Shushu Zhang, Bowen Tian

- [Trap and Replace: Defending Backdoor Attacks by Trapping Them into an Easy-to-Replace Subnetwork](#) Haotao Wang, Junyuan Hong, Aston Zhang, Jiayu Zhou, Zhangyang Wang
- [Reinforcement Learning with Logarithmic Regret and Policy Switches](#) Grigoris Velegkas, Zhuoran Yang, Amin Karbasi
- [Cluster and Aggregate: Face Recognition with Large Probe Set](#) Minchul Kim, Feng Liu, Anil K Jain, Xiaoming Liu
- [GLIPv2: Unifying Localization and Vision-Language Understanding](#) Haotian Zhang, Pengchuan Zhang, Xiaowei Hu, Yen-Chun Chen, Liunian Li, Xiyang Dai, Lijuan Wang, Lu Yuan, Jenq-Neng Hwang, Jianfeng Gao
- [Rethinking Alignment in Video Super-Resolution Transformers](#) Shuwei Shi, Jinjin Gu, Liangbin Xie, Xintao Wang, Yujiu Yang, Chao Dong
- [Robustness in deep learning: The good \(width\), the bad \(depth\), and the ugly \(initialization\)](#) Zhenyu Zhu, Fanghui Liu, Grigorios Chrysos, Volkan Cevher
- [Learning to Scaffold: Optimizing Model Explanations for Teaching](#) Patrick Fernandes, Marcos Treviso, Danish Pruthi, André Martins, Graham Neubig
- [AutoLink: Self-supervised Learning of Human Skeletons and Object Outlines by Linking Keypoints](#) Xingzhe He, Bastian Wandt, Helge Rhodin
- [Composite Feature Selection Using Deep Ensembles](#) Fergus Imrie, Alexander Norcliffe, Pietro Lió, Mihaela van der Schaar
- [Scaling Multimodal Pre-Training via Cross-Modality Gradient Harmonization](#) Junru Wu, Yi Liang, feng han, Hassan Akbari, Zhangyang Wang, Cong Yu
- [Bounded-Regret MPC via Perturbation Analysis: Prediction Error, Constraints, and Nonlinearity](#) Yiheng Lin, Yang Hu, Guannan Qu, Tongxin Li, Adam Wierman
- [AniFaceGAN: Animatable 3D-Aware Face Image Generation for Video Avatars](#) Yue Wu, Yu Deng, Jiaolong Yang, Fangyun Wei, Qifeng Chen, Xin Tong
- [Addressing Resource Scarcity across Sign Languages with Multilingual Pretraining and Unified-Vocabulary Datasets](#) Gokul NC, Manideep Ladi, Sumit Negi, Prem Selvaraj, Pratyush Kumar, Mitesh Khapra
- [Detection and Localization of Changes in Conditional Distributions](#) Lizhen Nie, Dan Nicolae
- [Refining Low-Resource Unsupervised Translation by Language Disentanglement of Multilingual Translation Model](#) Xuan-Phi Nguyen, Shafiq Joty, Kui Wu, Ai Ti Aw
- [A Mean-Field Game Approach to Cloud Resource Management with Function Approximation](#) Weichao Mao, Haoran Qiu, Chen Wang, Hubertus Franke, Zbigniew Kalbarczyk, Ravishankar Iyer, Tamer Basar
- [Regret Bounds for Risk-Sensitive Reinforcement Learning](#) Osbert Bastani, Jason Yecheng Ma, Estelle Shen, Wanqiao Xu
- [BILCO: An Efficient Algorithm for Joint Alignment of Time Series](#) Xuelong Mi, Mengfan Wang, Alex Chen, Jing-Xuan Lim, Yizhi Wang, Misha B Ahrens, Guoqiang Yu
- [Giving Feedback on Interactive Student Programs with Meta-Exploration](#) Evan Liu, Moritz Stephan, Allen Nie, Chris Piech, Emma Brunskill, Chelsea Finn
- [Left Heavy Tails and the Effectiveness of the Policy and Value Networks in DNN-based best-first search for Sokoban Planning](#) Dieqiao Feng, Carla P. Gomes, Bart Selman
- [Nonparametric Uncertainty Quantification for Single Deterministic Neural Network](#) Nikita Kotelevskii, Aleksandr Artemenkov, Kirill Fedyanin, Fedor Noskov, Alexander Fishkov, Artem Shelmanov, Artem Vazhentsev, Aleksandr Petiushko, Maxim Panov
- [Decoupling Features in Hierarchical Propagation for Video Object Segmentation](#) Zongxin Yang, Yi Yang
- [Computationally Efficient Horizon-Free Reinforcement Learning for Linear Mixture MDPs](#) Dongruo Zhou, Quanquan Gu
- [Counterfactual harm](#) Jonathan Richens, Rory Beard, Daniel H. Thompson
- [Chain of Thought Imitation with Procedure Cloning](#) Mengjiao (Sherry) Yang, Dale Schuurmans, Pieter Abbeel, Ofir Nachum
- [Revisiting Optimal Convergence Rate for Smooth and Non-convex Stochastic Decentralized Optimization](#) Kun Yuan, Xinmeng Huang, Yiming Chen, Xiaohan Zhang, Yingya Zhang, Pan Pan
- [Training with More Confidence: Mitigating Injected and Natural Backdoors During Training](#) Zhenting Wang, Hailun Ding, Juan Zhai, Shiqing Ma
- [Out-of-Distribution Detection via Conditional Kernel Independence Model](#) Yu Wang, Jingjing Zou, Jingyang Lin, Qing Ling, Yingwei Pan, Ting Yao, Tao Mei
- [Online Convex Optimization with Hard Constraints: Towards the Best of Two Worlds and Beyond](#) Hengquan Guo, Xin Liu, Honghao Wei, Lei Ying
- [ResT V2: Simpler, Faster and Stronger](#) Qinglong Zhang, Yu-Bin Yang
- [Object-Category Aware Reinforcement Learning](#) Qi Yi, Rui Zhang, shaohui peng, Jiaming Guo, Xing Hu, Zidong Du, xishan zhang, Qi Guo, Yunji Chen
- [Learning Partial Equivariances From Data](#) David W. Romero, Suhas Lohit
- [Photorealistic Text-to-Image Diffusion Models with Deep Language Understanding](#) Chitwan Saharia, William Chan, Saurabh Saxena, Lala Li, Jay Whang, Emily L. Denton, Kamyar Ghasemipour, Raphael Gontijo Lopes, Burcu Karagol Ayan, Tim Salimans, Jonathan Ho, David J. Fleet, Mohammad Norouzi
- [A Simple Decentralized Cross-Entropy Method](#) Zichen Zhang, Jun Jin, Martin Jagersand, Jun Luo, Dale Schuurmans
- [MSDS: A Large-Scale Chinese Signature and Token Digit String Dataset for Handwriting Verification](#) Peirong Zhang, Jiajia Jiang, Yuliang Liu, Lianwen Jin
- [Do Residual Neural Networks discretize Neural Ordinary Differential Equations?](#) Michael Sander, Pierre Ablin, Gabriel Peyré

- [Diffusion-based Molecule Generation with Informative Prior Bridges](#) Lemeng Wu, Chengyue Gong, Xingchao Liu, Mao Ye, Qiang Liu
- [Generative multitask learning mitigates target-causing confounding](#) Taro Makino, Krzysztof Geras, Kyunghyun Cho
- [Revisit last-iterate convergence of mSGD under milder requirement on step size](#) ruinan Jin, Xingkang He, Lang Chen, Difei Cheng, Vijay Gupta
- [Are You Stealing My Model? Sample Correlation for Fingerprinting Deep Neural Networks](#) Jiyang Guan, Jian Liang, Ran He
- [SecureFedYJ: a safe feature Gaussianization protocol for Federated Learning](#) Tanguy Marchand, Boris Muzellec, Constance Béguier, Jean Ogier du Terrail, Mathieu Andreux
- [When to Trust Your Simulator: Dynamics-Aware Hybrid Offline-and-Online Reinforcement Learning](#) Haoyi Niu, shubham sharma, Yiwen Qiu, Ming Li, Guyue Zhou, Jianming HU, Xianyuan Zhan
- [Data-Efficient Structured Pruning via Submodular Optimization](#) Marwa El Halabi, Suraj Srinivas, Simon Lacoste-Julien
- [Interpolation and Regularization for Causal Learning](#) Leena Chennuru Vankadara, Luca Rendsburg, Ulrike Luxburg, Debarghya Ghoshdastidar
- [A Direct Approximation of AIXI Using Logical State Abstractions](#) Samuel Yang-Zhao, Tianyu Wang, Kee Siong Ng
- [Learning Representations via a Robust Behavioral Metric for Deep Reinforcement Learning](#) Jianda Chen, Sinno Pan
- [Finding Second-Order Stationary Points in Nonconvex-Strongly-Concave Minimax Optimization](#) Luo Luo, Yujun Li, Cheng Chen
- [Contextual Squeeze-and-Excitation for Efficient Few-Shot Image Classification](#) Massimiliano Patacchiola, John Bronskill, Aliaksandra Shysheya, Katja Hofmann, Sebastian Nowozin, Richard Turner
- [Fuzzy Learning Machine](#) Junbiao Cui, Jiye Liang
- [Learning Structure from the Ground up---Hierarchical Representation Learning by Chunking](#) Shuchen Wu, Noemi Elteto, Ishita Dasgupta, Eric Schulz
- [AMOS: A Large-Scale Abdominal Multi-Organ Benchmark for Versatile Medical Image Segmentation](#) Yuanfeng Ji, Haotian Bai, Chongjian GE, Jie Yang, Ye Zhu, Ruimao Zhang, Zhen Li, Lingyan Zhanng, Wanling Ma, Xiang Wan, Ping Luo
- [The price of ignorance: how much does it cost to forget noise structure in low-rank matrix estimation?](#) Jean Barbier, TianQi Hou, Marco Mondelli, Manuel Saenz
- [Scalable Interpretability via Polynomials](#) Abhimanyu Dubey, Filip Radenovic, Dhruv Mahajan
- [DeVRF: Fast Deformable Voxel Radiance Fields for Dynamic Scenes](#) Jia-Wei Liu, Yan-Pei Cao, Weijia Mao, Wenqiao Zhang, David Junhao Zhang, Jussi Keppo, Ying Shan, Xiaohu Qie, Mike Zheng Shou
- [Hiding Images in Deep Probabilistic Models](#) Haoyu Chen, Linqi Song, Zhenxing Qian, Xinpeng Zhang, Kede Ma
- [ViewFool: Evaluating the Robustness of Visual Recognition to Adversarial Viewpoints](#) Yinpeng Dong, Shouwei Ruan, Hang Su, Caixin Kang, Xingxing Wei, Jun Zhu
- [Learning Superpoint Graph Cut for 3D Instance Segmentation](#) Le Hui, Linghua Tang, Yaqi Shen, Jin Xie, Jian Yang
- [Estimating the Arc Length of the Optimal ROC Curve and Lower Bounding the Maximal AUC](#) Song Liu
- [Acceleration in Distributed Sparse Regression](#) Marie Maros, Gesualdo Scutari
- [Latency-aware Spatial-wise Dynamic Networks](#) Yizeng Han, Zhihang Yuan, Yifan Pu, Chenhao Xue, Shiji Song, Guangyu Sun, Gao Huang
- [Towards Versatile Embodied Navigation](#) Hanqing Wang, Wei Liang, Luc V Gool, Wenguan Wang
- [Explain My Surprise: Learning Efficient Long-Term Memory by predicting uncertain outcomes](#) Artyom Sorokin, Nazar Buzun, Leonid Pugachev, Mikhail Burtsev
- [Polyhistor: Parameter-Efficient Multi-Task Adaptation for Dense Vision Tasks](#) Yen-Cheng Liu, CHIH-YAO MA, Junjiao Tian, Zijian He, Zsolt Kira
- [LAPO: Latent-Variable Advantage-Weighted Policy Optimization for Offline Reinforcement Learning](#) Xi Chen, Ali Ghadirzadeh, Tianhe Yu, Jianhao Wang, Alex Yuan Gao, Wenzhe Li, Liang Bin, Chelsea Finn, Chongjie Zhang
- [GAMA: Generative Adversarial Multi-Object Scene Attacks](#) Abhishek Aich, Calvin-Khang Ta, Akash Gupta, Chengyu Song, Srikanth Krishnamurthy, Salman Asif, Amit Roy-Chowdhury
- [Predicting Label Distribution from Multi-label Ranking](#) Yunan Lu, Xiuyi Jia
- [Trajectory of Mini-Batch Momentum: Batch Size Saturation and Convergence in High Dimensions](#) Kiwon Lee, Andrew Cheng, Elliot Paquette, Courtney Paquette
- [Transformers from an Optimization Perspective](#) Yongyi Yang, zengfeng Huang, David P Wipf
- [PDSketch: Integrated Domain Programming, Learning, and Planning](#) Jiayuan Mao, Tomás Lozano-Pérez, Josh Tenenbaum, Leslie Kaelbling
- [Amortized Projection Optimization for Sliced Wasserstein Generative Models](#) Khai Nguyen, Nhat Ho
- [GT-GAN: General Purpose Time Series Synthesis with Generative Adversarial Networks](#) Jinsung Jeon, JEONGHAK KIM, Haryong Song, Seunghyeon Cho, Noseong Park
- [Heterogeneous Skill Learning for Multi-agent Tasks](#) Yuntao Liu, Yuan Li, Xinhai Xu, Yong Dou, Donghong Liu
- [On Margin Maximization in Linear and ReLU Networks](#) Gal Vardi, Ohad Shamir, Nati Srebro
- [Unsupervised Causal Generative Understanding of Images](#) Titas Anciukevicius, Patrick Fox-Roberts, Edward Rosten, Paul Henderson
- [DART: Articulated Hand Model with Diverse Accessories and Rich Textures](#) Daiheng Gao, Yuliang Xiu, Kailin Li, Lixin Yang, Feng Wang, Peng Zhang, Bang Zhang, Cewu Lu, Ping Tan
- [BadPrompt: Backdoor Attacks on Continuous Prompts](#) Xiangrui Cai, Haidong Xu, Sihan Xu, Ying ZHANG, Yuan xiaojie
- [Off-Policy Evaluation with Policy-Dependent Optimization Response](#) Wenshuo Guo, Michael Jordan, Angela Zhou

- [Sampling in Constrained Domains with Orthogonal-Space Variational Gradient Descent](#) Ruqi Zhang, Qiang Liu, Xin Tong
- [Change-point Detection for Sparse and Dense Functional Data in General Dimensions](#) Carlos Misael Madrid Padilla, Daren Wang, Zifeng Zhao, Yi Yu
- [ClimbQ: Class Imbalanced Quantization Enabling Robustness on Efficient Inferences](#) Ting-An Chen, De-Nian Yang, Ming-syan Chen
- [Look Around and Refer: 2D Synthetic Semantics Knowledge Distillation for 3D Visual Grounding](#) Eslam Bakr, Yasmeen Alsaedy, Mohamed Elhoseiny
- [Nonstationary Dual Averaging and Online Fair Allocation](#) Luofeng Liao, Yuan Gao, Christian Kroer
- [Incentivizing Combinatorial Bandit Exploration](#) Xinyan Hu, Dung Ngo, Aleksandrs Slivkins, Steven Z. Wu
- [Transfer Learning on Heterogeneous Feature Spaces for Treatment Effects Estimation](#) Ioana Bica, Mihaela van der Schaar
- [Is Out-of-Distribution Detection Learnable?](#) Zhen Fang, Yixuan Li, Jie Lu, Jiahua Dong, Bo Han, Feng Liu
- [DMAP: a Distributed Morphological Attention Policy for learning to locomote with a changing body](#) Alberto Silvio Chiappa, Alessandro Marin Vargas, Alexander Mathis
- [On the role of overparameterization in off-policy Temporal Difference learning with linear function approximation](#) Valentin Thomas
- [Whitening Convergence Rate of Coupling-based Normalizing Flows](#) Felix Draxler, Christoph Schnörr, Ullrich Köthe
- [Neurosymbolic Deep Generative Models for Sequence Data with Relational Constraints](#) Halley Young, Maxwell Du, Osbert Bastani
- [Semi-Supervised Generative Models for Multiagent Trajectories](#) Dennis Fassmeyer, Pascal Fassmeyer, Ulf Brefeld
- [PhysGNN: A Physics--Driven Graph Neural Network Based Model for Predicting Soft Tissue Deformation in Image--Guided Neurosurgery](#) Yasmin Salehi, Dennis Giannacopoulos
- [Towards Diverse and Faithful One-shot Adaption of Generative Adversarial Networks](#) Yabo Zhang, mingshuai Yao, Yuxiang Wei, Zhilong Ji, Jinfeng Bai, Wangmeng Zuo
- [Deep Bidirectional Language-Knowledge Graph Pretraining](#) Michihiro Yasunaga, Antoine Bosselut, Hongyu Ren, Xikun Zhang, Christopher D Manning, Percy S. Liang, Jure Leskovec
- [Theoretical analysis of deep neural networks for temporally dependent observations](#) Mingliang Ma, Abolfazl Safikhani
- [A Theoretical Framework for Inference Learning](#) Nick Alonso, Beren Millidge, Jeffrey Krichmar, Emre O Neftci
- [Cross-modal Learning for Image-Guided Point Cloud Shape Completion](#) Emanuele Aiello, Diego Valsesia, Enrico Magli
- [Coreset for Line-Sets Clustering](#) Sagi Lotan, Ernesto Evgeniy Sanches Shayda, Dan Feldman
- [Robust On-Policy Sampling for Data-Efficient Policy Evaluation in Reinforcement Learning](#) Rujie Zhong, Duohan Zhang, Lukas Schäfer, Stefano Albrecht, Josiah Hanna
- [Pre-Trained Model Reusability Evaluation for Small-Data Transfer Learning](#) Yao-Xiang Ding, Xi-Zhu Wu, Kun Zhou, Zhi-Hua Zhou
- [Tree ensemble kernels for Bayesian optimization with known constraints over mixed-feature spaces](#) Alexander Thebelt, Calvin Tsay, Robert Lee, Nathan Sudermann-Merx, David Walz, Behrang Shafei, Ruth Misener
- [RLIP: Relational Language-Image Pre-training for Human-Object Interaction Detection](#) Hangjie Yuan, Jianwen Jiang, Samuel Albanie, Tao Feng, Ziyuan Huang, Dong Ni, Mingqian Tang
- [Skills Regularized Task Decomposition for Multi-task Offline Reinforcement Learning](#) Minjong Yoo, SangWoo Cho, Honguk Woo
- [Contextual Dynamic Pricing with Unknown Noise: Explore-then-UCB Strategy and Improved Regrets](#) Yiyun Luo, Will Wei Sun, Yufeng Liu
- [Adversarially Robust Learning: A Generic Minimax Optimal Learner and Characterization](#) Omar Montasser, Steve Hanneke, Nati Srebro
- [The Implicit Delta Method](#) Nathan Kallus, James McInerney
- [Singular Value Fine-tuning: Few-shot Segmentation requires Few-parameters Fine-tuning](#) Yanpeng Sun, Qiang Chen, Xiangyu He, Jian Wang, Haocheng Feng, Junyu Han, Errui Ding, Jian Cheng, Zechao Li, Jingdong Wang
- [On the relationship between variational inference and auto-associative memory](#) Louis Annabi, Alexandre Pitti, Mathias Quoy
- [Annihilation of Spurious Minima in Two-Layer ReLU Networks](#) Yossi Arjevani, Michael Field
- [A Closer Look at Weakly-Supervised Audio-Visual Source Localization](#) Shentong Mo, Pedro Morgado
- [Posterior Collapse of a Linear Latent Variable Model](#) Zihao Wang, Liu Ziyin
- [Accelerating SGD for Highly Ill-Conditioned Huge-Scale Online Matrix Completion](#) Jialun Zhang, Hong-Ming Chiu, Richard Y Zhang
- [Lifting Weak Supervision To Structured Prediction](#) Harit Vishwakarma, Frederic Sala
- [A Lagrangian Duality Approach to Active Learning](#) Juan Elenter, Navid Naderializadeh, Alejandro Ribeiro
- [Instance-optimal PAC Algorithms for Contextual Bandits](#) Zhaoqi Li, Lillian Ratliff, houssam nassif, Kevin G. Jamieson, Lalit Jain
- [xView3-SAR: Detecting Dark Fishing Activity Using Synthetic Aperture Radar Imagery](#) Fernando Paolo, Tsu-ting Tim Lin, Ritwik Gupta, Bryce Goodman, Nirav Patel, Daniel Kuster, David Kroodsma, Jared Dunnmon
- [Understanding the Failure of Batch Normalization for Transformers in NLP](#) Jiaxi Wang, Ji Wu, Lei Huang
- [Exploration via Elliptical Episodic Bonuses](#) Mikael Henaff, Roberta Raileanu, Minqi Jiang, Tim Rocktäschel
- [Efficient and Effective Multi-task Grouping via Meta Learning on Task Combinations](#) Xiaozhuang Song, Shun Zheng, Wei Cao, James Yu, Jiang Bian
- [Learning Fractional White Noises in Neural Stochastic Differential Equations](#) Anh Tong, Thanh Nguyen-Tang, Toan Tran, Jaesik Choi
- [CryptoGCN: Fast and Scalable Homomorphically Encrypted Graph Convolutional Network Inference](#) Ran Ran, Wei Wang, Quan Gang, Jieming Yin, Nuo Xu, Wujie Wen
- [UniGAN: Reducing Mode Collapse in GANs using a Uniform Generator](#) Ziqi Pan, Li Niu, Liqing Zhang
- [UMIX: Improving Importance Weighting for Subpopulation Shift via Uncertainty-Aware Mixup](#) Zongqbo Han, Zhipeng Lian, Fan Yang, Liu

[Report an Issue](#) | [Name Change Policy](#)

- [Exploit Reward Shifting in Value-Based Deep-RL: Optimistic Curiosity-Based Exploration and Conservative Exploitation via Linear Reward Shaping](#) Hao Sun, Lei Han, Rui Yang, Xiaoteng Ma, Jian Guo, Bolei Zhou
- [Stability and Generalization for Markov Chain Stochastic Gradient Methods](#) Puyu Wang, Yunwen Lei, Yiming Ying, Ding-Xuan Zhou
- [Degradation-Aware Unfolding Half-Shuffle Transformer for Spectral Compressive Imaging](#) Yuanhao Cai, Jing Lin, Haoqian Wang, Xin Yuan, Henghui Ding, Yulun Zhang, Radu Timofte, Luc V Gool
- [Deep Surrogate Assisted Generation of Environments](#) Varun Bhatt, Bryon Tjanaka, Matthew Fontaine, Stefanos Nikolaidis
- [Toward Equation of Motion for Deep Neural Networks: Continuous-time Gradient Descent and Discretization Error Analysis](#) Taiki Miyagawa
- [FLAIR: Federated Learning Annotated Image Repository](#) Congzheng Song, Filip Granqvist, Kunal Talwar
- [Self-Supervised Pretraining for Large-Scale Point Clouds](#) Zaiwei Zhang, Min Bai, Erran Li Li
- [Vision Transformers provably learn spatial structure](#) Samy Jelassi, Michael Sander, Yuezhi Li
- [Unsupervised Learning of Shape Programs with Repeatable Implicit Parts](#) Boyang Deng, Sumith Kulal, Zhengyang Dong, Congyue Deng, Yonglong Tian, Jiajun Wu
- [Detecting Abrupt Changes in Sequential Pairwise Comparison Data](#) Wanshan Li, Alessandro Rinaldo, Daren Wang
- [Rethinking Resolution in the Context of Efficient Video Recognition](#) Chuofan Ma, Qiushan Guo, Yi Jiang, Ping Luo, Zehuan Yuan, Xiaojuan Qi
- [The Effects of Regularization and Data Augmentation are Class Dependent](#) Randall Balestriero, Leon Bottou, Yann LeCun
- [Learning to Mitigate AI Collusion on Economic Platforms](#) Gianluca Brero, Eric Mibuari, Nicolas Lepore, David C. Parkes
- [The Query Complexity of Cake Cutting](#) Simina Branzei, Noam Nisan
- [PAC Prediction Sets for Meta-Learning](#) Sangdon Park, Edgar Dobriban, Insup Lee, Osbert Bastani
- [High-dimensional Asymptotics of Feature Learning: How One Gradient Step Improves the Representation](#) Jimmy Ba, Murat A. Erdogdu, Taiji Suzuki, Zhichao Wang, Denny Wu, Greg Yang
- [Pruning's Effect on Generalization Through the Lens of Training and Regularization](#) Tian Jin, Michael Carbin, Dan Roy, Jonathan Frankle, Gintare Karolina Dziugaite
- [Cluster Randomized Designs for One-Sided Bipartite Experiments](#) Jennifer Brennan, Vahab Mirrokni, Jean Pouget-Abadie
- [Deep Equilibrium Approaches to Diffusion Models](#) Ashwini Pokle, Zhengyang Geng, J. Zico Kolter
- [Inducing Neural Collapse in Imbalanced Learning: Do We Really Need a Learnable Classifier at the End of Deep Neural Network?](#) Yibo Yang, Shixiang Chen, Xiangtai Li, Liang Xie, Zhouchen Lin, Dacheng Tao
- [Pruning Neural Networks via Coresets and Convex Geometry: Towards No Assumptions](#) Murad Tukan, Loay Mualem, Alaa Maalouf
- [Intermediate Prototype Mining Transformer for Few-Shot Semantic Segmentation](#) YUANWEI LIU, Nian Liu, Xiwen Yao, Junwei Han
- [Long-Form Video-Language Pre-Training with Multimodal Temporal Contrastive Learning](#) Yuchong Sun, Hongwei Xue, Ruihua Song, Bei Liu, Huan Yang, Jianlong Fu
- [A new dataset for multilingual keyphrase generation](#) Frédéric Piedboeuf, Philippe Langlais
- [Model Preserving Compression for Neural Networks](#) Jerry Chee, Megan Flynn (née Renz), Anil Damle, Christopher M. De Sa
- [Neural Conservation Laws: A Divergence-Free Perspective](#) Jack Richter-Powell, Yaron Lipman, Ricky T. Q. Chen
- [Towards Effective Multi-Modal Interchanges in Zero-Resource Sounding Object Localization](#) Yang Zhao, Chen Zhang, Haifeng Huang, Haoyuan Li, Zhou Zhao
- [On the Convergence of Stochastic Multi-Objective Gradient Manipulation and Beyond](#) Shiji Zhou, Wenpeng Zhang, Jiyan Jiang, Wenliang Zhong, Jinjie GU, Wenwu Zhu
- [Decentralized Local Stochastic Extra-Gradient for Variational Inequalities](#) Aleksandr Beznosikov, Pavel Dvurechenskii, Anastasiia Koloskova, Valentin Samokhin, Sebastian U. Stich, Alexander Gasnikov
- [Model Zoos: A Dataset of Diverse Populations of Neural Network Models](#) Konstantin Schürholt, Diyar Taskiran, Boris Knyazev, Xavier Giró-i-Nieto, Damian Borth
- [Weakly-Supervised Multi-Granularity Map Learning for Vision-and-Language Navigation](#) Peihao Chen, Dongyu Ji, Kunyang Lin, Runhao Zeng, Thomas Li, Minghui Tan, Chuang Gan
- [Thinned random measures for sparse graphs with overlapping communities](#) Federica Zoe Ricci, Michele Guindani, Erik Sudderth
- [Fine-tuning language models to find agreement among humans with diverse preferences](#) Michiel Bakker, Martin Chadwick, Hannah Sheahan, Michael Tessler, Lucy Campbell-Gillingham, Jan Balaguer, Nat McAleese, Amelia Glaese, John Aslanides, Matt Botvinick, Christopher Summerfield
- [What is a Good Metric to Study Generalization of Minimax Learners?](#) Asuman Ozdaglar, Sarath Pattathil, Jiawei Zhang, Kaiqing Zhang
- [Sequencer: Deep LSTM for Image Classification](#) Yuki Tatsunami, Masato Taki
- [Double Check Your State Before Trusting It: Confidence-Aware Bidirectional Offline Model-Based Imagination](#) Jiafei Lyu, Xiu Li, Zongqing Lu
- [Learning on Arbitrary Graph Topologies via Predictive Coding](#) Tommaso Salvatori, Luca Pinchetti, Beren Millidge, Yuhang Song, Tianyi Bao, Rafal Bogacz, Thomas Lukasiewicz
- [Robustness Disparities in Face Detection](#) Samuel Dooley, George Z Wei, Tom Goldstein, John Dickerson
- [Marksman Backdoor: Backdoor Attacks with Arbitrary Target Class](#) Khoa D Doan, Yingjie Lao, Ping Li
- [Memorization Without Overfitting: Analyzing the Training Dynamics of Large Language Models](#) Kushal Tirumala, Aram Markosyan, Luke

- [Bandit Theory and Thompson Sampling-Guided Directed Evolution for Sequence Optimization](#) Hui Yuan, Chengzhuo Ni, Huazheng Wang, Xuezhou Zhang, Le Cong, Csaba Szepesvari, Mengdi Wang
- [Scalable and Efficient Training of Large Convolutional Neural Networks with Differential Privacy](#) Zhiqi Bu, Jialin Mao, Shiyun Xu
- [Weakly supervised causal representation learning](#) Johann Brehmer, Pim de Haan, Phillip Lippe, Taco S. Cohen
- [Zeroth-Order Negative Curvature Finding: Escaping Saddle Points without Gradients](#) Hualin Zhang, Huan Xiong, Bin Gu
- [Exposing and Exploiting Fine-Grained Block Structures for Fast and Accurate Sparse Training](#) Peng Jiang, Lihan Hu, Shihui Song
- [DABS 2.0: Improved Datasets and Algorithms for Universal Self-Supervision](#) Alex Tamkin, Gaurab Banerjee, Mohamed Owda, Vincent Liu, Shashank Rammoorthy, Noah Goodman
- [Operator Splitting Value Iteration](#) Amin Rakhsha, Andrew Wang, Mohammad Ghavamzadeh, Amir-massoud Farahmand
- [Enhanced Latent Space Blind Model for Real Image Denoising via Alternative Optimization](#) Chao Ren, Yizhong Pan, Jie Huang
- [Self-Explaining Deviations for Coordination](#) Hengyuan Hu, Samuel Sokota, David Wu, Anton Bakhtin, Andrei Lupu, Brandon Cui, Jakob Foerster
- [Communication Efficient Federated Learning for Generalized Linear Bandits](#) Chuanhao Li, Hongning Wang
- [Active Learning for Multiple Target Models](#) Ying-Peng Tang, Sheng-Jun Huang
- [Descent Steps of a Relation-Aware Energy Produce Heterogeneous Graph Neural Networks](#) Hongjoon Ahn, Yongyi Yang, Quan Gan, Taesup Moon, David P Wipf
- [Multi-agent Performative Prediction with Greedy Deployment and Consensus Seeking Agents](#) Qiang Li, Chung-Yiu Yau, Hoi-To Wai
- [Preservation of the Global Knowledge by Not-True Distillation in Federated Learning](#) Gihun Lee, Minchan Jeong, Yongjin Shin, Sangmin Bae, Se-Young Yun
- [Finite-Time Regret of Thompson Sampling Algorithms for Exponential Family Multi-Armed Bandits](#) Tianyuan Jin, Pan Xu, Xiaokui Xiao, Anima Anandkumar
- [Graph Reordering for Cache-Efficient Near Neighbor Search](#) Benjamin Coleman, Santiago Segarra, Alexander J. Smola, Anshumali Shrivastava
- [MetaMask: Revisiting Dimensional Confounder for Self-Supervised Learning](#) Jiangmeng Li, Wenwen Qiang, Yanan Zhang, Wenyi Mo, Changwen Zheng, Bing Su, Hui Xiong
- [On Feature Learning in the Presence of Spurious Correlations](#) Pavel Izmailov, Polina Kirichenko, Nate Gruver, Andrew G. Wilson
- [Sparse2Dense: Learning to Densify 3D Features for 3D Object Detection](#) Tianyu Wang, Xiaowei Hu, Zhengzhe LIU, Chi-Wing Fu
- [Exploring Length Generalization in Large Language Models](#) Cem Anil, Yuhuai Wu, Anders Andreassen, Aitor Lewkowycz, Vedant Misra, Vinay Ramasesh, Ambrose Slone, Guy Gur-Ari, Ethan Dyer, Behnam Neyshabur
- [Stability and Generalization Analysis of Gradient Methods for Shallow Neural Networks](#) Yunwen Lei, Rong Jin, Yiming Ying
- [ViTPose: Simple Vision Transformer Baselines for Human Pose Estimation](#) Yufei Xu, Jing Zhang, Qiming ZHANG, Dacheng Tao
- [Re-Analyze Gauss: Bounds for Private Matrix Approximation via Dyson Brownian Motion](#) Oren Mangoubi, Nisheeth Vishnoi
- [ASPiRe: Adaptive Skill Priors for Reinforcement Learning](#) Mengda Xu, Manuela Veloso, Shuran Song
- [Neural Differential Equations for Learning to Program Neural Nets Through Continuous Learning Rules](#) Kazuki Irie, Francesco Faccio, Jürgen Schmidhuber
- [MEMO: Test Time Robustness via Adaptation and Augmentation](#) Marvin Zhang, Sergey Levine, Chelsea Finn
- [Learning-Augmented Algorithms for Online Linear and Semidefinite Programming](#) Elena Grigorescu, Young-San Lin, Sandeep Silwal, Maoyuan Song, Samson Zhou
- [Text-Adaptive Multiple Visual Prototype Matching for Video-Text Retrieval](#) Chengzhi Lin, Ancong Wu, Junwei Liang, Jun Zhang, Wenhao Ge, Wei-Shi Zheng, Chunhua Shen
- [Asymptotically Unbiased Instance-wise Regularized Partial AUC Optimization: Theory and Algorithm](#) HuiYang Shao, Qianqian Xu, Zhiyong Yang, Shilong Bao, Qingming Huang
- [Universality of Group Convolutional Neural Networks Based on Ridgelet Analysis on Groups](#) Sho Sonoda, Isao Ishikawa, Masahiro Ikeda
- [Error Correction Code Transformer](#) Yoni Choukroun, Lior Wolf
- [Towards Improving Calibration in Object Detection Under Domain Shift](#) Muhammad Akhtar Munir, Muhammad Haris Khan, M. Sarfraz, Mohsen Ali
- [Renyi Differential Privacy of Propose-Test-Release and Applications to Private and Robust Machine Learning](#) Jiachen T. Wang, Saeed Mahloujifar, Shouda Wang, Ruoxi Jia, Prateek Mittal
- [A Transformer-Based Object Detector with Coarse-Fine Crossing Representations](#) Zhishan Li, Ying Nie, Kai Han, Jianyuan Guo, Lei Xie, Yunhe Wang
- [Beyond Adult and COMPAS: Fair Multi-Class Prediction via Information Projection](#) Wael Alghamdi, Hsiang Hsu, Haewon Jeong, Hao Wang, Peter Michalak, Shahab Asoodeh, Flavio Calmon
- [Explicit Tradeoffs between Adversarial and Natural Distributional Robustness](#) Mazda Moayeri, Kiarash Banihashem, Soheil Feizi
- [Learning Latent Seasonal-Trend Representations for Time Series Forecasting](#) Zhiyuan Wang, Xovee Xu, Weifeng Zhang, Goce Trajcevski, Ting Zhong, Fan Zhou
- [JAHS-Bench-201: A Foundation For Research On Joint Architecture And Hyperparameter Search](#) Archit Bansal, Danny Stoll, Maciej Janowski, Arber Zela, Frank Hutter
- [Capturing Graphs with Hypo-Elliptic Diffusions](#) Csaba Toth, Darrick Lee, Celia Hacker, Harald Oberhauser

- [FedSR: A Simple and Effective Domain Generalization Method for Federated Learning](#) A. Tuan Nguyen, Philip Torr, Ser Nam Lim
- [SIXO: Smoothing Inference with Twisted Objectives](#) Dieterich Lawson, Allan Raventós, andrew warrington, Scott Linderman
- [Explicable Policy Search](#) Ze Gong, Yu ("Tony") Zhang
- [Exploring evolution-aware & -free protein language models as protein function predictors](#) Mingyang Hu, Fajie Yuan, Kevin Yang, Fusong Ju, Jin Su, Hui Wang, Fei Yang, Qiuyang Ding
- [Breaking Bad: A Dataset for Geometric Fracture and Reassembly](#) Silvia Sellán, Yun-Chun Chen, Ziyi Wu, Animesh Garg, Alec Jacobson
- [Fair and Optimal Decision Trees: A Dynamic Programming Approach](#) Jacobus van der Linden, Mathijs de Weerd, Emir Demirović
- [Recommender Forest for Efficient Retrieval](#) Chao Feng, Wuchao Li, Defu Lian, Zheng Liu, Enhong Chen
- [Graph Few-shot Learning with Task-specific Structures](#) Song Wang, Chen Chen, Jundong Li
- [DASCO: Dual-Generator Adversarial Support Constrained Offline Reinforcement Learning](#) Quan Vuong, Aviral Kumar, Sergey Levine, Yevgen Chebotar
- [Beyond L1: Faster and Better Sparse Models with skglm](#) Quentin Bertrand, Quentin Klopfenstein, Pierre-Antoine Bannier, Gauthier Gidel, Mathurin Massias
- [You Can't Count on Luck: Why Decision Transformers and RvS Fail in Stochastic Environments](#) Keiran Paster, Sheila McIlraith, Jimmy Ba
- [Cost-efficient Gaussian tensor network embeddings for tensor-structured inputs](#) Linjian Ma, Edgar Solomonik
- [Neural Transmitted Radiance Fields](#) Chengxuan Zhu, Renjie Wan, Boxin Shi
- [Geoclidean: Few-Shot Generalization in Euclidean Geometry](#) Joy Hsu, Jiajun Wu, Noah Goodman
- [Enabling Detailed Action Recognition Evaluation Through Video Dataset Augmentation](#) Jihoon Chung, Yu Wu, Olga Russakovsky
- [Unsupervised Skill Discovery via Recurrent Skill Training](#) Zheyuan Jiang, Jingyue Gao, Jianyu Chen
- [Structural Kernel Search via Bayesian Optimization and Symbolical Optimal Transport](#) Matthias Bitzer, Mona Meister, Christoph Zimmer
- [Robust Models are less Over-Confident](#) Julia Grabinski, Paul Gavrikov, Janis Keuper, Margret Keuper
- [Near-Optimal No-Regret Learning Dynamics for General Convex Games](#) Gabriele Farina, Ioannis Anagnostides, Haipeng Luo, Chung-Wei Lee, Christian Kroer, Tuomas Sandholm
- [OTKGE: Multi-modal Knowledge Graph Embeddings via Optimal Transport](#) Zongsheng Cao, Qianqian Xu, Zhiyong Yang, Yuan He, Xiaochun Cao, Qingming Huang
- [ComMU: Dataset for Combinatorial Music Generation](#) Hyun Lee, Taehyun Kim, Hyolim Kang, Minjoo Ki, Hyeonchan Hwang, kwanho park, Sharang Han, Seon Joo Kim