

≡ Hide menu

Mechanistic Interpretability Concepts

Circuits and Superposition

Representation Learning & LLMs

Video: Mechanistic Interpretability Representation Learning

5 min

Video: Scaling Mechanistic Interpretability to LLMs

7 min

Discussion Prompt: Visualizing an LLM

10 min

Reading: Mechanistic Interpretability Resource Guide

10 min

End of Module Assignments

Mechanistic Interpretability Resource Guide

- Read online:** For immediate reading or to bookmark for later, access the [Mechanistic Interpretability Resource Guide](#).
- Download PDF:** For those who prefer to curate their own learning library or need offline access, download the paper as a PDF.

< 1 / 2 >

🔍

100%

🔍

⬇️

🖨️

Mechanistic Interpretability Resource Guide

Getting Started

[Mechanistic Interpretability Quickstart Guide](#) [Neel Nanda, 2023]

[Concrete Steps to Get Started in Transformer Mechanistic Interpretability](#) [Neel Nanda, 2022]

[A Comprehensive Mechanistic Interpretability Explainer and Glossary](#) [Neel Nanda, 2022]

Explanatory Papers

[Zoom In: An Introduction to Circuits](#) [Oliah, et al. 2020]

[Circuit Defectology](#) [Cammurata, et al. 2020]

[Toy Models of Superposition](#) [Elhage, et al. 2022]

Research Papers

[List of Mechanistic Interpretability Papers](#) [Neel Nanda]

[Multimodal Neurons in Artificial Neural Networks](#) [Goh, et al. 2021]

[Towards Monosemanticity: Decomposing Language Models with Dictionary Learning](#) [Bricken, et al. 2023]

[Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet](#) [Templeton and Conerly, et al. 2024]

[Scaling and Evaluating Sparse Autoencoders](#) [Gao, et al. 2024]

Code

[OpenAI Sparse Autoencoders GitHub](#) (sparse autoencoders trained on gpt2-small)

[Transformer Interpretability short course](#)

Videos

[Induction Head Circuits](#) [lecture by Chris Olah]

[Mechanistic Interpretability & Mathematics](#) [lecture by Neel Nanda]

Future Looking

[Interpretability Dreams](#) [Olah, et al. 2023]

[200 Concrete Open Problems in Mechanistic Interpretability](#) [Neel Nanda, 2022]

[Open Questions](#) (from Toy Models of Superposition paper)

Go to next item

✓ Completed

👍 Like

👎 Dislike

📄 Report an issue

