



Menu


External lab: Interact with hosted models

In this external lab, you will use Hugging Face to explore an Open Source generative AI model and interact with one of its hosted models.

Optionally, you can create an account, but this is not required to interact with the models.

Steps

1. Go to the main [Hugging Face landing page](#) and navigate to the "Models" section
2. Filter down the models by finding "Text Generation" capabilities, or by searching a specific model by name. For example you can search for "Phi" which is Microsoft's Open Source model for text generation
3. Click on the model and interact with the inference API. For example, [this is the Phi-2 model page](#) with its inference API.



The screenshot shows the Hugging Face model card for Phi-2. The top navigation bar includes tabs for 'Model card', 'Files and versions', and 'Community' (with 105 members). Action buttons for 'Train', 'Deploy', and 'Use in Transformers' are visible. The 'Model Summary' section describes Phi-2 as a Transformer with 2.7 billion parameters, trained on synthetic texts and filtered websites. It mentions that the model was assessed against benchmarks for common sense, language understanding, and logical reasoning. The 'How to Use' section states that Phi-2 has been integrated into the development version (4.37.0.dev) of transformers. On the right, a 'Downloads last month' graph shows 546,965 downloads. Below this, the 'Safetensors' section shows the model size as 2.78B params and the tensor type as FP16. The 'Inference API' section includes a 'Text Generation' dropdown with 'Examples' selected. A text input field contains 'My name is Teven and I am'. Below the input is a 'Compute' button and a 'Maximize' link. A note at the bottom states 'This model can be loaded on the Inference API on-demand.'

Reflection Questions

1. Do you think the text generation was fast or slow? What are some reasons you think this is the case?
2. What are some of the warnings displayed in the model card? Use the Phi-2 page if the model you selected doesn't display any warnings in its documentation
3. Can you find limitations when interacting with the model using the inference API? What are those limitations? Are these explained and highlighted in the documentation?