

 Menu

External lab: Create a RAG HTTP API

For this external lab, you will use your knowledge on the Qdrant vector database alongside with Llamafire for the Large Language Model solution to update an existing application.

The application is currently using Azure but your task will be to update the Azure components to use the Qdrant database as well as Llamafire for the OpenAI APIs.

You will have to apply the knowledge of previous lessons to replace the components.

Given the requirements of running this application you will need to be mindful in selecting a machine that has a performant CPU and over 8GB of memory.

Steps

1. Clone the [Azure RAG repository](#) 
2. Update the requirements.txt file to [include the Qdrant and Sentence Transformers dependencies using the requirements.txt file from the Vector Database examples](#) 
3. [Download the Phi-2 Llamafire](#)  and start it up in your system
4. Update the `.env` file in the repository to point to the local Llamafire instead of Azure
5. Create a helper function that will produce the embeddings and load them to the vector database when the web application starts

Finally, verify your work by starting the web application and navigating to the `/docs` URL and interacting with the `/ask` endpoint.

[Go to next item](#) Completed Like  Dislike  Report an issue