



nanogeek815

Jun 20 '12

```
Hello all,

So after a few hours of different tests, I found that my cudaMemcpy, although returning successful, is not actually copying data from the host to the device.

Here's the code:

#include "cuda_runtime.h"

#include "device_launch_parameters.h"

#include <stdio.h>

#include <cstdlib>

#include <fstream>

#include <windows.h>

using namespace std;

/*****

* Variables that may be changed before running

*****/

//TESTING TYPE

//Timing test:  true

//Centroid test:  false

bool timing

bool cudaTest

//Accuracy changing here

typedef float decimalAccuracy;

//Size of the image in pixels

const int height

const int width

//Search box specifications

const int numberOfSearches

const decimalAccuracy initialSearchBoxRadius

//Iterative Specifications

decimalAccuracy searchBoxMultiplier

int numberOfIterations

//Timing testing

int timing_repetitions

//Input file names:

//Data file

const char *fileName      = "blah.dat";

//Reference coordinates

const char *fileNameRef    = "blah2.dat";

//Output file names:

//Result output file name

const char *outputFileName = "out.dat";

//Timing data output file name

const char *timingFileName  = "time_out.dat";

__global__ void centroidKernel  (decimalAccura

                                decimalAccura

                                decimalAccura

                                decimalAccura

                                decimalAccura

                                int

                                int
```

tera Top Contributor

Jun 20 '12

So after a few hours of different tests, I found that my cudaMemcpy, although returning successful, is not actually copying data from the host to the device.

How do you know the cudaMemcpy() calls are successful? Nowhere in your code the return code is checked...

nanogeek815

Jun 20 '12

How do you know the cudaMemcpy() calls are successful? Nowhere in your code the return code is checked...

It's right under it. I use printf so I can check it each time it runs.

```
cudaMemcpy(device_image_to_centroid,      imag

          printf("Device Variable Copying:\t\t\t\n"
```

nanogeek815

Jun 20 '12

Just to add a little more detail, it simply fills the new array with 0's, no memory leaks or anything either.

A.Goude

Jun 20 '12

Since I don't have the input files, I cannot run the code, which makes it hard to determine what happens without any additional information, as I do not know the output of the code. I did notice that you are actually copying the same vector twice in your code

```
cudaMemcpy(working_reference_coordinates_x, dev

cudaMemcpy(working_reference_coordinates_y, dev
```

could this be your problem?

nanogeek815

Jun 20 '12

Since I don't have the input files, I cannot run the code, which makes it hard to determine what happens without any additional information, as I do not know the output of the code. I did notice that you are actually copying the same vector twice in your code

```
cudaMemcpy(working_reference_coordinates_x, dev

cudaMemcpy(working_reference_coordinates_y, dev
```

could this be your problem?

That part has no real issue. Its the image file that it takes in. I would give you input files but it would take far too long to explain them all.

It copied the data if i make them all something simple like 7.0 but if it's the data pulled from the file (which i printf'ed as well and it was correct) it just copies an entire array of 0's.

I'm not really sure where to look for errors at this point.

RezaRob3

Jun 20 '12

It's right under it. I use printf so I can check it each time it runs.

You're not error checking after every call! Why??

New & Unread Topics

Topic	Replies	Views	Activity
Questions about "globaltimer", functionality, accessing and configuring	2	11	2h
Program stucks on cudaErrorMemoryAllocation after failing a cudaMallocHost	12	14	2h
How to operate irregular gemm on tensor core?	8	8	5h
Delay before kernel runs?	6	24	9h
Why Are My CUDA Compilation Results Different with Direct and Separate Compilation?	4	11	9h
Why early exit will influence later __syncthreads()?	2	7	10h
Does TMA need to commit?	0	6	10h
Why will different numbers of thread blocks cause bank conflicts? cuda kernel ubuntu	1	13	Aug 22
CUDA Trigonometric Function Issues	1	13	Aug 22
Is VBIOS version significant to report for reproducibility? pytorch	1	11	Aug 22

There are 199 new topics remaining, or browse other topics in

