

You can login using your huggingface.co credentials.



This forum is powered by [Discourse](#) and relies on a [trust-level system](#). As a new user, you're temporarily limited in the number of topics and posts you can create. To lift those restrictions, just spend time reading other posts (to be precise, enter 5 topics, read through 30 posts and spend a total of 10 minutes reading).

Start with reading [this post](#). Then maybe someone already had that error that is bugging you check with a quick [search](#). Or you can read the latest [awesome paper](#) the team discussed.

How can I prompt Llama to only use my provided context?

Lahjonta

Jan 22

Hi I'm trying to implement a RAG system using kNN to retrieve relevant documents and generate an answer based on those documents for a school project. Is there an easy way to implement prompt engineering in Llama-2 so it actually only uses the provided context? The output of my code is very nonsensical. Thanks in advance!

```
model_llama = "/kaggle/input/llama-2/pytorch/7b-chat-hf/1"

tokenizer_llama = AutoTokenizer.from_pretrained(model_llama)
pipeline_llama = transformers.pipeline(
    "text-generation",
    model=model_llama,
    tokenizer=tokenizer_llama,
)

def generate_answer_llama(question, pipeline=pipeline_llama, tokeni:
    # Perform the kNN query
    top_docs = query_knn(query, index, df)['text']
    # Concatenate the top documents to create the context
    documents = "".join(top_docs)

    # Define the System Prompt to only use our documents
    input = f"""
    <<SYS>>
    Only respond with "Not in the text." if the information need
    Answer the question using only the information from the att
    Ensure that the questions are answered fully and effectively
    Respond in short and concise yet fully formulated sentences
    <</SYS>>
    [INST]
    User:{question}
    [/INST]\n
    [INST]
    User:{documents}
    [/INST]\n

    Assistant:
    """

    sequences = pipeline(
        input,
        do_sample=True,
        top_k=50,
        num_return_sequences=2,
```

📁 Topics

⋮ More

⌵ Categories

Beginners

Intermediate

Course

Research

Models

☰ All categories

```
eos_token_id=tokenizer.eos_token_id,
max_new_tokens=2048,
return_full_text=False,
temperature=0.1,
)

for seq in sequences:
    generated_text = seq['generated_text']
    # Find the start of the assistant's answer and return only :
    answer_start = generated_text.find("Assistant:") + len("Ass:
    return generated_text[answer_start:].strip()

# Example usage
query = 'What is NLP?'
results = generate_answer_llama(query)

print(results)
```



svshenoy

Mar 3

Hey did you ever get an answer, trying to do the same thing

Related Topics

Topic	Replies	Views	Activity
Trying to understand system prompts with Llama 2 and transformers interface	8	31.1k	Apr 29
Llama 2 repeats its prompt as output without answering the prompt	2	1.6k	Mar 20
How to set Llama-2-Chat prompt context	2	13.4k	Oct 2023
How to prompt Llama2 for text classification?	0	2.4k	Sep 2023
Llama-2 7B-hf repeats context of question directly from input prompt, cuts off with newlines	14	24.2k	Jun 8