

You can login using your huggingface.co credentials.



This forum is powered by [Discourse](#) and relies on a [trust-level system](#). As a new user, you're temporarily limited in the number of topics and posts you can create. To lift those restrictions, just spend time reading other posts (to be precise, enter 5 topics, read through 30 posts and spend a total of 10 minutes reading).

Start with reading [this post](#). Then maybe someone already had that error that is bugging you check with a quick [search](#). Or you can read the latest [awesome paper](#) the team discussed.

How to just get the answer from Llama-2 instead of repeating the whole prompt?

anon9969843

Sep 2023

Here is my script:

```
from transformers import AutoTokenizer, AutoModelForCausalLM

tokenizer = AutoTokenizer.from_pretrained("meta-llama/Llama-2-7b-chat-hf")
model = AutoModelForCausalLM.from_pretrained("meta-llama/Llama-2-7b-chat-hf")

prompt = """
CONTEXT: Harvard University is a private Ivy League research university
Founded in 1636 as Harvard College and named for its first benefactor,
John Harvard.

QUESTION: Which year was Harvard University founded?
"""

input_ids = tokenizer(prompt, return_tensors="pt").input_ids
outputs = model.generate(input_ids, max_new_tokens=200)
print(tokenizer.decode(outputs[0]))
```

Here is the output:

```
<S>
CONTEXT: Harvard University is a private Ivy League research university
Founded in 1636 as Harvard College and named for its first benefactor,
it is the oldest institution of higher learning in the United States
and rankings have made it one of the most prestigious universities in the world.

QUESTION: Which year was Harvard University founded?
ANSWER: Harvard University was founded in 1636.</S>
```

How to prompt in order to get just the answer instead of repeating the input prompt?

dblakely

Sep 2023

There's no way to do this. The model is stateless so each time you call it, it needs to see the prompt for context about what you want it to do.

craighagerman

Apr 16

In order to retrieve just the response you need to slice the output at the index of the input tokens. Like so using `input_token_len`:

Topics

More

Categories

Beginners

[Intermediate](#)[Course](#)[Research](#)[Models](#)[☰ All categories](#)

```
input_ids = tokenizer(prompt, return_tensors="pt").input_ids
input_token_len = input_ids.shape[-1]
outputs = model.generate(input_ids, max_new_tokens=200)
print(tokenizer.decode(outputs[0][input_token_len:]))
```

Related Topics

Topic	Replies	Views	Activity
How can I prompt Llama to only use my provided context?	1	1.3k	Mar 3
Llama 2 repeats its prompt as output without answering the prompt	2	1.6k	Mar 20
How to prompt Llama2 for text classification?	0	2.4k	Sep 2023
LLama2 trained on completions only repeating prompt during inference	0	191	Apr 2
Llama 2 don't reponse prompt invokes	0	356	Feb 10

