

How to just get the answer from Llama-2 instead of repeating the whole prompt?

Asked 9 months ago Modified 4 months ago Viewed 3k times



Here is my script:

5

```
from transformers import AutoTokenizer, AutoModelForCausalLM
```



```
tokenizer = AutoTokenizer.from_pretrained("meta-llama/Llama-2-7b-chat-hf")
model = AutoModelForCausalLM.from_pretrained("meta-llama/Llama-2-7b-chat-hf")
```



```
prompt = """
CONTEXT: Harvard University is a private Ivy League research university in Cambridge, Massachusetts.
Founded in 1636 as Harvard College and named for its first benefactor, the Puritan clergyman John Harvard, it is the
oldest institution of higher learning in the United States. Its influence, wealth, and rankings have made it one of the
most prestigious universities in the world.
```

```
QUESTION: Which year was Harvard University found?
"""
```

```
input_ids = tokenizer(prompt, return_tensors="pt").input_ids
outputs = model.generate(input_ids, max_new_tokens=200)
print(tokenizer.decode(outputs[0]))
```

Here is the output:

```
<s>
CONTEXT: Harvard University is a private Ivy League research university in Cambridge, Massachusetts.
Founded in 1636 as Harvard College and named for its first benefactor, the Puritan clergyman John Harvard,
it is the oldest institution of higher learning in the United States. Its influence, wealth,
and rankings have made it one of the most prestigious universities in the world.
```

```
QUESTION: Which year was Harvard University found?
ANSWER: Harvard University was founded in 1636.</s>
```

You can see Llama-2 includes the input prompt for the output. Is there any way to remove the input prompt from the output?

[prompt-design](#) [llm](#) [llama-2](#)

Share Improve this question Follow

edited Sep 23, 2023 at 6:21

asked Sep 22, 2023 at 5:13

[etang](#)
213 1 8

Insert "don't say anything else but the value" and it'll only give the value. This could work in LLama, I used it in ChatGPT. – [Alias Cartellano](#) Sep 23, 2023 at 20:19

1 Maybe stackoverflow.com/a/77409249/3659012 can help? – [wxh](#) Nov 2, 2023 at 11:38

4 Answers

Sorted by: Highest score (default)



what i do is if it contains "Answer:" it only prints what follows. but if you find a better solution please let me know:

1

```
if "Answer:" in output_text:
    start_index = output_text.index("Answer:")
    print(output_text[start_index:])
else:
    print(output_text)
```



Share Improve this answer Follow

edited Dec 7, 2023 at 10:05

answered Dec 7, 2023 at 10:05

[Bryan Koolman](#)
11 2



You can use the `generate()` method and pass the `input_ids` directly without including the prompt text. This will generate the continuation without repeating the prompt.

0

For example:



```
input_ids = tokenizer(prompt, return_tensors="pt").input_ids  
outputs = model.generate(input_ids, max_new_tokens=50)  
print(tokenizer.decode(outputs[0]))
```



This will generate just the continuation text without the original prompt. The key is to not include the prompt text when encoding the input. Pass just the `input_ids` to `generate()`.

Share Improve this answer Follow

answered Dec 16, 2023 at 15:46



devvy devvy
36 1



From [transformers documentation](#)

0

```
input_length = len(model_inputs[0])  
Output = Tokenizer.batch_decode(GeneratedIds[:, ModelInputsLen:], skip_special_tokens=True)[0]
```



Share Improve this answer Follow

edited Mar 14 at 18:49

answered Mar 13 at 18:46



pixol
1 1



you're making it print its thought process, essentially. it should only print the answer. changing the last bit to the below should work:

-1

```
answer = tokenizer.decode(outputs[0], skip_special_tokens=True)  
print(answer)
```



Share Improve this answer Follow

answered Sep 23, 2023 at 12:57



iiOxygen
29 2