

≡ Hide menu

Under the Hood with Llama2 and Python: Understanding How it Works

Lab: Llama 2 Workspace

1 min

Video: Llama 2 Introduction

5 min

Video: Building Our First Llama 2 Application with llama.cpp

9 min

Video: Choices/Parameters

8 min

Video: Llama 2 Chat and Conversations

7 min

Graded Assignment: Module 2 Quiz

Submitted

Reading: Optional Reading: Sampling Methods, Bindings, and Quantization

10 min

Your grade: 100%

Your latest: 100% • Your highest: 100% • Your previous: 100% • We keep your highest score.

Module 2 Quiz

Next item →

1. True or false: You generate **Python Learning Objectives** using C and C++

1 / 1 point

☐ True, llama.cpp applications are exclusively written in C and C++ due to their high performance and compatibility with the library's low-level APIs.

☒ False, while you can write llama.cpp applications using C and C++, you can also use a wide array of languages including Rust and, of course, **Python**.

Correct

Submitted Jul 14, 12:51 AM IST

Try again

2. True or false: When you use **llama.cpp**, it automatically quantizes your models to fit your hardware.

1 / 1 point

☐ True, llama.cpp integrates an intelligent auto-tuning feature that quantizes models to achieve **best** performance on the available hardware.

☒ False, you must quantize models yourself with llama.cpp, or obtain already quantized models from another source (such as HuggingFace).

Correct

Like

Dislike

Report an issue

See feedback

3. True or false: llama.cpp is a toolkit which only works for Llama 2 and the original Llama model.

1 / 1 point

☒ False, it works with a number of other open source models.

☐ True, llama.cpp is specifically designed to support only Llama 2 and the original Llama model, ensuring optimized performance for these models.

Correct

Correct: Llama.cpp is not limited to only Llama 2 and the original Llama model. It supports a variety of other open source models such as Mistral 7B among others.

4. What is the context length for Llama 2?

1 / 1 point

☐ 8,192 tokens

☒ 4,096 tokens

☐ 2,048 tokens

☐ 4,096 words

Correct

Correct.

5. True or false: Tokenization converts each English word into a single number which is used by Llama 2 to predict new words.

1 / 1 point

☒ False, rarely does each word map to a single token, often words are made up of multiple tokens, e.g. the word Llama is represented by three tokens, one for L, one for i and one for ama.

☐ True, tokenization simplifies the text processing by assigning a unique number to each distinct word, facilitating the prediction of new words by Llama 2.

Correct

6. Fill in the blank. The top_k parameter ____:

1 / 1 point

☐ Determine the "creativity" of the model by changing how likely the token with the best probability (logit) is chosen.

☒ Restricts the number of tokens as an integer value which llama will then choose from.

☐ Restricts the combined probability of candidate tokens which llama will then reweight and choose from.

☐ Adjusts the frequency threshold for token selection, filtering out tokens that appear below a certain usage rate in the training data

Correct

Correct.

7. The role options for Llama 2 chat models include ____:

1 / 1 point

☐ Client, Server, Observer

☐ Administrator, Moderator, Participant

☐ Speaker, Listener, Interpreter

☒ User, System, Agent

Correct

Correct. These are the role options for Llama 2 chat models.

8. True or false: Using llama.cpp with the Llama 2 chat model requires you to fill in the instruction prompt string using the [INST] and [/INST] tags.

1 / 1 point

☒ False, the llama.cpp python bindings abstract this away, and allow you to create chat messages where you set the roles (e.g. system) and text of the messages. This will be converted to the appropriate prompt format automatically for you.

☐ True, manual insertion of [INST] and [/INST] tags is mandatory in all instruction prompt strings.

Correct

https://www.coursera.org/learn/llama-for-python-programmers/assignment/submission/UQ1zr/module-2-quiz/view-feedback

1/1