



Hugging Face is way more fun with friends and colleagues! 🤗 [Join an organization](#)

Dismiss this message

TheBloke / **Llama-2-70B-Chat-GPTQ** like 257

Text Generation Transformers Safetensors PyTorch English
llama facebook meta llama-2 text-generation-inference 4-bit precision
gptq arxiv:2307.09288 License: llama2

Train Deploy Use this model

Model card Files Community 51



Why the input prompt is part of the output? #25

by neo-benjamin - opened Jul 27, 2023

Discussion

neo-benjamin Jul 27, 2023



When I ran the code:

```
prompt = "Tell me about AI"
system_message = "You are a helpful, respectful and honest assistant"
prompt_template=f'''[INST] <<SYS>>
{system_message}
<</SYS>>

{prompt} [/INST]
...

print("\n\n*** Generate:")

input_ids = tokenizer(prompt_template, return_tensors='pt').input_ids
output = model.generate(inputs=input_ids, temperature=0.7, max_new_tokens=
print(tokenizer.decode(output[0]))
```

output contains the prompt along with the output i.e. response from the model.
So how token is calculated? the output token is max_new_tokens and the
max_new_tokens + output tokens needs to be less than 4096?

TheBloke Owner Jul 27, 2023



That's how `model.generate()` works. You could use `model.pipeline()` instead with `return_full_text=False`

`max_new_tokens` is the number of tokens it will generate in response to your prompt. `prompt + max_new_tokens` must be less than 4096.

 3



 **Michael528** Aug 30, 2023  

Is there a way to remove the input prompt from the output? I use below but still gives me the input prompt.

```
out_obj[0]['generated_text']
```

 **which47** Nov 26, 2023 · edited Nov 26, 2023  

```
output_ids = model.generate(input_ids)
output = tokenizer.decode(output_ids[:, input_ids.shape[1]:])
```

Edit

Preview

Leave a comment

 Tap or paste here to upload images

 Comment



Company

TOS

Privacy

About

Jobs

Website

Models

Datasets

Spaces

Pricing

Docs

© Hugging Face