Hide menu

NLP models

- **Video:** Introduction 1 min
- **Video:** ANN 10 min
- **Video:** TensorFlow 6 min
- **Video:** DNN 6 min
- **Video:** RNN 6 min
- **Video:** LSTM 8 min
- **Video:** GRU 2 min
- **Video:** Lab introduction: Text Classification with Keras 2 min
- **Graded App Item:** Lab: Text Classification with Keras 2h
- **Video:** Summary 3 min
- **Quiz:** Quiz 6 min
- **Reading:** Reading list 10 min

Congratulations! You passed!

Grade received 100%

Latest Submission Grade 100%

To pass 80% or higher

Go to next item

coach

Ready to review what you've learned before taking the quiz? I'm here to help.

Help me practice Let's chat

Submit your assignment

Try again

Due Aug 28, 11:59 PM IST

Receive grade

To Pass 80% or higher

Your grade 100%

View Feedback

We keep your highest score

Like Dislike Report an issue

1. What is the key feature to enable a "memory" of an RNN (recurrent neural network)? **1 / 1 point**

- An RNN has one hidden layer.
- An RNN has multiple hidden layers.
- An RNN uses a mechanism to carry the previous information to the next learning iteration.
- An RNN has a single lambda layer.

Correct
This option is correct

2. What are the major gates in a standard LSTM (long short-term memory) cell? **1 / 1 point**

- A standard LSTM cell includes three gates: the forget gate to forget irrelevant information, the input gate to remember relevant information, and the update gate to update new information.
- A standard LSTM cell includes three gates: the input gate to input information, the hidden gate to remember information, and the output gate to output information.
- A standard LSTM cell includes two gates: the input gate to input information and the output gate to output information.
- A standard LSTM cell includes two gates: the remember gate to remember relevant information and the forget gate to forget irrelevant information.

Correct
This option is correct

3. What is the coding in Keras to build the hidden layer of a GRU (gated recurrent unit) model? **1 / 1 point**

- Dense(N_CLASSES, activation="softmax")
- Lambda(lambda x: tf.reduce_mean(x, axis=1))
- gru_model = build_gru_model(embed_dim=EMBED_DIM)
- GRU(units)

Correct
This option is correct