

- Hide menu
- Introduction
- Data extraction, analysis, and preparation
- Design Decisions
- Designing an Architecture from Scratch
- Introducing Vertex AI
- Video: Using Vertex AI8 min
- Video: Lab Introduction: Structured data prediction22 sec
- Video: Getting Started with Google Cloud and Qwiklabs4 min
- Graded App Item: Lab: Structured data prediction using Vertex AIStarted
- Quiz: Architecting production ML systemsSubmitted
- Reading: Architecting production ML systems10 min

🎉 Congratulations! You passed!

Grade received 100%

Latest Submission

To pass 80% or higher

Go to next item

Architecting production ML systems

COACH

1. What percent of system code is typically written before taking the quiz? I'm here to help.1 / 1 point

5%

25%

50%

90%

Correct

Correct!

Help me practice

Let's chat

Submit your assignment

Due Sep 6, 11:59 PM IST

Try again

Receive grade

To Pass: 80% or higher

Your grade

100%

View Feedback

We keep your highest score

2. Match the three types of data ingest with an appropriate source of training data.1 / 1 point

Streaming batch (Dataflow), structured batch (BigQuery), unstructured batch (App Engine)

Streaming (Pub/Sub), structured batch (BigQuery), unstructured batch (Cloud Storage)

Streaming (BigQuery), structured batch (Pub/Sub), unstructured batch (Cloud Storage)

Correct

Correct! On Google Cloud, the three types of data ingestion map to three different products. If you are ingesting streaming data, use Pub/Sub. If you are ingesting structured data directly into your ML model, use BigQuery, and if you are transforming data from training so that you can train on it later, read from Cloud Storage.

3. What is the responsibility of model evaluation and validation components?1 / 1 point

To ensure that the models are good after moving them into a production/staging environment.

To ensure that the models are good before moving them into a production/staging environment.

To ensure that the models are not good after moving them into a staging environment.

To ensure that the models are not good before moving them into a staging environment.

Correct

Correct!

4. Which type of training do you use if your data set doesn't change over time?1 / 1 point

Static training

Dynamic training

Online training

Real-time training

Correct

Correct!

5. Which type of logging should be enabled in the online prediction that logs the stderr and stdout streams from your prediction nodes to Cloud Logging and can be useful for debugging?1 / 1 point

Access logging

Request-response logging

Container logging

Cloud logging

Correct

Correct!

6. Vertex AI has a unified data preparation tool that supports image, tabular, text, and video content. Where are uploaded datasets stored in Vertex AI?1 / 1 point

A Google Cloud Storage bucket that acts as an input for both AutoML and custom training jobs.

A Google Cloud database that acts as an input for both AutoML and custom training jobs.

A Google Cloud Storage bucket that acts as an output for both AutoML, custom training jobs, serialized training jobs.

A Google Cloud database that acts as an output for both AutoML and custom training jobs.

Correct

Correct!

7. In the featurstore, the timestamps are an attribute of the feature values, not a separate resource type.1 / 1 point

True

False

Correct

Correct!

8. When you use the data to train a model, Vertex AI examines the source data type and feature values and infers how it will use that feature in model training. This is called the _____ for that feature.1 / 1 point

Duplication

Transmutation

Transformation

Translation

Correct

Correct!

https://www.coursera.org/learn/gcp-production-ml-systems/exam/09bWd/architecting-production-ml-systems/attempt/redirectToCover?due

1/1