

```
monitor.py
7 # Prepare your batch of prompts
8
9 prompts = ["who is the oldest president", " what is the name of a famous cave"]
10
11
12 # Record the start time
13 start_time = time.time()
14
15 # Make the API call
16 responses = openai.Completion.create(
17     engine="text-davinci-002",
18     prompt=prompts,
19     max_tokens=50,
20     n=len(prompts),
21     stop=None,
22     temperature=0.7
23 )
24
25 # Record the end time
26 end_time = time.time()
27
28 # Calculate response time
29 response_time = end_time - start_time
30
31 print(f"Response time: {response_time} seconds")
```

Guide

Collapse

<>⚙️☰

Monitoring and Managing Batch Processing Performance with the C...

response. Here's an example of how to extract usage information from the API response:

```
usage = responses['usage']
total_tokens = usage['total_tokens']
prompt_tokens = usage['prompt_tokens']
completion_tokens = usage['completion_tokens']

print(f"API Usage:")
print(f" Total tokens: {total_tokens}")
print(f" Prompt tokens: {prompt_tokens}")
print(f" Completion tokens: {completion_tokens}")
```

TRY IT!

Response time: 0.8815827369689941 seconds

Compare these values with the rate limits and quotas imposed by OpenAI to ensure your implementation stays within the allowed limits.

4. Error Rate:

To monitor the error rate, you need to track the number of errors or failed requests and calculate the ratio of errors to total requests. Here's an example of how to handle errors and calculate the error rate:

```
error_count = 0

try:
    responses = openai.Completion.create(
        engine="text-davinci-002",
        prompt=prompts,
        max_tokens=50,
        n=len(prompts),
        stop=None,
        temperature=0.7
    )
except Exception as e:
    error_count += 1
    print(f"Error encountered: {e}")

error_rate = error_count / num_requests

print(f"Error rate: {error_rate}")
```

TRY IT!

Response time: 0.8635814189910889 seconds

This example assumes a single batch request. In a real-world scenario, you would track errors over multiple API calls and calculate the error rate accordingly.

By monitoring key performance metrics such as response time, throughput, API usage, and error rate, you can ensure that your batch processing implementation with the ChatGPT API is efficient and reliable. Regularly tracking these metrics will help you identify and address potential issues or bottlenecks, resulting in a better user experience.

When monitoring the error rate, it's important to track the number of errors or failed requests and calculate the ratio of errors to total requests.

☒ True

☐ False

It is essential to track the number of errors or failed requests and calculate the ratio of errors to total requests. This helps ensure that the batch processing implementation is efficient and reliable.

Check It!

Next

1/1