

torch.autograd.grad

```
torch.autograd.grad(outputs, inputs, grad_outputs=None, retain_graph=None, create_graph=False, only_inputs=True, allow_unused=None, is_grads_batched=False, materialize_grads=False) [SOURCE]
```



Compute and return the sum of gradients of outputs with respect to the inputs.

`grad_outputs` should be a sequence of length matching `output` containing the “vector” in vector-Jacobian product, usually the pre-computed gradients w.r.t. each of the outputs. If an output doesn’t require `grad`, then the gradient can be `None`).

• NOTE

If you run any forward ops, create `grad_outputs`, and/or call `grad` in a user-specified CUDA stream context, see [Stream semantics of backward passes](#).

• NOTE

`only_inputs` argument is deprecated and is ignored now (defaults to `True`). To accumulate gradient for other parts of the graph, please use `torch.autograd.backward`.

Parameters

- **outputs** (sequence of *Tensor* or *GradientEdge*) – outputs of the differentiated function.
- **inputs** (sequence of *Tensor* or *GradientEdge*) – Inputs w.r.t. which the gradient will be returned (and not accumulated into `.grad`).
- **grad_outputs** (sequence of *Tensor*) – The “vector” in the vector-Jacobian product. Usually gradients w.r.t. each output. `None` values can be specified for scalar Tensors or ones that don’t require `grad`. If a `None` value would be acceptable for all `grad_tensors`, then this argument is optional. Default: `None`.
- **retain_graph** (*bool*, optional) – If `False`, the graph used to compute the grad will be freed. Note that in nearly all cases setting this option to `True` is not needed and often can be worked around in a much more efficient way. Defaults to the value of `create_graph`.
- **create_graph** (*bool*, optional) – If `True`, graph of the derivative will be constructed, allowing to compute higher order derivative products. Default: `False`.
- **allow_unused** (Optional[*bool*], optional) – If `False`, specifying inputs that were not used when computing outputs (and therefore their grad is always zero) is an error. Defaults to the value of `materialize_grads`.
- **is_grads_batched** (*bool*, optional) – If `True`, the first dimension of each tensor in `grad_outputs` will be interpreted as the batch dimension. Instead of computing a single vector-Jacobian product, we compute a batch of vector-Jacobian products for each “vector” in the batch. We use the `vmap` prototype feature as the backend to vectorize calls to the autograd engine so that this computation can be performed in a single call. This should lead to performance improvements when compared to manually looping and performing backward multiple times. Note that due to this feature being experimental, there may be performance cliffs. Please use `torch._C._debug_only_display_vmap_fallback_warnings(True)` to show any performance warnings and file an issue on github if warnings exist for your use case. Defaults to `False`.
- **materialize_grads** (*bool*, optional) – If `True`, set the gradient for unused inputs to zero instead of `None`. This is useful when computing higher-order derivatives. If `materialize_grads` is `True` and `allow_unused` is `False`, an error will be raised. Defaults to `False`.

Return type

*Tuple[*Tensor*, ...]*

Docs

Access comprehensive developer documentation for PyTorch

[View Docs >](#)

Tutorials

Get in-depth tutorials for beginners and advanced developers

[View Tutorials >](#)

Resources

Find development resources and get your questions answered

[View Resources >](#)



PyTorch

- [Get Started](#)
- [Features](#)
- [Ecosystem](#)
- [Blog](#)
- [Contributing](#)

Resources

- [Tutorials](#)
- [Docs](#)
- [Discuss](#)
- [Github Issues](#)
- [Brand Guidelines](#)

Stay up to date

- [Facebook](#)
- [Twitter](#)
- [YouTube](#)
- [LinkedIn](#)

PyTorch Podcasts

- [Spotify](#)
- [Apple](#)
- [Google](#)
- [Amazon](#)