

Hide menu

- Introduction
- Incorporating Context
- Introduction
- Quiz: Reinforcement Learning
- Quiz: Reinforcement Learning
- Submitted
- Reading: Reinforcement Learning
- 10 min
- Lab
- Summary

Reinforcement Learning

Submit your assignment

Due Mar 24, 11:59 PM IST

Attempts 3 every 8 hours

Receive grade

To Pass 75% or higher

Like

Dislike

Report an issue

Congratulations! You passed!

Grade received 100%

Latest Submission Grade 100%

To pass 75% or higher

Go to next item

Try again

1. In which scenarios is reinforcement learning preferable over supervised learning?

1 / 1 point

- ☐ When you have predictive modeling problems with an offline static dataset.
- ☒ When you have optimization or control problems where simulation trial and error is possible.
- ☐ When you have optimization or control problems with scarce data points and trial and error is impossible.
- ☐ When you have predictive modeling with scarce data and a differentiable metric to be optimized.

Correct

This is correct.

2. Which of the following is not a motivating rationale to use replay buffers?

1 / 1 point

- ☐ For de-correlating experience trajectories.
- ☐ For repeating rare experiences.
- ☐ For achieving data efficiency.
- ☒ To keep the model policy well aligned with the newest experience

Correct

This answer is correct.

3.

1 / 1 point

Which of the following would make for suitable good value functions? (Choose 2 correct Answers)

- ☐ Scenario: You want to train an agent to win a race. The reward is the total time taken to run the race.
- ☒ Scenario: You have a movie recommender system. The reward is the count of clicks.

Correct

This answer is correct!

☐ Scenario: You have a tennis video game. The reward is the negated value of the final score.

☒ Scenario: You want to train an agent to win a race. The reward is the negative value of the total time taken to run the race.

Correct

This answer is correct!

4.

1 / 1 point

In which of the following use cases is it recommended to go with a contextual bandit system?

- ☐ Forecasting demand for various products in a supermarket in a given time horizon.
- ☐ Training a robot to walk.
- ☐ Training two agents to cooperate with each other to win a multi-agent strategy game.
- ☒ Tailoring the results of a search engine to a specific user.

Correct

This answer is correct!

5.

1 / 1 point

You would like to train an agent to drive a car. The action space consists of the following variables: the acceleration (between 0 and 300), the angular degree of turn or tilt (between 0 and 180 degrees), and the direction (either forward or reverse). Select the three algorithms which are appropriate. Credit is given for selecting the correct three.

☒ Proximal Policy Optimization (PPO)

Correct

This answer is correct! All three correct answers must be selected to receive credit for the question.

☒ Deep deterministic policy gradient (DDPG)

Correct

This answer is correct! All three correct answers must be selected to receive credit for the question.

☐ Deep Q Networks

☒ REINFORCE

Correct

This answer is correct! All three correct answers must be selected to receive credit for the question.

Your grade

100%

View Feedback

We keep your highest score

