

A community index of packages for Apache Spark (<http://spark.apache.org>).

173 packages

spark-avro (/package/databricks/spark-avro)

Integration utilities for using Spark with Apache Avro data

from: @databricks (/user/databricks) / owner: @pwendell (/user/pwendell) / Latest release: 2.0.1-s_2.10 (2015-09-08) / Apache-2.0 / ★★★★★ (28)

4 sql 3 input 3 avro

spark-redshift (/package/databricks/spark-redshift)

Spark and Redshift integration

from: @databricks (/user/databricks) / owner: @pwendell (/user/pwendell) / Latest release: 0.6.0 (2016-01-06) / Apache-2.0 / ★★★★★ (3)

1 input 1 sql 1 redshift

kafka-spark-consumer (/package/dibbhatt/kafka-spark-consumer)

Receiver Based Low Level Kafka-Spark Consumer with built-in Back-Pressure Controller

@dibbhatt (/user/dibbhatt) / Latest release: 1.0.6 (2016-01-07) / Apache-2.0 / ★★★★★ (6)

4 streaming 3 kafka

thunder (/package/freeman-lab/thunder)

Large-scale neural data analysis with Spark

@freeman-lab (/user/freeman-lab) / Latest release: 0.4.1 (2014-11-27) / Apache-2.0 / ★★★★★ (4)

3 neuroscience 2 python 2 machine learning

spark_azure (/package/sigmoidanalytics/spark_azure)

Spark launch script for Microsoft Azure

@sigmoidanalytics (/user/sigmoidanalytics) / No release yet / ★★★★★★ (19)

1 Azure 1 spark 1 Microsoft

spork (/package/sigmoidanalytics/spork)

Pig on Apache Spark

@sigmoidanalytics (/user/sigmoidanalytics) / No release yet / ★★★★★★ (19)

1 streaming 1 spark 1 pig

spark_gce (/package/sigmoidanalytics/spark_gce)

Spark launch script for Google Compute Engine

@sigmoidanalytics (/user/sigmoidanalytics) / No release yet / ★★★★★★ (18)

1 Google Compute Engine 1 spark 1 GCE

spark-jobserver (/package/spark-jobserver/spark-jobserver)

REST job server for Spark

from: @spark-jobserver (/user/spark-jobserver) / owner: @velvia (/user/velvia) / No release yet / ★★★★★★ (12)

1 application 1 REST 1 Mesos

GMM (/package/FlytxtRnD/GMM)

Gaussian Mixture Model Implementation in Pyspark

@FlytxtRnD (/user/FlytxtRnD) / Latest release: 0.1 (2015-04-07) / EPL-1.0 / ★★★★★★ (15)

1 python

1 mllib

spark-csv (/package/databricks/spark-csv)

Spark SQL CSV data source

from: @databricks (/user/databricks) / owner: @falaki (/user/falaki) / Latest release: 1.3.0-s_2.10 (2015-11-20) / Apache-2.0 / ★★★★★ (7)

1 sql

1 DataSource

1 SparkSQL

spark-indexedrdd (/package/amplab/spark-indexedrdd)

An efficient updatable key-value store for Apache Spark

from: @amplab (/user/amplab) / owner: @ankurdave (/user/ankurdave) / Latest release: 0.3 (2015-09-10) / Apache-2.0 / ★★★★★ (1)

1 core

1 kv

spark-perf (/package/databricks/spark-perf)

Performance tests for Spark

from: @databricks (/user/databricks) / owner: @pwendell (/user/pwendell) / No release yet / ★★★★★ (1)

zeppelin (/package/NFLabs/zeppelin)

Zeppelin, a web-based notebook that enables interactive data analytics.

from: @NFLabs (/user/NFLabs) / owner: @Leemoonsoo (/user/Leemoonsoo) / No release yet / ★★★★★ (2)

1 Applications

1 notebook

1 interactive

spark-hbase (/package/haosdent/spark-hbase)

Integration utilities for using Spark with Apache HBase data

@haosdent (/user/haosdent) / No release yet / ★★★★★ (1)

1 hbase

spark_hbase (/package/GenTang/spark_hbase)

The example in Scala of reading data saved in hbase by Spark and the example of converter for python

@GenTang (/user/GenTang) / No release yet / ★★★★★ (3)

1 python

1 hbase

sparkling (/package/gorillalabs/sparkling)

A Clojure library for Apache Spark: fast, fully-features, and developer friendly

from: @gorillalabs (/user/gorillalabs) / owner: @chrisbetz (/user/chrisbetz) / Latest release: 1.0.0 (2014-12-31) / EPL-1.0 / ★★★★★ (3)

2 clojure

1 API

deep-spark (/package/Stratio/deep-spark)

Connecting Apache Spark with different data stores

from: @Stratio (/user/Stratio) / owner: @albertope (/user/albertope) / Latest release: 0.7.0-RC1 (2015-01-14) / Apache-2.0 / ★★★★★ (20)

6 database

6 mongo

6 cassandra

streaming-cep-engine (/package/Stratio/streaming-cep-engine)

Streaming CEP Engine Powered by Spark Streaming & Siddhi

from: @Stratio (/user/Stratio) / owner: @dmorales-stratio (/user/dmorales-stratio) / Latest release: 0.6.2 (2015-01-14) / Apache-2.0 / ★★★★★ (19)

5 spark streaming

5 cep

4 complex event processing

generalized-kmeans-clustering (/package/derrickburns/generalized-kmeans-clustering)

This project generalizes the Spark MLLIB K-Means clusterer to support arbitrary distance functions

@derrickburns (/user/derrickburns) / No release yet / ★★★★★ (2)

1 clustering 1 mllib 1 machine learning

sparkling-water (/package/h2oai/sparkling-water)

Sparkling Water provides H2O algorithms inside Spark cluster

from: @h2oai (/user/h2oai) / owner: @mmalohlava (/user/mmalohlava) / Latest release: 1.4.3 (2015-07-06) / Apache-2.0 / ★★★★★ (2)

1 h2o 1 algorithms 1 machine learning

spark-ml-streaming (/package/freeman-lab/spark-ml-streaming)

Visualize streaming machine learning in Spark

@freeman-lab (/user/freeman-lab) / No release yet / ★★★★★ (1)

1 streaming 1 machine learning 1 visualization

spark-hbase-connector (/package/nerdammer/spark-hbase-connector)

Connect Spark to HBase for reading and writing data with ease

from: @nerdammer (/user/nerdammer) / owner: @nibbio84 (/user/nibbio84) / Latest release: 1.0.0 (2015-05-11) / Apache-2.0 / ★★★★★ (1)

1 streaming 1 hbase 1 library

spark-testing-base (/package/holdenk/spark-testing-base)

Base classes to use when writing tests with Spark

@holdenk (/user/holdenk) / Latest release: 1.5.1_0.2.1 (2015-11-05) / Apache-2.0 / ★★★★★ (5)

3 testing

1 streaming

1 tools

spark-mongodb (/package/Stratio/spark-mongodb)

MongoDB data source for Spark SQL

from: @Stratio (/user/Stratio) / owner: @albertope (/user/albertope) / Latest release: 0.10.1 (2015-11-18) / Apache-2.0 / ★★★★★ (13)

5 MongoDB

5 Spark SQL

1 sql

pyspark-cassandra (/package/TargetHolding/pyspark-cassandra)

PySpark Cassandra brings back the fun in working with Cassandra data in PySpark.

from: @TargetHolding (/user/TargetHolding) / owner: @frensjan (/user/frensjan) / Latest release: 0.2.3 (2016-01-10) / Apache-2.0 / ★★★★★ (1)

1 python

1 spark

1 sql

demo-scala-python (/package/brkyvz/demo-scala-python)

A Spark Package Template

@brkyvz (/user/brkyvz) / Latest release: 1.0.1-s_2.11 (2015-08-04) / Apache-2.0 / ★★★★★ (1)

1 python

1 demo

1 template

spark-cassandra-connector (/package/datastax/spark-cassandra-connector)

Connects Spark to Cassandra

from: @datastax (/user/datastax) / owner: @pkolaczka (/user/pkolaczka) / Latest release: 1.5.0-RC1-s_2.11 (2016-01-07) / Apache-2.0 / ★★★★★★ (10)

3 spark

3 cassandra

2 nosql

spark-package-cmd-tool (/package/databricks/spark-package-cmd-tool)

A command line tool for Spark packages

from: @databricks (/user/databricks) / owner: @brkyvz (/user/brkyvz) / Latest release: 0.3.0 (2015-03-17) / Apache-2.0 / ★★★★★★ (1)

1 tools

meetup-stream (/package/actions/meetup-stream)

Spark Streaming, Machine Learning and meetup.com streaming API.

from: @actions (/user/actions) / owner: @fountainheadpro (/user/fountainheadpro) / No release yet / ★★★★★★ (1)

1 ml

1 example

1 streaming

spark-sequoiadb (/package/SequoiaDB/spark-sequoiadb)

Spark connector for SequoiaDB

from: @SequoiaDB (/user/SequoiaDB) / owner: @wangzhonnew (/user/wangzhonnew) / Latest release: 1.12-s_2.11 (2015-03-30) / Apache-2.0 / ★★★★★ (1)

1 sequoiadb 1 nosql 1 sql

spark-notebook (/package/andypetrella/spark-notebook)

Use Apache Spark straight from the Browser

@andypetrella (/user/andypetrella) / Latest release: v0.4.0 (2015-03-29) / Apache-2.0 / ★★★★★ (1)

1 notebook 1 charts 1 interactive

spark-cloudant (/package/cloudant/spark-cloudant)

Spark SQL IBM Cloudant External Datasource

from: @cloudant (/user/cloudant) / owner: @davis (/user/davis) / No release yet / ★★★★★ (1)

1 data source 1 sql

MLlib-dropout (/package/rakeshchhalasani/MLlib-dropout)

Package adding dropout regularization to Apache Spark MLlib project

@rakeshchhalasani (/user/rakeshchhalasani) / No release yet / ★★★★★ (1)

1 machine learning 1 mllib 1 scala

spark-infotheoretic-feature-selection (/package/sramirez/spark-infotheoretic-feature-selection)

Feature Selection framework based on Information Theory that includes: mRMR, InfoGain, JMI and other commonly used FS filters.

@sramirez (/user/sramirez) / Latest release: 1.2 (2015-12-18) / Apache-2.0 / ★★★★★ (5)

2 feature-selection 2 mllib 2 machine learning

spark-MDLP-discretization (/package/sramirez/spark-MDLP-discretization)

Spark implementation of Fayyad's discretizer based on Minimum Description Length Principle (MDLP)

@sramirez (/user/sramirez) / Latest release: 1.1 (2015-12-04) / Apache-2.0 / ★★★★★ (5)

2 discretization 2 mllib 1 machine learning

couchbase-spark-connector (/package/couchbaselabs/couchbase-spark-connector)

The Official Couchbase Server <-> Apache Spark Connector

from: @couchbaselabs (/user/couchbaselabs) / owner: @daschl (/user/daschl) / Latest release: 1.0.0 (2015-10-20) / Apache-2.0 / ★★★★★ (1)

1 streaming 1 library 1 sql

spark-pmml-exporter-validator (/package/selvinsource/spark-pmml-exporter-validator)

Using JPMML Evaluator to validate the PMML models exported from Spark

@selvinsource (/user/selvinsource) / No release yet / ★★★★★ (1)

1 ml 1 mllib 1 machine learning

spark-sas7bdat (/package/saurfang/spark-sas7bdat)

Splittable SAS (.sas7bdat) Input Format for Hadoop and Spark SQL

@saurfang (/user/saurfang) / Latest release: 1.1.4-s_2.10 (2015-10-07) / GPL-3.0 / ★★★★★★ (1)

[1 sas](#)[1 tools](#)[1 sql](#)

succinct (/package/amplab/succinct)

Spark and Spark SQL integration for Succinct

from: @amplab (/user/amplab) / owner: @anuragkh (/user/anuragkh) / Latest release: 0.1.6 (2015-11-23) / Apache-2.0 / ★★★★★★ (1)

[1 application](#)[1 data source](#)

sparkit-learn (/package/lensacom/sparkit-learn)

PySpark + Scikit-learn = Sparkit-learn

from: @lensacom (/user/lensacom) / owner: @kszucs (/user/kszucs) / No release yet / ★★★★★★ (2)

[1 python](#)[1 scikit-learn](#)[1 machine learning](#)

RabbitMQ-Receiver (/package/Stratio/RabbitMQ-Receiver)

RabbitMQ Spark Streaming receiver

from: @Stratio (/user/Stratio) / owner: @danielcsant (/user/danielcsant) / Latest release: 0.1.0-RELEASE (2015-05-21) / Apache-2.0 / ★★★★★★ (9)

[3 streaming](#)

streaming-matrix-factorization (/package/brkyvz/streaming-matrix-factorization)

Streaming Recommendation Engine using matrix factorization with user and

product bias

@brkyvz (/user/brkyvz) / Latest release: 0.1.0 (2015-05-26) / Apache-2.0 / ★★★★★★ (1)

1 streaming 1 ml 1 machine learning

zen (/package/cloudml/zen)

Zen aims to provide the largest scale and the most efficient machine learning platform on top of Spark, including but not limited to logistic regression, latent dirichlet allocation, factorization machines and DNN.

from: @cloudml (/user/cloudml) / owner: @witgo (/user/witgo) / No release yet / ★★★★★★ (2)

1 ml 1 mllib 1 machine learning

spark-es (/package/SHSE/spark-es)

ElasticSearch integration for Apache Spark

@SHSE (/user/SHSE) / Latest release: 1.0.4 (2015-08-18) / Apache-2.0 / ★★★★★★ (1)

1 analytics 1 search 1 elasticsearch

aerosolve (/package/airbnb/aerosolve)

A machine learning package built for humans.

from: @airbnb (/user/airbnb) / owner: @hectorgon (/user/hectorgon) / No release yet / ★★★★★★ (1)

1 machine learning

spark-streamingsql (/package/Intel-bigdata/spark-streamingsql)

Manipulate Apache Spark Streaming by SQL

from: @Intel-bigdata (/user/Intel-bigdata) / owner: @GraceH (/user/GraceH) / No release yet / ★★★★★★ (1)

1 streaming 1 sql

DDF (/package/ddf-project/DDF)

Distributed DataFrame: Productivity = Power x Simplicity For Scientists & Engineers, on any Data/Compute Engine

from: @ddf-project (/user/ddf-project) / owner: @dabaitu (/user/dabaitu) / No release yet / ★★★★★★ (11)

3 API

2 tools

2 machine learning

spark-datetime (/package/SparklineData/spark-datetime)

A library for exposing dateTime functions from the joda library as SQL functions. With a dsl to build dateTime catalyst expressions.

from: @SparklineData (/user/SparklineData) / owner: @hbutani (/user/hbutani) / Latest release: 0.0.2 (2015-10-29) / Apache-2.0 / ★★★★★★ (1)

1 spark

1 sql

1 dateTime

adam (/package/bigdatagenomics/adam)

A genomics processing engine and specialized file format built using Apache Avro, Apache Spark and Parquet. Apache 2 licensed.

from: @bigdatagenomics (/user/bigdatagenomics) / owner: @fnofhaft (/user/fnofhaft) / No release yet / ★★★★★★ (1)

elasticsearch-hadoop (/package/elastic/elasticsearch-hadoop)

Official integration between Apache Spark and Elasticsearch real-time search and analytics

from: @elastic (/user/elastic) / owner: @costin (/user/costin) / Latest release: 2.2.0-rc1 (2016-01-09) / Apache-2.0 / ★★★★★★ (3)

1 search

1 elasticsearch

1 sql

modelmatrix

(/package/collectivemedia/modelmatrix)

Alternative to Spark machine learning pipeline feature extractors, focused on building sparse feature vectors.

from: @collectivemedia (/user/collectivemedia) / owner: @ezhulenev (/user/ezhulenev) / No release yet / ★★★★★★ (1)

2 feature extraction

2 machine learning

spark-knn-graphs (/package/tdebatty/spark-knn-graphs)

Spark algorithms for building k-nn graphs

@tdebatty (/user/tdebatty) / Latest release: 0.7 (2015-12-13) / LGPL-3.0 / ★★★★★★ (1)

1 graph

1 machine learning

Spark-SQL-on-HBase (/package/Huawei-Spark/Spark-SQL-on-HBase)

Native, optimized access to HBase Data through Spark SQL/Dataframe Interfaces

from: @Huawei-Spark (/user/Huawei-Spark) / owner: @yzhou2001 (/user/yzhou2001) / Latest release: 1.0.0 (2015-07-17) / Apache-2.0 / ★★★★★★ (1)

magellan (/package/harsha2010/magellan)

Geo Spatial Data Analytics on Spark

@harsha2010 (/user/harsha2010) / Latest release: 1.0.3-s_2.10 (2015-08-20) / Apache-2.0 / ★★★★★★ (1)

2 geospatial

2 data source

2 sql

SparkTwitterAnalysis

(/package/prabeesh/SparkTwitterAnalysis)

An Apache Spark standalone application using the Spark API in Scala. The application uses Simple Build(SBT) for building the project.

@prabeesh (/user/prabeesh) / Latest release: 0.1.0 (2015-08-04) / Apache-2.0 / ★★★★★ (1)

pyspark-notebook (/package/prabeesh/pyspark-notebook)

Pyspark Notebook With Docker

@prabeesh (/user/prabeesh) / Latest release: 0.1.0 (2015-08-04) / Apache-2.0 / ★★★★★ (1)

spark-streaming-gnip (/package/knoldus/spark-streaming-gnip)

An Apache Spark utility for pulling Tweets from Gnip's PowerTrack in realtime

from: @knoldus (/user/knoldus) / owner: @vhazrati (/user/vhazrati) / No release yet / ★★★★★ (1)

spark-salesforce (/package/springml/spark-salesforce)

Spark Salesforce Wave Connector

@springml (/user/springml) / Latest release: 1.0.4 (2016-01-11) / Apache-2.0 / ★★★★★ (2)

1 salesforce

1 data source

spark-centrality (/package/webgeist/spark-centrality)

Library for computing centrality for graph nodes

@webgeist (/user/webgeist) / Latest release: 0.11 (2015-08-09) / LGPL-3.0 / ★★★★★ (1)

1 graph

spark-corenlp (/package/databricks/spark-corenlp)

a Stanford CoreNLP wrapper for Spark ML pipeline API

from: @databricks (/user/databricks) / owner: @mengxr (/user/mengxr) / No release yet /

★★★★★ (1)

1 NLP

1 machine learning

spark-tfocs (/package/databricks/spark-tfocs)

A Spark port of TFOCS: Templates for First-Order Conic Solvers
(cvxr.com/tfocs)

from: @databricks (/user/databricks) / owner: @mengxr (/user/mengxr) / No release yet /

★★★★★ (1)

1 machine learning

1 optimization

1 convex

aliyun-spark-sdk (/package/aliyun/aliyun-spark-sdk)

Spark on Aliyun, supporting interactions with Aliyun's base services.

from: @aliyun (/user/aliyun) / owner: @uncleGen (/user/uncleGen) / No release yet / ★★★★★ (

1)

1 streaming

1 data source

spark-druid-olap (/package/SparklineData/spark-druid-olap)

Spark Druid Package

from: @SparklineData (/user/SparklineData) / owner: @hbutani (/user/hbutani) / Latest release:

0.0.1 (2015-09-01) / Apache-2.0 / ★★★★★ (2)

TopicModeling (/package/intel-analytics/TopicModeling)

Distribtued Topic Modeling on Apache Spark

from: @intel-analytics (/user/intel-analytics) / owner: @jason-dai (/user/jason-dai) / No release yet / ★★★★★★ (1)

1 graph 1 LDA 1 machine learning

lazy-linalg (/package/brkyvz/lazy-linalg)

Linear algebra operators for Apache Spark MLlib's linalg package

@brkyvz (/user/brkyvz) / Latest release: 0.1.0 (2015-09-09) / Apache-2.0 / ★★★★★★ (1)

1 linear algebra 1 lazy 1 machine learning

spark-stochastic-outlier-selection (/package/rug-ds-lab/spark-stochastic-outlier-selection)

Implementation of Stochastic Outlier Selection (SOS) which is an unsupervised outlier selection algorithm.

from: @rug-ds-lab (/user/rug-ds-lab) / owner: @Fokko (/user/Fokko) / No release yet / ★★★★★★ (1)

1 outlier detection

spark-stochastic-outlier-selection (/package/Fokko/spark-stochastic-outlier-selection)

Implementation of Stochastic Outlier Selection (SOS) which is an unsupervised outlier selection algorithm.

@Fokko (/user/Fokko) / Latest release: 0.1.0 (2015-09-11) / Apache-2.0 / ★★★★★★ (1)

1 outlier detection

mCNN (/package/hhbyyh/mCNN)

A Spark machine learning package containing the implementation of classical Convolutional Neural Network

@hhbyyh (/user/hhbyyh) / No release yet / ★★★★★ (1)

1 app

1 machi

pipeline (/package/fluxcapacitor/pipeline)

Docker-based, End-to-End, Real-time, Advanced Analytics Big Data Reference Pipeline using Spark, Spark SQL, Spark Streaming, ML, MLlib, GraphX, Kafka, Cassandra, Redis, Apache Zeppelin, Spark-Notebook, iPython/Jupyter Notebook, Tableau, H2O Flow, Tachyon,

from: @fluxcapacitor (/user/fluxcapacitor) / owner: @cfregly (/user/cfregly) / No release yet / ★★★★★ (1)

1 streaming

1 kafka

1 machine learning

SpatialSpark (/package/syoummer/SpatialSpark)

Big Spatial Data Processing using Spark

@syoummer (/user/syoummer) / Latest release: 1.0 (2015-10-08) / Apache-2.0 / ★★★★★ (1)

1 geospatial

1 WKT

1 spatial

spark-cc (/package/SherlockYang/spark-cc)

Library for computing clustering coefficient

@SherlockYang (/user/SherlockYang) / Latest release: 0.1 (2015-10-22) / LGPL-3.0 / ★★★★★ (1)

1 graph

hue (/package/cloudera/hue)

Livy, a REST Spark Server for submitting jobs and code snippets

from: @cloudera (/user/cloudera) / owner: @romainr (/user/romainr) / No release yet /

★★★★★ (2)

1 application 1 REST 1 interactive

spark-xml (/package/HyukjinKwon/spark-xml)

XML data source for Spark SQL and DataFrames

@HyukjinKwon (/user/HyukjinKwon) / Latest release: 0.1.1-s_2.10 (2015-11-19) / Apache-2.0 /

★★★★★ (1)

1 sql 1 DataSource 1 SparkSQL

graphx-diameter (/package/Cecca/graphx-diameter)

A spark package to approximate the diameter of large graphs

@Cecca (/user/Cecca) / Latest release: 0.1.0-s_2.11 (2015-12-03) / Apache-2.0 / ★★★★★ (1)

1 graph

spark-ryft-connector (/package/getryft/spark-ryft-connector)

Spark connector for Ryft ONE

from: @getryft (/user/getryft) / owner: @demon-xxi (/user/demon-xxi) / Latest release: 0.4.0 (2016-01-06) / other license / ★★★★★ (1)

1 search 1 pyspark 1 scala

spark-sftp (/package/springml/spark-sftp)

Spark connector for SFTP

@springml (/user/springml) / Latest release: 1.0.0 (2016-01-11) / Apache-2.0 / ★★★★★ (1)

1 data source

flintrock (/package/nchammas/flintrock)

A command-line tool for launching Apache Spark clusters.

@nchammas (/user/nchammas) / No release yet / ★★★★★★ (1)

1 tools

1 ec2

1 deployment

spark-als (/package/mengxr/spark-als)

Another, hopefully better, implementation of ALS on Spark

@mengxr (/user/mengxr) / Latest release: 0.1.0 (2014-11-27) / Apache-2.0 / ★★★★★★ (1)

1 ml

1 mllib

1 recommendation

pyspark-csv (/package/seahboonsiew/pyspark-csv)

An external PySpark module that works like R's read.csv or Panda's read_csv, with automatic type inference and null value handling. Parses csv data into SchemaRDD. No installation required, simply include pyspark_csv.py via SparkContext.

@seahboonsiew (/user/seahboonsiew) / No release yet / ★★★★★★ (1)

2 python

2 csv

1 sql

sbt-spark-package (/package/databricks/sbt-spark-package)

Sbt plugin for Spark packages

from: @databricks (/user/databricks) / owner: @brkyvz (/user/brkyvz) / Latest release: 0.2.1 (2015-04-02) / Apache-2.0 / ★★★★★★ (2)

1 tools

1 sbt

pyspark-elastic (/package/TargetHolding/pyspark-elastic)

Pyspark support for Elastic Search

from: @TargetHolding (/user/TargetHolding) / owner: @frensjan (/user/frensjan) / Latest release: 0.3.1 (2015-12-15) / Apache-2.0 / ★★★★★ (1)

[1 python](#)[1 spark](#)[1 database](#)

spark-tsne (/package/saurfang/spark-tsne)

Distributed t-SNE via Apache Spark

@saurfang (/user/saurfang) / No release yet / ★★★★★ (1)

[1 machine learning](#)

DistML (/package/intel-machine-learning/DistML)

DistML provide a supplement to mllib to support model-parallel on Spark

from: @intel-machine-learning (/user/intel-machine-learning) / owner: @Yunlong-He (/user/Yunlong-He) / No release yet / ★★★★★ (1)

[1 parameter server](#)[1 machine learning](#)

dl4j-spark-ml (/package/deeplearning4j/dl4j-spark-ml)

Deep Learning for Spark ML

from: @deeplearning4j (/user/deeplearning4j) / owner: @agibsonccc (/user/agibsonccc) / Latest release: 0.4-rc3.4 (2015-10-02) / Apache-2.0 / ★★★★★ (1)

[1 Spark ML](#)[1 machine learning](#)

spark-google-spreadsheets (/package/potix2/spark-google-spreadsheets)

Google Spreadsheets datasource for SparkSQL and DataFrames

@potix2 (/user/potix2) / Latest release: 0.2.0-s_2.11 (2015-12-29) / Apache-2.0 / ★★★★★ (1)

1 sql

1 data source

1 scala

dllib (/package/Lewuathe/dllib)

dllib is a deep learning module running on Apache Spark.

@Lewuathe (/user/Lewuathe) / Latest release: 0.0.7 (2015-12-30) / Apache-2.0 / ★★★★★★ (1)

1 deep learning

1 machine learning

mllib-grid-search (/package/spark-ml/mllib-grid-search)

An example project for doing grid search in MLlib

from: @spark-ml (/user/spark-ml) / owner: @mengxr (/user/mengxr) / Latest release: 0.0.1 (2014-11-27) / Apache-2.0 / ★★★★★★ (1)

1 ml

1 example

1 examples

killrweather (/package/killrweather/killrweather)

KillrWeather is a reference application (in progress) showing how to easily leverage and integrate Apache Spark, Apache Cassandra, and Apache Kafka for fast, streaming computations on time series data in asynchronous Akka event-driven environments.

from: @killrweather (/user/killrweather) / owner: @pmcfadin (/user/pmcfadin) / No release yet / ★★★★★★ (0)

1 streaming

spark-hash (/package/mrsqueeze/spark-hash)

Locality Sensitive Hashing for Apache Spark

@mrsqueeze (/user/mrsqueeze) / No release yet / ★★★★★★ (0)

1 mllib

spark-dbf (/package/mraad/spark-dbf)

Spark SQL DBF Library

@mraad (/user/mraad) / No release yet / ★★★★★ (10)

1 sql

spark-kernel (/package/ibm-et/spark-kernel)

A kernel that enables applications to interact with Apache Spark.

from: @ibm-et (/user/ibm-et) / owner: @rcsenkbeil (/user/rcsenkbeil) / No release yet / ★★★★★ (10)

1 ipython

1 foundation

1 interactive

pyspark-pictures (/package/jkthompson/pyspark-pictures)

Learn the pySpark API through pictures and simple examples

@jkthompson (/user/jkthompson) / No release yet / ★★★★★ (10)

2 tutorials

2 examples

spark (/package/synsys/spark)

Apache Camel Streaming Consumer

from: @synsys (/user/synsys) / owner: @sbrewin (/user/sbrewin) / Latest release: 1.0.0 (2015-01-26) / Apache-2.0 / ★★★★★ (10)

1 streaming

1 consumer

1 camel

spark_dbscan (/package/alitouka/spark_dbscan)

DBSCAN clustering algorithm on top of Apache Spark

@alitouka (/user/alitouka) / No release yet / ★★★★★ (10)

1 clustering

1 ml

1 dbscan

spark-scalding (/package/tresata/spark-scalding)

Use Cascading Taps and Scalding DSL with Spark — Edit

from: @tresata (/user/tresata) / owner: @koertkuipers (/user/koertkuipers) / Latest release: 0.5.0-s_2.10 (2015-11-13) / Apache-2.0 / ★★★★★ (20)

spark-sorted (/package/tresata/spark-sorted)

Secondary sort and streaming reduce for Spark

from: @tresata (/user/tresata) / owner: @koertkuipers (/user/koertkuipers) / Latest release: 0.4.0-s_2.11 (2015-11-03) / Apache-2.0 / ★★★★★ (20)

1 core

spark-kafka (/package/tresata/spark-kafka)

Low level integration of Spark and Kafka

from: @tresata (/user/tresata) / owner: @koertkuipers (/user/koertkuipers) / Latest release: 0.6.0-s_2.10 (2015-11-13) / Apache-2.0 / ★★★★★ (20)

1 streaming

spawncamping-dds (/package/FRosner/spawncamping-dds)

Data-Driven Spark allows quick data exploration based on Apache Spark

@FRosner (/user/FRosner) / No release yet / ★★★★★ (20)

spark-power-bi (/package/granturing/spark-power-bi)

Power BI API adapter for Apache Spark

@granturing (/user/granturing) / Latest release: 1.5.0_0.0.7 (2015-09-13) / Apache-2.0 /

★★★★★ (10)

1 streaming

1 sql

1 realtime

demo-python-sp (/package/brkyvz/demo-python-sp)

Pure python package used for testing Spark Packages

@brkyvz (/user/brkyvz) / Latest release: 0.4.1 (2015-08-25) / Apache-2.0 / ★★★★★ (10)

spark-mrmr-feature-selection (/package/wxC3SC6OPm8M1HXboMy/spark-mrmr-feature-selection)

Feature selection based on information gain: maximum relevancy minimum redundancy

@wxC3SC6OPm8M1HXboMy (/user/wxC3SC6OPm8M1HXboMy) / No release yet / ★★★★★ (10)

1 mllib

spark-csv2sql (/package/wxC3SC6OPm8M1HXboMy/spark-csv2sql)

Hand routine to import csv files as tables in spark sql

@wxC3SC6OPm8M1HXboMy (/user/wxC3SC6OPm8M1HXboMy) / No release yet / ★★★★★ (10)

docker-spark (/package/epahomov/docker-spark)

Docker container for spark standalone cluster.

@epahomov (/user/epahomov) / No release yet / ★★★★★ (10)

1 tools

1 deployment

spark-libFM (/package/zhengruifeng/spark-libFM)

An implement of Factorization Machines (LibFM)

@zhengruifeng (/user/zhengruifeng) / No release yet / ★★★★★ (10)

1 ml

1 mllib

1 machine learning

spark-ml-class (/package/zinniasystems/spark-ml-class)

Coursera Machine Learning class examples in Spark

from: @zinniasystems (/user/zinniasystems) / owner: @phatak-dev (/user/phatak-dev) / No release yet / ★★★★★ (10)

2 ml

2 machine learning

1 example

spark-archetype-scala (/package/mbonaci/spark-archetype-scala)

Maven archetype used to bootstrap a Spark Scala project

@mbonaci (/user/mbonaci) / Latest release: 0.9 (2015-04-24) / MIT / ★★★★★ (10)

1 Maven

1 tools

1 scala

sbt-spark-ec2 (/package/pishen/sbt-spark-ec2)

SBT plugin for spark-ec2

@pishen (/user/pishen) / No release yet / ★★★★★ (10)

1 tools

1 sbt

1 deployment

spark-solr (/package/LucidWorks/spark-solr)

Tools for reading data from Solr as a Spark RDD and indexing objects from Spark into Solr using SolrJ.

from: @LucidWorks (/user/LucidWorks) / owner: @thelabdude (/user/thelabdude) / No release yet / ★★★★★★ (10)

mypackage (/package/EronWright/mypackage)

Test Project

@EronWright (/user/EronWright) / Latest release: 0.0.13 (2015-06-11) / Apache-2.0 / ★★★★★★ (10)

dissolve-struct (/package/dalab/dissolve-struct)

Distributed solver library for large-scale structured output prediction

from: @dalab (/user/dalab) / owner: @tribhuvanesh (/user/tribhuvanesh) / No release yet / ★★★★★★ (10)

1 Support Vector Machine

1 Structured Prediction

1 machine learning

mba (/package/mfawadalam/mba)

Two way association analysis

@mfawadalam (/user/mfawadalam) / No release yet / ★★★★★★ (10)

spark-deployer (/package/pishen/spark-deployer)

Deploy Spark cluster in an easy way.

@pishen (/user/pishen) / Latest release: 0.5.1 (2015-06-25) / Apache-2.0 / ★★★★★★ (10)

1 tools

1 sbt

1 deployment

sparkboost (/package/tizfa/sparkboost)

A distributed implementation of AdaBoost.MH and MP-Boost using Apache Spark

@tizfa (/user/tizfa) / Latest release: 0.6 (2015-07-01) / Apache-2.0 / ★★★★★ (10)

1 adaboost

1 classification

1 machine learning

hivemall-spark (/package/maropu/hivemall-spark)

A Hivemall wrapper for Spark

@maropu (/user/maropu) / Latest release: 0.0.5 (2015-10-20) / Apache-2.0 / ★★★★★ (10)

1 sql

1 hive

1 machine learning

patchwork (/package/thomastriple/patchwork)

Highly Scalable Grid-Density Clustering Algorithm for Spark MLlib

@thomastriple (/user/thomastriple) / No release yet / ★★★★★ (10)

1 clustering

1 spark

1 machine learning

spark-lda (/package/EntilZha/spark-lda)

Spark package with multiple LDA implementations

@EntilZha (/user/EntilZha) / No release yet / ★★★★★ (10)

jaws-spark-sql-rest (/package/Atigeo/jaws-spark-sql-rest)

Restful service for running Spark SQL/Shark queries on top of Spark, with Mesos and Tachyon support.

from: @Atigeo (/user/Atigeo) / owner: @Adrian-Henry (/user/Adrian-Henry) / No release yet / ★★★★★ (10)

spark-job-rest (/package/Atigeo/spark-job-rest)

Restful service that enables support for multiple spark contexts created from the same server.

from: @Atigeo (/user/Atigeo) / owner: @Adrian-Henry (/user/Adrian-Henry) / No release yet / ★★★★★ (10)

sp-demo (/package/brkyvz/sp-demo)

WIP Demo Package

@brkyvz (/user/brkyvz) / No release yet / ★★★★★ (10)

xpatterns-xframe (/package/Atigeo/xpatterns-xframe)

Simplified tabular data processing library for Spark

from: @Atigeo (/user/Atigeo) / owner: @Adrian-Henry (/user/Adrian-Henry) / No release yet / ★★★★★ (10)

sparrow (/package/ypg-data/sparrow)

Scala library for converting Spark rows to case classes

from: @ypg-data (/user/ypg-data) / owner: @suhailshergill (/user/suhailshergill) / Latest release: 0.1.1-s_2.11 (2015-08-03) / Apache-2.0 / ★★★★★ (10)

1 sql 1 library 1 scala

spark-json-relay (/package/hammerlab/spark-json-relay)

SparkListener that converts SparkListenerEvents to JSON and forwards them to an external service via RPC.

from: @hammerlab (/user/hammerlab) / owner: @ryan-williams (/user/ryan-williams) / Latest release: 2.0.1 (2015-10-12) / Apache-2.0 / ★★★★★ (10)

spark-on-hbase (/package/michal-harish/spark-on-hbase)

Generic solution for scanning, joining and mutating HBase tables to and from the Spark RDDs.

@michal-harish (/user/michal-harish) / No release yet / ★★★★★ (20)

infinispan-spark (/package/infinispan/infinispan-spark)

Infinispan Spark Connector

from: @infinispan (/user/infinispan) / owner: @gustavonalle (/user/gustavonalle) / Latest release: 0.2 (2015-12-09) / Apache-2.0 / ★★★★★ (20)

1 streaming 1 sql 1 scala

spark-skewjoin (/package/tresata/spark-skewjoin)

Joins for skewed datasets in Spark

from: @tresata (/user/tresata) / owner: @koertkuipers (/user/koertkuipers) / Latest release: 0.2.0-s_2.10 (2015-11-13) / Apache-2.0 / ★★★★★ (20)

1 core

sparksql-protobuf (/package/saurfang/sparksql-protobuf)

Read SparkSQL parquet file as RDD[Protobuf]

@saurfang (/user/saurfang) / Latest release: 0.1.2-s_2.10 (2015-08-18) / Apache-2.0 / ★★★★★ (20)

1 data source 1 protobuf 1 sql

sbt-spark-submit (/package/saurfang/sbt-spark-submit)

sbt plugin for spark-submit

@saurfang (/user/saurfang) / No release yet / ★★★★★★ (10)

1 tools 1 sbt 1 deployment

twitter-stream-ml (/package/giorgioinf/twitter-stream-ml)

Machine Learning over Twitter's stream. Using Apache Spark, Web Server and Lightning Graph server.

@giorgioinf (/user/giorgioinf) / No release yet / ★★★★★★ (10)

1 ml 1 example 1 streaming

bisecting-kmeans (/package/yu-iskw/bisecting-kmeans)

This is a prototype implementation of Bisecting K-Means Clustering on Spark.

@yu-iskw (/user/yu-iskw) / Latest release: 0.1.1 (2015-08-28) / Apache-2.0 / ★★★★★★ (10)

1 clustering 1 machine learning 1 scala

TopicModeling (/package/dding3/TopicModeling)

Distributed text modeling algorithms with spark

@dding3 (/user/dding3) / No release yet / ★★★★★★ (10)

spark-mainframe-connector (/package/Syncsort/spark-mainframe-connector)

Spark mainframe connector

@Syncsort (/user/Syncsort) / Latest release: 1.0.0 (2015-09-01) / Apache-2.0 / ★★★★★★ (10)

1 input

1 data source

1 sql

sparkling-ferns (/package/CeON/sparkling-ferns)

Implementation of Random Ferns for Apache Spark

from: @CeON (/user/CeON) / owner: @matfed (/user/matfed) / Latest release: 0.2.0 (2015-10-08) / Apache-2.0 / ★★★★★★ (10)

3 machine learning

SparkGIS (/package/drubbo/SparkGIS)

GIS extension for SparkSQL

@drubbo (/user/drubbo) / No release yet / ★★★★★★ (10)

spark-eventhubs (/package/hdinsight/spark-eventhubs)

EventHubs Receiver for Spark Streaming

from: @hdinsight (/user/hdinsight) / owner: @shanyu (/user/shanyu) / No release yet / ★★★★★★ (10)

1 Azure

1 streaming

1 eventhubs

cloudphylo (/package/xingjianxu/cloudphylo)

Phylogenetic tree inference tool

@xingjianxu (/user/xingjianxu) / No release yet / ★★★★★★ (10)

1 Phylogenetic

1 Bioinformatics

spark-ext (/package/collectivemedia/spark-ext)

Spark Extension : ML transformers, SQL aggregations, etc that are missing in Apache Spark

from: @collectivemedia (/user/collectivemedia) / owner: @ezhulenev (/user/ezhulenev) / No release yet / ★★★★★★ (10)

spark-on-hpc (/package/ekasitk/spark-on-hpc)

Script to submit spark jobs on a traditional HPC cluster

@ekasitk (/user/ekasitk) / No release yet / ★★★★★★ (10)

1 deployment

spark-redis (/package/RedisLabs/spark-redis)

A connector for Spark that allows reading and writing to/from Redis cluster

from: @RedisLabs (/user/RedisLabs) / owner: @itamarhaber (/user/itamarhaber) / Latest release: 0.1.0 (2015-09-12) / BSD 3-Clause / ★★★★★★ (10)

ScalaNetwork (/package/nearbydelta/ScalaNetwork)

A Neural network implementation with Scala

@nearbydelta (/user/nearbydelta) / No release yet / ★★★★★★ (10)

1 neural network

1 machine learning

keystone (/package/amplab/keystone)

Simplifying robust end-to-end machine learning on Apache Spark.

from: @amplab (/user/amplab) / owner: @etrain (/user/etrain) / No release yet / ★★★★★★ (10)

spark-xml-utils (/package/elsevierlabs-os/spark-xml-utils)

Spark-xml-utils provides the ability to filter documents based on an xpath expression, return specific nodes for an xpath/xquery expression, or transform documents using a xslt stylesheet.

from: @elsevierlabs-os (/user/elsevierlabs-os) / owner: @dmcbeath (/user/dmcbeath) / Latest release: 1.3.0 (2015-12-04) / Apache-2.0 / ★★★★★★ (10)

[1 xml](#)[1 xpath](#)[1 xquery](#)

spark-python-knn (/package/jakac/spark-python-knn)

Function for computing K-NN in Apache Spark

@jakac (/user/jakac) / Latest release: 0.0.3 (2015-10-06) / Apache-2.0 / ★★★★★★ (10)

smartframes (/package/dondrake/smartframes)

Enhanced Python Dataframes for Spark

@dondrake (/user/dondrake) / No release yet / ★★★★★★ (10)

[1 python](#)[1 sql](#)[1 pyspark](#)

spark-FM-parallelSGD (/package/blebreton/spark-FM-parallelSGD)

Implementation of Factorization Machines on Spark using parallel stochastic gradient descent (python and scala)

@blebreton (/user/blebreton) / No release yet / ★★★★★★ (10)

1 ml

1 mllib

1 machine learning

SMV (/package/TresAmigosSD/SMV)

Spark Modularized View

from: @TresAmigosSD (/user/TresAmigosSD) / owner: @AliTajeldin (/user/AliTajeldin) / No release yet / ★★★★★★ (10)

1 core

1 sql

soda (/package/elsevierlabs-os/soda)

Solr Dictionary Annotator (Microservice for Spark)

from: @elsevierlabs-os (/user/elsevierlabs-os) / owner: @sujitpal (/user/sujitpal) / No release yet / ★★★★★★ (10)

1 application

1 tools

Mean-Shift-LSH (/package/Kybe67/Mean-Shift-LSH)

Spark implementation of Nearest Neighbours Mean Shift using LSH

@Kybe67 (/user/Kybe67) / No release yet / ★★★★★★ (10)

1 machine learning

sparkplug (/package/springnz/sparkplug)

Create composable data processing pipelines in Spark, and execute them on a cluster using simple Scala code

from: @springnz (/user/springnz) / owner: @szoio (/user/szoio) / No release yet / ★★★★★★ (10)

1 application

1 testing

1 tools

sparkxgboost (/package/rotationsymmetry/sparkxgboost)

gradient boosting tree with arbitrary user-defined loss function

@rotationsymmetry (/user/rotationsymmetry) / Latest release: 0.2.1-s_2.10 (2015-11-01) /

Apache-2.0 / ★★★★★ (10)

1 machine learning

SparkRext (/package/saurfang/SparkRext)

SparkR extension for dplyr

@saurfang (/user/saurfang) / No release yet / ★★★★★ (10)

1 sparkr

1 r

1 tools

SparkCLR (/package/skaarthik/SparkCLR)

C# language binding and extensions to Apache Spark

@skaarthik (/user/skaarthik) / No release yet / ★★★★★ (10)

1 streaming

1 examples

1 sql

mandolin (/package/project-mandolin/mandolin)

Large-scale Machine Learning using Apache Spark

from: @project-mandolin (/user/project-mandolin) / owner: @wellner (/user/wellner) / No release yet / ★★★★★ (10)

1 machine learning

Estus (/package/EstusDev/Estus)

Estus Scientific Library

@EstusDev (/user/EstusDev) / No release yet / ★★★★★ (10)

[1 machine learning](#)

drunken-data-quality (/package/FRosner/drunken-data-quality)

Some utility classes for checking data quality in Spark

@FRosner (/user/FRosner) / No release yet / ★★★★★ (10)

wookie (/package/elyast/wookie)

Library for building data products

@elyast (/user/elyast) / No release yet / ★★★★★ (10)

spark-knn (/package/saurfang/spark-knn)

k-Nearest Neighbors algorithm on Spark

@saurfang (/user/saurfang) / Latest release: 0.1.0 (2015-11-23) / Apache-2.0 / ★★★★★ (10)

[2 ml](#)[2 machine learning](#)

cookie-datasets (/package/cookieai/cookie-datasets)

Popular ML Datasets for Spark ML (MNIST, IRIS, CIFAR)

from: @cookieai (/user/cookieai) / owner: @EronWright (/user/EronWright) / Latest release: 0.1.0 (2015-12-22) / Apache-2.0 / ★★★★★ (10)

[1 data source](#)[1 machine learning](#)

spark-kaggle-examples (/package/Lewuathe/spark-kaggle-examples)

Kaggle Job repository

@Lewuathe (/user/Lewuathe) / Latest release: 0.0.1 (2015-11-22) / Apache-2.0 / ★★★★★ (10)

[1 examples](#)

spark-google-analytics (/package/crealytics/spark-google-analytics)

A Spark package for retrieving data from Google Analytics

from: @crealytics (/user/crealytics) / owner: @nightscape (/user/nightscape) / Latest release: 0.8.1 (2015-12-14) / Apache-2.0 / ★★★★★ (10)

spark-netezza (/package/SparkTC/spark-netezza)

Spark Netezza Connector

from: @SparkTC (/user/SparkTC) / owner: @lresende (/user/lresende) / No release yet / ★★★★★ (10)

spark-igmulus (/package/IGNF/spark-igmulus)

Spark Package to read and write PLY, LAS and XYZ lidar point clouds using Spark SQL.

from: @IGNF (/user/IGNF) / owner: @mbredif (/user/mbredif) / Latest release: 0.1.0-s_2.10 (2015-12-08) / Apache-2.0 / ★★★★★ (10)

[1 geospatial](#)[1 data source](#)[1 sql](#)

graphx-citymap (/package/fancellu/graphx-citymap)

CityMap coding test plus 3 solutions, 1 with Spark/GraphX

@fancellu (/user/fancellu) / No release yet / ★★★★★ (10)

[1 graph](#)[1 example](#)

spark-benford-analysis (/package/dvgodoy/spark-benford-analysis)

Benford Analysis for Spark package.

@dvgodoy (/user/dvgodoy) / Latest release: v0.0.1 (2015-12-13) / Apache-2.0 / ★★★★★ (10)

spark-google-adwords (/package/crealytics/spark-google-adwords)

A library for querying Google AdWords data with Apache Spark, for Spark SQL and DataFrames

from: @crealytics (/user/crealytics) / owner: @nightscape (/user/nightscape) / Latest release: 0.8.2 (2015-12-14) / Apache-2.0 / ★★★★★ (10)

sparkpipe (/package/unchartedsoftware/sparkpipe)

Modular, non-linear data pipeline framework for Spark

from: @unchartedsoftware (/user/unchartedsoftware) / owner: @GhnuBerath (/user/GhnuBerath) / No release yet / ★★★★★ (10)

shc (/package/zhzhan/shc)

DataFrame support for HBase

@zhzhan (/user/zhzhan) / No release yet / ★★★★★ (10)

apache-spark-node (/package/henridf/apache-spark-node)

Node.js bindings for Apache Spark

@henridf (/user/henridf) / No release yet / ★★★★★ (10)

spark-statsd (/package/vidhyaarvind/spark-statsd)

StatsD Metrics Reporter for Spark metrics

@vidhyaarvind (/user/vidhyaarvind) / No release yet / ★★★★★★ (10)

sscheck (/package/juanrh/sscheck)

ScalaCheck for Spark

@juanrh (/user/juanrh) / No release yet / ★★★★★★ (10)

1 streaming

1 testing

1 tools

als-ngc (/package/mbhynes/als-ngc)

NGC acceleration of ALS for computing low rank matrix factorizations for Collaborative Filtering

@mbhynes (/user/mbhynes) / No release yet / ★★★★★★ (10)

Spark.statistics (/package/hhbyyh/Spark.statistics)

Assembly of fundamental statistics implemented based on Apache Spark

@hhbyyh (/user/hhbyyh) / No release yet / ★★★★★★ (10)

1 statistics

1 machine learning

spark-hdp (/package/tund/spark-hdp)

Implementations of Hierarchical Dirichlet Process (HDP) on Spark

@tund (/user/tund) / No release yet / ★★★★★★ (10)

sparkpipe-core

(/package/unchartedsoftware/sparkpipe-core)

Modular, non-linear data pipeline framework for Spark

from: @unchartedsoftware (/user/unchartedsoftware) / owner: @Ghnuberath (/user/Ghnuberath) /

Latest release: 0.9.4 (2016-01-11) / Apache-2.0 / ★★★★★ (10)

1 etl

1 data processing
