

A Note from Dr. Bent

While there is a single course instructor listed for this course, this course would not have been possible without a number of incredible contributors and collaborators:

[Megan Lancaster](#) , Senior Learning Experience Designer

From the inception of this course, Megan has been instrumental in the planning and development of the course content. Her expertise in creating learning experiences and her ability to keep everything organized is unmatched. Megan marched bravely through the course content - code bugs and all - to bring this course to you.

[Devon Henry](#) , Media Producer

Devon played a number of roles to bring this course to you: from director to producer to video editor to hair stylist (no, really!). Devon is the true mastermind behind all of the lecture videos you watched in the course. I'm still not quite sure how she managed to juggle all of the things. All I had to do was talk!

[Nick Janes](#) , Learning Media Designer

The video introductions and logos were the brainchild of Nick, who took on the challenge of bringing the course to life visually. Nick also provided invaluable feedback on the design and development of the slide decks used in the lecture videos.

[Prof. Jon Reifschneider](#) , Director of the Duke Artificial Intelligence MEng

Jon sparked the idea of putting together this Coursera course over Summer 2024. A Coursera instructor veteran himself, he provided feedback and advice throughout the inception and development of the course.

Others who helped bring this course to life include:

- [Dr. Brad Fox](#) , Executive Director of Professional Masters Programs
- [Quentin Ruiz-Esparza](#) , Director for Digital Product Strategy and Design
- **Mich Donovan**, Manager of Academic Media Production
- [Duke Learning Innovation and Lifetime Education \(LILE\)](#) 

Thank you to the contributors of the Interpretable Machine Learning Book (specifically primary author [Christoph Molnar](#) ) for several figures used in the course lectures. And thank you to [Neel Nanda](#)  for his work on Mechanistic Interpretability (his tutorial on TransformerLens is included in the course as an educational resource). Thank you to the authors of the TCAV ([Been Kim](#) , et.al.) for their tutorial that the TCAV lab is based on. Thank you to [Ryan Bolick](#)  for his contribution of a resource on UX for AI.

A special thanks to all of the incredible researchers who took on (and continue to take on) the challenge of interpretability and explainability in machine learning. I hope I have done your work justice with this course.

And THANK YOU for participating in this course. It is critical that we develop AI responsibly, and I am thrilled that you decided to take it upon yourself to learn more about building AI systems that are interpretable, explainable, and transparent. Please share what you have learned with others - we can all benefit from having more responsible AI.

Feel free to [connect with me on LinkedIn](#) . I am looking forward to hearing your thoughts on the course.

Cheers,

Brinnae Bent, PhD

[Go to next item](#)

✓ Completed

 Like  Dislike  Report an issue