


Dr. Bent's Responsible AI Reading List

What books to read if you want to learn more:



Unmasking AI by Joy Buolamwini

Why you should read it: An autobiographical journey towards uncovering algorithmic bias and standing up for what is right. A nonfiction that reads like a novel.

The Alignment Problem by Brian Christian

Why you should read it: Survey of AI Safety including history, science, philosophy, and limitations.

A Hacker's Mind by Bruce Schneier

Why you should read it: Adversarial attacks, defenses, and red teaming are all emerging trends that are closely tied to XAI. In particular, the chapter on Hacking AI is extremely relevant to XAI and Human-AI Interaction.