

Never miss a tutorial:

in

Twitter

on

Facebook

on

LinkedIn

on

Medium

on

YouTube

Machine Learning Mastery

Making Developers Awesome at Machine Learning

you:

A Tour of Machine Learning Algorithms

Click to Take the FREE Algorithms Crash-Course

Search...

Supervised and Unsupervised Machine Learning

Bagging and Random Forest Ensemble Algorithms for Machine Learning

by Jason Brownlee on December 3, 2020 in Machine Learning Algorithms

Machine Learning Algorithms Mini-Course

Tweet

Share

Random Forest is one of the most popular and most powerful machine learning algorithms. It is a type of ensemble machine learning algorithm called

Stacked Generalization (Blending)

Machine Learning

In this post you will discover the Bagging ensemble algorithm and the Random Forest algorithm for predictive modeling. After reading this post you will know about:

Logistic Regression for Machine Learning

The bootstrap method for estimating statistical quantities from samples.

The Bootstrap Aggregation algorithm for creating multiple different models from a single training dataset.

The Random Forest algorithm that makes a small tweak to Bagging and results in a very powerful classifier.

Loving the Tutorials?

This post was written for developers and assumes no background in statistics or mathematics. The post focuses on how the algorithm works and how to use it for predictive modeling problems. Book is where you'll find the *Really Good* stuff.

If you have any questions, leave a comment and I will do my best to answer.

SEE WHAT'S INSIDE

Kick-start your project with my new book *Master Machine Learning Algorithms*, including *step-by-step tutorials* and the *Excel Spreadsheet* files for all examples.

Let's get started.

Bootstrap Method

Before we get to Bagging, let's take a quick look at an important foundation technique called the bootstrap.

The bootstrap is a powerful statistical method for estimating a quantity from a data sample. This is easiest to understand if the quantity is a descriptive statistic such as a mean or a standard deviation.

Let's assume we have a sample of 100 values (x) and we'd like to get an estimate of the mean of the sample.

We can calculate the mean directly from the sample as:

$$\text{mean}(x) = 1/100 * \text{sum}(x)$$

We know that our sample is small and that our mean has error in it. We can improve the estimate of our mean using the bootstrap procedure:

- Create many (e.g. 1000) random sub-samples of our dataset with replacement (meaning we can select the same value multiple times).
- Calculate the mean of each sub-sample.
- Calculate the average of all of our collected means and use that as our estimated mean for the data.

For example, let's say we used 3 resamples and got the mean values 2.3, 4.5 and 3.3. Taking the average of these we could take the estimated mean of the data to be 3.367.

This process can be used to estimate other quantities like the standard deviation and even quantities used in machine learning algorithms, like learned coefficients.

Get your FREE Algorithms Mind Map

Stacked Generalization (Blending)

Gradient Boosted Regression Trees (GBRT)

Radial Basis Function Network (RBFN)

Perceptron

Back-Propagation

Hopfield Network

Ridge Regression

Lasso and Selection Operator (LASSO)

Elastic Net

Least Angle Regression (LARS)

Cubist

Neural Networks

Regularization

Mach

I've created a handy mind map of 60+ algorithms organized by type.

Download it, print it and use it.

Download For Free

Also get exclusive access to the machine learning algorithms email mini-course.

Sample of the handy machine learning algorithms mind map.

Bootstrap Aggregation (Bagging)

Bootstrap Aggregation (or Bagging for short), is a simple and very powerful ensemble method.

An ensemble method is a technique that combines the predictions from multiple machine learning algorithms together to make more accurate predictions than any individual model.

Bootstrap Aggregation is a general procedure that can be used to reduce the variance for those algorithm that have high variance. An algorithm that has high variance are decision trees, like classification and regression trees (CART).

Decision trees are sensitive to the specific data on which they are trained. If the training data is changed (e.g. a tree is trained on a subset of the training data) the resulting decision tree can be quite different and in turn the predictions can be quite different.

Bagging is the application of the Bootstrap procedure to a high-variance machine learning algorithm, typically decision trees.

Let's assume we have a sample dataset of 1000 instances (x) and we are using the CART algorithm. Bagging of the CART algorithm would work as follows.

- Create many (e.g. 100) random sub-samples of our dataset with replacement.
- Train a CART model on each sample.
- Given a new dataset, calculate the average prediction from each model.

For example, if we had 5 bagged decision trees that made the following class predictions for a in input sample: blue, blue, red, blue and red, we would take the most frequent class and predict blue.

When bagging with decision trees, we are less concerned about individual trees overfitting the training data. For this reason and for efficiency, the individual decision trees are grown deep (e.g. few training samples at each leaf-node of the tree) and the trees are not pruned. These trees will have both high variance and low bias. These are important characteristics of sub-models when combining predictions using bagging.

The only parameters when bagging decision trees is the number of samples and hence the number of trees to include. This can be chosen by increasing the number of trees on run after run until the accuracy begins to stop showing improvement (e.g. on a cross validation test harness). Very large numbers of models may take a long time to prepare, but will not overfit the training data.

Just like the decision trees themselves, Bagging can be used for classification and regression problems.

Random Forest

Random Forests are an improvement over bagged decision trees.

A problem with decision trees like CART is that they are greedy. They choose which variable to split on using a greedy algorithm that minimizes error. As such, even with Bagging, the decision trees can have a lot of structural similarities and in turn have high correlation in their predictions.

Combining predictions from multiple models in ensembles works better if the predictions from the sub-models are uncorrelated or at best weakly correlated.

Random forest changes the algorithm for the way that the sub-trees are learned so that the resulting predictions from all of the subtrees have less correlation.

It is a simple tweak. In CART, when selecting a split point, the learning algorithm is allowed to look through all variables and all variable values in order to select the most optimal split-point. The random forest algorithm changes this procedure so that the learning algorithm is limited to a random sample of features of which to search.

The number of features that can be searched at each split point (m) must be specified as a parameter to the algorithm. You can try different values and tune it using cross validation.

- For classification a good default is: m = sqrt(p)
- For regression a good default is: m = p/3

Where m is the number of randomly selected features that can be searched at a split point and p is the number of input variables. For example, if a dataset had 25 input variables for a classification problem, then:

- m = sqrt(25)
- m = 5

×

https://machinelearningmastery.com/bagging-and-random-forest-ensemble-algorithms-for-machine-learning/

1/7



A Tour of Machine Learning Algorithms

 Supervised and Unsupervised Machine Learning Algorithms

experiments, you could perform where some features are removed from the dataset

Learning

DREAM

Android 6.6 | Play Store | 40% off | WhatsApp | Sale is Live

Loving the Tutorials?

The Machine Learning Algorithms EBook is

- Do you have any questions about this post or the Bagging or Random Forest Ensemble algorithms?

Leave a comment and ask your question and I will do my best to answer it.

[Share](#)
[Tweet](#)
[Share](#)

Never miss a tutorial!
Sign Up



**Jason Brownlee**

January 6, 2020 at 7:10 am #

REPLY

↗



A Tour of Machine Learning Algorithms

It's bagging and boosting we typically use one algorithm type and traditionally this is a decision tree.

**Supervised Machine Learning Algorithms**

January 6, 2020 at 9:00 pm #

REPLY

↗



@Jason Brownlee can u Elaborate all concepts in machine learning with real time examples?

**Machine Learning Algorithms Mini-Course**

January 6, 2020 at 7:11 am #

REPLY

↗



@Jason Brownlee That's what the blog covers.

**Logistic Regression Tutorial for Machine Learning**

January 6, 2020 at 7:11 am #

REPLY

↗



Open Enroll in Machine Learning

Hey Jason,


You're doing a great job here. I didn't know anything about machine learning until I found your site. Believe it or not, I follow it pretty well. #LoveMath.

Loving the Tutorials?

The Machine Learning Algorithms eBook is available for free download

REPLY

↗



@Jason Brownlee

>> SEE WHAT'S INSIDE

**Shilpa Anil**

January 4, 2017 at 4:52 am #

REPLY

↗

Sir, I have to predict daily air temperature values using random forest regression and i have 5 input varibales. Could you please explain how splitting is performed in regression? exactly what is done at each split point?

**Jason Brownlee**

January 4, 2017 at 8:59 am #

REPLY

↗

Here is some advice on splitting time series data for machine learning:
<https://machinelearningmastery.com/backtest-machine-learning-models-time-series-forecasting/>

**Ola**

January 13, 2017 at 3:14 am #

REPLY

↗

Thank You for that post! It helps me to clarify decision about using Random Forest in my Master's Thesis analysis.

**Jason Brownlee**

January 13, 2017 at 9:14 am #

REPLY

↗

I'm glad to hear it.

**Joe Ellsworth**

January 24, 2017 at 1:06 pm #

REPLY

↗

Hi Jason, I liked your article. I am working on a Quantized classifier and would love to collaborate on an article. <https://bitbucket.org/joexdobs/ml-classifier-gesture-recognition>

**Jason Brownlee**

January 25, 2017 at 9:57 am #

REPLY

↗

Good luck with your project Joel!

**Chandrakant Pattekar**

May 28, 2017 at 6:49 pm #

REPLY

↗

Very well explained

**Jason Brownlee**

June 2, 2017 at 12:10 pm #

REPLY

↗

Thanks, I'm glad it helped.

**Mercy Peter**

June 23, 2017 at 3:26 pm #

REPLY

↗

Very crisp and clear explanations, nailed to the point.

**Jason Brownlee**

June 24, 2017 at 7:58 am #

REPLY

↗

Thanks Mercy.

**Paul C**

June 28, 2017 at 1:35 am #

REPLY

↗

Hi Jason, it's not true that bootstrapping a sample and computing the mean of the bootstrap sample means "improves the estimate of the mean." The standard MLE (I.e just the sample mean) is the best estimate of the population mean. Bootstrapping is great for many things but not for giving a better estimate of a mean.

**Jason Brownlee**

June 28, 2017 at 6:29 am #

REPLY

↗

Perhaps. Also, check this:
[https://en.wikipedia.org/wiki/Bootstrapping_\(statistics\)#Estimating_the_distribution_of_sample_mean](https://en.wikipedia.org/wiki/Bootstrapping_(statistics)#Estimating_the_distribution_of_sample_mean)

**Felipe**

July 30, 2017 at 1:49 am #

REPLY

↗

Hi, Jason! I'm reading your article and helped me understand the context about bagging. I need to implement a Bagging for Caltech 101 dataset and I do not know how can I start. I am programing something in Matlab but I dont know how can I create a file from Caltech101 to Matlab and studying the data to create Ensemble. I am little confusing! Do you have any consideration to help me?

**Jason Brownlee**

July 30, 2017 at 7:49 am #

REPLY

↗

Sorry, I do not have matlab examples. I may not the best person to give you advice.

**Sonal**

August 8, 2017 at 12:43 am #

REPLY

↗

Hi Jason, Can you recommend any C++ libraries (open source or commercially licensed) with an accurate implementation of decision trees and its variants(bagged, random forests)?

**Jason Brownlee**

August 8, 2017 at 7:50 am #

REPLY

↗

Not really. Perhaps xgboost – I think it is written in cpp.

**Himanshu**

November 17, 2017 at 2:24 am #

REPLY

↗

When label data is very less in my training how can I use bagging to validate performance on the full distribution of training?

**Jason Brownlee**

November 17, 2017 at 9:27 am #

REPLY

↗

Not sure I follow, can you elaborate?

**Himanshu**

November 17, 2017 at 9:23 pm #

REPLY

↗

I mean out of 100k training data I have 2k labeled, so can I use bagging to label rest of my unlabeled data in training data set, I will do cross validation before bagging within 2k labelled

**Jason Brownlee**

November 18, 2017 at 10:16 am #

REPLY

↗

You could build a model on the 2K and predict labels for the remaining 100k, and you will need to test a suite of methods to see what works best using cross validation.

**Rajit Nair**

December 4, 2017 at 5:08 pm #

REPLY

↗

Hi Jason, Your blogs are always very useful to me, but it will be more useful when you take an example and explain the whole process. But anyways you blogs are very new and interesting

**Jason Brownlee**

December 5, 2017 at 5:41 am #

REPLY

↗

Thanks Rajit.

**Simon Ting**

December 26, 2017 at 9:45 pm #

REPLY

↗

Nice tutorial, Jason! I'm a bit confuse about the "Variable Importance" part, which step in bagging algorithm do you need to calculate the importance of each variables by estimate the error function drops?

**Jason Brownlee**

December 27, 2017 at 5:19 am #

REPLY

↗

Thanks,

I'm not sure I follow, perhaps you can restate the question?

**sachindeo chavan**

January 14, 2018 at 10:33 am #

REPLY

↗

Dear Jason, I'm new to regression am a student of MSc Big Data Analytics University of Liverpool UK. My query is on Random Forest, is Random Forest non-parametric regression model? regards sachin

**Jason Brownlee**

January 15, 2018 at 6:55 am #

REPLY

↗

Yes.

**Alejandra**

January 23, 2018 at 6:56 pm #

REPLY

↗

Thanks for making it clear. I have not enough background (I am a journalist) and was easy to understand. Thanks for sharing your knowledge!

**Jason Brownlee**

January 24, 2018 at 9:52 am #

REPLY

↗

✕

<https://machinelearningmastery.com/bagging-and-random-forest-ensemble-algorithms-for-machine-learning/>

37

A split point uses one value for one feature. Different values for the same or different features can be reused, even the same value for the same feature – although I doubt it.



Picked for You by Popkin

February 10, 2019 at 2:36 am #

REPLY ↗


A little bit about Machine Learning Algorithms

A I understand this post, but I'm getting confused as I read up on ensembles. Specifically, is applying them...

option 1: as simple as just choosing to use an ensemble algorithm (I'm using Random Forest and AdaBoost)

So I think it's option 1, but as mentioned above some of the reading I've been doing is confusing me.

Machine Learning Algorithms Mini-Course

 **Jason Brownlee**

February 10, 2019 at 9:45 am #

REPLY ↗

What is confusing you exactly?

Logistic Regression for Machine Learning

 **Venu**

February 27, 2019 at 1:02 am #

REPLY ↗


Loving the Tutorials?

So it means each tree in the random forest will have low bias and high variance?

At the moment, I am using the **Really Good** stuff.

Thanks,

Venu >> SEE WHAT'S INSIDE

 **Jason Brownlee**

February 27, 2019 at 7:33 am #

REPLY ↗

Correct!

 **Aisha**

May 11, 2019 at 12:56 am #

REPLY ↗

Hello sir,

Thanks for your article.

I have a question about time series forecasting with bagging.

In fact my base is composed of 500 days, each day is a time series (database: 24 lines (hours), 500 columns (days))

I want to apply a bagging to predict the 501 day.

How can I apply this technique given it resamples the base into subsets randomly and each subset makes one-day forecasting at random.

Is the result of the aggregation surely the 501 day?

thank you

 **Jason Brownlee**

May 11, 2019 at 6:18 am #


REPLY ↗

This post will help to frame your data:

<https://machinelearningmastery.com/convert-time-series-supervised-learning-problem-python/>

And this:

<https://machinelearningmastery.com/time-series-forecasting-supervised-learning/>

 **Sanjeev Kumar TM**

July 23, 2019 at 10:59 pm #

REPLY ↗

Hello,

Actually I trained the model with 4 predictors and later based on predictor importance one variable is not at all impact on response so I removed that parameter and trained the model but I am getting error obtained during 3 predictors is less as compared with 4 predictor model. will u please help me out why I am getting this error difference if I removed the parameter if it is not at all related to the response variable is reducing error or the error is same please help me out.

 **Jason Brownlee**

July 24, 2019 at 7:55 am #

REPLY ↗

It is likely that the parameter that is "not useful" has nonlinear interactions with the other parameters and is in fact useful.

 **Sanjeev Kumar TM**

July 23, 2019 at 11:00 pm #

REPLY ↗


How to get the coefficient of the predictor weights in ensemble boosted tree model.

 **Jason Brownlee**

July 24, 2019 at 7:58 am #

REPLY ↗

You don't, they are not useful/interpretable.

 **Priyanka Dave**

July 26, 2019 at 9:37 pm #


REPLY ↗

Hi Jason, I have total 47 input columns and 15 output columns (all are continuous values). Currently I am working on Random forest regression model.

My question is:

- 1) Can we define input -> output correlation or output -> output correlation ?
- 2) Can we tell model that particular these set of inputs are more powerful ?
- 3) Can we do sample wise classification ?
- 4) It is giving 98% accuracy on training data but still I am not getting expected result. Because model can not identify change in that particular input.

Thanks.

 **Jason Brownlee**

July 27, 2019 at 6:09 am #

REPLY ↗


The algorithm will learn the relationships/correlations that are most relevant to making a prediction, no need to specify them.

You can make per-sample predictions, if you're using Python, here's an example:

<https://machinelearningmastery.com/make-predictions-scikit-learn/>

I recommend evaluating the model on a hold out test set, or better yet using cross validation:


<https://machinelearningmastery.com/k-fold-cross-validation/>

 **Briana T**

August 7, 2019 at 2:12 am #

REPLY ↗

I just wanted to say that this explanation is so good and easy to follow! Thank you for providing this.

 **Jason Brownlee**

August 7, 2019 at 8:02 am #

REPLY ↗

Thanks, I'm glad it helped!

 **Luis Angel**

November 4, 2019 at 5:27 am #

REPLY ↗

Very thankful for your post.


I only have a simple question. The imbalanced sample could affect the performance of the algorithm?

 **Jason Brownlee**

November 4, 2019 at 6:49 am #

REPLY ↗


Yes.

 **Neha Manpreet**

November 16, 2019 at 1:25 am #

REPLY ↗


If bagging uses the entire feature space then in python we have max_features option in BaggingClassifier. Why we have this option of max_features ?

 **Jason Brownlee**

November 16, 2019 at 7:26 am #

REPLY ↗

To limit the feature on each tree.

 **Daniel**

January 4, 2020 at 7:02 pm #

REPLY ↗


Hi Jason,

I am confused on bootstrapping: how can we have a 'better mean' than the calculated one?

Why do I want to estimate the mean instead of calculating it?

THanks,

Daniel

 **Jason Brownlee**

January 5, 2020 at 7:03 am #

REPLY ↗

A better estimate of the population mean from the data sample.

Recall that the population is all data, sample is a subset we actually have.

Perhaps see this tutorial:


<https://machinelearningmastery.com/a-gentle-introduction-to-the-bootstrap-method/>

 **Daniel**

January 6, 2020 at 8:04 am #

REPLY ↗


Got it, thanks.

 **Jason Brownlee**

January 6, 2020 at 1:30 pm #

REPLY ↗

You're welcome.

 **Anil**

January 9, 2020 at 9:03 am #

REPLY ↗

Hello, Jason,

My question is:


- Does the random forest algorithm include bagging by default? So when I use the random forest algorithm, do I actually do bagging?
- If the random forest algorithm includes bagging by default and I apply bagging to my data set first and then use the random forest algorithm, can I get a higher success rate or a meaningful result?

 **Jason Brownlee**

January 9, 2020 at 1:42 pm #

REPLY ↗

Yes and no. Think of it bagging by feature rather than by sample. You can also bag by sample by using a bootstrap sample for each tree.

 **Paula Garcia Llamas**

February 7, 2020 at 8:59 pm #

REPLY ↗

Hi Jason,

i am a bit confused with bagging in regression.

I run random forest with 1000 total observations, I set ntree to 1000 and I calculate the mean-squared error estimate and thus, the variance explained based on the out-of-bag.

If my ntree is 1000, that means that the number of bootstrap samples is 1000, each containing, by default, two thirds of the sampled points and one third is used to get predictions out-of-bag, is this correct?

